

Update 2015

Portfolio Optimization **with R/Rmetrics**

Diethelm Würtz
Tobias Setz
Yohan Chalabi
William Chen
Andrew Ellis



R/Rmetrics eBook Series

"R/Rmetrics eBooks" is a series of electronic books and user guides aimed at students, and practitioners entering the increasing field of using R/Rmetrics software in the analysis of financial markets.

Book Suite:

Basic R for Finance (2010),
Diethelm Würtz, Tobias Setz, Yohan Chalabi, Longhow Lam, Andrew Ellis

Chronological Objects with Rmetrics (2010),
Diethelm Würtz, Tobias Setz, Yohan Chalabi, Andrew Ellis

Financial Market Data for R/Rmetrics (2010)
Diethelm Würtz, Tobias Setz, Andrew Ellis, Yohan Chalabi

Portfolio Optimization with R/Rmetrics (2010),
Diethelm Würtz, Tobias Setz, William Chen, Yohan Chalabi, Andrew Ellis

Asian Option Pricing with R/Rmetrics (2010)
Diethelm Würtz

Indian Financial Market Data for R/Rmetrics (2010)
Diethelm Würtz, Mahendra Mehta, Andrew Ellis, Yohan Chalabi

Free Documents:

A Discussion of Time Series Objects for R in Finance (2009)
Diethelm Würtz, Yohan Chalabi, Andrew Ellis

Long Term Statistical Analysis of US Asset Classes (2011)
Diethelm Würtz, Haiko Bailer, Yohan Chalabi, Fiona Grimson, Tobias Setz

R/Rmetrics Workshop Meielisalp 2010
Proceedings of the Meielisalp Summer School and Workshop 2010

Editor: Diethelm Würtz

R/Rmetrics Workshop Singapore 2010

Proceedings of the Singapore Workshop 2010

Editors: Diethelm Würtz, Mahendra Mehta, David Scott, Juri Hinz

Contributed Authors:

tinn-R Editor (2010)

José Cláudio Faria, Philippe Grosjean, Enio Galinkin Jelihovschi and Riccardo Pietrobon

Topics in Empirical Finance with R and Rmetrics (2013)

Patrick Hénaff

Under Preparation:

Advanced Portfolio Optimization with R/Rmetrics (2014),

Diethelm Würtz, Tobias Setz, Yohan Chalabi

R/Rmetrics Meielisalp 2011

Proceedings of the Meielisalp Summer School and Workshop 2011

Editor: Diethelm Würtz

R/Rmetrics Meielisalp 2012

Proceedings of the Meielisalp Summer School and Workshop 2012

Editor: Diethelm Würtz

PORTFOLIO OPTIMIZATION WITH R/RMETRICS

DIETHELM WÜRTZ

TOBIAS SETZ

YOHAN CHALABI

WILLIAM CHEN

ANDREW ELLIS

JANUARY 2015

Series Editors:

Professor Dr. Diethelm Würtz
Institute for Theoretical Physics and
Curriculum for Computational Science
ETH - Swiss Federal Institute of Technology
Hönggerberg, HIT G 32.3
8093 Zurich

Dr. Martin Hanf
Finance Online GmbH
Zeltweg 7
8032 Zurich

Contact Address:

Rmetrics Association
Zeltweg 7
8032 Zurich
info@rmetrics.org

Publisher:

Finance Online GmbH
Swiss Information Technologies
Zeltweg 7
8032 Zurich

Authors and Contributors:

Diethelm Würtz, ETH Zurich
Tobias Setz, ETH Zurich
Yohan Chalabi, ETH Zurich
William Chen, University of Auckland
Andrew Ellis, Finance Online GmbH Zurich

ISBN: 978-3-906041-01-8 (Update 2015)

© 2010-2015, Rmetrics Association and Finance Online GmbH, Zurich

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions, except that this permission notice may be stated in a translation approved by the Rmetrics Association, Zurich.

Limit of Liability/Disclaimer of Warranty: While the publisher and authors have used their best efforts in preparing this book, they make no representations or warranties with respect to the accuracy or completeness of the contents of this book and specifically disclaim any implied warranties of merchantability or fitness for a particular purpose. No warranty may be created or extended by sales representatives or written sales materials. The advice and strategies contained herein may not be suitable for your situation. You should consult with a professional where appropriate. Neither the publisher nor authors shall be liable for any loss of profit or any other commercial damages, including but not limited to special, incidental, consequential, or other damages.

Trademark notice: Product or corporate names may be trademarks or registered trademarks, and are used only for identification and explanation, without intent to infringe.

DEDICATION

*This book is dedicated to all those who
have helped make Rmetrics what it is today:
The leading open source software environment in
computational finance and financial engineering.*

PREFACE

ABOUT THIS BOOK

This is a book about portfolio optimization from the perspective of computational finance and financial engineering. Thus the main emphasis is to briefly introduce the concepts and to give the reader a set of powerful tools to solve the problems in the field of portfolio optimization.

This book divides roughly into six parts. The first two parts, Chapters 1-10, are dedicated to the exploratory data analysis of financial assets, the third part, Chapters 11-14, to the framework of portfolio design, selection and optimization, the fourth part, Chapters 15-20, to the mean-variance portfolio approach, the fifth part, Chapters 21-24, to the mean-conditional value-at-risk portfolio approach, and the sixth part, Chapters 25-27, to portfolio backtesting and benchmarking.

COMPUTATIONS

In this book we use the statistical software environment R to perform our computations. R is an advanced statistical computing system with very high quality graphics that is freely available for many computing platforms. It can be downloaded from the [CRAN server](http://cran.r-project.org)¹ (central repository), and is distributed under the GNU Public Licence. The R project was started by Ross Ihaka and Robert Gentleman at the University of Auckland. The R base system is greatly enhanced by extension packages. R provides a command line driven interpreter for the S language. The dialect supported is very close to that implemented in S-Plus. R is an advanced system and provides powerful state-of-the-art methods for almost every application in statistics.

Rmetrics is a collection of several hundreds of R functions and enhances the R environment for computational finance and financial engineering. Source packages of Rmetrics and compiled MS Windows and Mac OS X

¹<http://cran.r-project.org>

binaries can be downloaded from CRAN and the development branch of Rmetrics can be downloaded from the [R-Forge repository](#)².

AUDIENCE BACKGROUND

The material presented in this book was originally written for my students in the areas of empirical finance and financial econometrics. However, the audience is not restricted to academia; this book is also intended to offer researchers and practitioners in the finance industry an introduction to using the statistical environment R and the Rmetrics packages for modelling and optimizing portfolios.

It is assumed that the reader has a basic familiarity with the R statistical environment. A background in computational statistics and finance and in financial engineering will be helpful. Most importantly, the authors assume that the reader is interested in analyzing and modelling financial data sets and in designing and optimizing portfolios.

Note that the book is not only intended as a user guide or as a reference manual. The goal is also that you learn to interpret and to understand the output of the R functions and, even more importantly, that you learn how to modify and how to enhance functions to suit your personal needs. You will become an R developer and expert, which will allow you to rapidly prototype your models with a powerful scripting language and environment.

GETTING HELP

There are various manuals available on the CRAN server as well as a list of frequently asked questions (FAQ). The [FAQ document](#)³ ranges from basic syntax questions to help on obtaining R and downloading and installing R packages. The [manuals](#)⁴ range from a basic introduction to R to detailed descriptions of the R language definition or how to create your own R packages. The manuals are described in more detail in [Appendix C](#).

We also suggest having a look at the [mailing lists](#)⁵ for R and reading the general instructions. If you need help for any kind of R and/or Rmetrics problems, we recommend consulting [r-help](#)⁶, which is R's main mailing list. R-help has become quite an active list with often dozens of messages per day. [r-devel](#)⁷ is a public discussion list for R 'developers' and 'pre-testers'. This list is for discussions about the future of R and pre-testing

²<http://r-forge.r-project.org/projects/rmetrics/>

³<http://cran.r-project.org/doc/FAQ/R-FAQ.html>

⁴<http://cran.r-project.org/manuals.html>

⁵<http://www.r-project.org/mail.html>

⁶<https://stat.ethz.ch/mailman/listinfo/r-help>

⁷<https://stat.ethz.ch/mailman/listinfo/r-devel>

of new versions. It is meant for those who maintain an active position in the development of R. Also, all bug reports are sent there. And finally, [r-sig-finance](https://stat.ethz.ch/mailman/listinfo/r-sig-finance)⁸ is the ‘Special Interest Group’ for R in finance. Subscription requests to all mailing lists can be made by using the usual confirmation system employed by the mailman software.

GETTING STARTED

R can be downloaded and installed from the [CRAN](http://cran.r-project.org/)⁹ (Comprehensive R Archive Network) web site. Contributed R packages can also be downloaded from this site. Alternatively, packages can be installed directly in the R environment. A list of R packages accompanied by a brief description can be found on the web site itself, or, for financial and econometrics packages, from the [CRAN Task View](http://cran.r-project.org/web/views/Finance.html)¹⁰ in finance and econometrics. This task view contains a list of packages useful for empirical work in finance and econometrics grouped by topic.

To install all packages required for the examples of this eBook we recommend that you install the packages `cluster`, `mvoutlier`, `pastecs` and `fPortfolio` including its dependencies. This can be done with the following command in the R environment.

```
> install.packages(c("cluster", "mvoutlier", "pastecs", "fPortfolio"),  
+                  repos = "http://cran.r-project.org")
```

To update your installed packages use:

```
> update.packages(repos = "http://cran.r-project.org")
```

If there is no binary package for your operating system, you can install the package from source by using the argument `type = "source"`. The [R Installation and Administration](http://cran.r-project.org/doc/manuals/R-admin.html)¹¹ manual has detailed instructions regarding the required tools to compile packages from source for different platforms.

⁸<https://stat.ethz.ch/mailman/listinfo/r-sig-finance>

⁹<http://cran.r-project.org>

¹⁰<http://cran.r-project.org/web/views/Finance.html>

¹¹<http://cran.r-project.org/doc/manuals/R-admin.html>

GETTING SUPPORT

This book was compiled using the following R and package versions:

- R version: 3.1.2 (2014-10-31)
- Packages (loaded): boot 1.3-14, cluster 1.15.3, fAssets 3011.83, fBasics 3011.87, fPortfolio 3011.81, mvoutlier 2.0.5, pastecs 1.3-18, sgeostat 1.0-25, timeDate 3012.100, timeSeries 3012.99
- Packages (loaded via namespace): bitops 1.0-6, colorspace 1.2-4, DEoptimR 1.0-2, digest 0.6.8, ecodist 1.2.9, energy 1.6.2, fCopulae 3011.81, fMultivar 3011.78, GGally 0.5.0, ggplot2 1.0.0, grid 3.1.2, gtable 0.1.2, kernlab 0.9-20, lattice 0.20-29, MASS 7.3-37, mnormt 1.5-1, munsell 0.4.2, mvnrmtest 0.1-9, mvtnorm 1.0-2, numDeriv 2012.9-1, parallel 3.1.2, pcaPP 1.9-60, pls 2.4-3, plyr 1.8.1, proto 0.3-10, quadprog 1.5-5, Rcpp 0.11.4, RCurl 1.95-4.5, reshape 0.8.5, reshape2 1.4.1, Rglpk 0.6-0, rneos 0.2-8, robCompositions 1.9.0, robustbase 0.92-3, rrcov 1.3-8, Rsolnp 1.15, Rsymphony 0.1-18, scales 0.2.4, slam 0.1-32, sn 1.1-2, stats4 3.1.2, stringr 0.6.2, tools 3.1.2, truncnorm 1.0-7, XML 3.98-1.1

Generally we give our best to make sure that the code examples within this book are also compatible with newer versions of R or packages. But we cannot guarantee functionality for other setups as the one described above. The reason for this is that we don't have any control for changes within the R core and base functionality as well as changes within dependant packages that are not maintained by us. In our experience there is not much to worry about this. The code snippets usually remain functionality with maybe minor changes for a couple of years.

Note that especially for Mac OS X the situation is not very satisfying for operating systems newer than Snow Leopard. This due to the extensive changes made to the Xcode environment. Many packages are not available as OS X binaries and installing them from source seems rather tricky. As long as this situation doesn't change we can not give any guarantee for this book to work for Mac. One solution for Mac users is to install Windows or Linux as a virtual machine. Internally we successfully compiled all the necessary packages for newer OS X operating systems.

If you need help in setting up an environment for Mac, porting the code within this book to newer systems or implementing your own models you can get support from the Rmetrics association.¹²

¹²Terms and conditions may apply.

ACKNOWLEDGEMENTS

This book would not be possible without the R environment developed by the R Development Core Team.

We are also grateful to many people who have read and commented on draft material and on previous manuscripts of this eBook. Thanks also to those who contribute to the R-sig-finance mailing list, helping us to test our software.

We cannot name all who have helped us but we would like to thank Enrique Bengoechea, Dirk Eddelbuettel, Alexios Ghalanos, Francisco Gochez, Oliver Greshake, Martin Hanf, Elmar Heeb, Kurt Hornik, Stephan Joeri, Dominik Locher, David Lüthi, Dominik Lutz, Martin Mächler, Mahendra Mehta, Brian Peterson, Bernhard Pfaff, Jeff Ryan, David Scott, Stefan Theussl, Achim Zeileis, and all other people who made this book possible, the Institute for Theoretical Physics at ETH Zurich for their continuous support, and the participants and sponsors of the R/Rmetrics Meielisalp workshops and summer schools.

This book is the first in a series of Rmetrics eBooks. These eBooks will cover the whole spectrum of basic R and the Rmetrics packages; from managing chronological objects, to dealing with risk, to portfolio design. In this eBook we introduce those Rmetrics packages that deal with the whole spectrum of portfolio analysis, selection, and optimization.

Enjoy it!

Diethelm Würtz
Zurich, May 2009

This book was written six years ago. Since then many changes have been made in the base R environment. Most of them had impact on our eBook and have been continuously updated. Now with R 3.X we have done a complete revision of the book. This refreshed version of the book should take care of all updates until the beginning of 2015.

Diethelm Würtz
Tobias Setz
Zurich, January 2015

CONTENTS

DEDICATION	III
PREFACE	V
<i>About this Book</i>	v
<i>Computations</i>	v
<i>Audience Background</i>	vi
<i>Getting Help</i>	vi
<i>Getting Started</i>	vii
<i>Getting Support</i>	viii
<i>Acknowledgements</i>	ix
CONTENTS	XI
LIST OF FIGURES	XIX
LIST OF TABLES	XXIII
INTRODUCTION	1
I Managing Data Sets of Assets	3
INTRODUCTION	5
1 GENERIC FUNCTIONS TO MANIPULATE ASSETS	7
1.1 <i>timeDate and timeSeries Objects</i>	7
1.2 <i>Loading timeSeries Data Sets</i>	9
1.3 <i>Sorting and Reversing Assets</i>	10
1.4 <i>Alignment of Assets</i>	13
1.5 <i>Binding and Merging Assets</i>	14
1.6 <i>Subsetting Assets</i>	17
1.7 <i>Aggregating Assets</i>	20
1.8 <i>Rolling Assets</i>	22

2	FINANCIAL FUNCTIONS TO MANIPULATE ASSETS	25
2.1	<i>Price and Index Series</i>	25
2.2	<i>Returns and Cumulated Returns Series</i>	26
2.3	<i>Drawdowns Series</i>	27
2.4	<i>Durations Series</i>	28
2.5	<i>How to Add Your Own Functions</i>	29
3	BASIC STATISTICS OF FINANCIAL ASSETS	33
3.1	<i>Summary Statistics</i>	33
3.2	<i>Sample Mean and Covariance Estimates</i>	36
3.3	<i>Estimates for Higher Moments</i>	37
3.4	<i>Quantiles and Related Risk Measures</i>	38
3.5	<i>Computing Column Statistics</i>	39
3.6	<i>Computing Cumulated Column Statistics</i>	40
4	ROBUST MEAN AND COVARIANCE ESTIMATES	41
4.1	<i>Robust Covariance Estimators</i>	41
4.2	<i>Comparisons of Robust Covariances</i>	43
4.3	<i>Minimum Volume Ellipsoid Estimator</i>	43
4.4	<i>Minimum Covariance Determinant Estimator</i>	44
4.5	<i>Orthogonalized Gnanadesikan-Kettenring Estimator</i>	46
4.6	<i>Nearest-Neighbour Variance Estimator</i>	47
4.7	<i>Shrinkage Estimator</i>	48
4.8	<i>Bagging Estimator</i>	49
4.9	<i>How to Add a New Estimator to the Suite</i>	50
4.10	<i>How to Detect Outliers in a Set of Assets</i>	51
II	Exploratory Data Analysis of Assets	55
	INTRODUCTION	57
5	FINANCIAL TIME SERIES AND THEIR PROPERTIES	59
5.1	<i>Financial Time Series Plots</i>	59
5.2	<i>Box Plots</i>	69
5.3	<i>Histogram and Density Plots</i>	71
5.4	<i>Quantile-Quantile Plots</i>	73
6	CUSTOMIZATION OF PLOTS	77
6.1	<i>Plot Labels</i>	77
6.2	<i>More About Plot Function Arguments</i>	79
6.3	<i>Selecting Colours</i>	83
6.4	<i>Selecting Character Fonts</i>	90
6.5	<i>Selecting Plot Symbols</i>	92

7	MODELLING ASSET RETURNS	93
7.1	<i>Testing Asset returns for Normality</i>	93
7.2	<i>Fitting Asset returns</i>	95
7.3	<i>Simulating Asset Returns from a given Distribution</i>	96
8	SELECTING SIMILAR OR DISSIMILAR ASSETS	99
8.1	<i>Functions for Grouping Similar Assets</i>	99
8.2	<i>Grouping Asset Returns by Hierarchical Clustering</i>	100
8.3	<i>Grouping Asset Returns by k-means Clustering</i>	102
8.4	<i>Grouping Asset Returns through Eigenvalue Analysis</i>	105
8.5	<i>Grouping Asset Returns by Contributed Cluster Algorithms</i>	105
8.6	<i>Ordering Data Sets of Assets</i>	108
9	COMPARING MULTIVARIATE RETURN AND RISK STATISTICS	109
9.1	<i>Star and Segment Plots</i>	109
9.2	<i>Segment Plots of Basic Return Statistics</i>	110
9.3	<i>Segment Plots of Distribution Moments</i>	112
9.4	<i>Segment Plots of Box Plot Statistics</i>	113
9.5	<i>How to Position Stars and Segments in Star Plots</i>	114
10	PAIRWISE DEPENDENCIES OF ASSETS	115
10.1	<i>Simple Pairwise Scatter Plots of Assets</i>	115
10.2	<i>Pairwise Correlations Between Assets</i>	118
10.3	<i>Tests of Pairwise Correlations</i>	122
10.4	<i>Image Plot of Correlations</i>	122
10.5	<i>Bivariate Histogram Plots</i>	124
III	Portfolio Framework	129
	INTRODUCTION	131
11	S4 PORTFOLIO SPECIFICATION CLASS	133
11.1	<i>Class Representation</i>	133
11.2	<i>The Model Slot</i>	137
11.3	<i>The Portfolio Slot</i>	141
11.4	<i>The Optim Slot</i>	145
11.5	<i>The Message Slot</i>	147
11.6	<i>Consistency Checks on Specifications</i>	147
12	S4 PORTFOLIO DATA CLASS	149
12.1	<i>Class Representation</i>	149
12.2	<i>The Data Slot</i>	152
12.3	<i>The Statistics Slot</i>	153

13	S4 PORTFOLIO CONSTRAINTS CLASS	155
13.1	<i>Class Representation</i>	155
13.2	<i>Long-Only Constraint String</i>	158
13.3	<i>Unlimited Short Selling Constraint String</i>	159
13.4	<i>Box Constraint Strings</i>	160
13.5	<i>Group Constraint Strings</i>	161
13.6	<i>Covariance Risk Budget Constraint Strings</i>	162
13.7	<i>Non-Linear Weight Constraint Strings</i>	163
13.8	<i>Case study: How To Construct Complex Portfolio Constraints</i>	165
14	PORTFOLIO FUNCTIONS	169
14.1	<i>S4 Class Representation</i>	169
IV	Mean-Variance Portfolios	175
	INTRODUCTION	177
15	MARKOWITZ PORTFOLIO THEORY	179
15.1	<i>The Minimum Risk Mean-Variance Portfolio</i>	179
15.2	<i>The Feasible Set and the Efficient Frontier</i>	180
15.3	<i>The Minimum Variance Portfolio</i>	181
15.4	<i>The Capital Market Line and Tangency Portfolio</i>	182
15.5	<i>Box and Group Constrained Mean-Variance Portfolios</i>	182
15.6	<i>Maximum Return Mean-Variance Portfolios</i>	183
15.7	<i>Covariance Risk Budgets Constraints</i>	183
16	MEAN-VARIANCE PORTFOLIO SETTINGS	185
16.1	<i>Step 1: Portfolio Data</i>	185
16.2	<i>Step 2: Portfolio Specification</i>	185
16.3	<i>Step 3: Portfolio Constraints</i>	186
17	MINIMUM RISK MEAN-VARIANCE PORTFOLIOS	187
17.1	<i>How to Compute a Feasible Portfolio</i>	187
17.2	<i>How to Compute a Minimum Risk Efficient Portfolio</i>	189
17.3	<i>How to Compute the Global Minimum Variance Portfolio</i>	190
17.4	<i>How to Compute the Tangency Portfolio</i>	192
17.5	<i>How to Customize a Pie Plot</i>	194
18	MEAN-VARIANCE PORTFOLIO FRONTIERS	197
18.1	<i>Frontier Computation and Graphical Displays</i>	197
18.2	<i>The 'long-only' Portfolio Frontier</i>	201
18.3	<i>The Unlimited 'short' Portfolio Frontier</i>	202
18.4	<i>The Box-Constrained Portfolio Frontier</i>	205
18.5	<i>The Group-Constrained Portfolio Frontier</i>	209

18.6	<i>The Box/Group-Constrained Portfolio Frontier</i>	211
18.7	<i>Creating Different 'Reward/Risk Views' on the Efficient Frontier</i>	213
19	CASE STUDY: DOW JONES INDEX	217
20	ROBUST PORTFOLIOS AND COVARIANCE ESTIMATION	221
20.1	<i>Robust Mean and Covariance Estimators</i>	222
20.2	<i>The MCD Robustified Mean-Variance Portfolio</i>	222
20.3	<i>The MVE Robustified Mean-Variance Portfolio</i>	224
20.4	<i>The OGK Robustified Mean-Variance Portfolio</i>	229
20.5	<i>The Shrunked Mean-Variance Portfolio</i>	233
20.6	<i>How to Write Your Own Covariance Estimator</i>	234
V	Mean-CVaR Portfolios	239
	INTRODUCTION	241
21	MEAN-CVaR PORTFOLIO THEORY	243
21.1	<i>Solution of the Mean-CVaR Portfolio</i>	244
21.2	<i>Discretization</i>	244
21.3	<i>Linearization</i>	246
22	MEAN-CVaR PORTFOLIO SETTINGS	247
22.1	<i>Step 1: Portfolio Data</i>	247
22.2	<i>Step 2: Portfolio Specification</i>	247
22.3	<i>Step 3: Portfolio Constraints</i>	248
23	MEAN-CVaR PORTFOLIOS	249
23.1	<i>How to Compute a Feasible Mean-CVaR Portfolio</i>	249
23.2	<i>How to Compute the Mean-CVaR Portfolio with the Lowest Risk for a Given Return</i>	251
23.3	<i>How to Compute the Global Minimum Mean-CVaR Portfolio</i>	252
24	MEAN-CVaR PORTFOLIO FRONTIERS	257
24.1	<i>The Long-only Portfolio Frontier</i>	257
24.2	<i>The Unlimited 'Short' Portfolio Frontier</i>	259
24.3	<i>The Box-Constrained Portfolio Frontier</i>	263
24.4	<i>The Group-Constrained Portfolio Frontier</i>	265
24.5	<i>The Box/Group-Constrained Portfolio Frontier</i>	268
24.6	<i>Other Constraints</i>	269
24.7	<i>More About the Frontier Plot Tools</i>	270

VI	Portfolio Backtesting	275
	INTRODUCTION	277
25	S4 PORTFOLIO BACKTEST CLASS	279
25.1	<i>Class Representation</i>	279
25.2	<i>The Windows Slot</i>	281
25.3	<i>The Strategy Slot</i>	285
25.4	<i>The Smoother Slot</i>	287
25.5	<i>Rolling Analysis</i>	290
26	CASE STUDY: SPI SECTOR ROTATION	299
26.1	<i>SPI Portfolio Backtesting</i>	299
26.2	<i>SPI Portfolio Weights Smoothing</i>	300
26.3	<i>SPI Portfolio Backtest Plots</i>	301
26.4	<i>SPI Performance Review</i>	301
27	CASE STUDY: GCC INDEX ROTATION	305
27.1	<i>GCC Portfolio Weights Smoothing</i>	306
27.2	<i>GCC Performance Review</i>	306
27.3	<i>Alternative Strategy</i>	307
VII	Appendix	311
A	PACKAGES REQUIRED FOR THIS EBOOK	313
A.1	<i>Rmetrics Package: fPortfolio</i>	313
A.2	<i>Rmetrics Package: timeDate</i>	314
A.3	<i>Rmetrics Package: timeSeries</i>	315
A.4	<i>Rmetrics Package: fBasics</i>	316
A.5	<i>Rmetrics Package: fAssets</i>	317
A.6	<i>Contributed R Package: quadprog</i>	317
A.7	<i>Contributed Package: Rglpk</i>	318
A.8	<i>Recommended Packages from base R</i>	319
A.9	<i>Contributed R Packages</i>	319
A.10	<i>Rmetrics Package: fPortfolioBacktest</i>	319
B	DESCRIPTION OF DATA SETS	321
B.1	<i>Data Set: SWX</i>	321
B.2	<i>Data Set: LPP2005</i>	321
B.3	<i>Data Set: SPISECTOR</i>	322
B.4	<i>Data Set: SMALLCAP</i>	323
B.5	<i>Data Set: GCCINDEX</i>	323
C	R MANUALS ON CRAN	325

D Rmetrics Association	327
E Rmetrics Terms of Legal Use	331
Bibliography	333
Index	339
About the Authors	347

LIST OF FIGURES

2.1	Smoothed SPI times series with turnpoints	31
4.1	Comparison of sample and MVE robust covariances	45
4.2	Comparison of sample and mcd/MCD robust covariances	46
4.3	Comparison of sample and OGK robust covariances	48
5.1	Time series plots of the Swiss pension fund benchmark	60
5.2	Time series plots of the LPP benchmark indices	61
5.3	Scatter plot of the SPI versus the SBI	62
5.4	Plots of the SPI index and the returns	64
5.5	Plots of major Swiss indices and pension fund benchmark	65
5.6	Tailored graphs for the SPI	66
5.7	Plots with tailored axis labelling for the SPI	68
5.8	Box and box percentile plots of Swiss pension fund assets	70
5.9	Histogram and density plots of Swiss pension fund assets	72
5.10	Quantile-Quantile Plots	76
6.1	Customized time series plots	80
6.2	Colour table of R's base colours	84
6.3	Colour locator with R's 657 named colours	85
6.4	Selected colour palettes.	89
6.5	Selected colour palettes, continued	90
6.6	Character font tables	91
6.7	Table of plot symbols	92
8.1	Hierarchical clustering of Swiss pension fund assets	102
8.2	Hierarchical clustering of Swiss pension fund assets	103
8.3	Grouping assets by eigenvalue analysis	106
8.4	Grouping assets by partitioning around medoids	107
9.1	Segment plots from return statistics	111
9.2	Star plots from distributional moments	112
9.3	Star Plots from box plot statistics	113

10.1	Scatter plots of assets from the Swiss pension fund index . . .	117
10.2	Customized pairwise scatter plots of assets with histograms	118
10.3	Correlation patterns for pairwise scatter plots	119
10.4	Display of sorted pairwise correlations between assets	120
10.5	Display of sorted pairwise correlations between assets	121
10.6	Scatter plots in combination with correlation tests	123
10.7	Image plots of pairwise correlations	125
10.8	Scatter plots in combination with correlation tests	126
10.9	Bivariate histogram plots	127
15.1	Risk versus return view of the mean-variance portfolio	181
17.1	Weights plots for equal weighted and min var portfolios	191
17.2	Weights plots for global min risk and tangency portfolios . .	193
17.3	Bar plots of weights	195
18.1	Long-only MV portfolio frontier plot	202
18.2	MV long-only weights plots along the frontier	203
18.3	MV LongOnly weights plots along the frontier	204
18.4	Unlimited short MV portfolio frontier plot	206
18.5	MV short constrained weights plot along the frontier	207
18.6	Box constrained MV portfolio frontier plot	208
18.7	Group constrained MV portfolio frontier plot	210
18.8	MV linear constrained weights plot along the frontier	211
18.9	Box/group constrained MV portfolio frontier plot	213
18.10	MV box/group constrained weights plot along the frontier . .	214
18.11	Different Risk/Reward views on the efficient frontier	215
19.1	Weights along the DJI MV frontier	219
20.1	The MCD robustified MV portfolio frontier	225
20.2	Weights plot for MCD and COV MV portfolios	226
20.3	The MVE robustified MV portfolio frontier	228
20.4	Weights plot for MVE and MCD MV portfolios	229
20.5	The OGK robustified MV portfolio frontier	231
20.6	Weights plot for OGK and MCD MV portfolios	232
20.7	The shrunk MV portfolio frontier	234
20.8	Weights plot for shrunk MV portfolios.	235
20.9	The Student's t MV portfolio frontier	237
20.10	Weights plot for Student's t robustified MV portfolio	238
21.1	Portfolio frontier with Var, CVaR and ES risk measures	245
21.2	Portfolio frontier with VaR constraint.	245
23.1	Weights plots for an equal-weights CVaR portfolio	251

23.2	Weights plots for a minimum risk CVaR portfolio	253
23.3	Weights plots for a global minimum risk CVaR portfolio	254
23.4	Weights plots for a max/risk ratio CVaR portfolio	256
24.1	Mean-CVaR Portfolio - Long Only Frontier Plot	260
24.2	Mean-CVaR Portfolio - Long Only Weights Plot	261
24.3	Mean-CVaR Portfolio - Short Frontier Plot	262
24.4	Mean-CVaR Portfolio - Short Weights Plot	263
24.5	Box constrained Min-CVaR portfolio frontier plot.	265
24.6	Mean-CVaR Portfolio - Box Constrained Weights Plot	266
24.7	Group constrained Min-CVaR portfolio frontier plot.	267
24.8	Mean-CVaR Portfolio - Group Constrained Weights Plot . . .	268
24.9	Box/Group constrained CVaR portfolio frontier plot.	270
24.10	Mean-CVaR Portfolio - Box/Group Constrained Weights Plot	271
24.11	Return vs. risk plots	272
25.1	Portfolio backtesting for major Swiss indices	295
25.2	Rolling Analysis of the Swiss Performance Index	297
26.1	Portfolio backtesting for the SPI Sector Indexes	302
27.1	Portfolio backtesting for the GCC Index	309
27.2	Backtesting for the GCC Index with alternative strategy . . .	310

LIST OF TABLES

1.1	Example data sets	8
1.2	timeDate and timeSeries objects	8
1.3	Functions for downloading data from the Internet	10
1.4	Functions to sort time series objects	11
1.5	Functions to align time series objects.	13
1.6	Functions to concatenate time series objects	14
1.7	Functions to subset time series objects	17
1.8	Function for aggregating time series objects	20
1.9	Utility functions for time series objects	21
1.10	Functions to generate time periods	22
2.1	Functions for computing and exploring financial returns . .	25
2.2	Functions for price/index values and returns	26
2.3	Functions to compute drawdowns	27
2.4	Function for intervals	28
2.5	User supplied function	29
3.1	Functions for basic statistics of financial return	33
3.2	Moments and related functions	36
3.3	Portfolio risk functions	38
3.4	Column statistics functions	39
3.5	Cumulated column statistics functions	40
4.1	Robust covariance estimators	41
4.2	Function to detect outliers	51
5.1	Plot and related functions	59
5.2	Tailored plot functions and their arguments	62
5.3	Example plot function for lowess smoother	67
5.4	Box plot functions	69
5.5	Histogram and density plot functions	71
5.6	Quantile-quantile plot functions	74

6.1	Main arguments for plot functions	78
6.2	Title and text functions	79
6.3	Plot decoration functions	79
6.4	Selected arguments for plot functions	80
6.5	Type argument for plot functions	81
6.6	Font arguments for plot functions	82
6.7	cex arguments for plot functions	82
6.8	las argument for plot functions	82
6.9	lty argument for plot functions	83
6.10	Colour palette functions	86
6.11	Undocumented colour functions	90
6.12	Function to display characters	91
7.1	Functions to model multivariate assets	93
8.1	Functions to select similar assets	99
8.2	Functions used in hierarchical clustering	100
8.3	Functions used in hierarchical clustering	101
8.4	Functions used in k-means clustering	103
8.5	Functions used in k-means clustering	104
8.6	Further functions for clustering	105
9.1	Functions for star and segment plots	110
10.1	Functions for pairwise assets plots	115
10.2	Diagonal panel functions	120
10.3	Off-diagonal panel functions	121
10.4	Correlation functions	123
10.5	Bivariate histogram functions	124
11.1	Arguments of the function <code>portfolioSpec()</code>	134
11.2	Extractor functions for the <code>@model</code> slot	137
11.3	Constructor functions for the <code>@model</code> slot	137
11.4	The type argument for the <code>@model</code> slot	138
11.5	The optimize argument for the <code>@model</code> slot	138
11.6	Model slot of function <code>portfolioSpec()</code>	139
11.7	Template for a custom estimator function	140
11.8	The <code>tailRisk</code> argument	140
11.9	The <code>params</code> argument	141
11.10	Extractor functions for the <code>@portfolio</code> slot	141
11.11	Assignment functions for the <code>@portfolio</code> slot	141
11.12	Arguments of the <code>@portfolio</code> slot	142
11.13	The <code>riskFreeRate</code> argument of the <code>@portfolio</code> slot	144
11.14	The <code>nFrontierPoints</code> argument of the <code>@portfolio</code> slot	144
11.15	The status argument	145

11.16	Extractor functions for the @optim slot	145
11.17	Constructor functions for the @optim slot	145
11.18	Solver arguments in the optim slot	146
11.19	The trace argument for the @optim slot	146
11.20	The objective argument for the @optim slot	146
11.21	The options argument for the @optim slot	147
11.22	The control argument for the @optim slot	147
11.23	The argument list for the @message slot	147
12.1	The @data slot	152
12.2	The @statistics slot	153
13.1	Arguments of the portfolioConstraints() function	156
13.2	Long-only constraints string	158
13.3	Short constraint string	159
13.4	Box constraints	160
13.5	Group constraints	161
13.6	Covariance risk budget constraints	162
13.7	Non-linear weight constraints	164
14.1	Functions for computing and optimizing portfolios	170
14.2	fPORTFOLIO slots	170
17.1	Mean Variance Portfolio Functions	187
17.2	Pie charts of weights	194
18.1	Mean Variance Portfolio Frontiers	198
18.2	Frontier plot functions	200
20.1	Robust Covariance Estimators	222
24.1	Mean CVaR Portfolio Frontiers	257
25.1	Extractors for the @windows slot	281
25.2	Constructors for the @windows slot	281
25.3	List entry in the @windows slot an fPFOLIOBACKTEST object	283
25.4	Example of a rolling windows function	284
25.5	Extractors for the @strategy slot	285
25.6	Constructors for the @strategy slot	285
25.7	Extractors for the @smoother slot	287
25.8	Constructors for the @smoother slot	287
25.9	Parameters of the emaSmoother function	289
25.10	Custom smoother function	290
25.11	Elements returned by portfolioBacktesting()	292
25.12	Elements returned by portfolioSmoothing()	293

25.13	Rolling statistics functions	296
-------	--	-----

INTRODUCTION

Portfolio analysis, selection and optimization is the practise of dividing resources among different investments. These may be for example between stocks in an equity portfolio or between asset classes, such as stocks, bonds, mutual funds, real estate, cash equivalents, and private equity in a broader sense.

Portfolio design and optimization with the Rmetrics `fPortfolio` package relies on four pillars:

- *Definition* of the portfolio input, writing specifications, loading the data of the assets, and setting up the constraints.
- *Optimization* of the portfolio, including the computation of single portfolios such as feasible, efficient, tangency (max reward/risk) or minimum variance (global minimum risk) portfolios, and the evaluation of the entire efficient frontier.
- *Generation* of portfolio reports: printing, plotting and summarizing the results.
- *Analysis* of portfolio performance, including rolling analysis, back-testing and benchmarking.

The book is divided into six parts.

In [Part I](#) and [Part II](#) we present several aspects of the process of *exploratory data analysis* of the financial asset returns. This includes a brief summary which describes how to modify data sets of financial assets, a description of how to measure their statistical properties, and how to plot the time series and display related properties. We present several examples of how these plots and graphs can be customized by the user. Furthermore, we show how to model the multivariate distribution of the returns, how to group and compare the time series returns included in the data set of the assets, and how investigate and explore pairwise correlations and dependencies.

In [Part III](#) we describe the Rmetrics framework used for portfolio selection, optimization and backtesting. This includes the specification of the

three S4 portfolio classes dealing with the *specification*, the *data*, and the portfolio *constraints*.

[Part IV](#) is dedicated to Markowitz mean-variance portfolio optimization. We give a brief introduction to the theoretical aspects of the model. We show how to optimize efficient portfolios, including the global minimum variance and the tangency portfolio. Furthermore, we show how to explore the portfolio's feasible set and the whole efficient frontier. Special emphasis is also given in this part to the aspects of robust covariance estimation.

In [Part V](#) we go through the same program for mean-CVaR portfolios. Again, we consider individual efficient portfolios, the feasible set, and the efficient frontier. CVaR is an alternative risk measure to the covariance which is also known as mean excess loss, mean shortfall or tail value at risk, VaR. We discuss the portfolio optimization as a convex optimization problem proposed in [Rockafellar & Uryasev \(2000\)](#). We briefly describe the mathematical formulation of mean-CVaR optimization problems which can be formulated as an equivalent linear programming problem and can be solved using standard linear programming solvers.

[Part VI](#) is dedicated to portfolio backtesting. We introduce the portfolio backtest class and show how to define rolling windows, how to define portfolio strategies, and how to re-balance the portfolios over time using for example a smoothing approach for the portfolio weights. To show how backtesting for portfolio explicitly works we present two detailed case studies: A Swiss Sector Rotation Portfolio, and a Gulf Country Index portfolio.

PART I

MANAGING DATA SETS OF ASSETS

INTRODUCTION

In [chapter 1](#) we start with the manipulation of data sets of financial assets. These are usually financial return series represented by an S4 `timeSeries` object. We show how to sort a time series by ascending or descending time, how to provide a time-reversed version of a time series and how to re-sample a time series either with or without replacement. Further R functions allow us to bind two or more time series by columns or rows, and to merge two time series objects by common columns and/or row names. Another kind of manipulation which is often required is to align a time series to unique date and time stamps.

Financial time series analysis is concerned with data from financial markets, which mainly consist of prices, indexes and derived values, such as returns, cumulated returns, volatilities, drawdowns and durations. In [chapter 2](#) we list and describe functions provided by `Rmetrics` to compute such derived series.

In [chapter 3](#) we discuss basic statistics of time series. This includes summary and basic statistics reports, as well as the computation of measures such as mean, standard deviation, covariance, quantiles, or risk estimates.

CHAPTER 1

GENERIC FUNCTIONS TO MANIPULATE ASSETS

```
> library(fPortfolio)
```

Portfolio optimization with R/Rmetrics and the acquisition and selection of financial data as input go hand in hand. Rmetrics has a very intuitive way of working with financial time series. A financial time series consists of the data themselves and date/time stamps, which tell us when the data were recorded. In the generic case, when we consider a multivariate data set of financial assets, the data, usually prices or index values, are represented by a numeric matrix, where each column belongs to the data of an individual asset and each row belongs to a specific time/date stamp. This is most easily represented by a position vector of character strings. Combining the string vector of positions and the numeric matrix of data records, we can generate `timeSeries` objects.

In Rmetrics, date/time stamps are used to create `timeDate` objects, which are composed of a position vector of character strings, and the information of the name of the financial centre where the data were recorded. The financial centre is related to a time zone and appropriate daylight saving time rules, so that we can use the data worldwide without any loss of information.

1.1 TIME_DATE AND TIME_SERIES OBJECTS

The chronological objects implemented by Rmetrics and used in portfolio optimization are described in detail in the ebook *Chronological Objects in R/Rmetrics*. We highly recommend consulting this ebook if you have any questions concerning creating, modifying, and qualifying financial data sets.

Several financial datasets, which are used throughout this book, are provided with the `fPortfolio` package. They are listed in [Listing 1.1](#). Datasets are available as price/index series and as financial (log)-returns. They are stored as S4 `timeSeries` objects and do not need to be explicitly loaded. If you want to use data sets in the form of CSV files, you can load them as data frames using the `data()` function, and then convert them into S4 `timeSeries` objects using the function `as.timeSeries()`¹. [Listing 1.2](#) gives a brief summary of time series functions in `Rmetrics`, with a short description of their function.

LISTING 1.1: RMETRICS EXAMPLE DATA SETS USED IN THIS EBOOK. THE DATA SETS ARE DESCRIBED IN DETAIL IN [APPENDIX B](#).

Data Sets:

SWX	Daily Swiss equities, bonds, and reits series
LPP2005	Daily Pictet Swiss pension fund benchmarks
SPISECTOR	Swiss Performance sector indexes
GCCINDEX	Gulf Cooperation Council equity indexes
SMALLCAP	Monthly selected US small capitalized equities
MSFT	Daily Microsoft open, high, low, close, volume

LISTING 1.2: RMETRICS BASIC FUNCTIONS TO WORK WITH 'TIMEDate' AND 'timeSeries' OBJECTS. FOR FURTHER INFORMATION WE REFER TO THE HELP PAGES AND TO THE RMETRICS EBOOK 'CHRONOLOGICAL OBJECTS IN R/RMETRICS'

Function:

<code>timeDate</code>	creates <code>timeDate</code> objects from scratch
<code>timeSeq, seq</code>	creates regularly spaced <code>timeDate</code> objects
<code>timeCalendar</code>	creates <code>timeDate</code> objects from calendar atoms
<code>as.timeDate</code>	coerces and transforms <code>timeDate</code> objects
<code>timeSeries</code>	creates a <code>timeSeries</code> object from scratch
<code>readSeries</code>	reads a <code>timeSeries</code> from a spreadsheet file.
<code>as.timeSeries</code>	coerces and transforms <code>timeSeries</code> objects
<code>print, plot</code>	generic <code>timeSeries</code> functions
<code>+, -, *, ...</code>	math operations on <code>timeSeries</code> objects
<code>>, < == ...</code>	logical operations on <code>timeSeries</code> objects
<code>diff, log, ...</code>	function operations on <code>timeSeries</code> objects

The functions in [Listing 1.2](#) can be used to create time series objects from scratch, to convert to and from different representations, to read the time series data from files, or to download data from the Internet. We assume that the reader is familiar with the basics of the `timeDate` and `timeSeries` classes in `Rmetrics`.

Often data sets of assets are not in the form required for portfolio design, analysis and optimization. If this is the case, we have to compose and mod-

¹You can also work with other time series objects, such as `ts` or `zoo` objects. Note that if `zoo` is required, you must load the package before the `Rmetrics` packages are loaded. In this case you have to coerce these objects into a '`timeSeries`' object using functions `as.timeSeries.foo()`, where `foo` is a placeholder for the alternative time series class.

ify the data sets. In the following we briefly present the most important functions for managing timeSeries objects.

1.2 LOADING TIMESERIES DATA SETS

How to load a demo file

The Rmetrics software environment comes with selected demo data sets, which can be used to execute and test examples. Demo data sets are provided as S4 timeSeries objects. Below, we show the returns from the daily SWX market indices:

```
> class(SWX.RET)
[1] "timeSeries"
attr(,"package")
[1] "timeSeries"

> colnames(SWX.RET)
[1] "SBI" "SPI" "SII" "LP25" "LP40" "LP60"

> head(SWX.RET[, 1:3])

GMT
      SBI      SPI      SII
2000-01-04 -0.00208812 -0.0343901 1.3674e-05
2000-01-05 -0.00010452 -0.0104083 -4.9553e-03
2000-01-06 -0.00135976  0.0121191  3.8129e-03
2000-01-07  0.00041859  0.0224617 -6.1621e-04
2000-01-10  0.00000000  0.0021077  2.3806e-03
2000-01-11 -0.00104679 -0.0027737 -2.9385e-04
```

In the third line, we have restricted the output to the first 6 lines of the Swiss Bond Index, the Swiss Performance Index and the Swiss Immofunds Index.

How to read data from CSV text files

timeSeries files can also be written to and read from CSV files; in the example given below, we first create a small data set using just the first 6 lines of the SBI, SPI and SII. Then, we can use the `write.csv()` function to write the data set to a CSV file². The file 'myData.csv' will be created in the current working directory.

```
> # create small data set
> data <- head(SWX.RET[, 1:3])
> # write data to a CSV file in the current directory
> write.csv(data, file = "myData.csv")
```

²For help on this function, see `?write.csv`

We can now read the data from our CSV file, using the function `readSeries()`. Note that we have to specify the separator, `sep = ","` because the default separator is `sep = ";"`.

```
> # write CSV file in current directory, specifying the separators
> data2 <- readSeries(file = "myData.csv", header = TRUE, sep = ",")
```

The arguments of the `readSeries()` function are:

```
> args(readSeries)

function (file, header = TRUE, sep = ";", zone = "", FinCenter = "",
         format, ...)
NULL
```

For details we refer to the ebook *Chronological objects with R/Rmetrics* and the `timeSeries` help files.

How to download data from the Internet

The Rmetrics `fImport` (Würtz, 2009c) provides several functions to download time series data from the Internet, for example from

LISTING 1.3: FUNCTIONS FOR DOWNLOADING DATA FROM THE INTERNET

Download Functions:	
<code>fredSeries</code>	imports market data from the US Federal Reserve
<code>oandaSeries</code>	imports FX market data from OANDA
<code>yahooSeries</code>	imports market data from Yahoo Finance

These functions are able to download CSV files and HTML files and then format the data and make the records available as an S4 `timeSeries` object.

For further information, please consult the user and reference guide of the package `fImport` (Würtz, 2009c).

1.3 SORTING AND REVERSING ASSETS

In this chapter we use for our examples the daily data sets `SWX` and `SWX.RET`. The `SWX` data set contains six financial time series. The first three are Swiss indexes from the Swiss Exchange in Zurich, the *Swiss Performance Index*, `SPI`, the *Swiss Bond Index*, `SBI`, and the *Swiss ImmoFund Index* (reits), `SII`. The remaining three time series, named `LP25`, `LP40`, `LP60`, are *Swiss Pension Fund Benchmarks* provided by Pictet, a Swiss private bank in Geneva. The data set starts on January 3rd, 2000, and ends on May 5th, 2007. The data set contains 1917 time series records. The second data set, `SWX.RET`, contains daily log-returns derived from the `SWX` data set.

```
> head(SWX)
```



```

GMT
      SBI    SPI    SII  LP25  LP40  LP60
2000-01-03 95.88 5022.9 146.26 99.81 99.71 99.55
2000-01-04 95.68 4853.1 146.27 98.62 97.93 96.98
2000-01-05 95.67 4802.8 145.54 98.26 97.36 96.11
2000-01-06 95.54 4861.4 146.10 98.13 97.20 95.88
2000-01-07 95.58 4971.8 146.01 98.89 98.34 97.53
2000-01-10 95.58 4982.3 146.36 99.19 98.79 98.21

> end(SWX)

GMT
[1] [2007-05-08]

> class(SWX)

[1] "timeSeries"
attr(,"package")
[1] "timeSeries"

```

Loading the example data set `SWX` returns an object of class `timeSeries` as required for portfolio optimization.

Sometimes the records in a data set of assets are not ordered in time, or they are in reverse order. In this case the time stamps can be rearranged so that the series of assets becomes ordered in the desired way. `Rmetrics` has generic functions to sort, `sort()`, and reverse, `rev()`, the time stamps of time series so that they appear in ascending or descending order. The function `sample()` samples a series in random order.

LISTING 1.4: FUNCTIONS FOR SORTING, REVERSING AND SAMPLING DATA SETS OF ASSETS

Functions:

<code>sort</code>	sorts a 'timeSeries' in ascending or descending order
<code>rev</code>	provides a time-reversed version of a 'timeSeries'
<code>sample</code>	generates a sample either with or without replacement

Arguments:

<code>x</code>	an object of class 'timeSeries'
----------------	---------------------------------

How to sample a time series randomly

The generic function `sample()` takes a random sample either with or without replacement.

In this example, we randomly take ten rows (without replacement) from the `SWX` data set:

```

> SAMPLE <- sample(SWX[1:10, ])
> SAMPLE

GMT
      SBI    SPI    SII  LP25  LP40  LP60
2000-01-10 95.58 4982.3 146.36 99.19 98.79 98.21

```

```

2000-01-14 95.65 5042.2 146.94 99.79 99.68 99.52
2000-01-05 95.67 4802.8 145.54 98.26 97.36 96.11
2000-01-13 95.51 4985.2 147.09 99.20 98.81 98.24
2000-01-04 95.68 4853.1 146.27 98.62 97.93 96.98
2000-01-11 95.48 4968.5 146.31 98.95 98.48 97.80
2000-01-06 95.54 4861.4 146.10 98.13 97.20 95.88
2000-01-12 95.47 4977.8 146.28 98.91 98.42 97.71
2000-01-03 95.88 5022.9 146.26 99.81 99.71 99.55
2000-01-07 95.58 4971.8 146.01 98.89 98.34 97.53

```

Notice that the records of the sampled time series are no longer ordered in time, and thus follow each other in a completely irregular fashion.

How to sort a series in ascending order

The generic function `sort()` sorts the records of a time series in ascending or descending order. This is shown in the following example:

```

> sort(SAMPLE)
GMT
      SBI    SPI    SII  LP25  LP40  LP60
2000-01-03 95.88 5022.9 146.26 99.81 99.71 99.55
2000-01-04 95.68 4853.1 146.27 98.62 97.93 96.98
2000-01-05 95.67 4802.8 145.54 98.26 97.36 96.11
2000-01-06 95.54 4861.4 146.10 98.13 97.20 95.88
2000-01-07 95.58 4971.8 146.01 98.89 98.34 97.53
2000-01-10 95.58 4982.3 146.36 99.19 98.79 98.21
2000-01-11 95.48 4968.5 146.31 98.95 98.48 97.80
2000-01-12 95.47 4977.8 146.28 98.91 98.42 97.71
2000-01-13 95.51 4985.2 147.09 99.20 98.81 98.24
2000-01-14 95.65 5042.2 146.94 99.79 99.68 99.52

```

How to reverse a series in time

A sorted `timeSeries` object is given either in an ascending or descending order. The time ordering of the records of a data set can be reversed using the generic function `rev()`. Alternatively, we can also use the function `sort(x, decreasing=FALSE)`, setting the argument `decreasing` either to `TRUE` or `FALSE`.

```

> rev(sort(SAMPLE))

```

and

```

> sort(SAMPLE, decreasing = TRUE)
GMT
      SBI    SPI    SII  LP25  LP40  LP60
2000-01-14 95.65 5042.2 146.94 99.79 99.68 99.52
2000-01-13 95.51 4985.2 147.09 99.20 98.81 98.24
2000-01-12 95.47 4977.8 146.28 98.91 98.42 97.71
2000-01-11 95.48 4968.5 146.31 98.95 98.48 97.80

```

```

2000-01-10 95.58 4982.3 146.36 99.19 98.79 98.21
2000-01-07 95.58 4971.8 146.01 98.89 98.34 97.53
2000-01-06 95.54 4861.4 146.10 98.13 97.20 95.88
2000-01-05 95.67 4802.8 145.54 98.26 97.36 96.11
2000-01-04 95.68 4853.1 146.27 98.62 97.93 96.98
2000-01-03 95.88 5022.9 146.26 99.81 99.71 99.55

```

produce the same output.

1.4 ALIGNMENT OF ASSETS

The alignment of `timeSeries` objects is an important aspect in managing assets. Due to holidays, we must expect missing data records for daily data sets. For example, around Easter, data records for Good Friday may be missing in most countries of the world, and it is likely that the markets are also closed on Easter Monday. Even so, we still want to align the series to a regular weekly calendar series. Missing records can then be coded in several ways; a straightforward way is to use the price of the previous day for the subsequent day. The function `align()` aligns the asset series on calendar dates by default, i.e. on every day of the week, or, more naturally, on the weekdays from Monday to Friday.

Let us align the `SWX` series of indices to daily dates, including holidays, and replace the missing values with the indices from the previous days:

```

> nrow(SWX)
[1] 1917

> ALIGNED <- align(x = SWX, by = "1d", method = "before", include.weekends = FALSE)
> nrow(ALIGNED)
[1] 1917

```

The returned number of rows shows that the original series does not have any missing data records due to holidays. The alignment function can be used not only to align daily series, but also to align a series to other time horizons. For example, we can align daily data sets by weekly time horizons starting on any arbitrary day of the week using the `offset` argument of the function.

LISTING 1.5: FUNCTION TO ALIGN A TIME SERIES RECORDS TO TIME AND CALENDAR ATOMS.

Function:

`align` aligns a 'timeSeries' object to calendar objects.

Arguments:

`x` an object of class 'timeSeries'
`by` a character string formed from an integer length and a period identifier. Valid values are "w", "d", "h",

	"m", "s", for weeks, days, hours, minutes and seconds
offset	For example, a bi-weekly period is expressed as "2w"
	a character string formed from an integer length and
	a period identifier in the same way as for 'by'
method	a character string, defining the alignment. Substitutes
	a missing record with the value of the previous
	("before") record, of the following ("after") record,
	interpolates ("interp") or fills with NAs ("NA")
include.weekends	should the weekend days (Saturdays and Sundays) be
	included?

1.5 BINDING AND MERGING ASSETS

In many cases we have to compose the desired assets from several univariate and/or multivariate time series. Then we have to bind different time series together. The functions available in Rmetrics are shown in [Listing 1.6](#), in order of increasing complexity:

LISTING 1.6: FUNCTIONS TO CONCATENATE DATA SETS OF ASSETS

Function:	
c	concatenates a 'timeSeries' object.
cbind	combines a 'timeSeries' by columns.
rbind	combines a 'timeSeries' by rows.
merge	merges two 'timeSeries' by common columns and/or rows.
Arguments:	
x, y	objects of class 'timeSeries'.

Before we start to interpret the results of binding and merging several time series objects, let us consider the following three time series examples to better understand how binding and merging works.

```
> set.seed(1953)
> charvec <- format(timeCalendar(2008, sample(12, 6)))
> data <- matrix(round(rnorm(6), 3))
> t1 <- sort(timeSeries(data, charvec, units = "A"))
> t1
```

```
GMT
      A
2008-02-01 0.236
2008-05-01 1.484
2008-06-01 0.231
2008-07-01 0.187
2008-10-01 -0.005
2008-11-01 1.099
```

```

> charvec <- format(timeCalendar(2008, sample(12, 9)))
> data <- matrix(round(rnorm(9), 3), 3)
> t2 <- sort(timeSeries(data, charvec, units = "B"))
> t2

GMT
      B
2008-01-01 -1.097
2008-03-01 -0.890
2008-04-01 -1.472
2008-05-01 -1.009
2008-06-01  0.983
2008-07-01 -0.068
2008-10-01 -2.300
2008-11-01  1.023
2008-12-01  1.177

> charvec <- format(timeCalendar(2008, sample(12, 5)))
> data <- matrix(round(rnorm(10), 3), ncol = 2)
> t3 <- sort(timeSeries(data, charvec, units = c("A", "C")))
> t3

GMT
      A      C
2008-02-01  0.620 -0.109
2008-03-01 -1.490  0.796
2008-04-01  0.210 -0.649
2008-05-01  0.654  0.231
2008-06-01 -1.603  0.318

```

The first series `t1` and second series `t2` are univariate series with 6 and 9 random records and column names "A" and "B", respectively. The third `t3` series is a bivariate series with 5 records per column and column names "A" and "C". Notice that the first column "A" of the third time series `t3` describes the same time series "A" as the first series "`t1`".

How to bind time series column- and row-wise

The functions `cbind()` and `rbind()` allow us to bind time series objects together either by column or by row. Let us bind series `t1` and series `t2` by columns

```

> cbind(t1, t2)

GMT
      A      B
2008-01-01  NA -1.097
2008-02-01  0.236   NA
2008-03-01  NA -0.890
2008-04-01  NA -1.472
2008-05-01  1.484 -1.009
2008-06-01  0.231  0.983
2008-07-01  0.187 -0.068
2008-10-01 -0.005 -2.300
2008-11-01  1.099  1.023
2008-12-01  NA  1.177

```

We obtain a bivariate time series with column names "A" and "B", where the gaps were filled with NAs. Binding series t1 and t3 together column by column

```
> cbind(t1, t3)
GMT
      A.1  A.2  C
2008-02-01 0.236 0.620 -0.109
2008-03-01  NA -1.490  0.796
2008-04-01  NA  0.210 -0.649
2008-05-01 1.484  0.654  0.231
2008-06-01 0.231 -1.603  0.318
2008-07-01 0.187   NA   NA
2008-10-01 -0.005   NA   NA
2008-11-01 1.099   NA   NA
```

we obtain a new time series with three columns and the names of the two series with identical column names "A", but they receive the suffixes ".1" and ".2" to distinguish them.

The function rbind() behaves similarly, but the number of columns must be the same in all time series to be bound by rows

```
> rbind(t1, t2)
GMT
      A_B
2008-02-01 0.236
2008-05-01 1.484
2008-06-01 0.231
2008-07-01 0.187
2008-10-01 -0.005
2008-11-01 1.099
2008-01-01 -1.097
2008-03-01 -0.890
2008-04-01 -1.472
2008-05-01 -1.009
2008-06-01  0.983
2008-07-01 -0.068
2008-10-01 -2.300
2008-11-01  1.023
2008-12-01  1.177
```

The column name is now "A_B" to illustrate that series named "A" and "B" were bound together. Note that binding the univariate series t1 and the bivariate series t3 would result in an error because they do not have the same number of columns.

How to merge time series column-wise and row-wise

Merging two data sets of assets is the most general case and will take the names of the individual columns. merge() combines the two series, which can be either univariate or multivariate, by column and by row, and,

additionally, intersects columns with identical column names. This is the most important point. To show this, let us merge the time series `t1` and `t2`, and then merge them with `t3`

```
> tM <- merge(merge(t1, t2), t3)
> tM
```

GMT	A	B	C
2008-01-01	NA	-1.097	NA
2008-02-01	0.236	NA	NA
2008-02-01	0.620	NA	-0.109
2008-03-01	-1.490	NA	0.796
2008-03-01	NA	-0.890	NA
2008-04-01	0.210	NA	-0.649
2008-04-01	NA	-1.472	NA
2008-05-01	0.654	NA	0.231
2008-05-01	1.484	-1.009	NA
2008-06-01	-1.603	NA	0.318
2008-06-01	0.231	0.983	NA
2008-07-01	0.187	-0.068	NA
2008-10-01	-0.005	-2.300	NA
2008-11-01	1.099	1.023	NA
2008-12-01	NA	1.177	NA

This gives us a 3-column time series with names "A", "B", and "C". Note that the records from time series `t1` and from the first column of time series `t3`, both named "A", were merged into the same first column of the new time series.

1.6 SUBSETTING ASSETS

Subsetting a data set of assets and replacing parts of a data set by other records is a very important issue in the management of financial time series.

There are several functions that are useful in this context. These include the "[" operator, which extracts or replaces subsets, the `window()` function, which cuts out a piece from a data set between two 'timeDate' objects, start and end, and the functions `start()` and `end()` themselves, which return the first and last record of a data set.

Subsetting by using the "[" operator can be done by simple counts, by date/time stamps, by instrument (column) names, or even by logical predicates, e.g. extracting all records before or after a given date.

LISTING 1.7: FUNCTIONS FOR SUBSETTING DATA SETS OF ASSETS

Function:	
[extracts or replaces subsets by indexes, column names, date/time stamps, logical predicates, etc
subset	returns subsets that meet specified conditions

window	extracts a piece between two 'timeDate' objects
start	extracts the first record
end	extracts the last record

Arguments:

x	an object of class 'timeSeries'
---	---------------------------------

How to subset by counts

Subsetting by counts allows us to extract desired records from the rows, and desired instruments from the columns of the data series matrix. The first example demonstrates how to subset a univariate or multivariate timeSeries by row, here the second to the fifth rows

```
> SWX[2:5, ]
GMT
      SBI    SPI    SII LP25 LP40 LP60
2000-01-04 95.68 4853.1 146.27 98.62 97.93 96.98
2000-01-05 95.67 4802.8 145.54 98.26 97.36 96.11
2000-01-06 95.54 4861.4 146.10 98.13 97.20 95.88
2000-01-07 95.58 4971.8 146.01 98.89 98.34 97.53

> SWX[2:5, 2]
GMT
      SPI
2000-01-04 4853.1
2000-01-05 4802.8
2000-01-06 4861.4
2000-01-07 4971.8
```

Note that in the first example we have to explicitly write `SWX[2:5,]` instead of `SWX[2:5]` since the data part is a two dimensional rectangular object.

How to find the first and last records

To extract the first and the last record of a timeSeries object we can use the functions `start()` and `end()`. The function `start()` sorts the assets in increasing time order and returns the first element of the time positions

```
> SWX[start(SWX), ]
GMT
      SBI    SPI    SII LP25 LP40 LP60
2000-01-03 95.88 5022.9 146.26 99.81 99.71 99.55

> SWX[start(sample(SWX)), ]
GMT
      SBI    SPI    SII LP25 LP40 LP60
2000-01-03 95.88 5022.9 146.26 99.81 99.71 99.55
```


`end()` behaves in the same way, but in the opposite order.

How to subset by column names

Instead of using counts, e.g. the 4th column, we can reference and extract columns by column names, which are usually the names of the financial instruments

```
> tail(SWX[, "SPI"])

GMT
      SPI
2007-05-01 7620.1
2007-05-02 7634.1
2007-05-03 7594.9
2007-05-04 7644.8
2007-05-07 7647.6
2007-05-08 7587.9
```

How to subset by date/time stamps

Subsetting by date vectors allows you to extract desired records from the rows for a specified date or dates. We first show an example for the univariate case where we extract a specific date:

```
> # Extract a specific date:
> SWX["2007-04-24", ]

GMT
      SBI    SPI    SII  LP25  LP40  LP60
2007-04-24 97.05 7546.5 214.91 129.65 128.1 124.34

> # Subset all records from the first and second quarter:
> round(window(SWX, start = "2006-01-15", end = "2006-01-21"), 1)

GMT
      SBI    SPI    SII  LP25  LP40  LP60
2006-01-16 101.3 5939.4 196.7 123.1 118.2 110.5
2006-01-17 101.4 5903.5 196.7 123.0 117.9 110.2
2006-01-18 101.5 5861.7 197.2 122.8 117.5 109.5
2006-01-19 101.3 5890.5 198.9 122.9 117.8 110.0
2006-01-20 101.2 5839.4 197.6 122.5 117.3 109.4
```

Here we have rounded the results to one digit in order to shorten the output using the generic function `round()`. Note that there are additional functions to round numbers in R. These include: `ceiling()`, `floor()`, `truncate()`, and `signif()`. For details we refer to the help pages.

1.7 AGGREGATING ASSETS

In finance we often want to aggregate time series, that is we want to go from a fine-grained resolution to a coarse-grained resolution. For example, we have collected data on a daily basis and now we want to display them on a weekly, monthly, or quarterly basis.

We can use the generic function `aggregate()` from R's base package `stats` to do this. The function splits the data set into individual subsets, and then computes summary statistics for each subset. Finally, the result is returned in a convenient form. `Rmetrics` provides a method for aggregating `timeSeries` objects. The function requires three input arguments, the time series itself, a sequence of date/time stamps defining the grouping, and the function that is to be applied.

LISTING 1.8: FUNCTION FOR AGGREGATING A DATA SET OF ASSETS

Function:

`aggregate` aggregates a 'timeSeries' object.

Arguments:

`x` is a uni- or multivariate 'timeSeries' object
`by` is a 'timeDate' sequence of grouping dates
`FUN` a scalar function to compute the summary statistics
to be applied to all data subsets

To be more specific, let us define an artificial monthly `timeSeries` that we want to aggregate on a quarterly base

```
> charvec <- timeCalendar()
> data <- matrix(round(runif(24, 0, 10)), 12)
> tS <- timeSeries(data, charvec)
> tS
```

GMT	TS.1	TS.2
2015-01-01	1	6
2015-02-01	4	4
2015-03-01	2	10
2015-04-01	10	9
2015-05-01	6	6
2015-06-01	1	6
2015-07-01	1	2
2015-08-01	10	0
2015-09-01	7	4
2015-10-01	0	2
2015-11-01	9	8
2015-12-01	0	10

Next, we create the quarterly breakpoints from the `charvec` vector searching for the last day in a quarter for each date. To suppress double dates we make the breakpoints unique

```
> by <- unique(timeLastDayInQuarter(charvec))
> by
GMT
[1] [2015-03-31] [2015-06-30] [2015-09-30] [2015-12-31]
```

and finally we create the quarterly series with the aggregated monthly sums and new units passed in by the dots argument.

```
> aggregate(tS, by, FUN = sum, units = c("TSQ.1", "TSQ.2"))
GMT
      TSQ.1 TSQ.2
2015-03-31    7   20
2015-06-30   17   21
2015-09-30   18    6
2015-12-31    9   20
```

Rmetrics also has many utility functions to manage special dates. These are shown in [Listing 1.9](#).

LISTING 1.9: UTILITY FUNCTIONS FOR MANAGING SPECIAL DATES

Function:	
timeLastDayInMonth	last day in a given month/year
timeFirstDayInMonth	first day in a given month/ year
timeLastDayInQuarter	last day in a given quarter/year
timeFirstDayInQuarter	first day in a given quarter/year
timeNdayOnOrAfter	date month that is a n-day ON OR AFTER
timeNdayOnOrBefore	date in month that is a n-day ON OR BEFORE
timeNthNdayInMonth	n-th occurrence of a n-day in year/month
timeLastNdayInMonth	last n-day in year/month

to determine date breakpoints, e.g. when the accounting is quarterly on the first Monday or working day of the following quarter. More examples are provided in the Rmetrics ebook 'Chronological Objects with R/Rmetrics'.

Now let us demonstrate a real-world example. We will aggregate the daily returns of the SPI index on monthly periods:

```
> tS <- 100 * LPP2005.RET[, "SPI"]
> by <- timeLastDayInMonth(time(tS))
> aggregate(tS, by, sum)
GMT
      SPI
2005-11-30  4.81406
2005-12-31  2.48124
2006-01-31  3.19682
2006-02-28  1.39536
2006-03-31  2.48262
2006-04-30  1.41992
2006-05-31 -5.37181
2006-06-30  0.52306
```

```

2006-07-31  3.61945
2006-08-31  2.91602
2006-09-30  3.24490
2006-10-31  2.00248
2006-11-30 -0.54849
2006-12-31  3.90591
2007-01-31  4.41255
2007-02-28 -3.76150
2007-03-31  2.95389
2007-04-30  2.04765

```

1.8 ROLLING ASSETS

Let us write a simple function named `rollapply()` that can compute rolling statistics using the function `applySeries()`, which can be found in the Rmetrics package `timeSeries`. The periods may be overlapping or not, and we even allow gaps between the periods.

```

> rollapply <- function(x, by, FUN, ...)
{
  ans <- applySeries(x, from = by$from, to = by$to, by = NULL,
    FUN = FUN, format = x@format,
    zone = finCenter(x), FinCenter = finCenter(x),
    title = x$title, documentation = x@documentation, ...)
  attr(ans, "by") <- data.frame(from = format(by$from), to = format(by$to) )
  ans
}

```

Here we also want to focus on the `periods` function from the `timeDate` package. This allows us to compute periods from time spans. The following example demonstrates how to compute the returns on a multivariate data set of assets subset in annual windows and shifted monthly:

LISTING 1.10: FUNCTION TO GENERATE SHIFTED TIME PERIODS (WINDOWS)

Function:	
<code>periods</code>	constructs equidistantly sized and shifted windows
Arguments:	
<code>period</code>	size (length) of the periods
<code>by</code>	shift (interval) of the periods, "m" monthly, "w" weekly, "d" daily, "H" by hours, "M" by minutes, "S" by seconds.

```

> DATA <- 100 * SWX.RET[, c(1:2, 4:5)]
> by <- periods(time(DATA), "12m", "6m")
> SWX.ROLL <- rollapply(DATA, by, FUN = "colSums")
> SWX.ROLL

```

```

GMT
      SBI      SPI      LP25      LP40
2000-12-31 -0.64874  11.2533  1.9643  0.80907
2001-06-30  3.33258  -5.5704  3.0018  0.50810
2001-12-31  0.48173 -24.8813 -1.5145 -4.68422
2002-06-30  1.41993 -18.8377 -3.7161 -8.66230
2002-12-31  6.37416 -30.0450 -2.1779 -8.77699
2003-06-30  4.48767 -18.7557  3.4043  0.45410
2003-12-31 -1.78011  19.9374  7.5090 10.13347
2004-06-30 -3.05718  19.2258  4.3348  6.71645
2004-12-31  0.98298  6.6636  4.7731  5.13190
2005-06-30  4.45444  13.1578  9.4868 10.38565
2005-12-31  0.16783  30.4600  9.9139 13.55689
2006-06-30 -5.98790  22.5691  2.1522  5.01789
2006-12-31 -3.01400  18.7862  3.9933  6.15426

```

```

> attr(SWX.ROLL, "by")
      from      to
1 2000-01-01 2000-12-31
2 2000-07-01 2001-06-30
3 2001-01-01 2001-12-31
4 2001-07-01 2002-06-30
5 2002-01-01 2002-12-31
6 2002-07-01 2003-06-30
7 2003-01-01 2003-12-31
8 2003-07-01 2004-06-30
9 2004-01-01 2004-12-31
10 2004-07-01 2005-06-30
11 2005-01-01 2005-12-31
12 2005-07-01 2006-06-30
13 2006-01-01 2006-12-31

```

The values "12m" and "1m" in the function `periods()` are called time spans.

Here are two more examples with regular periodic periods and by-shifts.

```

> by <- periods(time(SWX), period = "52w", by = "4w")
> by <- periods(time(SWX), period = "360d", by = "30d")

```

The first example rolls on a period of 52 weeks shifted by 4 weeks, the second rolls every 30 days on 360 calendar days.

Periods created by "12m" yield an annual rolling period for a given fixed shift, "6m" yields a semi-annual rolling period, "3m" a quarterly rolling period, "2m" a bi-monthly rolling period, and so on. With these unit identifiers we can create calendar-based rolling as well as regular periodical rolling, or even irregular rolling periods and by-shifts. The latter are useful for periods triggered by the volatility, for example, and a shift given by automated trading signals or by human decision-makers.

CHAPTER 2

FINANCIAL FUNCTIONS TO MANIPULATE ASSETS

```
> library(fPortfolio)
```

Financial time series analysis investigates and models data sets from financial markets. These are usually prices, indices and derived values such as returns, cumulated returns, volatilities, drawdowns and durations, amongst others.

In this chapter we describe functions provided by Rmetrics to compute derived financial time series and show several examples how to use them. Moreover we show examples how a user can add his own functions to the Rmetrics framework.

2.1 PRICE AND INDEX SERIES

Price and index series can be downloaded either from free sources, such as [Yahoo Finance](http://finance.yahoo.com)¹, the [Swiss Exchange](http://www.six-swiss-exchange.com)² or the [Federal Reserve Bank in St. Louis](http://www.stlouisfed.org)³, or can be obtained from commercial providers such as Bloomberg or IBrokers. Rmetrics provides interfaces for downloading free data from the Internet (Würtz, 2009c). R also has packages for downloading data from commercial sources; these include for example the RBloomberg (Sams, 2009) and IBrokers (Ryan, 2008) packages.

LISTING 2.1: FUNCTIONS FOR COMPUTING AND EXPLORING FINANCIAL RETURNS

¹<http://finance.yahoo.com>

²<http://www.six-swiss-exchange.com>

³<http://www.stlouisfed.org>

Function:	
returns	generates returns from a price/index series
cumulated	generates indexed values from a returns series
drawdowns	computes drawdowns from financial returns
lowess	smooths a price/index series
turnpoints	finds turnpoints for a smoothed price/index series

2.2 RETURNS AND CUMULATED RETURNS SERIES

To calculate compound or simple *returns* (Bacon, 2008), usually on daily or monthly records for portfolio analysis and optimization, we can call the function `returns()`.

LISTING 2.2: FUNCTIONS TO COMPUTE AND CONVERT PRICE/INDEX VALUES AND FINANCIAL RETURNS

Function:	
returns	generates returns from a price/index series
cumulated	generates indexed values from a returns series
Arguments:	
x	a price/index for a uni or multivariate series of class <code>timeSeries</code>
method	the method of computing the returns "continuous", "discrete", "compound", "simple"
percentage	a logical, should percentual returns be computed?

The method argument allows us to define how the returns are computed. The methods "continuous" and "discrete" are synonyms for the methods "compound" and "simple", respectively.

In the following example we first compute the compound returns for the LP25 benchmark from the SWX data set, and then we cumulate the returns to recover the price/index series. To do so, we need to index the series to 1 on the first day before we calculate the returns. By cumulating the returns, we can recover the indexed series:

```
> LP25 <- SWX[, "LP25"]/as.numeric(SWX[1, "LP25"])
> head(LP25, 5)

GMT
      LP25
2000-01-03 1.00000
2000-01-04 0.98808
2000-01-05 0.98447
2000-01-06 0.98317
2000-01-07 0.99078

> head(returns(LP25), 5)
```



```

GMT
      LP25
2000-01-04 -0.0119943
2000-01-05 -0.0036571
2000-01-06 -0.0013239
2000-01-07  0.0077150
2000-01-10  0.0030291

> head(cumulated(returns(LP25)), 5)

GMT
      LP25
2000-01-04 0.98808
2000-01-05 0.98447
2000-01-06 0.98317
2000-01-07 0.99078
2000-01-10 0.99379

> head(returns(cumulated(returns(LP25))), 4)

GMT
      LP25
2000-01-05 -0.0036571
2000-01-06 -0.0013239
2000-01-07  0.0077150
2000-01-10  0.0030291

```

2.3 DRAWDOWNS SERIES

Drawdown measures describe the decline from a historical peak in some price or index variable. This is typically in cumulated return series of a financial trading strategy.

LISTING 2.3: FUNCTION TO COMPUTE DRAWDOWNS FROM FINANCIAL RETURNS

```

Function:
drawdowns          computes drawdowns from financial returns

Arguments:
x                  a 'timeSeries' of financial returns

```

The maximum drawdown up to a given time is the maximum of the drawdown over the history of the price or index variable and can be considered as an indicator of risk. A `drawdowns()`⁴ series can be computed from a return series as follows:

```
> head(drawdowns(SWX.RET[, 1:4]))
```

⁴The functions `drawdowns()` and `drawdownStats()` are reimplemented from the contributed R package `PerformanceAnalytics`, based on code written by [Carl & Peterson \(2008\)](#).

GMT	SBI	SPI	SII	LP25
2000-01-04	-0.0020881	-0.0343901	0.00000000	-0.0119943
2000-01-05	-0.0021924	-0.0444404	-0.00495531	-0.0156075
2000-01-06	-0.0035492	-0.0328598	-0.00116130	-0.0169107
2000-01-07	-0.0031321	-0.0111363	-0.00177680	-0.0093262
2000-01-10	-0.0031321	-0.0090521	0.00000000	-0.0063254
2000-01-11	-0.0041756	-0.0118006	-0.00029385	-0.0087326

The function returns a univariate time series object of `timeSeries`. Multiplying the series by a factor of 100 gives us the returns in percentages.

2.4 DURATIONS SERIES

The *duration* is the interval between time series records, and can be computed using the function `durations()`.

LISTING 2.4: FUNCTION TO COMPUTE INTERVALS FROM A FINANCIAL SERIES

Function:
`durations` computes intervals from a financial series

Arguments:
`x` a 'timeSeries' of financial returns

Let us consider 10 randomly selected records in ascending order:

```
> SPI <- SWX[, "SPI"]
> SPI10 <- SPI[c(4, 9, 51, 89, 311, 513, 756, 919, 1235, 1648),
+ ]
> SPI10
```

GMT	SPI
2000-01-06	4861.4
2000-01-13	4985.2
2000-03-13	4693.7
2000-05-04	5133.6
2001-03-12	5116.3
2001-12-19	4231.5
2002-11-25	3579.1
2003-07-10	3466.8
2004-09-24	4067.7
2006-04-26	6264.4

Then we compute their intervals in units of days.

```
> durations(SPI10)/(24 * 3600)
```

GMT	Duration
2000-01-06	NA

2000-01-13	7
2000-03-13	60
2000-05-04	52
2001-03-12	312
2001-12-19	282
2002-11-25	341
2003-07-10	227
2004-09-24	442
2006-04-26	579

Here we have divided the returned value by the length of one day, i.e. 24×3600 seconds so that the intervals are given in days.

Intervals are especially of interest when we consider irregular time series and we want to know the time period between consecutive records, or between consecutive events, such as turnpoints.

2.5 HOW TO ADD YOUR OWN FUNCTIONS

It is very easy to add new generic functions operating on `timeSeries` objects to Rmetrics. In this section we show how to add a function to smooth price and index series, and a function to find the turnpoints in a time series.

LISTING 2.5: USER SUPPLIED FUNCTIONS TO SMOOTH A SERIES AND TO FIND TURNING POINTS

Function:	
<code>lowess</code>	a locally-weighted polynomial regression smoother
<code>turnpoints</code>	finds turnpoints in a financial series
Arguments:	
<code>x</code>	a 'timeSeries' of financial returns
<code>f</code>	the smoother span for lowess
<code>iter</code>	the number of robustifying iterations for lowess

How to smooth a time series with `lowess()`

We will demonstrate this by writing a smoother for financial time series built on top of the `lowess()` function from the R `stats` package. `lowess()` is a smoother based on robust locally weighted regression (Cleveland, 1979, 1981). Using the function `setMethod()` from the R `methods` package we can create and save a formal method for `lowess()`.

```
> setMethod("lowess", "timeSeries", function(x, y = NULL, f = 2/3,
  iter = 3) {
  stopifnot(isUnivariate(x))
  ans <- stats::lowess(x = as.vector(x), y, f, iter)
  series(x) <- matrix(ans$y, ncol = 1)
  x
})
```

```
  })
[1] "lowess"
```

We first extract the SPI from the SWX data set and then we smooth the index. The argument `f` determines the smoother span. This gives the proportion of points which influence the smooth at each value. Larger values give more smoothness and smaller values result in less smoothness keeping more of the original structure of the curve. The graph in [Figure 2.1](#) shows the result.

```
> SPI <- SWX[, "SPI"]
> SPI.LW <- lowess(SPI, f = 0.08)
> plot(SPI)
> lines(SPI.LW, col = "brown", lwd = 2)
```

How to find the turnpoints of a time series

If we are interested in the turnpoints of the smoothed SPI index, we can use the `turnpoints()` function from the contributed R package `pastecs` ([Ibanez, Grosjean & Etienne, 2009](#)). The function determines the number and the positions of extrema, i.e. the turning points, either peaks or pits, in a regular time series. Writing an S4 method is straightforward:

```
> library(pastecs)
> setMethod("turnpoints", "timeSeries",

  function(x)
  {
    stopifnot(isUnivariate(x))
    tp <- suppressWarnings(pastecs::turnpoints(as.ts(x)))
    recordIDs <- data.frame(tp$peaks, tp$pits)
    rownames(recordIDs) <- rownames(x)
    colnames(recordIDs) <- c("peaks", "pits")
    timeSeries(data = x, charvec = time(x),
                units = colnames(x), zone = finCenter(x),
                FinCenter = finCenter(x),
                recordIDs = recordIDs, title = x$title,
                documentation = x@documentation)
  }
)
[1] "turnpoints"
```

Using the function `isUnivariate()`, we first check if the input time series is univariate, then we compute the turnpoints, converting the `timeSeries` into an object of class `ts` as expect by the underlying function. Then we extract the peaks and pits and save them as record identification codes in the data.frame `recordIDs`, which is represented by a slot in the S4 `timeSeries` object. The result is given back as a `timeSeries` object.




FIGURE 2.1: The graph shows the Swiss performance index SPI overlayed by a smoothed curve with turnpoints. For the smoother we used the R function `lowess()`, and for the turnpoints the contributed R function `turnpoints()` from the contributed R package `pastecs`.

Now let us compute the turnpoints for the smoothed SPI. We plot the original index series and the smoothed series. On top we put points for the peaks and pits in green and red, respectively.

```
> SPI.TP <- turnpoints(SPI.LW)
> SPI.PEAKS <- SPI.TP[SPI.TP@recordIDs[, "peaks"] == TRUE, ]
> SPI.PITS <- SPI.TP[SPI.TP@recordIDs[, "pits"] == TRUE, ]
> plot(SPI)
> lines(SPI.LW, col = "brown", lwd = 2)
> points(SPI.PEAKS, col = "green3", pch = 24)
> points(SPI.PITS, col = "red", pch = 25)
```

The turnpoints are added to [Figure 2.1](#).

CHAPTER 3

BASIC STATISTICS OF FINANCIAL ASSETS

```
> library(fPortfolio)
```

Rmetrics provides several functions and methods to compute basic statistics of financial time series from S4 `timeSeries` objects. These include summary and basic statistics, drawdown statistics, sample mean and covariance estimation, and quantile and risk estimation, amongst others. Moreover, we have functions to compute column statistics and cumulated column statistics, which are very useful tools if we are interested in the statistical properties of each column of a data set of assets.

3.1 SUMMARY STATISTICS

Three functions are available to compute basic statistics from a univariate or multivariate data set of assets, the generic `summary()` function and the functions `basicStats()` and `drawdownsStats()`.

Information on the size of a data set can be obtained from the functions `nrow`, `ncol`, `NROW`, `NCOL`, and `dim`. `nrow` and `ncol` return the number of rows or columns present in the `timeSeries` object `x`; `NCOL` and `NROW` do the same, but treat a univariate time series as 1-column multivariate time series.

LISTING 3.1: FUNCTIONS FOR COMPUTING BASIC STATISTICS OF FINANCIAL RETURNS

Function:	
<code>summary</code>	generates summary statistics of assets
<code>basicStats</code>	generates a basic statistics summary of assets
<code>drawdownsStats</code>	computes drawdown statistics from returns
<code>mean, cov</code>	computes sample mean and covariance of assets
<code>skewness</code>	computes sample skewness of assets

kurtosis	computes sample kurtosis of assets
quantile	computes quantiles of assets
colStats	computes column statistics of a data set of assets
colCumStats	computes cumulative column statistics of assets
covRisk	computes covariance portfolio risk
varRisk	computes value-at-risk for a portfolio
cvarRisk	computes conditional value-at-risk for a portfolio

How to create summary statistics

The `summary()` function for `timeSeries` objects behaves in the same way as for numerical matrices. The function returns the minimum and maximum values for each series, the first and third quartiles, and the mean and median values. The following example computes summary statistics for the log-returns of the SWX data set.

```
> summary(SWX.RET)
```

SBI	SPI	SII
Min. : -6.87e-03	Min. : -0.069039	Min. : -1.59e-02
1st Qu.: -7.24e-04	1st Qu.: -0.004794	1st Qu.: -1.40e-03
Median : 0.00e+00	Median : 0.000293	Median : 4.87e-05
Mean : 4.70e-06	Mean : 0.000215	Mean : 2.03e-04
3rd Qu.: 7.85e-04	3rd Qu.: 0.005681	3rd Qu.: 1.85e-03
Max. : 5.76e-03	Max. : 0.057860	Max. : 1.54e-02

LP25	LP40	LP60
Min. : -0.013154	Min. : -0.019720	Min. : -0.028106
1st Qu.: -0.001248	1st Qu.: -0.001940	1st Qu.: -0.002916
Median : 0.000247	Median : 0.000351	Median : 0.000430
Mean : 0.000139	Mean : 0.000135	Mean : 0.000123
3rd Qu.: 0.001587	3rd Qu.: 0.002283	3rd Qu.: 0.003326
Max. : 0.013287	Max. : 0.021178	Max. : 0.032057

How to create a basic statistics report

The function `basicStats()` behaves similarly to `summary()` but returns a broader spectrum of statistical measures.

```
> args(basicStats)
function (x, ci = 0.95)
NULL
```

The argument `ci` specifies the confidence interval for calculating standard errors.

The following example computes daily basic statistics for the percentual log-returns of the three SWX indices, SPI, SBI, SII, and the LP25 benchmark index from the SWX data set.

```
> basicStats(SWX.RET[, 1:4])
```


	SBI	SPI	SII	LP25
nobs	1916.000000	1916.000000	1916.000000	1916.000000
NAs	0.000000	0.000000	0.000000	0.000000
Minimum	-0.006868	-0.069039	-0.015867	-0.013154
Maximum	0.005757	0.057860	0.015411	0.013287
1. Quartile	-0.000724	-0.004794	-0.001397	-0.001248
3. Quartile	0.000785	0.005681	0.001851	0.001587
Mean	0.000005	0.000215	0.000203	0.000139
Median	0.000000	0.000293	0.000049	0.000247
Sum	0.008930	0.412553	0.389689	0.266111
SE Mean	0.000030	0.000248	0.000069	0.000058
LCL Mean	-0.000054	-0.000270	0.000069	0.000025
UCL Mean	0.000063	0.000701	0.000338	0.000253
Variance	0.000002	0.000118	0.000009	0.000006
Stdev	0.001298	0.010843	0.003005	0.002542
Skewness	-0.313206	-0.221507	0.084294	-0.134810
Kurtosis	1.516963	5.213489	2.592051	2.893592

The `basicStats()` function returns a data frame with the following entries and row names: `nobs`, `NAs`, `Minimum`, `Maximum`, `1. Quartile`, `3. Quartile`, `Mean`, `Median`, `Sum`, `SE Mean`, `LCL Mean`, `UCL Mean`, `Variance`, `Stdev`, `Skewness`, `Kurtosis`.

How to compute drawdown statistics

To compute the drawdowns statistics for the LPP25 benchmark index we use the `drawdownsStats()` function

```
> args(drawdownsStats)
function (x, ...)
NULL
```

which requires a univariate `timeSeries` object as input.

The example

```
> LP25 <- SWX.RET[, "LP25"]
> drawdownsStats(LP25)[1:10, ]
```

	From	Trough	To	Depth	Length	ToTrough	Recovery
1	2001-05-23	2001-09-21	2003-08-22	-0.084709	588	88	500
2	2006-02-23	2006-06-13	2006-09-04	-0.038749	138	79	59
3	2001-02-07	2001-03-22	2001-05-17	-0.031139	72	32	40
4	2004-03-09	2004-06-14	2004-11-12	-0.030972	179	70	109
5	2000-09-06	2000-10-12	2001-02-06	-0.021031	110	27	83
6	2005-10-04	2005-10-28	2005-11-24	-0.018214	38	19	19
7	2003-09-19	2003-09-30	2003-11-03	-0.017436	32	8	24
8	2000-03-23	2000-05-22	2000-07-11	-0.017321	79	43	36
9	2000-01-04	2000-01-06	2000-01-17	-0.016911	10	3	7
10	2000-01-18	2000-02-22	2000-03-17	-0.016687	44	26	18

returns the first ten drawdowns from the function value, which is a data.frame. The data frame lists the depth of the drawdown, the from (start) date, the

trough period, the to (end) date, the length of the period, the peak-to-trough, and the recovery periods. Note that lengths are measured in units of time series events.

3.2 SAMPLE MEAN AND COVARIANCE ESTIMATES

How to compute the sample mean

A fundamental task in many statistical analyses is to estimate a location parameter for the distribution, that is to find a typical or central value that best describes the data.

LISTING 3.2: FUNCTIONS TO ESTIMATE MOMENTS AND RELATED QUANTITIES

Function:	
mean	computes sample mean
var	computes sample variance
cov	computes sample covariance
skewness	computes sample skewness
kurtosis	computes sample kurtosis
Arguments:	
x	a 'timeSeries' object.

Sample means can be computed using R's base functions `mean()`. Note that calling the function `mean()` on a multivariate time series will return the grand mean, as if the time series were a numeric matrix. To obtain the column means, which is what you usually require for your financial time series, you have to apply the function `colMeans()`.

```
> mean(100 * SWX.RET)
[1] 0.013664

> colMeans(100 * SWX.RET)
      SBI      SPI      SII      LP25      LP40      LP60
0.00046605 0.02153198 0.02033869 0.01388886 0.01349041 0.01226859
```

How to compute the sample variance and covariance

Sample variance and covariance can be computed using the R base functions `var()` and `cov()`. Note that R's base function `cov()` operates in the same way on a `timeSeries` object as on a numeric matrix.

```
> Covariance <- round(cov(100 * SWX.RET), digits = 4)
> Covariance
      SBI      SPI      SII      LP25      LP40      LP60
SBI    0.0169 -0.0415 0.0014 -0.0011 -0.0094 -0.0206
SPI   -0.0415  1.1757 0.0066  0.2204  0.3617  0.5464
```

```

SII    0.0014  0.0066 0.0903  0.0027  0.0041  0.0062
LP25 -0.0011  0.2204 0.0027  0.0646  0.0993  0.1464
LP40 -0.0094  0.3617 0.0041  0.0993  0.1578  0.2372
LP60 -0.0206  0.5464 0.0062  0.1464  0.2372  0.3609

```

Here, we have rounded the output to four digits.

3.3 ESTIMATES FOR HIGHER MOMENTS

How to compute the sample skewness

Skewness is a measure of symmetry, or more precisely, the lack of symmetry. A distribution, or data set, is symmetric if it looks the same to the left and right of the centre point.

```

> args(skewness)
function (x, ...)
NULL

> SPI <- SWX[, "SPI"]
> skewness(SPI)
[1] 0.51945
attr(,"method")
[1] "moment"

```

How to compute the sample kurtosis

Kurtosis is a measure of whether the data are peaked or flat relative to a normal distribution. That is, data sets with high kurtosis tend to have a distinct peak near the mean, decline rather rapidly, and have heavier tails. Data sets with low kurtosis tend to have a flat top near the mean rather than a sharp peak¹.

```

> args(kurtosis)
function (x, ...)
NULL

> kurtosis(SPI)
[1] -0.31378
attr(,"method")
[1] "excess"

```

Note that R comes with the base functions `mean()` and `cov()`, but does not provide functions to compute skewness and kurtosis. The functions `skewness()` and `kurtosis()` are added by Rmetrics.

¹A distribution with high kurtosis is known as leptokurtic, whereas a distribution with low kurtosis is platykurtic.

3.4 QUANTILES AND RELATED RISK MEASURES

How to compute quantiles

Quantiles of assets can be calculated using R's base generic function `quantile()`. This function produces sample quantiles corresponding to the given probabilities. The smallest observation corresponds to a probability of 0 and the largest to a probability of 1. Note that according to [Hyndman & Fan \(1996\)](#) there are different ways to compute quantiles. The method used in finance for calculating the (conditional) Value at Risk is `type=1`, which is not the default setting.

```
> quantile(SWX.RET, probs = seq(0, 1, 0.25), type = 1)
      0%      25%      50%      75%     100%
-0.06903905 -0.00162222  0.00019352  0.00202488  0.05786042
```

As you can see, the function concatenates the columns of all assets in the data set to one vector (the same as for the `mean()`) and then computes the quantiles. To compute the quantiles for each column, use the function `colQuantiles()`, and do not forget to specify the proper `type=1`.

```
> colQuantiles(SWX.RET, prob = 0.05, type = 1)
      SBI      SPI      SII      LP25      LP40      LP60
-0.0022108 -0.0175881 -0.0044958 -0.0041020 -0.0064546 -0.0098262
```

Portfolio risk measures

To compute the three major risk measures for portfolios `Rmetrics` provides the functions `covRisk`, `varRisk`, and `cvarRisk`.

LISTING 3.3: FUNCTIONS TO COMPUTE PORTFOLIO RISK MEASURES

Function:	
<code>covRisk</code>	computes covariance portfolio risk
<code>varRisk</code>	computes Value at Risk for a portfolio
<code>cvarRisk</code>	computes conditional Value at Risk
Arguments:	
<code>x</code>	a 'timeSeries' object of asset returns
<code>weights</code>	vector of portfolio weights
<code>alpha</code>	the VaR and CVaR confidence level

The example shows the three risk measures for an equally weighted portfolio composed of Swiss equities, SPI, Swiss bonds, SBI, and Swiss reits, SII. For the sample covariance risk we obtain

```
> SWX3 <- 100 * SWX.RET[, 1:3]
> covRisk(SWX3, weights = c(1, 1, 1)/3)
```

```
Cov
0.36755
```

and for the sample VaR and Conditional VaR we obtain

```
> varRisk(SWX3, weights = c(1, 1, 1)/3, alpha = 0.05)
VaR.5%
-0.56351
> cvarRisk(SWX3, weights = c(1, 1, 1)/3, alpha = 0.05)
CVaR.5%
-0.87832
```

How to detect extreme values and outliers

The Rmetrics package fExtremes allows us to investigate univariate time-Series objects from the point of view of extreme value theory. The package provides functions to investigate extreme values in a time series using peak over threshold and block methods. From this we can estimate *Value-at-Risk* and *Conditional-Value at-Risk* much more reliably than is possible using sample estimates.

For a detailed description of the statistical approaches and algorithms for the analysis of extreme values in financial time series we refer to the Rmetrics ebook *Managing Risk with R/Rmetrics*.

3.5 COMPUTING COLUMN STATISTICS

Rmetrics implements several functions to compute column and row statistics of univariate and multivariate timeSeries objects. The functions return a numeric vector of the same length as the number of columns of the timeSeries.

Amongst the column statistics functions are

LISTING 3.4: COLUMN STATISTICS FUNCTIONS

Functions:	
colStats	calculates arbitrary column statistics
colSums	returns column sums
colMeans	returns column means
colSds	returns column standard deviations
colVars	returns column variances
colSkewness	returns column skewness
colKurtosis	returns column kurtosis
colMaxs	returns maximum values in each column
colMins	returns minimum values in each column
colProds	returns product of all values in each column
colQuantiles	returns quantiles of each column

Arguments:

x a 'timeSeries' object.

```
> 100 * colMeans(returns(SWX))
      SBI      SPI      SII      LP25      LP40      LP60
0.00046605 0.02153198 0.02033869 0.01388886 0.01349041 0.01226859

> 100 * colQuantiles(returns(SWX))
      SBI      SPI      SII      LP25      LP40      LP60
-0.21941 -1.74757 -0.44678 -0.40828 -0.64300 -0.98188
```

You can also define your own statistical functions and execute them with the function `colStats()`. If, for example, you want to know the column medians of the `timeSeries`, you can simply write

```
> round(colStats(returns(SWX, percentage = TRUE), FUN = "median"),
      digits = 4)
      SBI      SPI      SII      LP25      LP40      LP60
0.0000 0.0293 0.0049 0.0247 0.0351 0.0430
```

3.6 COMPUTING CUMULATED COLUMN STATISTICS

Functions to compute cumulated column statistics are also available in `Rmetrics`. These are

LISTING 3.5: FUNCTIONS FOR CUMULATED COLUMN STATISTICS

Functions:	
<code>colCumstats</code>	returns user-defined column statistics
<code>colCumsums</code>	returns column-cumulated sums
<code>colCummaxs</code>	returns column-cumulated maximums
<code>colCummins</code>	returns column-cumulated minimums
<code>colCumprods</code>	returns column-cumulated products
<code>colCumreturns</code>	returns column-cumulated returns

Arguments:

x a 'timeSeries' object.

The function `colCumstats()` allows you to define your own functions to compute cumulated column statistics, in the same way as for the function `colStats()`.

CHAPTER 4

ROBUST MEAN AND COVARIANCE ESTIMATES

Robust statistics provides an alternative approach to classical statistical methods. The idea behind robust statistics is to produce estimators that are not unduly affected by small departures from model assumptions. In Rmetrics we have included robust estimators for the mean and covariance of financial assets from several R packages.

Additionally, we have implemented a covariance ellipse plot, which visualizes the difference between two or more covariance matrices. It is intended to compare different methods of covariance estimation. We also show how to detect multivariate outliers.

Robust covariance estimators of a data set of asset returns are of great interest for the optimization of robust mean-covariance portfolios, where we replace the sample covariance estimate with a robust covariance estimate. From many investigations, we know that the use of robust covariances instead of the sample covariance achieves a much better diversification of the mean-variance portfolio weights.

4.1 ROBUST COVARIANCE ESTIMATORS

The function `assetsMeanCov()` provides a collection of several robust estimators. The functions have their origin in several contributed R packages for robust estimation.

LISTING 4.1: THE ROBUST FUNCTION ESTIMATORS FUNCTIONS

Function:	
<code>assetsMeanCov</code>	returns robustified covariance estimates
<code>getCenterRob</code>	extracts the robust centre estimate
<code>getCovRob</code>	extracts the robust covariance estimate
<code>covEllipsesPlot</code>	creates a covariance ellipses plot

`assetsOutliers` detects multivariate outliers in assets

Arguments:

<code>x</code>	a univariate 'timeSeries' object
<code>method</code>	the method of robustification:
"cov"	uses the sample covariance estimator from [base]
"mve"	uses the "mve" estimator from [MASS]
"mcd"	uses the "mcd" estimator from [MASS]
"MCD"	uses the "MCD" estimator from [robustbase]
"OGK"	uses the "OGK" estimator from [robustbase]
"nnve"	uses the "nnve" estimator from [covRobust]
"shrink"	uses "shrinkage" estimator from [corpcor]
"bagged"	uses "bagging" estimator from [corpcor]

First, let us have a look at the argument list of the function `assetsMeanCov()` and the methods provided.

```
> args(assetsMeanCov)
function (x, method = c("cov", "mve", "mcd", "MCD", "OGK", "nnve",
  "shrink", "bagged"), check = TRUE, force = TRUE, baggedR = 100,
  sigmamu = scaleTau2, alpha = 1/2, ...)
NULL
```

Through the argument `method`, we can select the desired estimator. The function `assetsMeanCov()` returns a named list with four entries `center` (the estimated mean), `cov` (the estimated covariance matrix), and `mu` and `Sigma` which are just synonyms for `center` and `cov`. In addition the returned value of the function has a control attribute `attr("control")`, a character vector which holds the name of the method of the estimator, the size (number) of assets, and two flags. If the covariance matrix was positive definite then the `posdef` flag is set to `TRUE`, and if not, then the flag is set to `FALSE`.

```
> assetsMeanCov(100 * SWX.RET[, c(1:2, 4:5)], method = "cov")
$center
      SBI      SPI      LP25      LP40
0.00046605 0.02153198 0.01388886 0.01349041

$cov
      SBI      SPI      LP25      LP40
SBI    0.0168515 -0.041468 -0.001076 -0.0094419
SPI   -0.0414682  1.175681  0.220376  0.3616704
LP25  -0.0010760  0.220376  0.064639  0.0992780
LP40  -0.0094419  0.361670  0.099278  0.1577805

$mu
      SBI      SPI      LP25      LP40
0.00046605 0.02153198 0.01388886 0.01349041

$Sigma
      SBI      SPI      LP25      LP40
```



```

SBI    0.0168515 -0.041468 -0.001076 -0.0094419
SPI    -0.0414682  1.175681  0.220376  0.3616704
LP25   -0.0010760  0.220376  0.064639  0.0992780
LP40   -0.0094419  0.361670  0.099278  0.1577805

```

```

attr("control")
  method    size posdef forced forced
  "cov"      "4"  "TRUE" "FALSE" "TRUE"

```

The functions `getCenterRob()` and `getCovRob()` can be used to extract the robust mean, center, and the robust covariance, `cov`, from an object as returned by the function `assetsMeanCov()`.

4.2 COMPARISONS OF ROBUST COVARIANCES

The function `covEllipsesPlot()` visualizes the differences between two covariance matrices. This allows us to compare the sample estimate with robust estimates, or to compare robust estimators with each other.

How to display ellipses plots

```

> args(covEllipsesPlot)
function (x = list(), ...)
NULL

```

The list argument has at least two covariance matrices as input, the dots argument allows us to pass optional arguments to the underlying plot, lines and text functions. Several examples how to use the `covEllipsesPlot()` function are shown in the following sections when we compare different robust covariance estimates.

4.3 MINIMUM VOLUME ELLIPSOID ESTIMATOR

The method "mve" is the minimum volume ellipsoid estimator as implemented in R's recommended package MASS (Venables & Ripley, 2008). This is called internally with the argument `method="mve"`

```

> args(MASS::cov.rob)
function (x, cor = FALSE, quantile.used = floor((n + p + 1)/2),
  method = c("mve", "mcd", "classical"), nsamp = "best", seed)
NULL

```

The following example shows how to call the estimator from the function suite `assetsMeanCov()` and how to extract the robust `$center` and `cov`

estimate. We investigate Swiss, SPI, and foreign equity indexes, MPI, and Swiss, SBI, and foreign bond indexes, LMI, which are stored in the columns 1, 2, 4, 5 of the LPP2005.RET data set

```
> set.seed(1954)
> lppData <- 100 * LPP2005.RET[, c(1:2, 4:5)]
> ans.mve <- assetsMeanCov(lppData, method = "mve")
> getCenterRob(ans.mve)

      SBI      SPI      LMI      MPI
-0.0041001 0.1406959 0.0013634 0.1246928

> getCovRob(ans.mve)

      SBI      SPI      LMI      MPI
SBI 0.0134075 -0.0114161 0.0083589 -0.017509
SPI -0.0114161 0.3380508 -0.0064112 0.194845
LMI 0.0083589 -0.0064112 0.0127780 -0.015432
MPI -0.0175086 0.1948451 -0.0154321 0.295309

> attr(ans.mve, "control")

method size posdef forced forced
"mve"   "4"  "TRUE" "FALSE" "TRUE"
```

Note that an attribute called "control" is returned, which allows us to extract additional information from the selected estimator.

With the help of the function `covEllipsesPlot()`, we can now compare the sample covariances with the "mve" robustified covariances

```
> covEllipsesPlot(list(cov(lppData), ans.mve$cov))
> title(main = "Sample vs. MVE Covariances")
```

The result is shown in [Figure 4.1](#).

4.4 MINIMUM COVARIANCE DETERMINANT ESTIMATOR

Two methods, called "mcd" and "MCD", are available to estimate a robust mean and covariance by the minimum covariance determinant estimator (Rousseeuw, 1985; Rousseeuw & Van Driessen, 1999). The first method uses the function `cov.rob()` from the MASS package (Venables & Ripley, 2008).

```
> args(MASS::cov.rob)

function (x, cor = FALSE, quantile.used = floor((n + p + 1)/2),
  method = c("mve", "mcd", "classical"), nsamp = "best", seed)
NULL
```

and the second method uses the function `covMcd()` from the contributed package `robustbase` (Rousseeuw, Croux, Todorov, Ruckstuhl, Salibián-Barrera, Verbeke & Maechler, 2008).




FIGURE 4.1: Comparison of sample and MVE robust Covariances.

```
> args(robustbase::covMcd)
function (x, cor = FALSE, raw.only = FALSE, alpha = control$alpha,
  nsamp = control$nsamp, nmini = control$nmini, kmini = control$kmini,
  scalefn = control$scalefn, maxcsteps = control$maxcsteps,
  inithsets = NULL, save.hsets = FALSE, names = TRUE, seed = control$seed,
  tolSolve = control$tolSolve, trace = control$trace, use.correction = control$use.correction,
  wgtFUN = control$wgtFUN, control = rrcov.control())
NULL
```

You can call the estimators for the two methods as

```
> ans.mcd <- assetsMeanCov(lppData, "mcd")
> ans.MCD <- assetsMeanCov(lppData, "MCD")
```

and compare them

```
> getCovRob(ans.mcd)
      SBI      SPI      LMI      MPI
SBI 0.0134909 -0.0107596 0.0083536 -0.013324
SPI -0.0107596 0.3377521 -0.0094661 0.205287
```




FIGURE 4.2: Comparison of sample and mcd/MCD robust covariances.

```

LMI 0.0083536 -0.0094661 0.0126455 -0.015428
MPI -0.0133236 0.2052866 -0.0154282 0.322434

> getCovRob(ans.MCD)

      SBI      SPI      LMI      MPI
SBI 0.0159968 -0.011929 0.0097842 -0.016060
SPI -0.0119294 0.403244 -0.0101378 0.242938
LMI 0.0097842 -0.010138 0.0149387 -0.018262
MPI -0.0160600 0.242938 -0.0182619 0.370622

> covEllipsesPlot(list(ans.mcd$cov, ans.MCD$cov))
> title(main = "mcd vs. MCD Covariances")

```

The result is shown in Figure 4.2.

4.5 ORTHOGONALIZED GNANADESIKAN-KETTENRING ESTIMATOR

The "OGK" method¹ computes the orthogonalized pairwise covariance matrix estimate described in [Maronna & Zamar \(2002\)](#). The pairwise

¹The "OGK" method is very efficient for large covariance matrices.

proposal goes back to [Gnanadesikan & Kettenring \(1972\)](#). The estimator is implemented in the contributed R package `robustbase` ([Rousseeuw et al., 2008](#)).

```
> args(robustbase::covOGK)
function (X, n.iter = 2, sigmamu, rcov = covGK, weight.fn = hard.rejection,
  keep.data = FALSE, ...)
NULL
```

We can use the estimator in the same way as the others

```
> ans.ogk <- assetsMeanCov(lppData, "OGK")
> getCovRob(ans.ogk)

      SBI      SPI      LMI      MPI
SBI  0.016585 -0.010687  0.010368 -0.012045
SPI -0.010687  0.420980 -0.010904  0.275285
LMI  0.010368 -0.010904  0.015081 -0.017319
MPI -0.012045  0.275285 -0.017319  0.407193

> covEllipsesPlot(list(cov(lppData), ans.ogk$cov))
> title(main = "Sample vs. OGK Covariances")
```

The result is shown in [Figure 4.3](#).

4.6 NEAREST-NEIGHBOUR VARIANCE ESTIMATOR

The method "nnve" provides robust covariance estimation by the nearest neighbour variance estimation method of [Wang & Raftery \(2002\)](#). The function `cov.nnve()` is implemented in the contributed R package `covRobust` ([Wang, Raftery & Fraley, 2008](#)) and is available as an internal function built into `Rmetrics`' `cov.nnve()`² function.

```
> args(fAssets:::cov.nnve)
function (datamat, k = 12, pnoise = 0.05, emconv = 0.001, bound = 1.5,
  extension = TRUE, devsm = 0.01)
NULL
```

and again

```
> ans.nnve <- assetsMeanCov(lppData, "nnve")
> getCenterRob(ans.nnve)

      SBI      SPI      LMI      MPI
4.0663e-05 8.4175e-02 5.5315e-03 5.9052e-02

> getCovRob(ans.nnve)

      SBI      SPI      LMI      MPI
SBI  0.0115245 -0.0092243  0.0068053 -0.010618
SPI -0.0092243  0.2978305 -0.0080592  0.180658
LMI  0.0068053 -0.0080592  0.0111839 -0.014251
MPI -0.0106177  0.1806576 -0.0142508  0.270602
```

²This means that the loading of the R package `covRobust` is not required.




FIGURE 4.3: Comparison of sample and OGK robust covariances.

```
> attr(ans.nnve, "control")
method size posdef forced forced
"nnve"  "4"  "TRUE" "FALSE" "TRUE"

> covEllipsesPlot(list(cov(lppData), ans.nnve$cov))
```

4.7 SHRINKAGE ESTIMATOR

The "shrink" method provides robust covariance estimation by the shrinkage method.

The `shrink()` function is implemented in the contributed R package `corpcor` (Schaefer, Opgen-Rhein & Strimmer, 2008) and is available as an internal function built into the `Rmetrics`' `cov.nnve()`³ function.

```
> ans.shrink <- assetsMeanCov(lppData, "shrink")
> getCenterRob(ans.shrink)
```

³This means that the loading of the Rpackage `covRobust` is not required.

```

      SBI      SPI      LMI      MPI
4.0663e-05 8.4175e-02 5.5315e-03 5.9052e-02

> getCovRob(ans.shrink)

      SBI      SPI      LMI      MPI
SBI  0.0158996 -0.012387  0.0095313 -0.015447
SPI -0.0123872  0.584612 -0.0136834  0.400155
LMI  0.0095313 -0.013683  0.0149511 -0.022674
MPI -0.0154466  0.400155 -0.0226738  0.535033
attr(,"lambda")
[1] 0.027802

> attr(ans.shrink, "control")

  method      size posdef forced forced
"shrink"      "4"  "TRUE" "FALSE" "TRUE"

> covEllipsesPlot(list(cov(lppData), ans.shrink$cov))

```

4.8 BAGGING ESTIMATOR

The "bagged" method provides a variance-reduced estimator of the covariance matrix using *bootstrap aggregation*, hence the name "bagged". "bagged" was implemented in a previous version of the contributed R package `corpcor` and is available as an internal function built into `Rmetrics`^{4 5}.

```

> ans.bagged <- assetsMeanCov(lppData, "bagged")
> getCenterRob(ans.bagged)

      SBI      SPI      LMI      MPI
4.0663e-05 8.4175e-02 5.5315e-03 5.9052e-02

> getCovRob(ans.bagged)

      SBI      SPI      LMI      MPI
SBI  0.0160430 -0.012019  0.0099105 -0.015864
SPI -0.0120185  0.578644 -0.0137355  0.409646
LMI  0.0099105 -0.013736  0.0150352 -0.023484
MPI -0.0158645  0.409646 -0.0234837  0.535081

> attr(ans.bagged, "control")

  method      R      size posdef forced forced
"bagged"    "100"    "4"  "TRUE" "FALSE" "TRUE"

> covEllipsesPlot(list(cov(lppData), ans.bagged$cov))

```

⁴This means that the loading of the R package `corpcor` is not required.

⁵"bagged" is built into `Rmetrics`' `cov.bagged()` function

4.9 HOW TO ADD A NEW ESTIMATOR TO THE SUITE

It is very simple to add your own estimators to the suite. To do this, you have to define an estimator function, which takes as its first argument a `timeSeries` object, and has a `dots` argument which allows you to pass optional arguments to the underlying estimator. This function has to return a list, with at least two entries, called `$center` and `$cov`.

How to add a multivariate Student's t estimator

The recommended MASS (Venables & Ripley, 2008) package implements a covariance estimator for the multivariate Student's t Distribution. The assumption that the data come from a multivariate Student's t distribution provides some degree of robustness to outliers without giving a high breakdown point. The breakdown point of an estimator is intuitively the proportion of incorrect large observations an estimator can handle before giving an arbitrarily large result.

To add this estimator to the function `assetsMeanCov()`, proceed as follows:

```
> # library(MASS)
> covt <- function(x, ...) MASS::cov.trob(x, ...)
> ans.covt <- assetsMeanCov(lppData, method = "covt")
> getCenterRob(ans.covt)
      SBI      SPI      LMI      MPI
-0.0012241  0.1187899  0.0017686  0.0920948
> getCovRob(ans.covt)
      SBI      SPI      LMI      MPI
SBI  0.0117503 -0.0091287  0.0073747 -0.011440
SPI -0.0091287  0.3581261 -0.0088391  0.239793
LMI  0.0073747 -0.0088391  0.0111483 -0.015084
MPI -0.0114398  0.2397926 -0.0150840  0.335681
```

How to write an adaptive re-weighted estimator

The contributed R package `mvoutlier` (Gschwandtner & Filzmoser, 2009) implements an adaptive re-weighted estimator for multivariate location and scatter with hard-rejection weights. The multivariate outliers are defined according to the supremum of the difference between the empirical distribution function of the robust Mahalanobis distance and the theoretical distribution function.

```
> # library(mvoutlier)
> arw <- function(x, ...) {
  ans <- mvoutlier::arw(as.matrix(x), colMeans(x), cov(x), ...)
  list(center = ans$m, cov = ans$c)
}
> ans.arw <- assetsMeanCov(lppData, method = "arw")
```



```

> getCenterRob(ans.arw)
      SBI      SPI      LMI      MPI
0.0019389 0.1503808 0.0060246 0.1121257

> getCovRob(ans.arw)
      SBI      SPI      LMI      MPI
SBI  0.0147643 -0.0092163  0.0090874 -0.013993
SPI -0.0092163  0.4045949 -0.0078659  0.267712
LMI  0.0090874 -0.0078659  0.0136485 -0.017386
MPI -0.0139928  0.2677117 -0.0173858  0.398549

```

4.10 HOW TO DETECT OUTLIERS IN A SET OF ASSETS

The function `assetsOutliers()` allows us to detect outliers by analyzing the estimates for the mean (center) and for the covariance matrix (cov) as described by [Filzmoser, Garrett & Reimann \(2005\)](#).

LISTING 4.2: FUNCTION TO DETECT OUTLIERS IN A MULTIVARIATE DATA SET OF ASSET RETURNS

Function:	
<code>assetsOutliers</code>	detects multivariate outliers in assets.
Values:	
<code>center</code>	mean vector as given by the input
<code>cov</code>	covariance matrix as given by the input
<code>cor</code>	correlation matrix computed from the covariances
<code>...</code>	optional arguments to be passed in
<code>quantile</code>	quantile
<code>outliers</code>	vector of outliers
<code>series</code>	return series of outliers

```

> args(assetsOutliers)
function (x, center, cov, ...)
NULL

```

The `assetsOutliers()` function expects as input a multivariate time-Series object and returns a named list with the mean vector, the covariance and correlation matrices, the quantile, and the outliers. From the vector of outliers we can relate position numbers to dates, and from the series we obtain the corresponding outlier returns.

The following example shows outlier detection for the mcd estimator

```

> outliers <- assetsOutliers(lppData, ans.mcd$center, ans.mcd$cov)
> outliers
$center
      SBI      SPI      LMI      MPI
-0.00017813 0.12976675 0.00043613 0.09916002

```

```
$cov
```

	SBI	SPI	LMI	MPI
SBI	0.0139246	-0.0093237	0.0084616	-0.014067
SPI	-0.0093237	0.3580961	-0.0085850	0.230081
LMI	0.0084616	-0.0085850	0.0126282	-0.017116
MPI	-0.0140674	0.2300810	-0.0171158	0.335830

```
$cor
```

	SBI	SPI	LMI	MPI
SBI	1.00000	-0.13204	0.63810	-0.20571
SPI	-0.13204	1.00000	-0.12766	0.66347
LMI	0.63810	-0.12766	1.00000	-0.26282
MPI	-0.20571	0.66347	-0.26282	1.00000

```
$quantile
```

```
[1] 12.439
```

```
$outliers
```

2005-11-04	2005-11-16	2005-12-14	2006-01-23	2006-03-28	2006-04-18
4	12	32	60	106	121
2006-05-11	2006-05-12	2006-05-17	2006-05-18	2006-05-22	2006-05-23
138	139	142	143	145	146
2006-05-24	2006-05-26	2006-05-30	2006-06-02	2006-06-05	2006-06-06
147	149	151	154	155	156
2006-06-08	2006-06-13	2006-06-15	2006-06-29	2006-06-30	2006-07-19
158	161	163	173	174	187
2006-07-24	2006-08-22	2006-09-22	2006-12-11	2007-02-27	2007-03-14
190	211	234	290	346	357

```
$series
```

```
GMT
```

	SBI	SPI	LMI	MPI
2005-11-04	-0.323575	-0.070276	-0.119853	1.167956
2005-11-16	0.299966	-0.718750	0.277342	0.387121
2005-12-14	-0.322309	-0.728358	0.224813	-0.647012
2006-01-23	-0.083813	0.008050	0.030172	-1.557646
2006-03-28	-0.285880	-0.072786	-0.336518	-0.824240
2006-04-18	-0.031136	-0.503126	0.111080	0.966265
2006-05-11	-0.377003	-0.109661	-0.133194	-1.142449
2006-05-12	-0.094473	-1.798973	-0.106245	-2.357418
2006-05-17	-0.196379	-2.840692	-0.187847	-1.400989
2006-05-18	0.180683	-0.971142	0.329356	-1.354839
2006-05-22	0.305009	-2.599776	0.350507	-3.009080
2006-05-23	0.000000	1.897068	0.179534	0.691444
2006-05-24	0.132662	-1.111559	-0.206778	0.119183
2006-05-26	0.038985	2.584213	-0.074680	2.192910
2006-05-30	-0.062373	-1.984241	-0.021746	-2.788328
2006-06-02	0.233973	0.699306	0.367916	-0.080858
2006-06-05	0.000000	0.000000	-0.051341	-1.457172
2006-06-06	-0.109119	-2.232654	-0.131750	-0.478857
2006-06-08	0.116900	-2.737931	0.191942	-0.923718
2006-06-13	0.139762	-2.399230	0.234377	-2.270706
2006-06-15	-0.209522	2.156939	-0.302316	2.407531
2006-06-29	0.062671	1.404176	0.182883	1.997279
2006-06-30	-0.054835	1.447381	0.152170	-0.237578

2006-07-19	0.085907	1.841958	0.184291	1.563777
2006-07-24	0.085760	1.923981	-0.025710	2.024162
2006-08-22	0.330960	0.313653	0.153983	0.976584
2006-09-22	0.198504	-0.918960	0.279948	-1.745647
2006-12-11	-0.258418	0.621642	0.058515	0.759359
2007-02-27	0.160030	-3.574624	0.258065	-3.375456
2007-03-14	0.114203	-2.820549	0.065135	-1.747682

PART II

EXPLORATORY DATA ANALYSIS OF ASSETS

INTRODUCTION

In [chapter 5](#) we show how to create and display graphs and plots of financial time series and their properties. We show how to use the generic plot function to produce univariate and multivariate graphs of assets. Several hints and recipes are given to customize the plots to the user's needs. In addition to the time series plots, we show how to display box plots, histograms and density plots, and quantile-quantile plots.

In [chapter 6](#) we present hints and tricks to customize graphs. This concerns plot labels, axis labels, the use of optional plot function arguments and how to select colours, fonts and plot symbols.

In [chapter 7](#) we show how to estimate the parameters of a data set of assets to a normal or Student's t distribution and how to simulate artificial data sets with the same statistical properties. In order to test whether the empirical asset returns are multivariate normally distributed we can perform hypothesis tests.

[chapter 8](#) deals with portfolio selection. We try to find which assets in a portfolio are similar, and thus grouped together in clusters. We introduce approaches which are suggested from a statistical point of view. We consider two approaches which group the asset returns by hierarchical or, alternatively, by k -means clustering. In addition, we show how we can group similar assets by an eigenvalue decomposition of the asset returns series. As a visual approach to detect similarities or dissimilarities we discuss star plots.

In [chapter 9](#) we discuss star and segment plots.

In [chapter 10](#) we concentrate on the pairwise comparison of assets. To make dependencies among the assets visible we display the correlation between two assets as scatter plots. In addition, we present alternative views displaying min/max panels, histogram panels, pie (or pac man) panels, shaded square panels, coloured ellipse panel, correlation test panels, and lowess fit panels. As an alternative, image correlation plots and bivariate hexagonal binned histogram plots are available.

CHAPTER 5

FINANCIAL TIME SERIES AND THEIR PROPERTIES

```
> library(fPortfolio)
```

Rmetrics offers several kinds of plot functions for quick and efficient exploratory data analysis of financial assets. These include *financial time series plots* for prices/indices, returns and their cumulated values, and plots for displaying their distributional properties: *box plots*, *histogram and density plots*, and *quantile-quantile plots*.

5.1 FINANCIAL TIME SERIES PLOTS

How to use the generic plot functions

The `plot()` function is a generic function to plot univariate and multivariate `timeSeries` objects. Furthermore, the two generic functions `lines()` and `points()` allow us to add lines and points to an already existing plot. The `plot()` function is implemented in the same spirit as the function `plot.ts()` for regular time series objects, `ts`, in R's base package `stats`. The function comes with the same arguments and some additional arguments, for user-specified "axis" labelling, and for modifying the plot "layout". As for `ts`, three different types of plots can be displayed: a multiple plot, a single plot, and a scatter plot.

LISTING 5.1: PLOT AND RELATED FUNCTIONS

Function:	
<code>plot</code>	displays a plot of a <code>timeSeries</code> object.
<code>lines</code>	adds lines to an already existing plot.
<code>points</code>	adds lines to an already existing plot.
<code>seriesPlot</code>	displays a time series plot given by its input.




FIGURE 5.1: Time series plots of the Swiss pension fund benchmark: The generic plot function creates a graph for each individual series where up to 10 subplots can be produced on one sheet of paper. The series of graphs shows the logarithmic returns of six asset classes and the three benchmark series included in the LPP2005 benchmark index.

<code>returnPlot</code>	displays returns given the price or index series.
<code>cumulatedPlot</code>	displays a cumulated series given the returns.

How to generate multiple plots

If the input argument `x` is a multivariate `timeSeries` object then the generic plot function creates a graph for each individual series. Up to ten subplots can be produced on one page.

```
> colnames(LPP2005.RET)
[1] "SBI"  "SPI"  "SII"  "LMI"  "MPI"  "ALT"  "LPP25" "LPP40" "LPP60"
> plot(LPP2005.RET, main = "LPP Pension Fund", col = "steelblue")
```




FIGURE 5.2: Time series plots of the LPP benchmark indices: The series of the three graphs show the logarithmic returns of the LPP benchmark indices, LPP25, LPP40, are part of the LPP2005 pension fund benchmark index family.

How to generate single plots

If the input argument `x` is a multivariate `timeSeries` object and the argument `plot.type` is set to "single" then the generic plot function creates a plot, where all curves are drawn in one plot on the same page.

```
> plot(SWX[, 6:4], plot.type = "single", , col = 2:4, xlab = "Date",
      ylab = "LP Index Family")
> title(main = "LP25 - LP40 - LP60")
> hgrid()
```

How to generate scatter plots

If two arguments `x` and `y` are specified, the generic plot function generates a scatter plot of two univariate `timeSeries` objects.

```
> SBI.RET <- 100 * SWX.RET[, "SBI"]
```




FIGURE 5.3: Scatter plot of the SPI versus the SBI: The generic plot function can also create scatter plots which show the values of the first versus the second time series. This figure shows the scatter plot for logarithmic returns of the SPI versus SBI indices in percentages.

```
> SPI.RET <- 100 * SWX.RET[, "SPI"]
> plot(SBI.RET, SPI.RET, xlab = "SBI", ylab = "SPI", pch = 19,
      cex = 0.4, col = "brown")
> grid()
```

This plot is useful if we want to compare the daily returns of two time series day by day. It gives an impression of the strength of correlations.

How to use tailored plot functions

Rmetrics comes with three major types of tailored plots to display a financial time series. We can display the price or index series given either the series itself or the returns, and we can also display the financial returns given the returns themselves or the price or index series. A third option allows us to plot the cumulated series when financial returns are given.

LISTING 5.2: TAILORED PLOT FUNCTIONS AND THEIR ARGUMENTS

Functions:

<code>seriesPlot</code>	generates an index plot
<code>returnPlot</code>	generates a financial returns plot
<code>cumulatedPlot</code>	generates a cumulative series plot

Arguments:

<code>labels</code>	a logical flag. Should the plot be returned with default labels? By default TRUE
<code>type</code>	determines type of plot. By default we use a line plot, <code>type="l"</code> . An alternative plot style which produces nice figures is for example <code>type="h"</code>
<code>col</code>	the colour for the series. In the univariate case, use just a colour name. The default is <code>col="steelblue"</code> . In the multivariate case we recommend selecting the colours from a colour palette, e.g. <code>col=heat.colors(ncol(x))</code>
<code>title</code>	a logical flag, by default TRUE. Should a default title be added to the plot?
<code>grid</code>	a logical flag. Should a grid be added to the plot? By default TRUE
<code>box</code>	a logical flag. Should a box be added to the plot? By default TRUE
<code>rug</code>	a logical flag. By default TRUE. Should a rug representation of the data added to the plot?

`seriesPlot()` displays the financial time series as given by its input. In most cases this may be either a price or index series when the prices or index values are given as input, or a return series when the values are given as financial returns. If the input values represent returns and we want to plot their cumulated values over time, we use the function `cumulatedPlot()`, and, in the opposite case, if we have a cumulated series and want to display the returns, we use the function `returnPlot()`.

Let us consider some examples. The example data file *SWX* contains in its columns the index values for the *Swiss Bond Index*, for the *Swiss Performance Index*, and for the *Swiss ImmoFunds Index*, *SII*. In the following code snippet the first line loads the example data file and converts it into a time series object, the second line extracts the *SPI* column, and the last line computes logarithmic returns from the index.

```
> SPI <- SWX[, "SPI"]
> SPI.RET <- SWX.RET[, "SPI"]
```

To create default plots we just call the functions `seriesPlot()`, `returnPlot()` and `cumulatedPlot()`

```
> seriesPlot(SPI)
> returnPlot(SPI)
> cumulatedPlot(SPI.RET)
```

The three graphs for the Swiss Performance Index are shown in [Figure 5.4](#).




FIGURE 5.4: Plots of the SPI index and the returns: The three graphs show the index, the logarithmic returns, and the cumulated returns indexed to 100. The plot options used are the default options.

The functions `seriesPlot()`, `returnPlot()` and `cumulatedPlot()` also allow for multivariate plots on one or more sheets. To create a two-column plot for the three SWX indices and the three LPP benchmarks on one sheet we proceed as follows:

```
> par(mfcol = c(3, 2))
> seriesPlot(SWX)
```

The indices for the SBI, SPI, SII, as well as for the three Pension Funds indices LPP25, LPP40, and LPP60 are shown in [Figure 5.5](#).

Notice that the arguments of the three plot functions

```
> args(seriesPlot)
function (x, labels = TRUE, type = "l", col = "steelblue", title = TRUE,
  grid = TRUE, box = TRUE, rug = TRUE, ...)
NULL
> args(returnPlot)
```

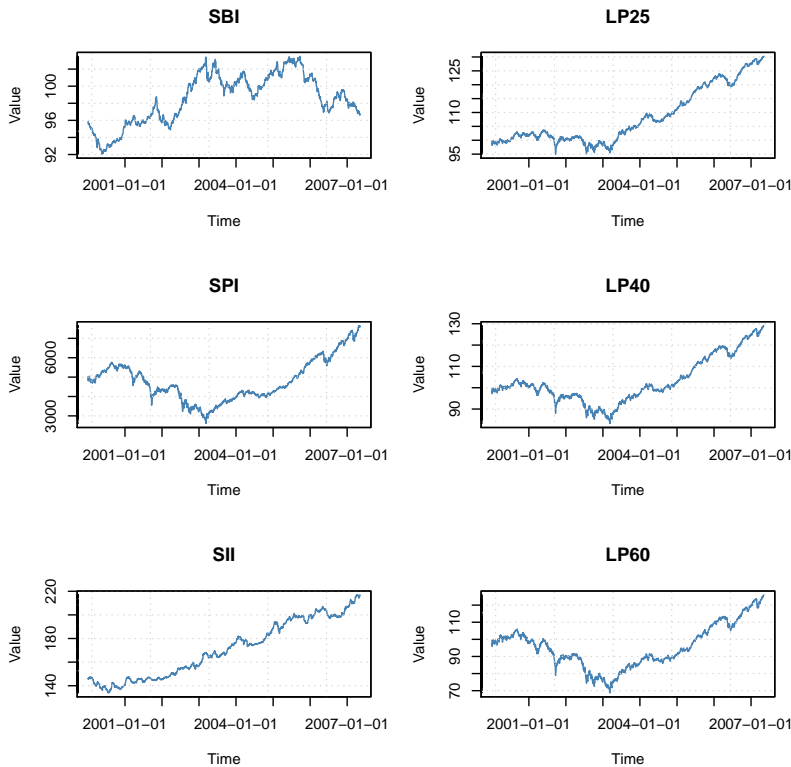


FIGURE 5.5: Plots of major Swiss indices and pension fund benchmark: The six graphs show to the left three SWX indices, the SBI, SPI and SII, as well as to the right the three Pictet Benchmark indices LLP25, LPP40 and LPP60 from Pictet's LPP2000 series.

```
function (x, labels = TRUE, type = "l", col = "steelblue", title = TRUE,
  grid = TRUE, box = TRUE, rug = TRUE, ...)
NULL

> args(cumulatedPlot)

function (x, index = 100, labels = TRUE, type = "l", col = "steelblue",
  title = TRUE, grid = TRUE, box = TRUE, rug = TRUE, ...)
NULL
```

allow you to adapt the plots according to your own requirements.
The following example shows a tailored graph:

```
> par(mfrow = c(1, 1))
> seriesPlot(SPI, labels = FALSE, type = "h", col = "brown",
  title = FALSE, grid = FALSE, rug = FALSE)
> lines(SPI, col = "orange")
> title(main = "Swiss Performance Index")
> hgrid()
> box_()
```




FIGURE 5.6: Tailored graphs for the SPI: The upper plot shows a tailored graph for the SPI, the middle plot shows the drawdowns of the SPI, and the lower plot a smoothed series of the SPI returns.

```
> copyright()
> mtext("SPI", side = 3, line = -2, adj = 1.02, font = 2)
```

In the `seriesPlot()` we suppress the labels, the title, the grid, and the rug, and change the type of the plot to histogram-like vertical lines. Finally, for `col` we choose a brown colour. Then we add an orange line on top of the plot. Then the main title is added calling the function `title()`. Horizontal grid lines are created by calling the Rmetrics function `hgrid()`, and a bottom lined box is created by calling the Rmetrics function `box_()`. Rmetrics also provides functions to easily add vertical grid lines, `vgrid()`, and L-shaped box frames, `boxL()`. The Rmetrics copyright is added by the function `copyright()`.

Derived Series Plots

For the future we plan to add several plots for derived series. Currently one such plot is available in Rmetrics for displaying the drawdowns of a financial time series. The function `drawdownsPlot()` takes as input a financial time series of returns and plots the drawdowns. For price or index series, we first have to compute the returns.

```
> drawdownPlot(returns(SPI, method = "discrete"))
```

User Generated Series

You can also plot any other derived series, or you can add your own plot functions using the function `seriesPlot()`. For example, let us smooth the SPI series using the `lowess()` function. This function performs the smoothing using locally-weighted polynomial regression.

```
> x <- SPI
> series(x) <- lowess(x = 1:nrow(SPI), y = as.vector(SPI),
  f = 0.1)$y
> seriesPlot(x, rug = FALSE, col = "red", ylim = c(2500, 8000),
  lwd = 2)
> lines(SPI)
```

Here is a recipe of how to create a user-generated plot function for the lowess smoother function:

LISTING 5.3: EXAMPLE PLOT FUNCTION FOR LOWESS SMOOTHER

```
lowessSeriesPlot <-
function(x, labels = TRUE, type = "l", col = "steelblue", title =
  TRUE, grid = TRUE, box = TRUE, rug = TRUE, add = TRUE, ...)
{
  stopifnot(isUnivariate(x))
  series(x) <- lowess(x = 1:nrow(x), y = as.vector(x), ...)$y
  seriesPlot(x = x, labels = labels, type = type, col = col,
    title = title, grid = grid, box = box, rug = rug, ...)
  invisible(x)
}
```

And now, let us run it:

```
> lowessSeriesPlot(SPI, rug = FALSE, col = "red", f = 0.1)
```

Bear in mind that you can easily generalize your tailored series plot. For instance, you can include the multivariate case (inspect the function `seriesPlot()`) and, optionally, add the unsmoothed series. Another possible use case is a volatility plot.




FIGURE 5.7: Plots with tailored axis labelling for the SPI: The upper plot shows the axis labelling using default settings. The plot in the middle shows "month-year" formatted axis labels. The lower plot labels the ticks from the beginning of the year.

How to tailor axis labelling

The graphs for the series and related plot functions use by default ISO8601 date/time formatted labels for the x-axis labelling. This can be modified by specifying the arguments `format` and `at` in the generic `plot()` function.

```
> plot(SPI, xlab = "", col = "steelblue")
> plot(SPI, format = "%b-%Y", xlab = "", col = "steelblue")
> plot(SPI, format = "%b-%Y", at = paste(2000:2007, 1, 1, sep = "-"),
  xlab = "", col = "steelblue")
```

How to display data from different time zones

The Rmetrics generic plot function can also display `timeSeries` objects recorded in different time zones in the same plot. For details we refer to the ebook *Chronological Objects with R/Rmetrics*.

5.2 BOX PLOTS

Box plots are an excellent tool for conveying location and variation information in data sets, particularly for detecting and illustrating location and variation changes between different groups of data (Chambers, Cleveland, Kleiner & Tukey, 1983).

The R base package `graphics` provides the `boxplot()` function, which takes as input a numeric vector. `Rmetrics` has added the functions `boxPlot()` and `boxPercentilePlot()` for `timeSeries` objects of financial returns. These allow two different views on distributional data summaries. Both functions are built on top of R's `boxplot()` function.

LISTING 5.4: BOX AND BOX PERCENTILE PLOT FUNCTIONS

Function:	
<code>boxPlot</code>	creates a side-by-side standard box plot
<code>boxPercentilePlot</code>	creates a side-by-side box-percentile plot
Arguments:	
<code>x</code>	a 'timeSeries' object
<code>col</code>	colours specified by a colour palette

How to display a box plot

Tukey (1977) introduced box plots as an efficient method for displaying a five-number data summary. The graph summarizes the following statistical measures: The median, upper and lower quartiles, and minimum and maximum data values. The box plot is interpreted as follows: The box itself contains the middle 50% of the data. The upper edge (hinge) of the box indicates the 75th percentile of the data set, and the lower hinge indicates the 25th percentile. The range of the middle two quartiles is known as the inter-quartile range. The line in the box indicates the median value of the data. If the median line within the box is not equidistant from the hinges, then the data is skewed. The ends of the vertical lines, the so called whiskers, indicate the minimum and maximum data values, unless outliers are present, in which case the whiskers extend to a maximum of 1.5 times the inter-quartile range. The points outside the ends of the whiskers are outliers or suspected outliers.

```
> args(boxPlot)

function (x, col = "steelblue", title = TRUE, ...)
NULL
```

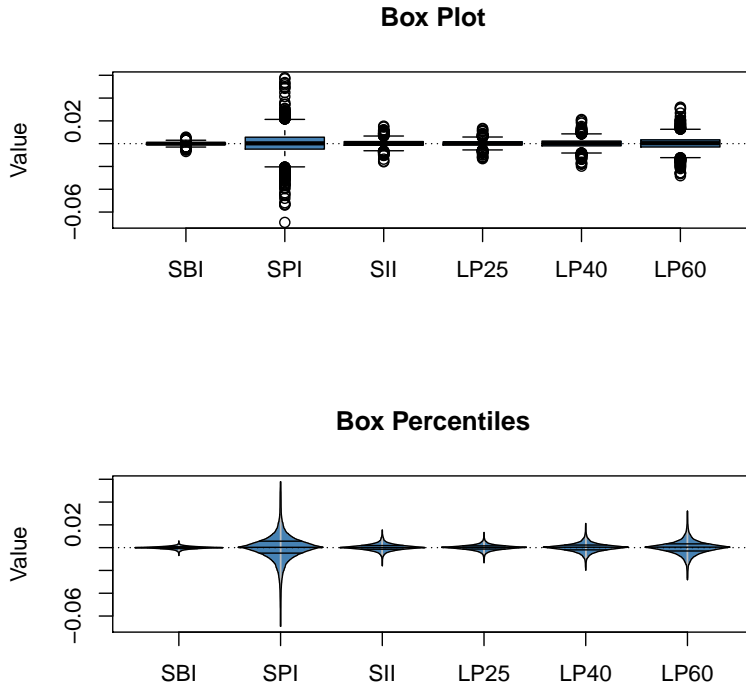


FIGURE 5.8: Box and box percentile plots of Swiss pension fund assets: The upper graph shows a box plot and the lower graph a box percentile plot. The presented data are the three Swiss assets classes SPI, SBI, SII, and Pictet's pension fund benchmark indices from the LPP2000 benchmark series.

The dot argument `...` allows us to pass optional parameters to the underlying `boxplot()` function from the `graphics` package¹.

```
> args(boxplot)
function (x, ...)
NULL

> boxPlot(returns(SWX))
```

How to display a box percentile plot

Unlike the box plot, which uses width only to emphasize the middle 50% of the data, the box-percentile plot uses width to encode information

¹`boxPlot()` is provided by `fBasics`, while `boxplot()` is from the `graphics` package

about the distribution of the data over the entire range of data values. Box-percentile plots convey the same graphical information as box plots. In addition, they also contain information about the shape of the distributions.

```
> args(boxPercentilePlot)
function (x, col = "steelblue", title = TRUE, ...)
NULL

> boxPercentilePlot(returns(SWX))
```

5.3 HISTOGRAM AND DENSITY PLOTS

To display a histogram or density plot for a univariate `timeSeries` object we can use R's base functions `hist()` and `density()`. In addition to these plots, Rmetrics offers three tailored plots, `histPlot()`, `densityPlot()` and `logDensityPlot()`, which allow different views on density functions².

LISTING 5.5: HISTOGRAM AND DENSITY PLOT FUNCTIONS. NOTE THAT THE INTERNAL RMETRICS FUNCTION `.hist()` ALLOWS FOR DEVELOPERS TO CREATE HISTOGRAMS WITH CONTROLLABLE FIXED BIN SIZES.

Function:	
<code>histPlot</code>	returns a tailored histogram plot
<code>densityPlot</code>	returns a kernel density estimate plot
<code>logDensityPlot</code>	returns a log kernel density estimate plot
<code>.hist</code>	creates histograms with a fixed bin size

Arguments:	
<code>x</code>	a 'timeSeries' object

How to display a histogram plot

The histogram is presumably the most pervasive of all graphical plots of financial returns. A histogram can be viewed as a graphical summary of distributional properties. On the other hand, we can consider it as a non-parametric estimator of a density function. The histogram is constructed by grouping the (return) data into equidistant bins or intervals and plotting the relative frequencies (or probabilities) falling in each interval. The `histPlot()` function plots a tailored histogram. By default, the probability is shown on the y-axis. Furthermore, the mean is added as an orange vertical line. For a comparison with a normal distribution with the same

²`help()` and `density()` are from R's base package, `fooPlot()` and the internal utility function `.help()` are from Rmetrics.

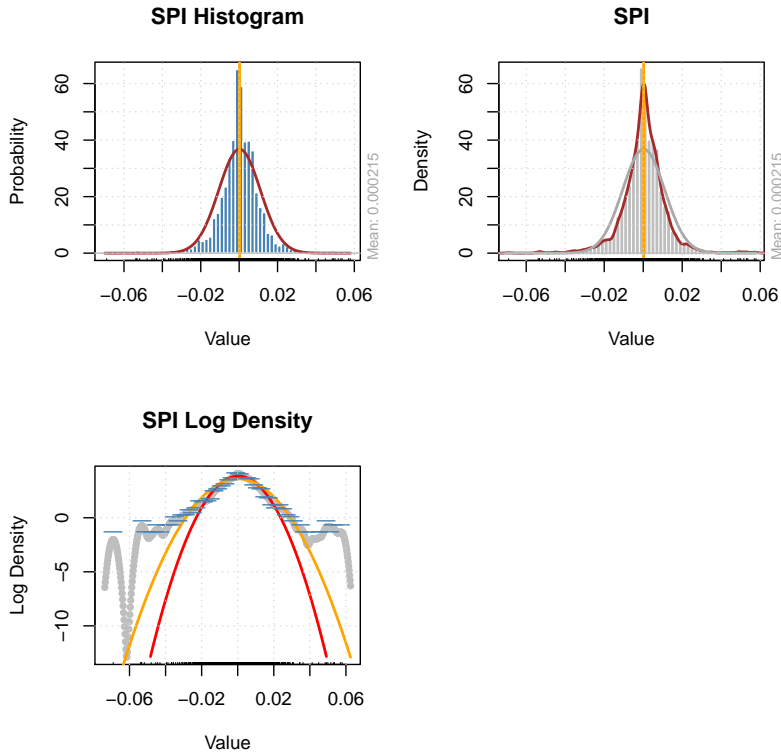


FIGURE 5.9: Histogram and density plots of Swiss pension fund assets: Upper Left: Histogram plot of the log returns of the Swiss Performance Index, SPI. The blue bins display the probability for the returns, the orange line the mean value, and the brown curve a normal density estimate with the same mean and variance as the empirical returns. Upper right: Kernel density estimate. Lower Left: Log Density Plot with a sample estimator and a robust estimate for the normal density fit.

mean and variance as the empirical data, a brown normal density line is added. The rugs on the x-line provide further helpful information about the density in the tails.

```
> histPlot(SPI.RET)
```

How to display a density plot and kernel density estimates

The function `densityPlot()` computes a kernel density estimate by calling the `density()` function, and then displays it graphically. The algorithm used disperses the mass of the empirical distribution function over a regular grid of at least 512 points and then uses the fast Fourier transform to convolve this approximation with a discretized version of the kernel. It

then uses linear approximation to evaluate the density at the specified points.

```
> args(densityPlot)
function (x, labels = TRUE, col = "steelblue", fit = TRUE, hist = TRUE,
         title = TRUE, grid = TRUE, rug = TRUE, skip = FALSE, ...)
NULL
```

The default plot adds a histogram, `hist=TRUE`, and overlays the density with a fitted normal distribution function, `fit=TRUE`.

```
> densityPlot(SPI.RET)
```

Optional dot arguments are passed to the `density()` function. This allows us to adapt the bandwidth, or to select an alternative smoothing kernel (the default kernel is Gaussian). For details we refer to the help page of the `density()` function.

How to display a log-density plot

The function `logDensityPlot()` creates a further view of the distributional properties of financial returns.

```
> args(logDensityPlot)
function (x, labels = TRUE, col = "steelblue", robust = TRUE,
         title = TRUE, grid = TRUE, rug = TRUE, skip = FALSE, ...)
NULL
```

The function displays the distribution on a logarithmic scale. Thus, in the case of normally distributed returns, we expect a parabolic shape, and heavy tails will be displayed as straight lines or are even bended upwards. The graph displays ...

```
> logDensityPlot(SPI.RET)
```

5.4 QUANTILE-QUANTILE PLOTS

The quantile-quantile plot, or qq-plot, is a graphical technique for determining if two data sets come from populations with a common distribution. A qq-plot is a plot of the quantiles of the first data set against the quantiles of the second data set. A 45-degree reference line is also plotted. If the two sets come from a population with the same distribution, the points should fall approximately along this reference line. The greater the departure from this reference line, the greater the evidence for the conclusion that the two data sets come from populations with different distributions.

To display a quantile-quantile plot for `timeSeries` objects we can use R's base functions `qqnorm()`, `qqline()`, and `qqplot()`. `qqnorm()` is a generic

function, the default method of which produces a normal quantile-quantile plot. `qqline()` adds a line to a normal quantile-quantile plot which passes through the first and third quartiles. `qqplot()` produces a quantile-quantile plot of two data sets. In addition to these plots, Rmetrics offers three tailored plots to display distributional properties of financial returns fitted by a normal, a normal inverse Gaussian, and a generalized hyperbolic Student's *t* distribution. These distributions are heavily used in modelling financial returns³.

LISTING 5.6: QUANTILE-QUANTILE PLOT FUNCTIONS

Function:
`qqnormPlot` returns a normal quantile-quantile plot
`qqnigPlot` returns a NIG quantile-quantile plot
`qqghtPlot` returns a GHT quantile-quantile plot

Arguments:
`x` a 'timeSeries' object

A qq-plot helps to answer the following questions:

Do two data sets come from populations with a common distribution?

Do two data sets have common location and scale?

Do two data sets have similar distributional shapes?

Do two data sets have similar tail behaviour?

How to display a normal quantile-quantile plot

The normal quantile-quantile plot

```
> args(qqnormPlot)
function(x, labels = TRUE, col = "steelblue", pch = 19, title = TRUE,
        mtext = TRUE, grid = FALSE, rug = TRUE, scale = TRUE, ...)
NULL
```

displays the empirical data points versus the quantiles of a normal distribution function. By default, the empirical data are scaled by their mean and standard deviation. If a non-scaled view is desired we have to set the argument `scale=FALSE`. If the empirical data points are drawn from a normal distribution then we expect them to all lie on the diagonal line added to the plot. In addition, the plot shows the 95% confidence intervals.

³The first three functions are from R's base package, the `fooPlot()` functions are from Rmetrics.


```
> set.seed(1953)
> x <- rnorm(250)
> qqnormPlot(x)
```

How to display a NIG quantile-quantile plot

In the case of the normal inverse Gaussian distribution, NIG, the empirical data are fitted by a log-likelihood approach. The `qqnigPlot()` function returns an invisible list with the plot coordinates `$x` and `$y`.

```
> y <- rnig(250)
> qqnigPlot(y)
```

How to display a GHT quantile plot

In the case of the generalized hyperbolic Student's t distribution, GHT, the empirical data are fitted by a log-likelihood approach to the generalized hyperbolic Student's t distribution function. The function returns an invisible list with the plot coordinates `$x` and `$y`.

```
> z <- rght(250)
> qqghtPlot(z)
```

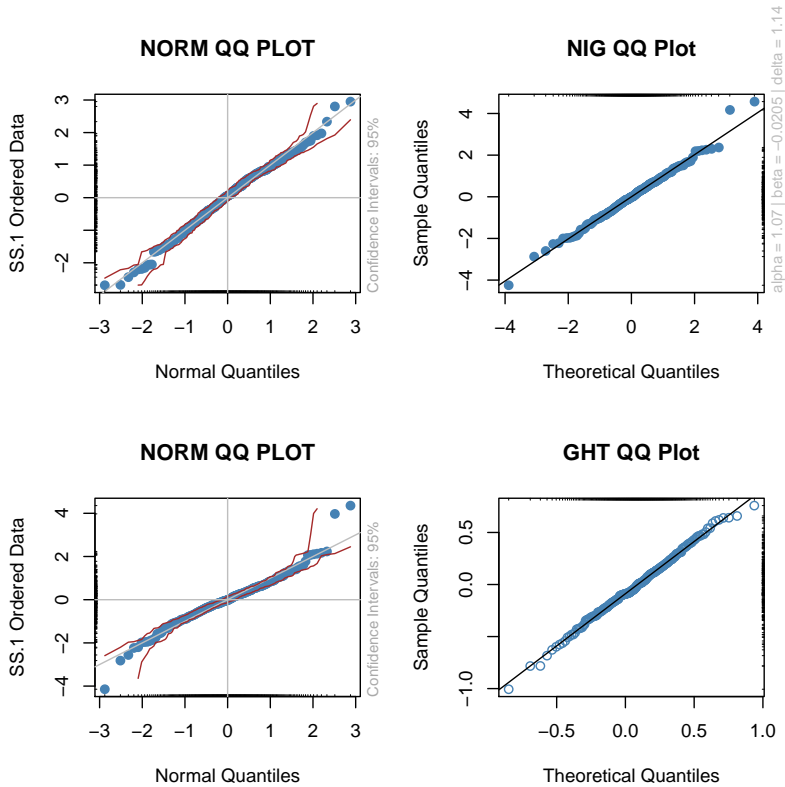


FIGURE 5.10: Quantile-Quantile Plots of simulated returns from a normal (top left), a normal inverse Gaussian (NIG) (top right and bottom left), and a generalized hyperbolic Student's t (GHT) (bottom right) distribution.

CHAPTER 6

CUSTOMIZATION OF PLOTS

```
> library(fPortfolio)
```

Rmetrics comes with several kinds of customized plots to display financial time series and their statistical properties. These plots can be adapted in many ways. The layout of the plot labels, including titles, labels and additional text information, can be modified by changing the content, the types of the fonts and the size of characters. Plot elements, such as lines and symbols, can be modified by changing their style, size and colours. In the following we give a brief overview of how to customize plot labels, and how to select colours, fonts and plot symbols.

6.1 PLOT LABELS

Most of the Rmetrics tailored plots, such as `seriesPlot()`, have common arguments to customize their layout.

```
> args(seriesPlot)
function (x, labels = TRUE, type = "l", col = "steelblue", title = TRUE,
  grid = TRUE, box = TRUE, rug = TRUE, ...)
NULL
```

The main arguments for customization a plot are summarized in the following function listing.

LISTING 6.1: MAIN ARGUMENTS FOR PLOT, POINTS AND LINES FUNCTIONS

Function:	
plot	generic plot function
points	adds points to a plot
lines	adds connected line segments to a plot
abline	adds straight lines through a plot
Arguments:	
type	determines the type of plot
col	colour or colour palette for lines or symbols
title	should a default title be added?
grid	should a grid be added to the plot?
box	should a box be added to the plot?
rug	should rugs be added?
...	optional arguments to be passed

For details we refer to the help functions for the `plot()` and `par()` functions. In the following we present some examples of how to customize a univariate time series plot.

How to create a series plot with default labels

The first graph shows a time series plot for the Swiss performance index created with default settings

```
> SPI <- SWX[, "SPI"]
> seriesPlot(SPI)
```

How to create a series plot with user-specified labels

The second graph shows the same plot but now with user-specified labels. Setting the argument `title=FALSE`

```
> seriesPlot(SPI, title = FALSE)
> title(main = "Swiss Performance Index", xlab = "", ylab = "SPI Index")
> text(as.POSIXct("2006-11-25"), rev(SPI)[1], as.character(rev(SPI)[1]),
      font = 2)
> mtext("Source: SWX", side = 4, col = "grey", adj = 0, cex = 0.7)
```

displays an untitled plot. Thus we can use the R base function `title()` to add a main title, subtitle, as well as x and y labels. Further text attributes can be added using R's base functions `text()` and `mtext()`.

For details please consult the help functions.

LISTING 6.2: TITLE, TEXT AND MARGIN TEXT FUNCTIONS

Function:	
title	adds a title, a subtitle, and axis labels
text	adds text string(s) to the plot
mtext	adds margin text string(s) to the plot

How to create a series plot with decorations

The third graph shows how to decorate the plot. By setting `rug = FALSE`, we remove the rug. Next, we add a horizontal grid and replace the framed box with an L-shaped box. Finally, we add a copyright string as margin text, and add an orange horizontal line on index level 5000.

```
> seriesPlot(SPI, grid = FALSE, box = FALSE, rug = FALSE)
> hgrid()
> boxL()
> copyright()
> abline(h = 5000, col = "orange")
```

For details of decoration functions we refer to the help pages of the following functions:

LISTING 6.3: PLOT DECORATION FUNCTIONS

Function:	
decor	simple decoration function
hgrid	creates horizontal grid lines
vgrid	creates vertical grid lines
boxL	creates a L-shaped box
box_	creates a bottom line box
copyright	adds Rmetrics copyright to a plot

How to create a series plot with optional dot arguments

The fourth graph demonstrates how to use optional plot parameters through the dot `...` arguments. Here we have modified the plotting point symbol and changed the orientation of the axis label style. Further arguments are shown in [Listing 6.4](#). For a complete list of all plot parameters we refer to `help(par)`.

```
> seriesPlot(SPI, grid = FALSE, rug = FALSE, type = "o", pch = 19,
  las = 1)
```

6.2 MORE ABOUT PLOT FUNCTION ARGUMENTS

Here are some of the arguments you might want to specify for plots:




FIGURE 6.1: Customized time series plots: The first graph shows the default plot, the second a user entitled and labelled plot, the third a user decorated plot, and the fourth plot modifications by adding optional parameters through the dot argument.

LISTING 6.4: SELECTED ARGUMENTS FOR PLOT FUNCTIONS

Function:	
plot	generic plot function
Arguments:	
type	what type of plot should be created?
axes	draw or suppress to plot the axes
ann	draw or suppress to add title and axis labels
pch	select the type of plotting symbol
cex	select the size of plotting symbol and text
xlab, ylab	names of the labels for the x and y axes
main	the (main) title of the plot
xlim, ylim	the range of the x and y axes
log	names of the axes which are to be logarithmic
col, bg	select colour of lines, symbols, background
lty, lwd	select line type, line width
las	select orientation of the text of axis labels

Notice that some of the relevant parameters are documented in `help(plot)` or `plot.default()`, but many only in `help(par)`. The function `par()` is for setting or querying the values of graphical parameters in traditional R graphics.

How to modify the plot type

Settings for the plot type can be modified using the following identifiers:

LISTING 6.5: TYPE ARGUMENT SPECIFICATIONS FOR PLOT FUNCTIONS

Function:	
<code>plot</code>	generic plot function
Argument:	
<code>type</code>	specifies the type of plot
<code>"p"</code>	point plot (default)
<code>"l"</code>	line plot
<code>"b"</code>	both points and lines
<code>"o"</code>	overplotted points and lines
<code>"h"</code>	histogram like
<code>"s"</code>	steps
<code>"n"</code>	no plotting

Note that by default, the `type` argument is set to `"p"`. If you want to draw the axes first and add points, lines and other graphical elements later, you should use `type="n"`.

How to select a font

With the `font` argument, an integer in the range from 1 to 5, we can select the type of fonts:

LISTING 6.6: FONT ARGUMENTS FOR PLOT FUNCTIONS

```

Function:
plot                generic plot function

Arguments:
font                integer specifying which font to use for text
font.axis           font number to be used for axis annotation
font.lab            font number to be used for x and y labels
font.main           font number to be used for plot main titles
font.sub            font number to be used for plot sub-titles

```

If possible, device drivers arrange so that 1 corresponds to plain text (the default), 2 to bold face, 3 to italic and 4 to bold italic. Also, font 5 is expected to be the symbol font, in Adobe symbol encoding.

How to modify the size of fonts

With the argument `cex`, a numeric value which represents a multiplier, we can modify the size of fonts

LISTING 6.7: CEX ARGUMENTS FOR PLOT FUNCTIONS

```

Function:
plot                generic plot function

Arguments:
cex                magnification of fonts/symbols relative to default
cex.axis           magnification for axis annotation relative to cex
cex.lab            magnification for x and y labels relative to cex
cex.main           magnification for main titles relative to cex
cex.sub            magnification for sub-titles relative to cex

```

How to orient axis labels

The argument `las`, an integer value ranging from 0 to 3, allows us to determine the orientation of the axis labels

LISTING 6.8: CEX PARAMETERS FOR PLOT FUNCTIONS

```

Function:
plot                generic plot function

Arguments:
las                orientation
0                  always parallel to the axis [default]
1                  always horizontal
2                  always perpendicular to the axis
3                  always vertical

```

Note that other string/character rotation (via argument `srt` to `par`) does not affect the axis labels.

How to select the line type

The argument `lty` sets the line type. Line types can either be specified as an integer, or as one of the character strings "blank", "solid", "dashed", "dotted", "dotdash", "longdash", or "twodash", where "blank" uses invisible lines, i.e. does not draw them.

LISTING 6.9: LTY ARGUMENT FOR PLOT FUNCTIONS

Function:	
<code>plot</code>	generic plot function
Arguments:	
<code>lty</code>	sets line type to
0	blank
1	solid (default)
2	dashed
3	dotted
4	dotdash
5	longdash
6	twodash

6.3 SELECTING COLOURS

Rmetrics provides tools and utilities to select individual colours by code numbers and sets of colours from colour palettes.

How to print the colour coding numbers

The function `colorTable()` displays a table of R's base colours together with their code numbers.

```
> colorTable()
```

Note that the colours are repeated cyclically.

How to use the colour name locator

The function `colorLocator()` displays R's 657 named colours for selection and optionally returns R's colour names. The idea and implementation of the colour locator originates in the contributed package `epitools` written and maintained by [Aragon \(2008\)](#).

Usually this function is used interactively, setting the argument to `TRUE`. Use the left mouse button to locate one or more colours and the stop the

Table of Color Codes




FIGURE 6.2: Colour table of R's base colours: The colours are shown together with their code numbers. Note that number 1 is white (invisible on white background), number 2 is black, and the next colours are red, green, blue, cyan, magenta, yellow, grey, then the cycle repeats with number 9 being black again.

locator, using the method appropriate to the screen device you are using. Here is what you get on the console:

```
> colorLocator(TRUE)

  x  y  colour.names
1 15 15  lightslateblue
2 18 17   orange
3 27 22  yellowgreen
```

To return all colour names in alphabetical order you can call the function `colorMatrix()`. The first 20 are:

```
> head(sort(colorMatrix()), 20)

[1] "aliceblue"      "antiquewhite"   "antiquewhite1"  "antiquewhite2"
[5] "antiquewhite3"  "antiquewhite4"  "aquamarine"     "aquamarine1"
[9] "aquamarine2"    "aquamarine3"    "aquamarine4"    "azure"
```




FIGURE 6.3: Colour locator with R's 657 named colours: The display shows the colour matrix. If you use the colour locator interactively, clicking on a colour square will return the name of the colour in the console window.

```
[13] "azure1"      "azure2"      "azure3"      "azure4"
[17] "beige"       "bisque"      "bisque1"     "bisque2"
```

The associated colour locator display is shown in [Figure 6.3](#).

Which colour palettes are available

To display a multivariate plot we would often want to use a personal set of colours. To this end, Rmetrics provides dozens of pre-implemented colour palettes taken from several contributed R packages. The functions and their arguments have been modified, so that we have a common usage for all palettes. Here is a list of the palettes provided with Rmetrics:

LISTING 6.10: COLOUR PALETTE FUNCTIONS

Function:	
rainbowPalette	Contiguous rainbow colour palette
heatPalette	Contiguous heat colour palette
terrainPalette	Contiguous terrain colour palette
topoPalette	Contiguous topo colour palette
cmPalette	Contiguous cm colour palette
grayPalette	R's gamma-corrected gray palette
timPalette	Tim's MATLAB-like colour palette
rampPalette	Colour ramp palettes
seqPalette	Sequential colour brewer palettes
divPalette	Diverging colour brewer palettes
qualiPalette	Qualified colour brewer palettes
focusPalette	Red, green and blue focus palettes
monoPalette	Red, green and blue mono palettes

All Rmetrics' colour sets are named as `fooPalette`, where the prefix `foo` denotes the name of the underlying colour set.

R's Contiguous Colour Palettes

Palettes for n contiguous colours are implemented in the `grDevices` package. To conform with Rmetrics' naming convention for colour palettes, we have built a wrapper around the underlying functions. These are the `rainbowPalette()`, `heatPalette()`, `terrainPalette()`, `topoPalette()`, and the `cmPalette()`. Conceptually, all of these functions actually use (parts of) a line cut out of the 3-dimensional colour space, parametrized by the function `hsv(h, s, v, gamma)`, where `gamma=1` for the `fooPalette()` function, and hence, equispaced hues in RGB space tend to cluster at the red, green and blue primaries. Some applications, such as contouring, require a palette of colours which do not wrap around to give a final colour close to the starting one. If you want to pass additional arguments to the underlying functions, please consult `help(rainbow)`. With `rainbow`, the parameters `start` and `end` can be used to specify particular subranges of hues. Synonymous function calls are `rainbow()`, `heat.colors()`, `terrain.colors()`, `topo.colors()`, and `cm.colors()`.

```
> pie(rep(1, 12), col = rainbowPalette(12), xlab = "rainbowPalette")
> pie(rep(1, 12), col = heatPalette(12), xlab = "heatPalette")
> pie(rep(1, 12), col = terrainPalette(12), xlab = "terrainPalette")
> pie(rep(1, 12), col = topoPalette(12), xlab = "topoPalette")
> pie(rep(1, 12), col = cmPalette(12), xlab = "cmPalette")
```

R's Gamma-Corrected Gray Palette

The function `grayPalette()` chooses a series of n gamma-corrected grey levels. The range of the grey levels can be optionally monitored through

the ... arguments, for details we refer to `help(gray.colors)`, which is a synonymous function call used in the `grDevices` package.

```
> pie(rep(1, 12), col = greyPalette(12), xlab = "greyPalette")
```

Tim's MATLAB-Like Colour Palette

The function `timPalette()` creates a colour set ranging from blue to red, and passes through the colours cyan, yellow, and orange. It is an implementation of the MATLAB jet colour palette, originally used in fluid dynamics simulations. The function here is a copy from R's contributed package `fields`, doing a spline interpolation on $n=64$ colour points.

```
> pie(rep(1, 12), col = timPalette(12), xlab = "timPalette")
```

Colour Ramp Palettes

The function `rampPalette()` creates several colour ramps. The function is implemented in the contributed R package `colorRamps` (Keitt, 2007). The following colour ramp palettes are supported through the `name` argument: "blue2red", "green2red", "blue2green", "purple2green", "blue2yellow", and "cyan2magenta".

```
> palettes <- c("blue2red", "green2red", "blue2green", "purple2green",
  "blue2yellow", "cyan2magenta")
> for (palette in palettes) pie(rep(1, 12), col = rampPalette(12,
  palette), xlab = paste("rampPalette -", palette))
```

Colour Brewer Palettes

The functions `seqPalette()`, `divPalette()`, and `qualiPalette()` create colour sets according to R's contributed `RColorBrewer` package. The first letter in the function name denotes the type of the colour set: "s" for sequential palettes, "d" for diverging palettes, and "q" for qualitative palettes.

Sequential palettes are suited to ordered data that progress from low to high. Lightness steps dominate the look of these schemes, with light colours for low data values to dark colours for high data values. The sequential palettes names are: Blues, BuGn, BuPu, GnBu, Greens, Greys, Oranges, OrRd, PuBu, PuBuGn, PuRd, Purples, RdPu, Reds, YlGn, YlGnBu, YlOrBr, YlOrRd.

```
> palettes <- c("Blues", "BuGn", "BuPu", "GnBu", "Greens",
  "Greys", "Oranges", "OrRd", "PuBu", "PuBuGn", "PuRd",
  "Purples", "RdPu", "Reds", "YlGn", "YlGnBu", "YlOrBr",
  "YlOrRd")
```

```
> for (palette in palettes) pie(rep(1, 12), col = seqPalette(12,
  palette), xlab = paste("seq -", palette))
```

Diverging palettes put equal emphasis on mid-range critical values and extremes at both ends of the data range. The critical class or break in the middle of the legend is emphasized with light colours and low and high extremes are emphasized with dark colours that have contrasting hues. The diverging palettes names are: "BrBG", "PiYG", "PRGn", "PuOr", "RdBu", "RdGy", "RdYlBu", "RdYlGn", "Spectral".

```
> palettes <- c("BrBG", "PiYG", "PRGn", "PuOr", "RdBu", "RdGy",
  "RdYlBu", "RdYlGn", "Spectral")
> for (palette in palettes) pie(rep(1, 12), col = divPalette(12,
  palette), names = NULL, xlab = paste("div -", palette))
```

Qualitative palettes do not imply magnitude differences between legend classes, and hues are used to create the primary visual differences between classes. Qualitative schemes are best suited to representing nominal or categorical data. The qualitative palettes names are: "Accent", "Dark2", "Paired", "Pastel1", "Pastel2", "Set1", "Set2", "Set3".

```
> palettes <- c("Accent", "Dark2", "Paired", "Pastel1", "Pastel2",
  "Set1", "Set2", "Set3")
> for (palette in palettes) pie(rep(1, 12), col = qualiPalette(12,
  palette), xlab = paste("quali -", palette))
```

In contrast to the original colour brewer palettes, the palettes here are created by spline interpolation from the colour variation with the most different values, i.e. for the sequential palettes these are 9 values, for the diverging palettes these are 11 values, and for the qualitative palettes these are between 8 and 12 values, depending on the colour set.

The brewer colour palettes are originally from the contributed R package RColorBrewer, written by [Neuwirth \(2007\)](#).

Graph Colour Palettes

The functions `focusPalette()` and `monoPalette()` create colour sets inspired by R's contributed package `PerformanceAnalytics` ([Carl & Peterson, 2008](#)). These colour palettes have been designed to create readable, comparable line and bar graphs with specific objectives.

Focused Colour Palettes: Colour sets designed to provide focus on the data graphed as the first element. This palette is best used when there is clearly an important data set for the viewer to focus on, with the remaining data being secondary, tertiary, etc. Later elements graphed in diminishing values of grey. The focus palette names are: "redfocus", "greenfocus", "bluefocus".

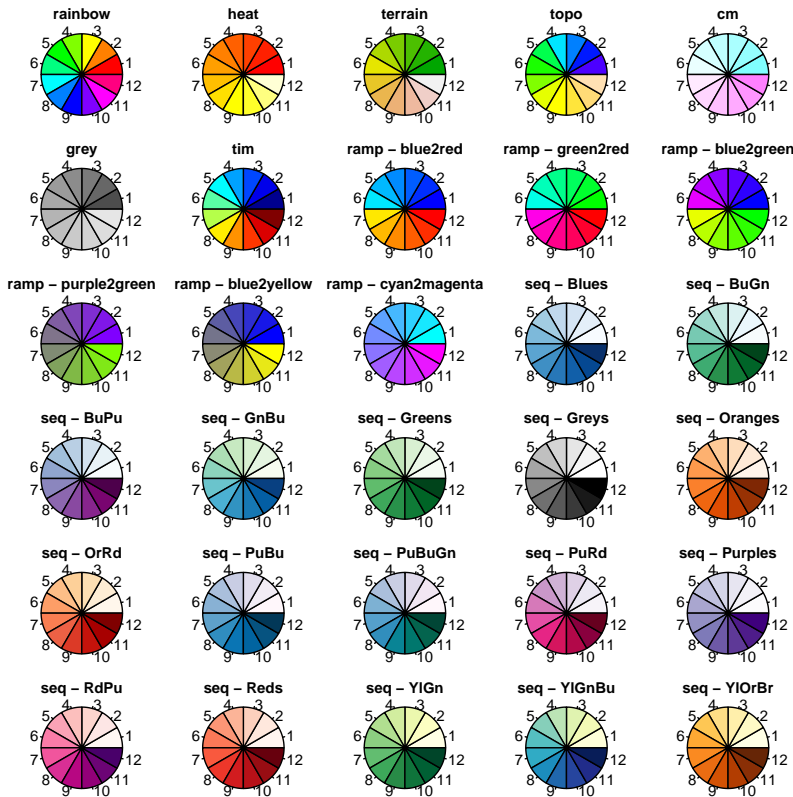


FIGURE 6.4: Selected colour palettes: The colour palettes provided by Rmetrics include R's base palettes, a grey palette, the MATLAB-like colour palette, colour ramp palettes, sequential colour brewer palettes, diverging colour brewer palettes, qualified colour brewer palettes, red/green/blue focus palettes, and red/green/blue mono palettes.

Monochrome Colour Palettes: These include colour sets for monochrome colour displays. The mono palette names are: "redmono", "greenmono", "bluemono".

Inspect the functions for the colour palettes and feel free to add your own palettes. For the developer we would like to mention the following undocumented functions:

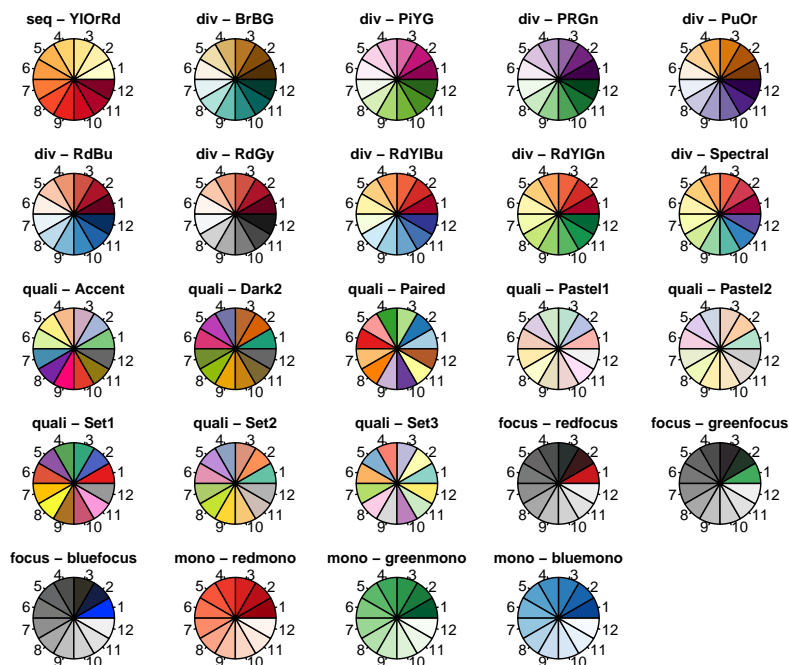


FIGURE 6.5: Selected colour palettes, continued: The remaining colour palettes from the previous figure.

LISTING 6.11: UNDOCUMENTED COLOUR FUNCTIONS

Function:	
<code>.asRGB</code>	converts any R colour to RGB (red/green/blue)
<code>.chcode</code>	changes from one to another number system
<code>.hex.to.dec</code>	converts heximal numbers do decimal numbers
<code>.dec.to.hex</code>	converts decimal numbers do heximal numbers

6.4 SELECTING CHARACTER FONTS

The function `characterTable()` displays the character for a given font. The font is specified by an integer number ranging from 1 to 5. This integer specifies which font to use for text. If possible, device drivers arrange the fonts in the following sequence:

	0	1	2	3	4	5	6	7
4		!	Ų	#	Ɔ	%	&	3
5	()	*	+	,	-	.	/
6	0	1	2	3	4	5	6	7
7	8	9	:	;	<	=	>	?
10	≡	Λ	B	X	Δ	E	Φ	Γ
11	H	I	θ	K	Λ	M	N	O
12	Π	Θ	P	Σ	T	Y	ς	Ω
13	Ξ	Ψ	Z	[∴]	⊥	-
14	—	α	β	χ	δ	ε	φ	γ
15	η	ι	φ	κ	λ	μ	ν	ο
16	π	θ	ρ	σ	τ	υ	ϖ	ω
17	ξ	ψ	ζ	{		}	~	
20								
21								
22								
23								
24	€	Υ	ˆ	≤	/	∞	f	♣
25	♦	♥	♠	↔	←	↑	→	↓
26	◊	±	″	≥	×	∞	∂	•
27	+	≠	≡	=	…		—	┐
30	×	§	℥	∅	⊗	⊕	∅	∩
31	∪	∩	⊇	∝	∞	∞	∞	∞
32	∠	∇	⊗	⊗	™	∏	√	∞
33	∟	^	∇	↔	↔	∏	⇒	∞
34	∅	(⊗	⊗	™	Σ	(∞
35	\							
36)						
37	J							

FIGURE 6.6: Character font tables: This table shows the characters for font number 5.

LISTING 6.12: FUNCTION TO DISPLAY CHARACTERS FOR A GIVEN FONT

```
Function:
characterTable      displays a table of characters

Arguments:
font                specifies font number
  1                plain text (the default)
  2                bold face
  3                italic
  4                bold italic
  5                symbol font in Adobe symbol encoding
```

To display a specific font in a graphics display we can use the command

```
> characterTable(font = 5)
```

Table of Plot Characters

□	0	▽	25	2	50	K	75	d	100	}	125	-	150	-	175	Ê	200	á	225	û	250
○	1		26	3	51	L	76	e	101	~	126	—	151	•	176	É	201	â	226	ü	251
△	2		27	4	52	M	77	f	102	•	127	-	152	±	177	Ê	202	ã	227	Û	252
+	3		28	5	53	N	78	g	103	€	128	™	153	²	178	Ê	203	ä	228	ý	253
×	4		29	6	54	O	79	h	104	•	129	§	154	³	179	î	204	å	229	þ	254
◇	5		30	7	55	P	80	i	105	•	130	›	155	•	180	í	205	æ	230	ÿ	255
▽	6		31	8	56	Q	81	j	106	f	131	œ	156	µ	181	ï	206	ç			
✖	7		32	9	57	R	82	k	107	•	132	•	157	¶	182	î	207	è	232		
*	8	!	33	:	58	S	83	l	108	•	133	ž	158	•	183	ð	208	é	233		
◆	9	"	34	;	59	T	84	m	109	†	134	ÿ	159	•	184	ñ	209	ê	234		
⊙	10	#	35	<	60	U	85	n	110	‡	135		160	•	185	ò	210	ë	235		
✖	11	\$	36	=	61	V	86	o	111	-	136	ı	161	•	186	ó	211	ì	236		
⊠	12	%	37	>	62	W	87	p	112	‰	137	¢	162	•	187	ô	212	í	237		
✖	13	&	38	?	63	X	88	q	113	§	138	£	163	¼	188	õ	213	î	238		
✖	14	•	39	@	64	Y	89	r	114	•	139	¤	164	½	189	ö	214	ï	239		
■	15	(40	A	65	Z	90	s	115	œ	140	¥	165	¾	190	×	215	ð	240		
•	16)	41	B	66	[91	t	116	•	141	ı	166	¿	191	ø	216	ñ	241		
▲	17	*	42	C	67	\	92	u	117	ž	142	§	167	À	192	Ú	217	ò	242		
•	18	+	43	D	68]	93	v	118	•	143	-	168	Á	193	Û	218	ó	243		
•	19	•	44	E	69	^	94	w	119	•	144	⊙	169	Â	194	Ü	219	ô	244		
•	20	-	45	F	70	-	95	x	120	•	145	a	170	Ã	195	Ý	220	õ	245		
•	21	•	46	G	71	•	96	y	121	•	146	«	171	Ä	196	Ÿ	221	ö	246		
□	22	/	47	H	72	a	97	z	122	•	147	-	172	Å	197	Þ	222	÷	247		
◇	23	0	48	ı	73	b	98	{	123	•	148	-	173	Æ	198	ß	223	ø	248		
△	24	1	49	J	74	c	99		124	•	149	⊙	174	Ç	199	à	224	ù	249		

FIGURE 6.7: Table of plot symbols: Displayed are the plot symbols for the current font.

6.5 SELECTING PLOT SYMBOLS

Plot symbols are set within the `plot()` function by setting the `pch` parameter, equal to an integer between 0 and usually 25. Since it is hard to remember what symbol each integer represents, [Figure 6.7](#) may serve as a reminder. The function `symbolTable()` displays the plot symbol for a given code.

```
> # The following example use latin1 characters: these may not
> # appear correctly (or be omitted entirely).
> symbolTable()
```

CHAPTER 7

MODELLING ASSET RETURNS

```
> library(fPortfolio)
```

In many cases we want to generate artificial data sets of assets which have the same statistical properties as a given set of empirical returns. In the simple case of the multivariate normal and Student's t as well as their skewed versions Rmetrics provides functions to fit the parameters for this family of elliptical distributions and to generate new random data sets from these parameters. To find out if a set of empirical financial asset returns is multivariate normally distributed we can perform an hypothesis test.¹

7.1 TESTING ASSET RETURNS FOR NORMALITY

The function `assetsTest()` is a suite of (currently two) functions to test whether or not a set of asset returns is (multivariate) normally distributed. The implemented tests are the *multivariate Shapiro test* (Royston, 1982) from the R package `mvnrmtest` contributed by Jarek (2009), and the *non-parametric E-statistics test*, also called energy test (Szekely, 1989; Rizzo, 2002; Szekely & Rizzo, 2005; Szekely, Rizzo & Bakirov, 2007) from the contributed R package `energy` (Rizzo & Szekely, 2008)².

LISTING 7.1: FUNCTIONS TO TEST A MULTIVARIATE DATA SET OF RETURNS FOR NORMALITY, TO FIT THE MODEL PARAMETERS, AND TO SIMULATE ARTIFICIAL DATA SETS WITH THE SAME STATISTICAL PROPERTIES AS THE EMPIRICAL DATA SET

¹An alternative way to model multivariate assets sets and their dependency structure uses copulae. Rmetrics functions for testing, fitting, and simulating copulae are described in the ebook *Managing Risk with R/Rmetrics*.

²These are implemented as built-in functions in Rmetrics

Function:	
assetsTest	tests for multivariate normal assets
assetsFit	estimates the parameters of a set of assets
assetsSim	simulates artificial data sets of assets

```
> args(assetsTest)
function (x, method = c("shapiro", "energy"), Replicates = 99)
NULL
```

The function `assetsTest()` requires a multivariate `timeSeries` object `x` as input and performs the test specified by the `method` argument, by default the multivariate Shapiro test. An object of S4 class `fHTEST` is returned.

Let us now investigate whether the Swiss bond returns, Swiss equities, and Swiss Reits in the LPP2005 data are normally distributed or not, and let us compare the results of the two methods, `shapiro` and `energy`.

How to perform a multivariate Shapiro test

To perform a multivariate Shapiro test, we call the function `assetsTest()` with `method="shapiro"`. This is the default setting, the specification of the argument `method` is therefore optional.

```
> shapiroTest <- assetsTest(LPP2005.RET[, 1:3], method = "shapiro")
> print(shapiroTest)

Shapiro-Wilk normality test

data:  Z
W = 0.9521, p-value = 1.018e-09
```

The function returns an object of class `fHTEST`, with the following slots

```
> #slotNames(shapiroTest)
> names(shapiroTest)

[1] "statistic" "p.value"  "method"   "data.name"
```

The slot named `@test` of the result returned by the Shapiro test returns a list with all entries from the original test as implemented in the R package `mvnrmtest`. The printout from the Shapiro test tells us that the hypothesis of a multivariate normal return distribution is rejected.

How to perform a multivariate E-Statistics test

Alternatively, we can perform a multivariate E-Statistics Test. To do this, we have to set the argument `method="energy"` explicitly.

```
> assetsTest(LPP2005.RET[, 1:3], method = "energy")
```

```
Energy test of multivariate normality: estimated parameters
```

```
data: x, sample size 377, dimension 3, replicates 99
E-statistic = 3.0382, p-value < 2.2e-16
```

Again, the slot named `@test` of the result returned by the E-statistics test gives a list with all entries from the original test as implemented in the R package *energy* (Rizzo & Szekely, 2008). The energy test also rejects the hypothesis that the returns are multivariate normally distributed.

7.2 FITTING ASSET RETURNS

Rmetrics also provides functions to fit a data set of asset returns to the most common multivariate distribution functions for financial returns, the normal and the Student's *t* distributions.

```
> args(assetsFit)
function (x, method = c("st", "sn", "sc"), title = NULL, description = NULL,
        fixed.df = NA, ...)
NULL
```

The function `assetsFit()` expects a multivariate `timeSeries` object of asset returns `x` as input and estimates the parameters of the specified distribution by the argument `method`. The choice can be a multivariate skew normal distribution, `method="sn"`, a multivariate skew-Student's *t* distribution, `"st"`, or a multivariate skew-Cauchy distribution, `"sc"`.

How to fit a normal or Student's t distribution

The most common multivariate distribution functions for financial returns include the *normal distribution* and the *Student's t distribution* together with their skewed versions. The `method` argument determines which distribution should be fitted to the asset returns, by default the skewed Student's *t* distribution. The following example shows how to fit a skew Student's *t* distribution to the set of Swiss asset returns including the SPI, SBI, and SII.

```
> fit <- assetsFit(LPP2005.RET[, 1:3], method = "st")
> print(fit)
Title:
  Student-t Parameter Estimation

Call:
  FUN(x = x, trace = FALSE)

Model:
  Skew Student-t Distribution

Estimated Parameter(s):
```

```

$beta
      SBI      SPI      SII
[1,] 9.143e-05 0.002414 -0.00020946

$Omega
      SBI      SPI      SII
SBI 1.2362e-06 -8.2959e-07 1.2688e-07
SPI -8.2959e-07 3.9630e-05 1.6155e-06
SII 1.2688e-07 1.6155e-06 6.1579e-06

$alpha
      SBI      SPI      SII
-0.14448 -0.32575 0.25228

$nu
[1] 6.5043

```

Description:

Tue Jan 27 13:37:15 2015 by user: Rmetrics

The multivariate skew-normal distribution is implemented as discussed by [Azzalini & Dalla Valle \(1996\)](#); the (Omega, alpha) parametrization is the one used by [Azzalini & Capitanio \(1999\)](#).

The family of multivariate skew-t distributions is an extension of the multivariate Student's t family, via the introduction of a shape parameter which regulates skewness. In the symmetric case the skew-t distribution reduces to the regular symmetric t-distribution. When the number of degrees becomes infinity the distribution reduces to the multivariate skew-normal one ([Azzalini & Capitanio, 2003](#)).

The fits are done using maximum likelihood estimation ([Azzalini & Capitanio, 1999, 2003](#)). The function `assetsFit` returns an object of S4 class `fDISTFIT`. The `@fit$d` slot provides the estimated parameters.

How to fit distributions using the copula approach

Rmetrics also offers functions to analyze, to model and to fit distribution functions using the copula approach. The relevant functions are available in the Rmetrics package `fCopulae` ([Würtz, 2009b](#)). For details we refer to the ebook *Managing Risk with R/Rmetrics*.

7.3 SIMULATING ASSET RETURNS FROM A GIVEN DISTRIBUTION

From the fitted parameters we can generate assets sets with the same distributional properties as the empirical data. The simulation of artificial data sets is performed by the function `assetsSim()`.

```
> args(assetsSim)
```

```
function (n, method = c("st", "sn", "sc"), model = list(beta = rep(0,
  2), Omega = diag(2), alpha = rep(0, 2), nu = 4), assetNames = NULL)
NULL
```

The first argument, `n`, sets the number of records to be simulated, the second, `method`, sets the method to be used and the third, `model`, sets the list of parameters. Alternatively, we can use as input the fitted model from the parameter estimation as returned by the function `assetsFit()`.

```
> #slotNames(fit)
> model <- fit@fit$dp
> SIM <- LPP2005.RET[,1:3]
> X <- assetsSim(n=nrow(SIM), method="st", model=model)
```

This simulation has generated random records of asset returns with the same distributional properties as the fitted empirical SPI equity, SBI bond, and SII real estate indices. The object returned `X` is a matrix, which can be easily transformed into a `timeSeries` object with the same instrument names (column names) and date stamps (row names), using the function `series()`.

```
> series(SIM) <- as.matrix(X)
> head(SIM)
GMT
      SBI      SPI      SII
2005-11-01 -6.1061e-04 -0.00129129 0.00182322
2005-11-02 6.9127e-05 0.00094257 -0.00033684
2005-11-03 4.3658e-04 0.00260074 0.00259501
2005-11-04 7.0820e-05 0.00069746 0.00200235
2005-11-07 -1.9384e-03 -0.00933346 -0.00542237
2005-11-08 2.8888e-03 0.00153843 -0.00159319
```


CHAPTER 8

SELECTING SIMILAR OR DISSIMILAR ASSETS

```
> library(fPortfolio)
```

In many cases we want to select in a data pre-processing step the most dissimilar assets in a large data set of assets to reduce the number of assets in portfolio design. This can be done using statistical approaches which sort out the assets in groups with similar behaviour and similar properties. To those approaches belong cluster algorithms, like hierarchical clustering or k-means clustering, which group similar assets together and separate dissimilar ones (Kaufman & Rousseeuw, 1990). Another popular approach uses eigenvalue analysis. In addition we show how to add additional cluster methods, such as a robust k-means approach, to the cluster approaches already supported by Rmetrics.

8.1 FUNCTIONS FOR GROUPING SIMILAR ASSETS

The Rmetrics function for selecting similar or dissimilar assets from a data set is the function `assetsSelect()`.

```
> args(assetsSelect)
function (x, method = c("hclust", "kmeans"), control = NULL,
  ...)
NULL
```

LISTING 8.1: FUNCTIONS TO SELECT SIMILAR OR DISSIMILAR ASSETS FROM A DATA SET BY CLUSTERING METHODS

Function:	
<code>assetsSelect</code>	selects assets from a cluster approach
<code>assetsCorEigenPlot</code>	performs eigenvalue analysis of assets

<code>assetsArrange</code>	rearranges columns in a data set of assets
----------------------------	--

The `method` argument of the `assetsSelect()` function allows you to select the type of grouping, either according to *hierarchical clustering* `method="hclust"` or according to *k-means clustering* `method="kmeans"`.

8.2 GROUPING ASSET RETURNS BY HIERARCHICAL CLUSTERING

Setting `method="hclust"` performs a hierarchical cluster analysis on a set of dissimilarities and methods for analyzing it.

LISTING 8.2: FUNCTIONS TO SELECT SIMILAR OR DISSIMILAR ASSETS FROM A DATA SET BY HIERARCHICAL CLUSTERING. THE ENTRY `method` IN THE `control` LIST DETERMINES THE METHOD USED FOR THE HIERARCHICAL CLUSTERING, AND THE ENTRY `measure` DETERMINE THE MEASURE FOR THE DISTANCE MATRIX

Function:	
<code>assetsSelect</code>	for hierarchical clustering of dissimilarities
Arguments:	
<code>x</code>	a 'timeSeries' object
<code>method</code>	"hclust"
<code>control</code>	list of optional cluster controls,
	<code>method</code> - the name of the clustering method
	"ward", "single", "complete", "average",
	"mcquitty", "median", "centroid",
	<code>measure</code> - the name of the distance measure
	"euclidean", "maximum", "manhattan",
	"canberra", "binary", "minkowski"
<code>...</code>	optional arguments

To perform a hierarchical clustering on asset returns, we call the function `assetsSelect()` with the argument `method="hclust"`. The underlying function is the R functions `hclust()` from the `stats` package. Arguments to this function can be passed to the `assetsSelect()` function through the `control` argument and the dots `...` argument.

Internally, the function `hclust()` performs a clustering using a set of dissimilarities for the n objects being clustered. Initially, each object is assigned to its own cluster and then the algorithm proceeds iteratively, at each stage joining the two most similar clusters, continuing until there is just a single cluster. At each stage distances between clusters are recomputed by the Lance-Williams dissimilarity update formula according to the particular clustering method being used.

Let us start to group the assets and the benchmark series from the data set `LPP2005.RET`. We use the default settings, the *complete linkage method* with an *euclidean distance measure*.

```

> lppData <- LPP2005.RET
> hclustComplete <- assetsSelect(lppData, method = "hclust")
> hclustComplete

Call:
hclust(d = dist(t(x), method = measure), method = method)

Cluster method   : complete
Distance         : euclidean
Number of objects: 9

> plot(hclustComplete, xlab = "LPP2005 Assets")
> mtext("Distance Metric: Euclidean", side = 3)

```

A number of different clustering methods can be provided through the dots argument. Ward's minimum variance method aims at finding compact, spherical clusters. The complete linkage method finds similar clusters. The single linkage method (which is closely related to the minimal spanning tree) adopts a friends-of-friends clustering strategy. The other methods can be regarded as aiming for clusters with characteristics somewhere between the single and complete link methods. Note however, that methods "median" and "centroid" do not lead to a monotone distance measure, or equivalently the resulting dendrograms can have so called inversions (which are hard to interpret). For details we refer to the help page of the function `hclust()`.

The next example shows how to set up an alternative hierarchical clustering with the *euclidean distance measure*, but now with Ward's method of *minimum variance agglomeration*

```

> hclustWard <- assetsSelect(lppData, method = "hclust", control = c(measure = "euclidean",
  method = "ward"))
> hclustWard

Call:
hclust(d = dist(t(x), method = measure), method = method)

Cluster method   : ward.D
Distance         : euclidean
Number of objects: 9

> plot(hclustWard)
> mtext("Distance Metric: Euclidean", side = 3)

```

The two plots have created dendrograms which are shown in [Figure 8.1](#) and [Figure 8.2](#).

Further information can be extracted from the results returned by the `hclust` function. This function returns a list of S3 class `hclust` with the following values:

LISTING 8.3: RETURNED VALUES FROM HIERARCHICAL CLUSTERING

Function:

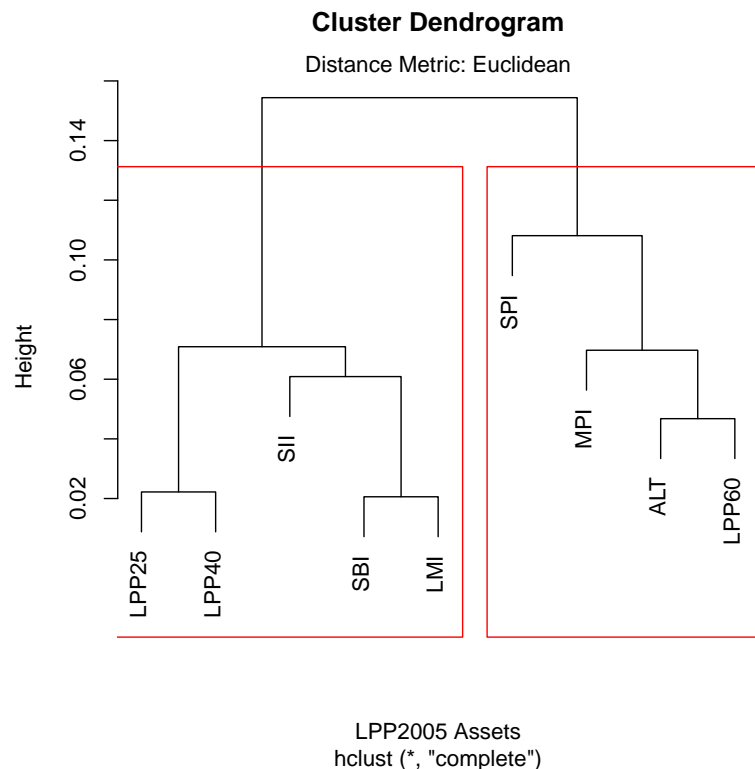


FIGURE 8.1: Hierarchical clustering of Swiss pension fund assets: Dendrogram plot (as obtained from default settings) for the Swiss pension fund assets set SBI, SBI, SII, LMI, MPI, and ALT, including the three benchmarks LPP25, LPP40, and LPP60.

hclust	hierarchical clusters of dissimilarities
Values:	
merge	merging process of clusters
height	the clustering height
order	permutation of the original observations
labels	labels for the objects being clustered
call	the call which produced the result
method	the cluster method that has been used
dist.method	the distance that has been used

8.3 GROUPING ASSET RETURNS BY K-MEANS CLUSTERING

If we set `method="kmeans"` then the function `assetsSelect()` performs a *k-means clustering* on the financial time series, i.e. the time series of




FIGURE 8.2: Hierarchical clustering of Swiss pension fund assets: Dendrogram plot as obtained using an euclidean distance measure and Ward's method for clustering for the Swiss pension fund assets set SBI, SBI, SII, LMI, MPI, and ALT, including the three benchmarks LPP25, LPP40, and LPP60.

financial returns.

LISTING 8.4: FUNCTIONS TO SELECT SIMILAR OR DISSIMILAR ASSETS BY K-MEANS CLUSTERING. THE ENTRY `center` IN THE `control` LIST SETS THE NUMBER OF CLUSTERS, AND THE ENTRY `algorithm` SELECTS THE NAME OF THE CLUSTERING ALGORITHM TO BE USED

```
Function:
assetsSelect      for k-means cluster analysis

Arguments:
x                a 'timeSeries' object
method          "kmeans"
control         list of cluster controls:
                  center - the number of clusters
                  algorithm - name of the algorithm
                        "Hartigan-Wong", "Lloyd",
                        "Forgy", "MacQueen"
```

The transposed data $t(x)$ given by the `@data` slot of the time series x is clustered by the k-means method, which aims to partition the points into k groups such that the sum of squares from points to the assigned cluster centres is minimized. At the minimum, all cluster centres are at the mean of their Voronoi sets, i.e. the set of data points which are nearest to the cluster centre.

The algorithm of [Hartigan & Wong \(1979\)](#) is used by default. Note that some authors use k-means to refer to a specific algorithm rather than the general method: most commonly the algorithm given by [MacQueen \(1967\)](#) but sometimes that given by [Lloyd \(1982\)](#) and [Forgy \(1965\)](#). The Hartigan-Wong ([Hartigan & Wong, 1979](#)) algorithm generally does a better job than either of those, but trying several random starts is often recommended. Except for the Lloyd-Forgy method, k clusters will always be returned if a number is specified. If an initial matrix of centres is supplied, it is possible that no point will be closest to one or more centres, which is currently an error in the Hartigan-Wong method.

The number of centres and the name of the desired algorithm can be passed by the `control` argument,

```
> control <- c(centers = 2, algorithm = "H")
```

The list of algorithms includes "Hartigan-Wong", "Lloyd", "Forgy", and "MacQueen". Note that the names can be abbreviated in the `control` argument. The default settings for the number of centres is three, and the default algorithm is the algorithm of Hartigan and Wong. Let us consider the case with two groups

```
> kmeans <- assetsSelect(lppData, method = "kmeans", control <- c(centers = 2,
  algorithm = "Hartigan-Wong"))
> sort(kmeans$cluster)
```

SBI	SII	LMI	LPP25	LPP40	SPI	MPI	ALT	LPP60
1	1	1	1	1	2	2	2	2

The result shows us that the assets are clustered in two groups, one with lower risk and the other with higher risky assets. In group 1 (lower risk) we find the assets SBI, SII, LMI as well as the L25 and LPP40 benchmarks), in group 2 (higher risk) we find the assets SPI, MPI, ALT, and the benchmark LPP60, a quite natural grouping which we would have expected.

Further information can be extracted from the results returned from `kmeans` clustering. The function returns a list of S3 class `kmeans` with the following values:

LISTING 8.5: RETURNED VALUES FROM K-MEANS CLUSTERING

Function:

`kmeans` performs k-means cluster analysis

Values:	
cluster	integer vector indicating the cluster to which each point is allocated
centers	matrix of cluster centres
withinss	within-cluster sum of squares for each cluster
size	number of points in each cluster

8.4 GROUPING ASSET RETURNS THROUGH EIGENVALUE ANALYSIS

A third approach groups the individual assets according to an eigenvalue analysis. This can be done calling the function `assetsCorEigenPlot()`.

```
> args(assetsCorEigenPlot)
function (x, labels = TRUE, title = TRUE, box = TRUE, method = c("pearson",
"kendall", "spearman"), ...)
NULL
```

The function takes a multivariate `timeSeries` object `x` as input and performs according to the specified method for the computation of the correlation matrix an eigenvalue analysis.

The function calculates the first two eigenvectors of the correlation matrix and plot their components against the `x` and `y` directions. This results in a grouping of the assets. Three methods are available to compute the correlation matrix, "person", "kendall", and "spearman". If method is "kendall" or "spearman", Kendall's tau or Spearman's rho statistic is used to estimate a rank-based measure of association. These are more robust and have been recommended if the data do not necessarily come from a bivariate normal distribution. If method is "pearson", the usual correlation coefficient will be returned.

```
> assetsCorEigenPlot(lppData, method = "kendall")
```

8.5 GROUPING ASSET RETURNS BY CONTRIBUTED CLUSTER ALGORITHMS

There are several other contributed R packages on the CRAN server which provide alternative cluster algorithms or alternative implementations. One of the most prominent packages is the R package `cluster` contributed by [Maechler, Rousseeuw, Struyf & Hubert \(2009\)](#).

LISTING 8.6: FURTHER FUNCTIONS FOR CLUSTERING FROM R'S CONTRIBUTED `cluster` PACKAGE

Function:	
	Hierarchical Clustering:
diana	divisive hierarchical clustering
mona	divisive hierarchical clustering with binary
variables	

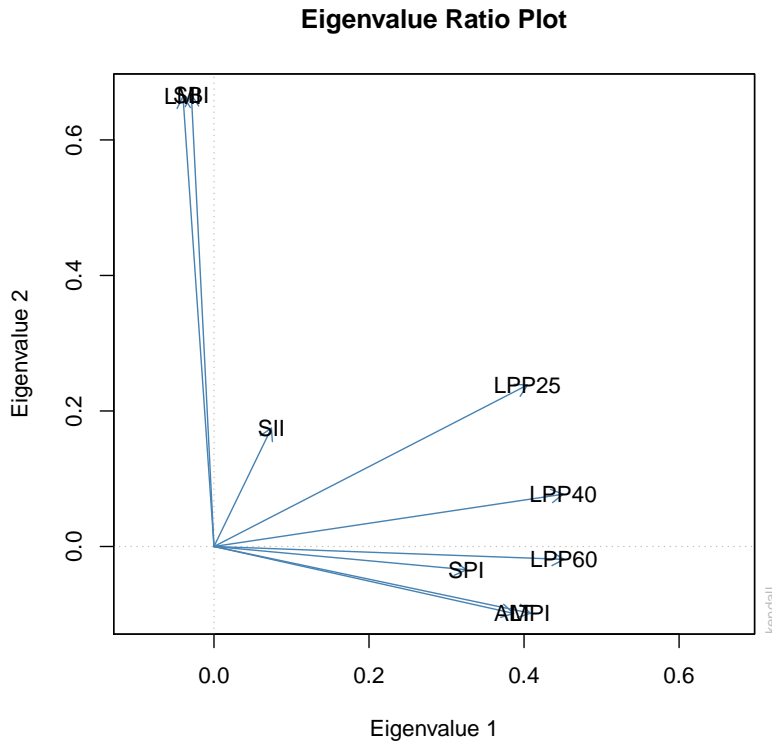


FIGURE 8.3: Grouping assets by eigenvalue analysis: The closeness of the arrows is a measure for similarities between individual assets.


agnes	agglomerative hierarchical clustering
pam	Partitioning Methods: partitioning into clusters around medoids, a more robust version of k-means
clara	partitioning method for much larger data sets
fanny	fuzzy clustering of the data into k clusters
daisy	Pairwise Dissimilarities: pairwise dissimilarities between observations

For the grouping of financial returns, one can use these functions with the transposed data matrix directly, `t(series(x))`, or one can add additional methods to the `assetsSelect()` function. This can be done in the following way, for example using the more robust k-means algorithm implemented in `pam` from the `cluster` package:

```
> library(cluster)
> .pamSelect <- function(x, control = NULL, ...) {
  if (is.null(control))
```



```
clusplot(pam(x = x <- as.matrix(x), k = k, metric = metric))
```



These two components explain 80.12 % of the point variability.

FIGURE 8.4: Grouping assets by partitioning around medoids: The graph shows a two-dimensional representation of the observations, in which the clusters are indicated by ellipses.

```
control = c(k = 2, metric = "euclidean")
k <- as.integer(control[1])
metric <- control[2]
pam(x <- as.matrix(x), k = k, metric = metric, ...)
}
```

Now apply Rmetrics' `assetSelect()` function:

```
> pam <- assetSelect(LPP2005.RET, method = "pam", control <- c(k = 2,
  metric = "euclidean"))
> plot(pam, which.plots = 1)
```

Note that the plot of a cluster partition consists of a two-dimensional representation of the observations, in which the clusters are indicated by ellipses.

8.6 ORDERING DATA SETS OF ASSETS

Pairwise correlations in financial data sets give important information on the pairwise dependencies of the individual asset returns. Graphical displays help us to visualize the correlations. This can be happen in many different ways depending on the ordering of the columns of the multivariate time series.

There are several ways to order a set of assets column by column. The most obvious ordering of assets may be the sorting in alphabetical order. From a statistical point of view we can use more sophisticated schemes, for example ordering by a PCA analysis (the default) or by hierarchical clustering. For this we can use the function `assetsArrange()`

```
> args(assetsArrange)
function (x, method = c("pca", "hclust", "abc"), ...)
NULL
```

with the following options for rearranging the data set of assets: `method="pca"` returns PCA correlation ordered column names, `"hclust"` returns hierarchical clustered column names, and `"abc"` returns alphabetically sorted column names.

In the following investigation of the pairwise correlations we sort the LPP2005 data sets by hierarchical clustering:

```
> colnames(lppData[, 1:6])
[1] "SBI" "SPI" "SII" "LMI" "MPI" "ALT"

> Assets <- assetsArrange(lppData[, 1:6], method = "hclust")
> LPP2005HC <- 100 * lppData[, Assets]
> head(round(LPP2005HC, 5))

GMT
```

	SII	SBI	LMI	SPI	MPI	ALT
2005-11-01	-0.31909	-0.06127	-0.11089	0.84146	0.15481	-0.25730
2005-11-02	-0.41176	-0.27620	-0.11759	0.25193	0.03429	-0.11416
2005-11-03	-0.52094	-0.11531	-0.09925	1.27073	1.05030	0.50074
2005-11-04	-0.11272	-0.32358	-0.11985	-0.07028	1.16796	0.94827
2005-11-07	-0.17958	0.13110	0.03604	0.62052	0.27096	0.47240
2005-11-08	0.21034	0.05393	0.23270	0.03293	0.03468	0.08536

CHAPTER 9

COMPARING MULTIVARIATE RETURN AND RISK STATISTICS

```
> library(fPortfolio)
```

The star and segment plots introduced by [Chambers et al. \(1983\)](#) allows you to display multivariate data sets. Each star in a star plot represents a single observation. Typically, star and segment plots are generated in a multi-plot format with many stars or segments on each page and each star or segment representing one observation. Star plots are used to examine the relative values for a single data point and to locate similar or dissimilar points.

9.1 STAR AND SEGMENT PLOTS

For the investigation of financial assets star plots can be used to answer the following questions:

- Which assets are dominant for a given observation?
- Which observations are most similar, i.e., are there clusters of observations?
- Are there outliers in the data set of assets?

Star plots are helpful for small-to-moderately-sized multivariate data sets. Their primary weakness is that their effectiveness is limited to data sets with less than a few hundred points. With data sets comprising more data points, they tend to be overwhelming. Rmetrics has implemented star plots to investigate several aspects of data sets of assets.

LISTING 9.1: FUNCTIONS FOR STAR AND SEGMENT PLOTS

Function:

<code>stars</code>	Star/Segment plots of a multivariate data
<code>assetsStarsPlot</code>	Segment/star diagrams of multivariate data sets
<code>assetsBasicStatsPlot</code>	Segment plot of basic return statistics
<code>assetsMomentsPlot</code>	Segment plot of distribution moments
<code>assetsBoxStatsPlot</code>	Segment plot of box plot statistics

R's basic function to create star and segment plots is named `stars()`. It has the following arguments:

```
> args(stars)

function (x, full = TRUE, scale = TRUE, radius = TRUE, labels = dimnames(x)[[1L]],
  locations = NULL, nrow = NULL, ncol = NULL, len = 1, key.loc = NULL,
  key.labels = dimnames(x)[[2L]], key.xpd = TRUE, xlim = NULL,
  ylim = NULL, flip.labels = NULL, draw.segments = FALSE, col.segments = 1L:n.seg,
  col.stars = NA, col.lines = NA, axes = FALSE, frame.plot = axes,
  main = NULL, sub = NULL, xlab = "", ylab = "", cex = 0.8,
  lwd = 0.25, lty = par("lty"), xpd = FALSE, mar = pmin(par("mar"),
  1.1 + c(2 * axes + (xlab != ""), 2 * axes + (ylab !=
  "")), 1, 0)), add = FALSE, plot = TRUE, ...)

NULL
```

The general Rmetrics star plot, `assetsStarsPlot()` is just a synonym for the basic function `stars()`

```
> args(assetsStarsPlot)

function (x, method = c("segments", "stars"), locOffset = c(0,
  0), keyOffset = c(0, 0), ...)

NULL
```

The specific star plots `assetsBasicStatsPlot()`, `assetsBoxStatsPlot()`, and `assetsMomentsPlot()` are built on top of the function `assetsStarsPlot()`. All three functions have the same list of arguments as input.

```
> args(assetsBasicStatsPlot)

function (x, par = TRUE, oma = c(0, 0, 0, 0), mar = c(4, 4, 4,
  4), keyOffset = c(-0.65, -0.5), main = "Assets Statistics",
  title = "Assets", titlePosition = c(3, 3.65), description = "Basic Returns Statistics",
  descriptionPosition = c(3, 3.5), ...)

NULL
```

9.2 SEGMENT PLOTS OF BASIC RETURN STATISTICS

Let us consider the returns of the Swiss Pension Fund Index LPP2005.RET and let us compute the basic statistics as returned by the function `basicStats()`. The following statistics are considered: minimum and maximum values, first and third quartiles, mean and median, sum, SE mean, LCL

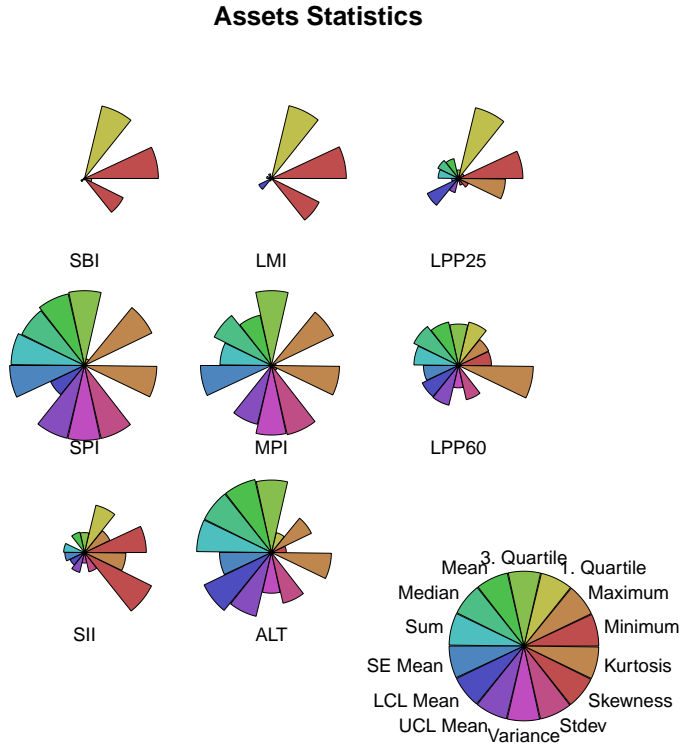


FIGURE 9.1: Segment plots based on the comparison of return statistics. 14 statistics are taken into account which represent the distributional properties of the six asset classes SBI, LMI, SII, SPI, MPI, and ALT as well as the two benchmarks LPP26, and LPP60.

mean and UCL mean, variance, standard deviation, skewness, and kurtosis. These observations are displayed as a segment plot. Which assets look similar and which look dissimilar?

```
> lppData <- LPP2005.RET
> assetsBasicStatsPlot(lppData[, -8], title = "", description = "")
```

This question is answered by [Figure 9.1](#). The SBI and LMI are similar to each other, as are the SPI, MPI, and ALT. The SII seems similar neither to the bonds nor to the equities. LPP25 can be interpreted as representing the bond assets, and the LPP60 can be understood to represent the equities and the alternative investment asset class.

Assets Statistics




FIGURE 9.2: Star plots from distributional moments including the mean, standard deviation, skewness and kurtosis. The segments for the Swiss, SBI, and foreign bonds, LMI, look similar, and also the segments for the equity investments, the SPI, MPI, and ALT. The Swiss ImmoFunds Index, SII, differs from the remaining assets. We can also say, that the LPP60 benchmark is dominated by equities.

9.3 SEGMENT PLOTS OF DISTRIBUTION MOMENTS

This segment plot displays four distributional sample estimates from the empirical asset returns including the mean, standard deviation, skewness, and kurtosis. When it is sufficient to characterize a distribution by its first four moments, then this plot allows for a simple comparison of the assets.

```
> assetsMomentsPlot(lppData[, -8], title = "", description = "")
```

Figure 9.2 shows the same results as already observed in the segment plots for the basic return statistics and the box plot statistics.

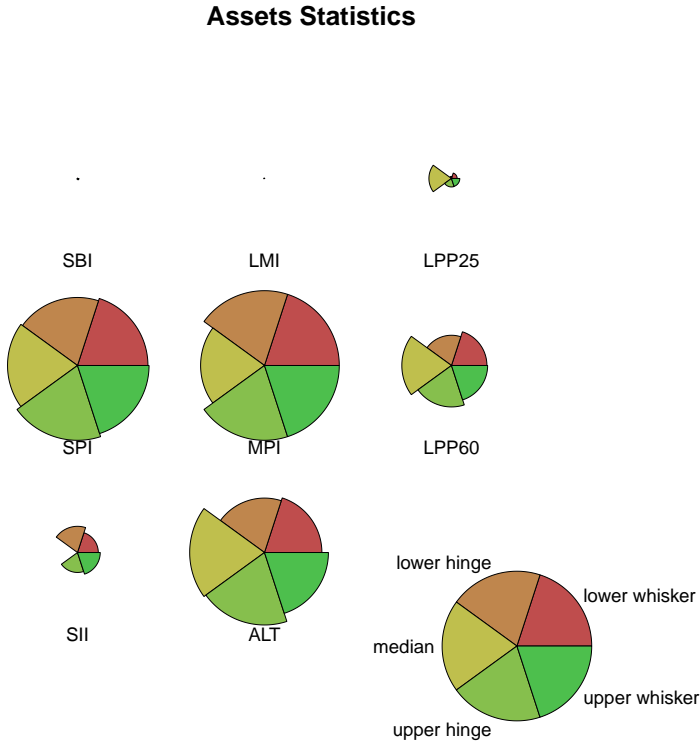


FIGURE 9.3: Star Plots from box plot statistics grouping the assets with respect to lower and upper hinge, lower and upper whisker, and the median of the assets and LPP benchmark series.

9.4 SEGMENT PLOTS OF BOX PLOT STATISTICS

These segment plot uses as observations the values returned by the function `boxPlot()`, i.e. lower and upper hinge, lower and upper whisker, and the median.

```
> assetsBoxStatsPlot(lppData[, -8], title = "", description = "")
```

This segment graph in [Figure 9.3](#) allows us to compare the assets from the view of the box plot statistics. Similarities are obvious between the Swiss and foreign bonds, and the Swiss and foreign equities together with the alternative investments. The Swiss Immo Funds Index is in between. The LPP25 benchmark is closer to the bonds, and the LPP60 benchmark is closer to the equities and alternative investments.

9.5 HOW TO POSITION STARS AND SEGMENTS IN STAR PLOTS

The default positions of stars or segments in a star plot are tailored for up to eight assets and a colour wheel legend. The positions for any other numbers of assets have to be adjusted individually by hand, and requires some process of trial and error. In most cases it is sufficient to modify the arguments `oma`, `mar`, `locOffset` and `keyOffset`. Another alternative is to use the low level function `stars()` directly.

It is also worth noting that star plots are most helpful for small to moderate sized multivariate data sets. Their primary weakness is that their effectiveness is limited to data sets with fewer than a few dozens points. For larger data sets, they tend to be overwhelming.

CHAPTER 10

PAIRWISE DEPENDENCIES OF ASSETS

```
> library(fPortfolio)
> library(fMultivar)
```

To display dependencies, similarities or correlations between two individual assets R offers the function `pairs()`. If you use the default settings, this function produces a scatter plot. Rmetrics adds customized plots to provide different views on pairs of assets. This allows us to judge on different aspects of pairwise correlations and dependencies. The views include bivariate scatterplots, correlation tests, image plots, and histogram binning, among other exploratory data analysis techniques.

10.1 SIMPLE PAIRWISE SCATTER PLOTS OF ASSETS

The function `assetsPairsPlot()` is a simple wrapper for R's base function `pairs()`. The function just transforms the data set of assets into a matrix object and plots `pairs(series(x), ...)`. Usually, the input `x` is given in form of a multivariate timeSeries object, then the `@data` slot is extracted by the function `series()`, and finally, a scatter plot of the data matrix is displayed.

```
> args(pairs.default)
function (x, labels, panel = points, ..., lower.panel = panel,
  upper.panel = panel, diag.panel = NULL, text.panel = textPanel,
  label.pos = 0.5 + has.diag/3, line.main = 3, cex.labels = NULL,
  font.labels = 1, rowlattice = TRUE, gap = 1, log = "")
NULL
```

LISTING 10.1: FUNCTIONS FOR PAIRWISE ASSETS PLOTS

Function:	
<code>pairs</code>	displays pairs of scatterplots of assets
<code>assetsPairsPlot</code>	displays pairs of scatterplots of assets
<code>assetsCorgramPlot</code>	displays correlations between assets
<code>assetsCorTestPlot</code>	displays and tests pairwise correlations
<code>assetsCorImagePlot</code>	displays an image plot of correlations
<code>squareBinning</code>	does a square binning of data points,
<code>hexBinning</code>	does a hexagonal binning of data points

How to create a simple scatter plot

In the following example we create a simple scatter plot for all pairwise asset returns using the function `assetsPairsPlot()`.

```
> args(assetsPairsPlot)
function (x, ...)
NULL
```

We will rearrange the assets as suggested by hierarchical clustering. This yields a nicer arrangement and view of the off-diagonal scatterplot panels of the graph.

```
> Assets <- assetsArrange(LPP2005.RET[, 1:6], method = "hclust")
> LPP2005HC <- 100 * LPP2005.RET[, Assets]
> assetsPairsPlot(LPP2005HC, pch = 19, cex = 0.5, col = "royalblue4")
```

We have tailored the plot layout of the graph, using small ($cex=0.5$) full dots ($pch=19$) and the colour `royalblue4`.

The optional dot arguments which are allowed to be passed, are the same as those for R's `pairs` function, see `help(pairs.default)`.

How to add a diagonal histogram panel

The possibility to modify and to define new panels makes these functions quite powerful. For example, to add histograms of the asset returns to the diagonal panels, we proceed as follows: First, we define the diagonal histogram panels,

```
> histPanel <- function(x, ...) {
  usr <- par("usr")
  on.exit(par(usr))
  par(usr = c(usr[1:2], 0, 1.5))
  h <- hist(x, plot = FALSE)
  breaks <- h$breaks
  nB <- length(breaks)
  y <- h$counts
  y <- y/max(y)
  rect(breaks[-nB], 0, breaks[-1], y, ...)
}
```

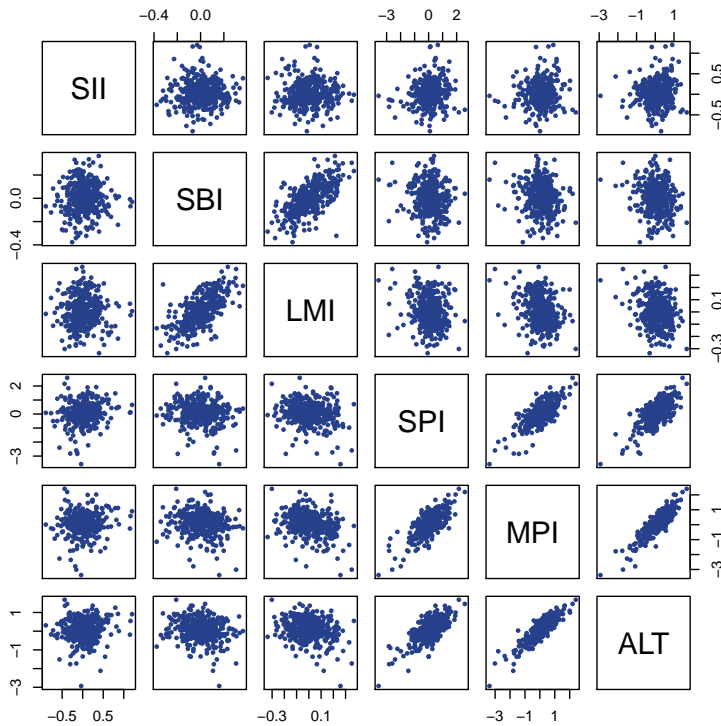


FIGURE 10.1: Pairwise scatter plots of assets from the Swiss pension fund index: The graph shows scatterplots for financial asset returns with default panels. In the default case both the lower and upper off-diagonal panels show a scatter plot, the diagonal panel is a text panel showing the names of the assets.

and then we plot the correlations together with the histograms:

```
> assetsPairsPlot(LPP2005HC, diag.panel = histPanel, pch = 19,
  cex = 0.5, col = "red4", tick = 0, col.axis = "white")
```

The result is shown in [Figure 10.2](#).

How to remove the axis labelling from a pairs plot

The two additional arguments `tick=0` and `col.axis="white"` just have to be added to suppress the ticks and the tick labels on the individual panel graphs.

Note that the scatter plots can be further customized by setting panel functions to appear as something completely different. The off-diagonal panel functions are passed the appropriate columns, the diagonal panel

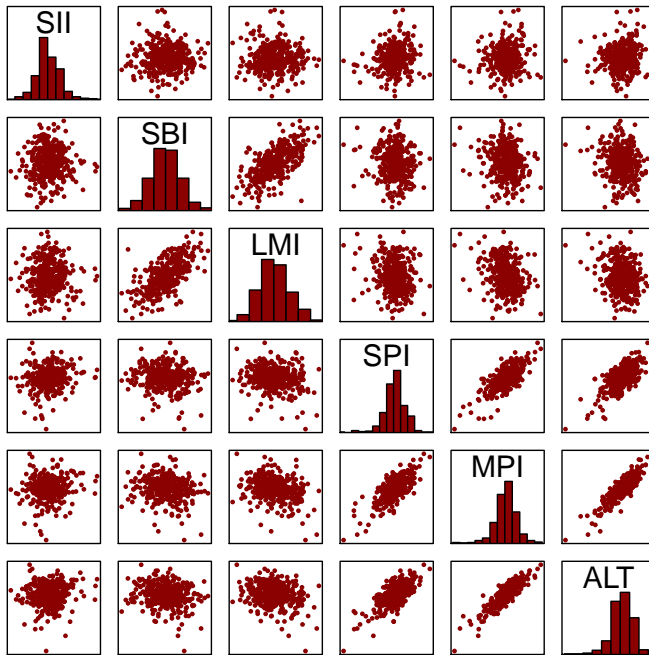


FIGURE 10.2: The plot shows customized pairwise scatter plots of assets with histograms in the diagonal panels. In addition, we have removed the axis labels on the panels.

function (if any) is passed a single column, and the text panel function is passed a single location and the column name.

10.2 PAIRWISE CORRELATIONS BETWEEN ASSETS

The function `assetsCorgramPlot()` displays a view of correlations as introduced by [Friendly \(2002\)](#), and is based on the implementations of the contributed R package `corrgram`¹ ([Wright, 2009](#)). This plot is also called a *correlogram plot*.

```
> args(assetsCorgramPlot)
function (x, method = c("pie", "shade"), ...)
NULL
```

The Rmetrics implementation works seamlessly with time series objects allows for two different correlogram views, a `method="pie"` and a `method="shade"`

¹The `corrgram` functions are available as built-ins in Rmetrics.




FIGURE 10.3: Correlation patterns for pairwise scatter plots: The patterns show how a matrix of correlations can be displayed schematically in the following forms: as numbers, as shaded squares, as bars, as ellipses, or as circular pac-man symbols. Source: [Friendly \(2002\)](#)

display of the off-diagonal panels. It is left to the user to extend the function to other off-diagonal displays as shown in [Figure 10.3](#).

In the implementation of [Friendly \(2002\)](#), a matrix of correlations can be displayed schematically in a variety of forms: as numbers, shaded squares, bars, ellipses, or as circular pac-man symbols². These schemes all attempt to show both the sign and magnitude of the correlation value, using a colour mapping of two hues in varying lightness, where the intensity of colour increases uniformly as the correlation value moves away from 0. Colour (blue for positive values, red for negative values) is used to encode the sign of the correlation, but the renderings are designed so that the sign may still be discerned when reproduced in black and white.

In the shaded row, each cell is shaded blue or red depending on the sign of the correlation, and with the intensity of colour scaled in proportion to the magnitude of the correlation. Such scaled colours are easily computed using RGB coding from red through white to blue. For simplicity, we ignore the non-linearities of colour reproduction and perception, but note that these are easily accommodated in the colour mapping function. White diagonal lines are added so that the direction of the correlation may still be discerned in black and white. This bipolar scale of colour was chosen to leave correlations near 0 empty (white), and to make positive and negative values of equal magnitude approximately equally intensely shaded. Gray scale and other colour schemes are implemented in the software, but not illustrated here.

The function `assetsCorgramPlot()` offers two options for the display of correlations. The lower panel is either a pie (pac-man) panel or a shaded panel overlayed by the scatter points, and the upper panel displays correlations as ellipses overlayed by a smooth fit of the data. Internally, smoothing is done by locally-weighted polynomial regression using the function `lowess()`. The first example displays a pac-man view

²Rmetrics has implemented the pac-man scheme, `method="pie"`, which is the default, and the shaded squares scheme, `method="shade"`.

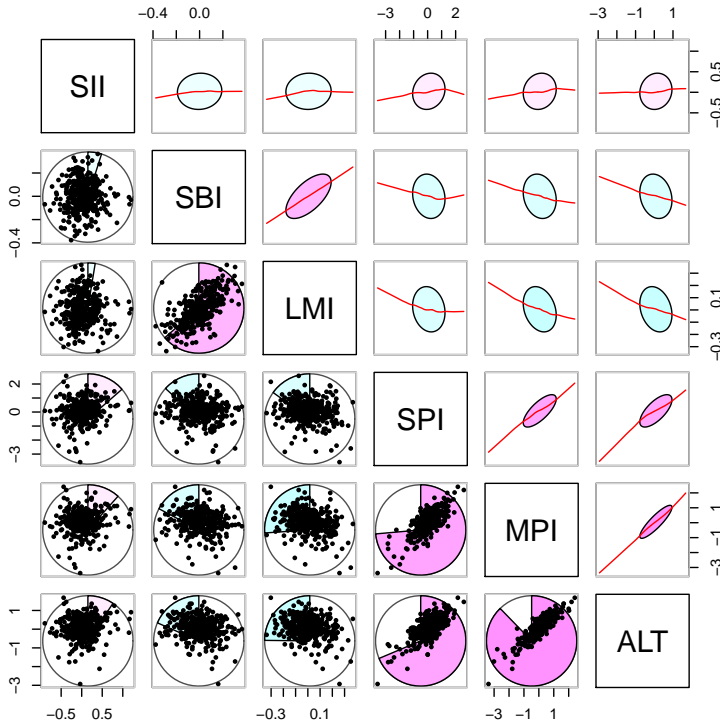


FIGURE 10.4: Display of sorted pairwise correlations between assets: The assets are sorted according to the grouping as obtained from hierarchical clustering. The lower off-diagonal panel returns a combination of the `piePanel()` together with the scatter plot as returned from the `pointsPanel()`. The upper off-diagonal panel returns a combination of the `ellipsePanel()` together with `lowess()` as returned from the `lowessPanel()` function. In the diagonal panel the names of the assets are shown.

```
> assetsCorgramPlot(LPP2005HC, pch = 19, cex = 0.5)
```

and the graph uses shaded off-diagonal panels.

```
> assetsCorgramPlot(LPP2005HC, method = "shade", pch = 19,
  cex = 0.5)
```

If you want to develop your own correlogram panels, several panel functions are available as hidden functions (within the `fAssets` package). They can be used to generate alternative views on correlogram plots with panel displays customized by the developer. These include diagonal panel functions:

LISTING 10.2: DIAGONAL PANEL FUNCTIONS

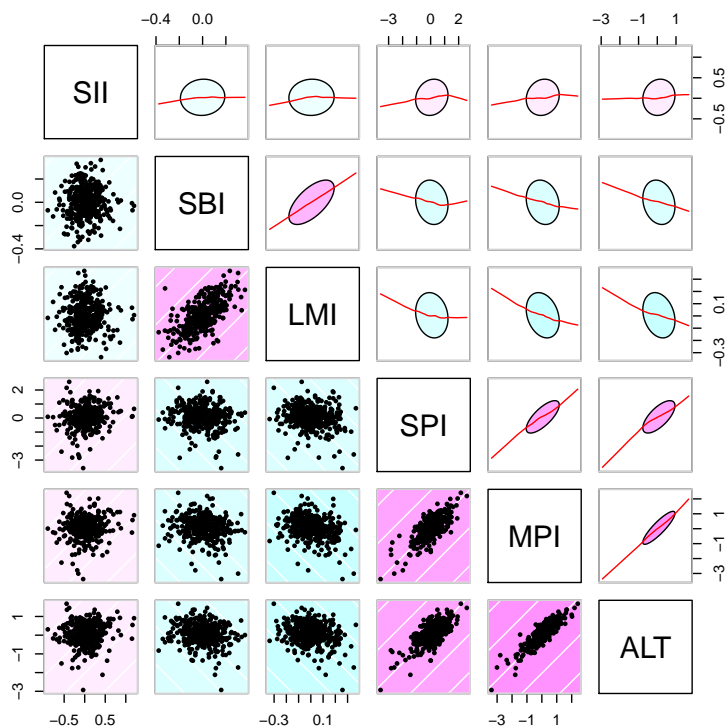


FIGURE 10.5: Display of sorted pairwise correlations between assets: The assets are sorted according to the grouping as obtained from hierarchical clustering. The lower off-diagonal panel returns a combination of the `piePanel()` together with the a scatter plot as returned from the `pointsPanel()`. The upper off-diagonal panel returns a combination of the `ellipsePanel()` together with `lowess()` as returned from the `lowessPanel()` function. In the diagonal panel the names of the assets are shown.

Function:

<code>.txtPanel</code>	displays a text panel with asset names
<code>.minmaxPanel</code>	displays min and max values
<code>.histPanel</code>	displays a histogram plot

and off-diagonal panel functions:

LISTING 10.3: OFF-DIAGONAL PANEL FUNCTIONS

Function:

<code>.ptsPanel</code>	displays a scatter plot panel
<code>.piePanel</code>	displays a pie (pac man) panel
<code>.shadePanel</code>	displays a shaded panel

<code>.ellipsePanel</code>	displays a coloured ellipse panel
<code>.cortestPanel</code>	displays a correlation test panel
<code>.lowessPanel</code>	displays a lowess fit panel
<code>.numberPanel</code>	displays correlations as numbers
<code>.piePtsPanel</code>	overlays a pie with a points panel

If you require further information, we recommend inspecting the source code of the provided hidden panel functions.

10.3 TESTS OF PAIRWISE CORRELATIONS

The function `assetsCorTestPlot()`

```
> args(assetsCorTestPlot)
function (x, ...)
NULL
```

combines a graphical view of the correlations together with the results from correlation tests.

```
> assetsCorTestPlot(LPP2005HC)
```

The lower off-diagonal panel displays the results from a scatter plot combined with a `lowess()` fit and the upper off diagonal panel shows the results returned from the function `cor.test()` which performs a test for association between paired samples, using one of Pearson's product moment correlation coefficient, Kendall's tau or Spearman's rho.

In the upper off-diagonal panels the numbers for the correlation coefficients are displayed as obtained from the function `cor()` together with the symbolic cutpoints "***" for <0.001, "**" for <0.01, "*" for <0.05, and "." for 0.1 for the p-values as obtained from the function `symnum()`. This function symbolically encodes a given numeric or logical vector or array and is thus particularly useful for the visualization of structured matrices, such as correlations.

The diagonal panel shows the names of the assets.

10.4 IMAGE PLOT OF CORRELATIONS

The image plot of correlations `assetsCorImagePlot()` can be used for a larger number of assets in data sets, mainly of financial returns. It gives another alternative view. The function `assetsCorImagePlot()`

```
> args(assetsCorImagePlot)
function (x, labels = TRUE, show = c("cor", "test"), use = c("pearson",
  "kendall", "spearman"), abbreviate = 3, ...)
NULL
```

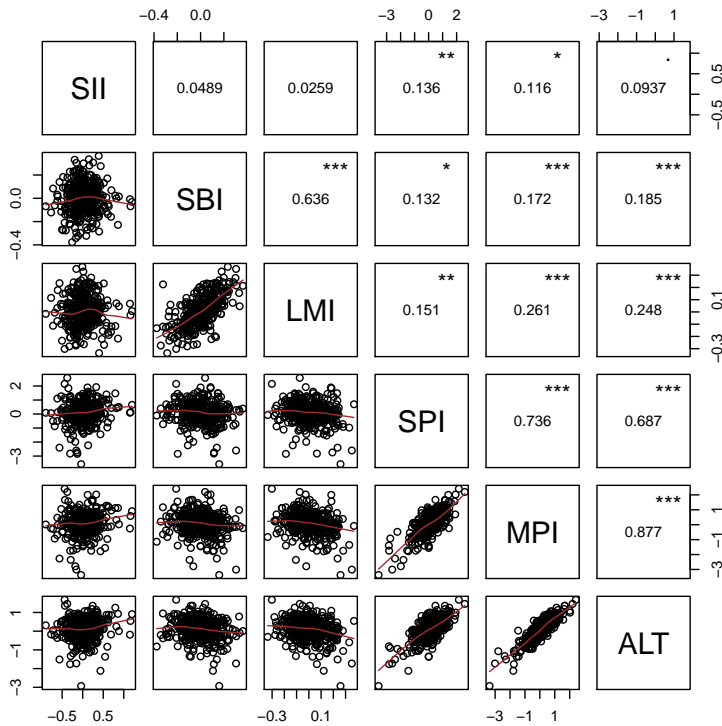



FIGURE 10.6: Scatter plots in combination with correlation tests: The plot represents a graphical view of correlations in combination with pairwise correlation tests. The lower off-diagonal panel displays a scatter plot combined with a `lowess()` fit. The upper off-diagonal panel shows the results returned from the correlation test. The diagonal panel shows the names of the assets.

returns a quadratic plot of squared images coloured according to the computed values either of the correlation coefficient, `show="cor"`, or of the correlation tests, `show="test"`.

LISTING 10.4: CORRELATION FUNCTIONS

```
Function:
cor          computes correlations
test         computes correlation test

Arguments:
show         specifies the image to be used, "cor" or "test"
use          specifies the correlation coefficient to be used,
              "pearson", "kendall" or "spearman"
abbreviate   abbreviates labels to specified length
```

If we have many assets, we can specify the argument `abbreviate` which allows to abbreviate assets name strings to the specified number of characters, such that they remain unique, if they were.

The following two graphs display different views on correlation images for the assets and benchmark series of the Swiss pension fund data set. In the first graph, [Figure 10.7](#), the assets are ordered according to the degree of similarity as suggested by hierarchical clustering,

```
> assetsCorImagePlot(LPP2005HC)
```

and in the second graph, [Figure 10.8](#), the columns of the data set are selected at random

```
> set.seed(1953)
> index <- sample(1:ncol(LPP2005HC))
> assetsCorImagePlot(LPP2005HC[, index])
```

Instead of using the the sample correlations, one can also think to use robust estimates for the correlation. The function `assetsCorImagePlot()` can easily be extended in this direction and is left as an example to the reader.

10.5 BIVARIATE HISTOGRAM PLOTS

Hexagon binning is useful for visualizing the structure in bivariate data sets of assets with a large number of records. The underlying concept is extremely simple, the plane of bivariate returns is tessellated by a regular grid of hexagons. Then the counts of points falling in each hexagon are counted, and finally the hexagons with at least one and more counts are plotted underlying a colour palette to the hexagons in proportion to the counts. If the size of the grid and the cuts in the colour palette are chosen in a proper fashion than the structure inherent in the data should emerge in the binned plots. Alternatively we can use squares instead of hexagons. To plot histograms of pairs of assets we can use the functions `squareBinning()` and `hexBinning()`. The functions return an S3 object either of class `squareBinning` or `hexBinning`. For both objects generic plot functions exist.

LISTING 10.5: BIVARIATE HISTOGRAM FUNCTIONS

Function:	
<code>squareBinning</code>	does a square binning of data points
<code>hexBinning</code>	does a hexagonal binning of data points
<code>plot</code>	generic bivariate binning plot function

The following example creates a bivariate histogram of two assets composed by hexagonal bins calling the function `hexBinning()`.




FIGURE 10.7: Image plots of pairwise correlations: The plot shows a symmetric coloured image with default settings: The numbers represent values for Pearson's correlation coefficient. Alternatively we can compute correlation tests. In both cases the underlying algorithms can use either Pearson's correlation coefficient, Kendall's rank correlation coefficient, or Spearman's rank correlation coefficient.

```
> args(hexBinning)
function (x, y = NULL, bins = 30)
NULL
```

The input can be either a bivariate timeSeries object `x`, or univariate timeSeries objects `x` and `y`. In the first case we set `y=NULL`, the default setting. In the next example we show the hexagonal binned histogram for Swiss bond, SBI, and Swiss performance index, SPI.

```
> hexHist <- hexBinning(SWX.RET[, c("SBI", "SPI")], bin = 20)
> plot(hexHist, xlab = "SBI", ylab = "SPI", col = rev(greyPalette(20)))
> title(main = "Bivariate Histogram Plot")
```

It is left to the reader to write his own panel functions with pairwise binned off-diagonal panels.




FIGURE 10.8: Image plots of pairwise correlations: The plot shows a symmetric coloured image with default settings: The numbers represent values for Pearson's correlation coefficient. Alternatively we can compute correlation tests. In both cases the underlying algorithms can use either Pearson's correlation coefficient, Kendall's rank correlation coefficient, or Spearman's rank correlation coefficient. In contrast to the previous graph, here the ordering of the assets is randomly chosen.




FIGURE 10.9: Bivariate histogram plots: The plots show bivariate histogram plots of the Swiss Bond, SBI, and Swiss Performance Index, SPI, expressed by hexagonal bins. The small dots in each bin express the centre of mass. The colours, here from a grey palette, indicate the frequency (counts) or probability of the counts.

PART III

PORTFOLIO FRAMEWORK

INTRODUCTION

Rmetrics provides functions to optimize a portfolio of assets. The goal of these functions is to determine the asset weights that minimize the risk for a desired return or, alternatively, that maximize the return for a given risk. After determining the optimal weights it is advisable to conduct a performance analysis on the optimal portfolio.

However, before you can optimize a portfolio, you have to create an environment which specifies a portfolio from the beginning and defines all the parameters and values that are required to perform the optimization. In the following chapters we define three S4 classes that describe the portfolio environment including (i) the specification of all parameters describing the portfolio, (ii) the selection and description of the assets data set for which we want to optimize the portfolio, and (iii) to set the constraints under which the portfolio will be optimized.

In [chapter 11](#) we introduce the S4 portfolio specification class. This process consists of three parts: First, we have to decide what kind of portfolio model we want to apply, e.g. an MV portfolio or a mean-CVaR portfolio. Secondly, we have to set the required portfolio parameters; these include, for example, the weights, the target return and risk, the risk-free rate, the number of frontier points and the status of the solver. Thirdly, we have to deal with the optimization parameter options. This involves setting the name of the solver to be used in optimization, e.g. linear, quadratic or nonlinear. Further, we can set the logical flag that tells us if the optimization should be traced. We explicitly show how to set, how to extract and how to modify these settings.

In [chapter 12](#) we discuss the S4 portfolio data class. In Rmetrics, data sets are represented by S4 timeSeries objects. These objects are used to represent financial asset returns for portfolio optimization. We describe the definition of the data, and the computation of the required statistical measures, including measures for the expected return and risk.

In [chapter 13](#) we describe the S4 portfolio constraints class. This is the most complex of the three classes. Constraints are defined by character strings or vectors of character strings. These strings can be used like a language to express lower and upper bounds on the ranges of weights that

have to be satisfied for box, group, covariance risk budgets and general non-linear constraint settings. We give detailed examples of how to specify these constraints.

In [chapter 14](#) we give a brief overview of the functions that are available to optimize portfolios. This includes the case of single portfolios as well as the case of the whole portfolio frontier.

CHAPTER 11

S4 PORTFOLIO SPECIFICATION CLASS

```
> library(fPortfolio)
```

To compose and optimize a portfolio of assets we first have to specify it. This process includes choosing the kind of portfolio model we want to investigate, choosing the required portfolio parameter settings, and choosing which type of programming solver (linear, quadratic, nonlinear) should be applied.

In this chapter we introduce the S4 portfolio specification class `fPFOLIOSPEC` and describe each of its slots. These are the `@model`, the `@portfolio`, the `@optim`¹ and the `@messages` slot. All four slots are represented by lists. In addition, we show how to extract and modify individual entries from these lists. A short discussion concerning consistency of the input parameters closes this chapter.

11.1 CLASS REPRESENTATION

All settings that specify a portfolio of assets are represented by an S4 class called `fPFOLIOSPEC`.

```
> showClass("fPFOLIOSPEC")  
  
Class "fPFOLIOSPEC" [package "fPortfolio"]  
  
Slots:  
  
Name:      model portfolio      optim messages      ampl  
Class:     list      list      list      list      list
```

¹optimization

An object of class `fPFOLIOSPEC` has four slots, named `@model`, `@portfolio`, `@optim`, and `@messages`. The first slot, `@model`, holds the model information, the second slot, `@portfolio`, the portfolio information and results, the `@optim` slot contains the information about the solver used for optimization, and the last slot, named `@messages`, holds a list of optional messages.

How to create a portfolio specification object

The function `portfolioSpec()` allows us to define specification settings from scratch. The default settings are for a mean-variance portfolio. To show the arguments of the function `portfolioSpec()`, you can use the function `formals()`, which prints an easy-to-read summary of the formal arguments.

```
> formals(portfolioSpec)
$model
list(type = "MV", optimize = "minRisk", estimator = "covEstimator",
      tailRisk = list(), params = list(alpha = 0.05, a = 1))

$portfolio
list(weights = NULL, targetReturn = NULL, targetRisk = NULL,
      riskFreeRate = 0, nFrontierPoints = 50, status = NA)

$optim
list(solver = "solveRquadprog", objective = c("portfolioObjective",
      "portfolioReturn", "portfolioRisk"), options = list(meq = 2),
      control = list(), trace = FALSE)

$messages
list(messages = FALSE, note = "")

$ampl
list(ampl = FALSE, project = "ampl", solver = "ipopt", protocol = FALSE,
      trace = FALSE)
```

The settings are created specifying the values for the model list, for the portfolio list, for the optimization list `optim`, and for the messages list². A more comprehensive listing of the arguments for the default settings is shown below³:

LISTING 11.1: ARGUMENTS OF THE FUNCTION `portfolioSpec()`

Arguments:

model slot

<code>type = "MV"</code>	a string value
<code>optimize = "minRisk"</code>	a string value

²A much more convenient way is to update an existing specification and to modify one or more of its parameters. In the next sections we present this in more detail.

³Note that when an argument is set to `NULL`, there is no default setting available and it is not required for the default portfolio.

```

estimator = "covEstimator"    a function name
tailRisk = list()             a list
params =
  list(alpha=0.05, a=1, ...) a list

portfolio slot
  weights = NULL              a numeric vector
  targetReturn = NULL         a numeric value
  targetRisk = NULL           a numeric value
  riskFreeRate = 0            a numeric value
  nFrontierPoints = 50        an integer value
  status = NA                 a integer value

optim slot
  solver = "solveRquadprog"   a function names
  objective = NULL            function names
  options = list()            a list with parameters
  control = list()            a list with controls
  trace = FALSE               a logical

messages slot:
  list = list()               a list

```

We can create the default settings for a mean-variance portfolio by calling the function `portfolioSpec()` without arguments⁴.

```
> defaultSpec <- portfolioSpec()
```

If we want to create a CVaR portfolio, we have to specify at least the model type, and the solver for the optimization.

```

> cvarSpec <- portfolioSpec(
  model = list(type = "CVaR", optimize = "minRisk",
    estimator = "covEstimator", tailRisk = list(),
    params = list(alpha = 0.05)),
  portfolio = list(weights = NULL,
    targetReturn = NULL, targetRisk = NULL,
    riskFreeRate = 0, nFrontierPoints = 50,
    status = 0),
  optim = list(solver = "solveRglpk", objective = NULL,
    params = list(), control = list(), trace = FALSE))

```

How to display the structure of a portfolio specification object

To look inside a portfolio's specification structure you can call the function `str()`. This function compactly displays the internal structure of the portfolio specification object⁵. It can be considered as a diagnostic function

⁴In this case, all arguments will just be set to their default value

⁵The entry `tailRisk` is only effective for portfolios constrained by tail risk budgets. How to use tail risk budgets in portfolio optimization is discussed in the ebook *Advanced Portfolio Optimization with R/Rmetrics*. The parameter `params$a=1` in the model slot is used as a risk aversion measure in mean-LPM portfolios. Lower partial moment, LPM, portfolios are also considered in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

and as a simple way to summarize the internal structure of the object. Let us inspect the structure of the default settings:

```
> str(defaultSpec)
Formal class 'fPFOLIOSPEC' [package "fPortfolio"] with 5 slots
 ..@ model      :List of 5
 .. ..$ type      : chr "MV"
 .. ..$ optimize   : chr "minRisk"
 .. ..$ estimator  : chr "covEstimator"
 .. ..$ tailRisk   : list()
 .. ..$ params     :List of 2
 .. .. ..$ alpha: num 0.05
 .. .. ..$ a      : num 1
 ..@ portfolio:List of 6
 .. ..$ weights      : NULL
 .. ..$ targetReturn  : NULL
 .. ..$ targetRisk    : NULL
 .. ..$ riskFreeRate  : num 0
 .. ..$ nFrontierPoints: num 50
 .. ..$ status        : logi NA
 ..@ optim          :List of 5
 .. ..$ solver       : chr "solveRquadprog"
 .. ..$ objective: chr [1:3] "portfolioObjective" "portfolioReturn" "portfolioRisk"
 .. ..$ options      :List of 1
 .. .. ..$ meq: num 2
 .. ..$ control      : list()
 .. ..$ trace        : logi FALSE
 ..@ messages :List of 2
 .. ..$ messages: logi FALSE
 .. ..$ note       : chr ""
 ..@ ampl         :List of 5
 .. ..$ ampl       : logi FALSE
 .. ..$ project    : chr "ampl"
 .. ..$ solver     : chr "ipopt"
 .. ..$ protocol   : logi FALSE
 .. ..$ trace      : logi FALSE
```

How to print a portfolio specification object

A nicely printed output of the same information can be obtained by using the generic `print()` function. Let us do this for the above-specified mean-CVaR portfolio.

```
> print(cvarSpec)
Model List:
Type:              CVaR
Optimize:          minRisk
Estimator:         covEstimator
Params:            alpha = 0.05

Portfolio List:
Target Weights:    NULL
Target Return:     NULL
```

```

Target Risk:           NULL
Risk-Free Rate:        0
Number of Frontier Points: 50
Status:                0

Optim List:
Solver:                solveRglpk
Objective:              list()
Trace:                 FALSE

```

11.2 THE MODEL SLOT

The `@model` slot covers all settings to specify a model portfolio. This includes the type of the portfolio, the objective function to be optimized, the estimators for mean and covariance, the tail risk⁶, and optional model parameters. To extract the current model specification we can use the extractor functions

LISTING 11.2: EXTRACTOR FUNCTIONS FOR THE `@model` SLOT

```

Model Slot - Extractor Functions:
getType           Extracts portfolio type from specification
getOptimize       Extracts what to optimize from specification
getEstimator      Extracts type of covariance estimator
getTailRisk       Extracts list of tail dependency risk matrices
getParams         Extracts parameters from specification

```

and to modify the settings from a portfolio specification we can use the following assignment functions:

LISTING 11.3: CONSTRUCTOR FUNCTIONS FOR THE `@model` SLOT

```

Model Slot - Constructor Functions:
setType           Sets type of portfolio optimization
setOptimize       Sets what to optimize, min risk or max return
setEstimator      Sets names of mean and covariance estimators
setParams         Sets optional model parameters

```

How to modify the type of the portfolio model

The list entry `$type` from the `@model` slot describes the type of the desired portfolio. In the current implementation, `type` can take different values to represent the type of portfolios⁷, such as

⁶see footnote 2

⁷The portfolio types "QLPM", "MAD", "SPS" are described in the ebook *Advanced Portfolio Optimization with R/Rmetrics*

LISTING 11.4: THE TYPE ARGUMENT FOR THE @model SLOT

Model Slot - Argument: type

Values:

"MV"	mean-variance (Markowitz) portfolio
"CVAR"	mean-conditional Value at Risk portfolio
"QLPM"	mean-quadratic-lower-partial-moment portfolio
"SPS"	Stone, Pedersen and Satchell type portfolios
"MAD"	mean-absolute-deviance Portfolio

One can now use the function `getType()` to retrieve the current setting and the assignment function `setType()` to modify this selection, e.g.

```
> mySpec <- portfolioSpec()
> getType(mySpec)
[1] "MV"

> setType(mySpec) <- "CVAR"
> getType(mySpec)
[1] "CVAR"
```

In this example we changed the specification from a mean-variance portfolio to a mean-conditional value-at-risk portfolio⁸.

Which objective to optimize

The list entry `$optimize` from the @model slot describes which objective function should be optimized. Possible choices are

LISTING 11.5: THE OPTIMIZE ARGUMENT FOR THE @model SLOT

Model Slot - Argument: optimize

Values:

"minRisk"	minimizes the risk for a given target return
"maxReturn"	maximizes the return for a given target risk
"objRisk"	gives the name of an alternative objective function

The first two options consider the most common choices; these are either minimizing the portfolio's risk for a given target return or maximizing the portfolio's return for a given target risk. In the default case of the mean-variance portfolio, the target risk is calculated from the sample covariance (or an alternative measure, e.g. a robust covariance estimate). The target return is computed by the sample mean of the assets if not otherwise specified. The third option leaves the user with the possibility to define any other portfolio objective function, such as maximizing the Sharpe ratio, for example⁹. You can use the function `getOptimize()` to retrieve and the assignment function `setOptimize()` to modify the current settings.

⁸Note that we now also have to modify the solver, since for CVAR portfolios, a linear solver is required. It is currently up to the user to make sure that the specification contains no conflicts.

⁹For examples of user defined objective functions for specific risk measure we refer to the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

How to estimate mean and covariance of asset returns

The list entry `$estimator` from the `@model` slot requires a string denoting the function name of the covariance estimator that should be used for estimating risk. In Markowitz' mean-variance portfolio model, `type="MV"`, the default function `covEstimator()` is used, which computes the sample column means and the sample covariance matrix of the multivariate assets data series. Alternative estimators include Kendall's and Spearman's rank based covariance estimators, robust estimators, and furthermore, a shrinkage and a bagged estimator.

The minimum covariance determinant estimator `mcdEstimator()` and the minimum volume ellipsoid estimator `mveEstimator()` are based on the robust covariance estimators from R's recommended MASS package (Venables & Ripley, 2008), which is part of R's base environment.

The estimators `covMcdEstimator()` and `covOGKEstimator()` require functions to be loaded from the contributed R package `robustbase` (Rousseeuw et al., 2008). `covMcdEstimator()` is an alternative implementation of the MCD estimator, and is faster than the one implemented in the MASS package. The Orthogonalized Gnanadesikan-Kettenring estimator, OGK, can be used when the dimensionality of the covariance matrix becomes large. The covariance estimators `shrinkEstimator()`, and `baggedEstimator()` use functions from the contributed R package `corpcor` (Schaefer et al., 2008). The shrinkage estimator computes the empirical variance of each considered random variable, and shrinks them towards their median (Schäfer & Strimmer, 2005; Opgen-Rhein & Strimmer, 2007). The bagged estimator uses bootstrap aggregating. This is a meta-algorithm to improve models in terms of stability and accuracy (Kotsiantis & Pintelas, 2004). Note that the R package `corpcor` does not have to be loaded explicitly, the required functions are available as built-ins¹⁰.

The function `nveEstimator()` performs robust covariance estimation by the nearest neighbour variance estimation, NNVE, method of Wang & Raftery (2002). The function is built-in from the contributed package `covRobust` (Wang et al., 2008).

LISTING 11.6: MODEL SLOT OF FUNCTION `portfolioSpec()`

```
Function:
portfolioSpec          specifies a portfolio

Model Slot:           specifies the type of estimator
  List Entry:
    estimator
      "covEstimator"    Covariance sample estimator
```

¹⁰Built-in functions are often modified or customized functions copied from external sources, usually from contributed packages. Rmetrics uses built-ins when only a small part of the code is required, or if the functions require slight modifications to work seamlessly in the Rmetrics environment.

"kendallEstimator"	Kendall's rank estimator
"spearmanEstimator"	Spearman's rank estimator
"mcdEstimator"	Minimum covariance determinant estimator
"mveEstimator"	Minimum volume ellipsoid estimator
"covMcdEstimator"	Minimum covariance determinant estimator
"covOGKEstimator"	Orthogonalized Gnanadesikan-Kettenring
"shrinkEstimator"	Shrinkage estimator
"baggedEstimator"	Bagged Estimator
"nnveEstimator"	Nearest neighbour variance estimator

You can add you own functions to estimate the mean and covariance of the multivariate assets data series. If you want to do so, you have to write a function, e.g. named

LISTING 11.7: TEMPLATE FOR A CUSTOM ESTIMATOR FUNCTION

```
myEstimator <- function(x, spec)
{
  <...>
  list(mu = <...> , Sigma = <...>)
}
```

where `x` is the multivariate time series object of assets, and `spec` is the portfolio specification. The argument `spec` allows additional parameters to be passed in. To be more specific, these arguments can usually be passed in through the `list@model$param`. Note that `myEstimator()` must return a named list, with at least the following two named entries: `$mu` and `$Sigma`. They represent the estimated values for the mean and covariance, respectively.

You can use the function `getEstimator()` to retrieve the current setting and the assignment function `setEstimator()` to modify the name of the estimator function to be used.

What is the tail risk list?

The list entry `tailRisk` from the `@model` slot is an empty list. It can be used to add tail risk budget constraints to the optimization. In this case a square matrix of pairwise tail dependence coefficients has to be specified as list entry. Usually, the matrix contains bivariate tail risk measures estimated via a copulae approach¹¹.

LISTING 11.8: THE TAILRISK ARGUMENT

```
Model Slot - Argument: tailRisk
List Entries:
...          a numeric matrix of tail dependence coefficients
```

You can use the function `getTailRisk()` to inspect the current setting and `setTailRisk()` to assign a tail risk matrix.

¹¹Modelling tail dependence coefficients using a copula approach is presented in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

How to set and modify model parameters

The list entry `$params` from the `@model` slot is a list with additional parameters used in different situations. It can be enhanced by the user if needed.

LISTING 11.9: THE PARAMS ARGUMENT

Model Slot - Argument: `params`

List Entries:

<code>alpha</code>	a numeric value, the VaR significance level α
<code>a</code>	a numeric value, the LPM risk measure exponent
<code>...</code>	optional parameters added by the user

By default, it contains the confidence level for "CVaR" portfolio optimization, `alpha=0.05`, and the exponent `a=1`, the parameter needed for portfolio optimization based on quadratic lower partial moments¹². Note that you can add additional parameters. For example, you could write your own robust covariance estimator, and if this function requires some parameters, you can pass them in through the model parameter list. Use the function `getModelParams()` and `setModelParams()` to inspect the current parameter settings, and to modify the values.

11.3 THE PORTFOLIO SLOT

The `@portfolio` slot covers all settings to specify the parameters for a portfolio. This includes the weights, the target return and risk, the risk-free rate, the number of frontier points and the status of the solver. Again, we can use the extractor functions to retrieve the current settings of the portfolio slot

LISTING 11.10: EXTRACTOR FUNCTIONS FOR THE `@portfolio` SLOT

Portfolio Slot - Extractor Functions:

<code>getWeights</code>	Extracts weights from a portfolio object
<code>getTargetReturn</code>	Extracts target return from specification
<code>getTargetRisk</code>	Extracts target risk from specification
<code>getRiskFreeRate</code>	Extracts risk-free rate from specification
<code>getNFrontierPoints</code>	Extracts number of frontier points
<code>getStatus</code>	Extracts the status of optimization

The assignment functions can be used to modify these settings

LISTING 11.11: ASSIGNMENT FUNCTIONS FOR THE `@portfolio` SLOT

Portfolio Slot - Assignment Functions:

¹²Optimizing portfolios based on the quadratic lower partial moment approach is discussed in the ebook *Advance Portfolio Optimization with R/Rmetrics*.

setWeights	Sets weights vector
setTargetReturn	Sets target return value
setTargetRisk	Sets target risk value
setRiskFreeRate	Sets risk-free rate value
setNFrontierPoints	Sets number of frontier points
setStatus	Sets status value

How to set the values of weights, target return and risk

The list entries `$weights`, `$targetReturn` and `$targetRisk` from the `@portfolio` slot have to be considered collectively.

LISTING 11.12: ARGUMENTS OF THE `@portfolio` SLOT

Portfolio slot - Arguments:

weights	a numeric vector of weights
targetReturn	a numeric value of the target return
targetRisk	a numeric value of the target risk

For example, if the weights for a portfolio are given, then the target return and target risk are determined, i.e. they are no longer free. As a consequence, if we set the weights to a new value, then the target return and risk also take new values, determined by the portfolio optimization. Since we do not know these values in advance, i.e. when we reset the weights, the values for the target return and risk are both set to NA. The same holds if we assign a new value to the target return or target risk; both of the other values are set to NA. By default, all three values are set to NULL. If this is the case, then it is assumed that an equal-weights portfolio should be calculated.

In summary, if only one of the three values is different from NULL, then the following procedure will be started:

1. If the weights are specified, it is assumed that a feasible portfolio should be considered.
2. If the target return is fixed, it is assumed that the efficient portfolio with the minimal risk should be considered.
3. And finally if the risk is fixed, the return should be maximized.

Use the functions `setWeights()`, `setTargetReturn()`, and `setTargetRisk()` to modify this selection. Note that a change in one of the three functions will influence the settings of the other two.

Let us look at an example of how to set the weights. First, let us display the default settings:

```
> mySpec <- portfolioSpec()
> getWeights(mySpec)
```

```
NULL
> getTargetReturn(mySpec)
NULL
> getTargetRisk(mySpec)
NULL
```

None of the three settings are available, therefore the extractor functions return NULL. Now, let us define a set of new weights, for example an equal-weights setting for four assets:

```
> setWeights(mySpec) <- c(1, 1, 1, 1)/4
> getWeights(mySpec)
[1] 0.25 0.25 0.25 0.25
> getTargetReturn(mySpec)
[1] NA
> getTargetRisk(mySpec)
[1] NA
> getOptimize(mySpec)
[1] "minRisk"
```

Now the target return and risk are set to NA, since we do not know the return and risk of the equal weights portfolio. On the other hand, if we want to fix the target return, for example to 2.5%, we can proceed as follows:

```
> setTargetReturn(mySpec) <- 0.025
> getWeights(mySpec)
[1] NA
> getTargetReturn(mySpec)
[1] 0.025
> getTargetRisk(mySpec)
[1] NA
> getOptimize(mySpec)
[1] "minRisk"
```

The weights and the target risk are now set to NA, since they are not known. In addition, the `getOptimize()` function returns "minRisk", since we have specified the target return. If we set the target risk, for example to 30%, then we obtain the following settings:

```
> setTargetRisk(mySpec) <- 0.3
> getWeights(mySpec)
[1] NA
> getTargetReturn(mySpec)
[1] NA
```

```

> getTargetRisk(mySpec)
[1] 0.3
> getOptimize(mySpec)
[1] "maxReturn"

```

Note that the `getOptimize()` function now returns the value "maxReturn". The name of the optimizer also has to be changed, since we are now dealing with quadratic constraints.

How to set the risk-free rate

The risk-free rate is the theoretical rate of return of an asset with zero risk. Its value, `riskFreeRate=0`, is stored in the `@portfolio` slot and set to zero by default.

LISTING 11.13: THE `RISKFREERATE` ARGUMENT OF THE `@portfolio` SLOT

Portfolio Slot - Argument:
`riskFreeRate` a numeric value of the risk-free rate

You can use the function `setRiskFreeRate()` to change the value of the risk-free rate, and the function `getRiskFreeRate()` to inspect its current value.

How to set the number of frontier points

The number of frontier points required by the calculation of the `portfolioFrontier` is obtained from the value of `nFrontierPoints` held in the portfolio slot. `nFrontierPoints` is set to 50 by default. You can change this with the function `setNFrontierPoints()`. The function `setNFrontierPoints()` returns the current setting for the number of frontier points.

LISTING 11.14: THE `NFRONTIERPOINTS` ARGUMENT OF THE `@portfolio` SLOT

Portfolio Slot - Argument:
`nFrontierPoints` an integer value specifying the number of
 frontier points

Bear in mind that if when considering a single portfolio, e.g. the tangency portfolio, the minimum-variance portfolio or any other efficient portfolio, the setting for the number of frontier points will be ignored.

How to obtain the solver status information

The final status of portfolio optimization is returned and stored in the `@portfolio` slot. Before optimization, the value is unset to NA, after optimization a value of `status=0` indicates a successful termination. For other values, we recommend that you inspect the help page of the selected solver. The name of the solver can be returned by the function `getSolver()`.

LISTING 11.15: THE STATUS ARGUMENT

```
Portfolio Slot - Argument:
status           an integer value of the status returned
                  by a portfolio optimization function
```

Note that the function `setStatus()` should only be used internally in solver functions to save and report the exit status.

11.4 THE OPTIM SLOT

The `@optim` slot deals with the solver settings, the name of the solver to be used in optimization, the logical flag which tells us if the optimization should be traced, and the message list.

For the optimization slot we have the following extractor functions

LISTING 11.16: EXTRACTOR FUNCTIONS FOR THE `@optim` SLOT

```
Optim slot - Extractor functions:
getSolver           Extracts the name of the solver
getTrace            Extracts solver's trace flag
getObjective        Extracts the name of the objective function
getOptions          Extracts optional solver parameters
getControl          Extracts the control list of the solver
```

and assignment functions to modify these settings

LISTING 11.17: CONSTRUCTOR FUNCTIONS FOR THE `@optim` SLOT

```
Portfolio slot - Constructor functions:
setSolver           sets the name of the solver
setTrace            sets solver's trace flag
setObjective        sets the name of the objective function
setOptions          sets optional solver parameters
setControl          sets the control list of the solver
```

How to select an appropriate solver

The name of the default solver used for the optimization of the mean-variance Markowitz portfolio, which is the default portfolio, is a quadratic programming (QP) solver, named `solveRquadprog()` in Rmetrics. This solver implements the approach of [Goldfarb & Idnani \(1982\)](#).

For mean-CVaR portfolio optimization, we use a linear programming (LP) solver, named `solveRglpk()` in Rmetrics. This solver uses R's interface to the GNU linear programming kit (GLPK) ([Makhorin, 2008](#)). Rmetrics provides a wide range of additional solvers:

LISTING 11.18: SOLVER ARGUMENTS IN THE OPTIM SLOT

```

Optim Slot - Argument: solver
Values:
"solveRquadprog"    Rmetrics default QP solver
"solveRglpk"        Rmetrics default LP solver
"solveRshortExact"  analytical short selling QP solver
"solveRipop"        alternative QP solver
"solveRlpSolveAPI"  alternative LP solver
"solveRsymphony"    alternative LP solver
"solveRsocp"        QP solver for quadratic constraints
"solveRdonlp2"      NL solver for non-linear constraints
...                for additional solvers

```

If you change the type of portfolio, remember to check whether you have specified a solver that is compatible with that type of portfolio . You can also choose the solver by calling the function `setSolver()` .

How to trace the iteration path

The logical flag `trace` in the `@optim` slot allows (most) solvers to trace the portfolio optimization process. By default, this will not be the case, i.e. `trace=FALSE`.

LISTING 11.19: THE TRACE ARGUMENT FOR THE @optim SLOT

```

Optim Slot - Argument:
trace                a logical flag to trace or not optimization
                    diagnostics from portfolio optimization

```

Tracing the process of portfolio optimization may be especially useful if we run into problems with the solver. By setting `trace=TRUE`, we can usually find out where the problems arise. You can use the function `setTrace()` to set or reset the selection.

How to add a user-defined objective function

When we optimize a portfolio for which the objective function to be optimized is neither the mean return nor the covariance risk, or any other predefined return or risk measure, then we can use a user-defined objective function, which we can pass in through the objective list entry of the `@optim` slot.

LISTING 11.20: THE OBJECTIVE ARGUMENT FOR THE @optim SLOT

```

Optim Slot - Argument:
objective            a character vector of three strings,
                    the objective function, the return,
                    and the risk function to be used.

```

You can use the function `setObjective()` to set or reset the selection.

How to add optional parameters

If a user-defined objective function requires additional options, you can pass them in through the options list entry of the @optim slot.

LISTING 11.21: THE OPTIONS ARGUMENT FOR THE @optim SLOT

```
Optim Slot - Argument:
options           a list of optional user supplied parameters
                  for the portfolio solvers
```

You can use the function `setOptions()` to set or reset the selection.

How to add control parameters for the solver

The argument `control` in the @optim slot allows you to control the parameters of the solvers. These are quantities such as the maximum number of iteration steps, or relative and absolute tolerances. Note that the entries in the list depend on which solver is used. An empty `list()` takes the default settings, which is what we recommend to control the parameters of the solvers.

LISTING 11.22: THE CONTROL ARGUMENT FOR THE @optim SLOT

```
Optim Slot - Argument:
control          a list of control parameters of the
                  portfolio solvers
```

Not all solvers allow you to modify their control settings, since these settings may be hard-coded. The QP quadprog solver, for instance, is one such solver. You can use the function `setControl()` to set or reset the selection, and the function `getControl()` to see the current control list.

11.5 THE MESSAGE SLOT

The message slots holds a list, into which you can save messages during the process of portfolio optimization. This option is especially helpful if you want to add your own portfolio models and solvers to the Rmetrics environment.

LISTING 11.23: THE ARGUMENT LIST FOR THE @message SLOT

```
Messages Slot - Argument: list
list           an optional list of messages added during
               the process of portfolio optimization
```

11.6 CONSISTENCY CHECKS ON SPECIFICATIONS

It is very important to be careful when modifying specification settings, because there are settings that are incompatible with certain other settings.

For example, if you want to minimize the covariance risk for a mean-variance portfolio, you cannot assign a linear programming solver. Currently we are working on implementing more consistency checks for the specification settings, so that you do not have to worry about creating conflicting settings. However, this has not been fully implemented in the current version of `fPortfolio`.

CHAPTER 12

S4 PORTFOLIO DATA CLASS

```
> library(fPortfolio)
```

In Rmetrics, data sets are represented by S4 `timeSeries` objects. These objects are used to represent financial returns series for portfolio optimization. Returns for a price or index series can be computed using the function `returns()`. The data summary information is stored in an object of class `fPFOLIODEATA`. An object of this class holds all the information about the data set of assets that is required for portfolio optimization. In this chapter, we introduce the S4 portfolio data class and describe each of the slots. These are the `@data` slot, the `@statistics` slot and the `@tailRisk` slot. In addition, we show how to extract and modify individual slots.

12.1 CLASS REPRESENTATION

An S4 `timeSeries` object only contains information on the series data themselves and the information on the date/time positions. Therefore, Rmetrics creates an S4 object of class `fPFOLIODEATA` with the function `portfolioData()`.

```
> showClass("fPFOLIODEATA")
```

```
Class "fPFOLIODEATA" [package "fPortfolio"]
```

```
Slots:
```

Name:	data	statistics	tailRisk
Class:	list	list	list

This S4 object holds additional information about the timeSeries data¹.

How to create a portfolio data object

The function `portfolioData()` allows you to define data settings for use in portfolio functions. The arguments of the function are

```
> args(portfolioData)
function (data, spec = portfolioSpec())
NULL
```

The settings are created by specifying the values for the time series data set and for the portfolio spec. First, we choose a subset of the `LPP2005.RET` returns data set, i.e. the "SBI", "SPI", "LMI" and "MPI"² columns. Then we create a portfolio object using the specified data and the default portfolio specification:

```
> lppAssets <- 100 * LPP2005.RET[, c("SBI", "SPI", "LMI", "MPI")]
> lppData <- portfolioData(data = lppAssets, spec = portfolioSpec())
```

How to display the structure of a portfolio data object

To look inside a portfolio's data structure you can call the function `str()`. This function compactly displays the internal structure of the portfolio data object. As in the case of the portfolio specification, the output can be considered as a diagnostic output and as a simple way to summarize the internal data structure.

```
> str(lppData, width = 65, strict.width = "cut")
Formal class 'fPFOLIODATA' [package "fPortfolio"] with 3 slots
  ..@ data      :List of 3
  .. ..$ series :Time Series:
    Name:      object
  Data Matrix:
    Dimension:      377 4
    Column Names:    SBI SPI LMI MPI
    Row Names:      2005-11-01 ... 2007-04-11
  Positions:
    Start:          2005-11-01
    End:            2007-04-11
  With:
    Format:         %Y-%m-%d
    FinCenter:      GMT
    Units:          SBI SPI LMI MPI
    Title:          Time Series Object
    Documentation:  Tue Jan 20 17:49:06 2009 by user:
```

¹The `@tailRisk` slot is only effective for portfolios constrained by tail risk budgets. We discuss how to use tail risk budgets in portfolio optimization in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

²For more information on the LPP2005, see [section B.2](#)

```

.. ..$ nAssets: int 4
.. ..$ names : chr [1:4] "SBI" "SPI" "LMI" "MPI"
..@ statistics:List of 5
.. ..$ mean : Named num [1:4] 4.07e-05 8.42e-02 5.53e-03 ..
.. ..$ attr(*, "names")= chr [1:4] "SBI" "SPI" "LMI" "MPI"
.. ..$ Cov : num [1:4, 1:4] 0.0159 -0.0127 0.0098 -0.015..
.. ..$ attr(*, "dimnames")=List of 2
.. ..$ : chr [1:4] "SBI" "SPI" "LMI" "MPI"
.. ..$ : chr [1:4] "SBI" "SPI" "LMI" "MPI"
.. ..$ estimator: chr "covEstimator"
.. ..$ mu : Named num [1:4] 4.07e-05 8.42e-02 5.53e-03 ..
.. ..$ attr(*, "names")= chr [1:4] "SBI" "SPI" "LMI" "MPI"
.. ..$ Sigma : num [1:4, 1:4] 0.0159 -0.0127 0.0098 -0.015..
.. ..$ attr(*, "dimnames")=List of 2
.. ..$ : chr [1:4] "SBI" "SPI" "LMI" "MPI"
.. ..$ : chr [1:4] "SBI" "SPI" "LMI" "MPI"
..@ tailRisk : list()

```

The internal structure shows us that we have three slots; the first is the `@data` slot, the second is the `@statistics` slot, and the third is the `@tail-Risk` slot.

How to print a portfolio data object

A nicely printed output of the same information can be obtained by using the generic `print()` function. Let us do this for the LPP2005 portfolio data object specified above.

```

> print(lppData)
Head/Tail Series Data:

GMT
      SBI      SPI      LMI      MPI
2005-11-01 -0.061275 0.84146 -0.110888 0.154806
2005-11-02 -0.276201 0.25193 -0.117594 0.034288
2005-11-03 -0.115309 1.27073 -0.099246 1.050296
GMT
      SBI      SPI      LMI      MPI
2007-04-09 0.000000 0.00000 -0.10324 0.817915
2007-04-10 -0.068900 0.63294 -0.00315 -0.142829
2007-04-11 0.030628 -0.10442 -0.00909 -0.099106

Statistics:

$mean
      SBI      SPI      LMI      MPI
4.0663e-05 8.4175e-02 5.5315e-03 5.9052e-02

$Cov
      SBI      SPI      LMI      MPI
SBI 0.0158996 -0.012741 0.0098039 -0.015888
SPI -0.0127414 0.584612 -0.0140747 0.411598
LMI 0.0098039 -0.014075 0.0149511 -0.023322

```

```

MPI -0.0158884  0.411598 -0.0233222  0.535033

$estimator
[1] "covEstimator"

$mu
      SBI      SPI      LMI      MPI
4.0663e-05 8.4175e-02 5.5315e-03 5.9052e-02

$Sigma
      SBI      SPI      LMI      MPI
SBI  0.0158996 -0.012741 0.0098039 -0.015888
SPI -0.0127414  0.584612 -0.0140747  0.411598
LMI  0.0098039 -0.014075  0.0149511 -0.023322
MPI -0.0158884  0.411598 -0.0233222  0.535033

```

The output displays the first and last three lines of the data set from the `@data` slot, and the sample mean and covariance estimates from the `@statistics` slot. Alternative or robust mean and covariance estimates, which are computed by the specified `covEstimator` function, are also printed. The name of the alternative estimator function has to be defined in the portfolio specification using the function `setEstimator()`. Note that the tail risk is only shown if it is not an empty list.

12.2 THE DATA SLOT

The `@data` slot keeps the S4 time series object, the number of assets, and their names in a list.

LISTING 12.1: THE `@data` SLOT

```

Data Slot - List Elements:
series          S4 timeSeries object
nAssets         number of assets
names           names of the assets

```

The contents of the `@data` slot can be extracted with the help of the function `getData()`.

```

> Data <- portfolioData(lppData)
> getData(Data)[-1]

$nAssets
[1] 4

$names
[1] "SBI" "SPI" "LMI" "MPI"

```

Since the time series in the first list element is quite long, we have excluded it from being printed.

12.3 THE STATISTICS SLOT

The `@statistics` slot holds information on the mean and covariance matrix of the timeSeries in a list.

LISTING 12.2: THE `@statistics` SLOT

Statistics Slot - List Elements:

mean	sample mean estimate
Cov	sample covariance estimate
estimator	name of alternative estimator function
mu	alternative mean estimate
Sigma	alternative covariance estimate

To be more precise, the `@statistics` slot holds the sample mean, `$mean`, and sample covariance matrix, `$Cov`, and additionally alternative measures for these two statistical measures, e.g. a robust estimate for the mean, `$mu`, and for the covariance matrix, `$Sigma`. The name of the estimator function used for the mean and covariance estimation can be retrieved from the character variable `$estimator`. A list of alternative mean and covariance estimators is given in chapter 4.

The contents of the `@statistics` slot can be extracted with the help of the function `getStatistics()`.

```
> getStatistics(Data)

$mean
      SBI      SPI      LMI      MPI
4.0663e-05 8.4175e-02 5.5315e-03 5.9052e-02

$Cov
      SBI      SPI      LMI      MPI
SBI  0.0158996 -0.012741 0.0098039 -0.015888
SPI -0.0127414 0.584612 -0.0140747 0.411598
LMI  0.0098039 -0.014075 0.0149511 -0.023322
MPI -0.0158884 0.411598 -0.0233222 0.535033

$estimator
[1] "covEstimator"

$mu
      SBI      SPI      LMI      MPI
4.0663e-05 8.4175e-02 5.5315e-03 5.9052e-02

$Sigma
      SBI      SPI      LMI      MPI
SBI  0.0158996 -0.012741 0.0098039 -0.015888
SPI -0.0127414 0.584612 -0.0140747 0.411598
LMI  0.0098039 -0.014075 0.0149511 -0.023322
MPI -0.0158884 0.411598 -0.0233222 0.535033
```


CHAPTER 13

S4 PORTFOLIO CONSTRAINTS CLASS

```
> library(fPortfolio)
```

Constraints define restrictions and boundary conditions on the weights and functional measures, depending on, or derived from, the portfolio weights. Constraints are defined by a character string or a vector of character strings. The formal style of these strings can be used like a language to express lower and upper bounds on the ranges of weights that have to be satisfied for box, group, covariance risk budgets and general non-linear constraint settings.

In this chapter we introduce the rules to express constraints as strings and introduce the functions used to create a summary for all the constraints. The function `portfolioConstraints()` creates such a summary, which is an S4 object of class `fPFOLIOCON`. We describe each of the slots that hold the constraint strings, the box and group constraints, the quadratic covariance risk budget constraints, and general non-linear constraints.

13.1 CLASS REPRESENTATION

In Rmetrics, portfolio constraints are represented by S4 `fPFOLIOCON` objects. These objects are used to represent all constraints settings for portfolio optimization. The function `portfolioConstraints()` creates the default settings and reports on all constraints in a compact form.

An S4 object of class `fPFOLIOCON` has the following representation:

```
> showClass("fPFOLIOCON")
Class "fPFOLIOCON" [package "fPortfolio"]

Slots:
```

Name:	stringConstraints	minWConstraints	maxWConstraints
Class:	character	numeric	numeric
Name:	eqsumWConstraints	minsumWConstraints	maxsumWConstraints
Class:	matrix	matrix	matrix
Name:	minBConstraints	maxBConstraints	listFConstraints
Class:	numeric	numeric	list
Name:	minFConstraints	maxFConstraints	minBuyinConstraints
Class:	numeric	numeric	numeric
Name:	maxBuyinConstraints	nCardConstraints	minCardConstraints
Class:	numeric	integer	numeric
Name:	maxCardConstraints		
Class:	numeric		

How to create a portfolio constraints object

The function `portfolioConstraints()` takes as arguments

```
> args(portfolioConstraints)
function (data, spec = portfolioSpec(), constraints = "LongOnly",
  ...)
NULL
```

the data set of assets as an object of `timeSeries`, the portfolio specification object `spec`, and the constraints, a vector of strings. The result is returned as an object of class `fPFOLIOCON`. The data set must be explicitly defined, whereas the last two arguments have default settings. The `spec=portfolioSpec()` argument takes as its default setting the mean-variance portfolio specification with long-only constraints, i.e. `constraints="LongOnly"`. In general, the argument `constraints` takes a character vector of constraints strings. The string vector can be composed from the following individual constraints:

LISTING 13.1: ARGUMENTS OF THE `portfolioConstraints()` FUNCTION

Argument: `constraints`

Values:

"LongOnly"	long-only constraints [0,1]
"Short"	unlimited short selling, [-Inf,Inf]
"minW[<...>]=<...>"	lower box bounds
"maxW[<...>]=<...>"	upper box bounds
"minsumW[<...>]=<...>"	lower group bounds
"maxsumW[<...>]=<...>"	upper group bounds
"minB[<...>]=<...>"	lower covariance risk budget bounds
"maxB[<...>]=<...>"	upper covariance risk budget bounds

```
"listF=list(<...>)"      list of non-linear functions
"minf[<...>]=<...>"     lower non-linear function bounds
"maxf[<...>]=<...>"     upper covariance risk budget
```

The returned S4 object of class `fPFOLIOCON` has eleven slots, the `@stringConstraints` slot (a character value or vector of constraints), and ten further slots for the individual constraints components. The following example creates the "LongOnly" default settings for the returns from the three Swiss assets from the LPP2005 Swiss Pension Fund Benchmark¹.

```
> Data <- 100 * LPP2005.RET[, 1:3]
> Spec <- portfolioSpec()
> setTargetReturn(Spec) <- mean(Data)
> Constraints <- "LongOnly"
> defaultConstraints <- portfolioConstraints(Data, Spec, Constraints)
```

How to display the structure of a portfolio constraints object

To explore the whole structure of an S4 `fPFOLIOCON`, type

```
> str(defaultConstraints, width = 65, strict.width = "cut")
Formal class 'fPFOLIOCON' [package "fPortfolio"] with 16 slots
..@ stringConstraints : chr "LongOnly"
..@ minWConstraints   : Named num [1:3] 0 0 0
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
..@ maxWConstraints   : Named num [1:3] 1 1 1
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
..@ eqsumWConstraints : num [1:2, 1:4] 3.60e-02 -1.00 4.07e-..
.. .. attr(*, "dimnames")=List of 2
.. .. ..$ : chr [1:2] "Return" "Budget"
.. .. ..$ : chr [1:4] "ceq" "SBI" "SPI" "SII"
..@ minsumWConstraints : logi [1, 1] NA
..@ maxsumWConstraints : logi [1, 1] NA
..@ minBConstraints    : Named num [1:3] -Inf -Inf -Inf
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
..@ maxBConstraints     : Named num [1:3] 1 1 1
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
..@ listFConstraints    : list()
..@ minFConstraints     : num(0)
..@ maxFConstraints     : num(0)
..@ minBuyinConstraints : Named num [1:3] 0 0 0
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
..@ maxBuyinConstraints : Named num [1:3] 1 1 1
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
..@ nCardConstraints    : int 3
..@ minCardConstraints : Named num [1:3] 0 0 0
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
..@ maxCardConstraints : Named num [1:3] 1 1 1
.. .. attr(*, "names")= chr [1:3] "SBI" "SPI" "SII"
```

¹Because the target return is not defined in the default spec settings, we set it to the grand mean of the assets data set.

How to print a portfolio constraints object

A nicely printed output of the same information can be obtained using the generic `print()` function. Let us do this for the above default mean-CVaR portfolio².

```
> print(defaultConstraints)
Title:
Portfolio Constraints

Lower/Upper Bounds:
      SBI SPI SII
Lower  0   0   0
Upper  1   1   1

Equal Matrix Constraints:
      ceq      SBI      SPI      SII
Return 0.036037 4.0663e-05 0.084175 0.023894
Budget -1.000000 -1.0000e+00 -1.000000 -1.000000

Cardinality Constraints:
      SBI SPI SII
Lower  0   0   0
Upper  1   1   1
```

13.2 LONG-ONLY CONSTRAINT STRING

Let us consider the settings for long-only constraints.

LISTING 13.2: LONG-ONLY CONSTRAINTS STRING

```
Constraints Settings:
"LongOnly"          long-only constraints, sets lower and upper bounds of
                    weights as box constraints
```

The long-only constraints generate the following set of weights:

```
> longConstraints <- "LongOnly"
> portfolioConstraints(Data, Spec, longConstraints)
Title:
Portfolio Constraints

Lower/Upper Bounds:
      SBI SPI SII
Lower  0   0   0
Upper  1   1   1

Equal Matrix Constraints:
      ceq      SBI      SPI      SII
Return 0.036037 4.0663e-05 0.084175 0.023894
Budget -1.000000 -1.0000e+00 -1.000000 -1.000000
```

²Non-defined constraints are not printed.

```
Cardinality Constraints:
      SBI SPI SII
Lower   0   0   0
Upper   1   1   1
```

This is also the default setting for the constraints if not otherwise defined. Alternatively, you can use `longConstraints=NULL`. Long-only positions reflect the fact that all weights are allowed to be between zero and one. Do not confuse this setting with box constraints, where the weights are restricted by arbitrary negative and/or positive lower and upper bounds. The output from the function `portfolioConstraints()` generates

1. the lower and upper bounds for each asset between 0 and 1 (100%),
2. two equal group constraints, where the first sums up to the target return, and the second sums up to 1 (i.e. we are fully invested), and
3. the lower and upper bounds for the covariance risk budgets between $-\infty$ and 1.

13.3 UNLIMITED SHORT SELLING CONSTRAINT STRING

Let us consider the settings for unlimited short constraints.

LISTING 13.3: SHORT CONSTRAINT STRING

```
Constraints Settings:
"Short"           short constraints, sets lower and upper bounds of
                  weights as box constraints ranging between minus
                  and plus infinity for all assets
```

The unlimited short constraints generate the following set of weights:

```
> shortConstraints <- "Short"
> portfolioConstraints(Data, Spec, shortConstraints)

Title:
Portfolio Constraints

Lower/Upper Bounds:
      SBI SPI SII
Lower -Inf -Inf -Inf
Upper  Inf  Inf  Inf

Equal Matrix Constraints:
      ceq      SBI      SPI      SII
Return 0.036037 4.0663e-05 0.084175 0.023894
Budget -1.000000 -1.0000e+00 -1.000000 -1.000000

Cardinality Constraints:
      SBI SPI SII
Lower   0   0   0
Upper   1   1   1
```

The setting allows for unlimited negative and positive weights. In this case, the mean-variance portfolio optimization problem can be solved analytically³.

13.4 BOX CONSTRAINT STRINGS

For arbitrary short-selling, use box constraints and set the lower bounds to negative values for those assets that are allowed for short selling. Weight-constrained portfolios, where the weights are limited by lower and upper bounds, are specified by the two character strings `minW` and `maxW`.

LISTING 13.4: CONSTRAINTS SETTING FOR BOX CONSTRAINTS

Constraints Settings:

```
"minW"          lower bounds of weights for box constraints
"maxW"          upper bounds of weights for box constraints
```

These character strings have to be used with indices between one and the number of assets, and appropriate values. If these values are all positive and between zero and one then we have constrained long-only portfolios. If they are allowed to become negative, then we have constrained (or limited) short portfolios.

Constraints are given as a vector composed of individual strings. For example, the constraints settings with the following strings form a box-constrained portfolio for a set of three assets:

```
> box.1 <- "minW[1:3] = 0.1"
> box.2 <- "maxW[c(1,3)] = c(0.5, 0.6)"
> box.3 <- "maxW[2] = 0.4"
> boxConstraints <- c(box.1, box.2, box.3)
> boxConstraints
[1] "minW[1:3] = 0.1"          "maxW[c(1,3)] = c(0.5, 0.6)"
[3] "maxW[2] = 0.4"
```

and the portfolio constraints become

```
> portfolioConstraints(Data, Spec, boxConstraints)
Title:
Portfolio Constraints

Lower/Upper Bounds:
SBI SPI SII
Lower 0.1 0.1 0.1
Upper 0.5 0.4 0.6

Equal Matrix Constraints:
          ceq          SBI          SPI          SII
Return 0.036037 4.0663e-05 0.084175 0.023894
```

³In order to solve the unlimited short-selling portfolio analytically, you have to set `setSolver(Spec)<-solveRshortExact`.

```
Budget -1.000000 -1.0000e+00 -1.000000 -1.000000
```

```
Cardinality Constraints:
```

```
      SBI SPI SII
Lower   0   0   0
Upper   1   1   1
```

The constraints tell us that we want to invest at least 10% in each asset, and no more than 50% in the first asset, a maximum of 40% in the second asset, and a maximum of 60% in number 3. Notice that we can repeat constraints settings, as for `maxW` in the example above. In this case, the previous settings for the assets will be overwritten if they are multiply defined. The variable `nAssets` is automatically recognized and the value of the total number of assets will be assigned.

13.5 GROUP CONSTRAINT STRINGS

Group constraints define the value of the total weight of a group of assets or lower and upper bounds on such groups. For this we can make use of the following strings:

LISTING 13.5: CONSTRAINTS SETTING FOR GROUP CONSTRAINTS

Constraints Settings:

```
"eqsumW"           equality group constraints
"minsumW"          lower bounds group constraints
"maxsumW"          upper bounds group constraints
```

Here, "eqsumW" sets the total amount of an investment in a group of assets to a fixed value, whereas "minsumW" and "maxsumW" set lower and upper bounds, e.g.

```
> group.1 <- "eqsumW[c(1,3)]=0.6"
> group.2 <- "minsumW[c(2,3)]=0.2"
> group.3 <- "maxsumW[c(1,2)]=0.7"
> groupConstraints <- c(group.1, group.2, group.3)
> groupConstraints
[1] "eqsumW[c(1,3)]=0.6" "minsumW[c(2,3)]=0.2" "maxsumW[c(1,2)]=0.7"
```

The first string means that we should invest exactly 60% of our money in the group consisting of assets one and three. The second string tells us that we should invest at least 20% in assets number two and three, and the third means that we want to invest no more than 70% of our money in assets one and two.

The portfolio constraints become

```
> portfolioConstraints(Data, Spec, groupConstraints)
```

```

Title:
Portfolio Constraints

Lower/Upper Bounds:
  SBI SPI SII
Lower   0   0   0
Upper   1   1   1

Equal Matrix Constraints:
      ceq      SBI      SPI      SII
Return  0.036037  4.0663e-05  0.084175  0.023894
Budget -1.000000 -1.0000e+00 -1.000000 -1.000000
eqsumW  0.600000  1.0000e+00  0.000000  1.000000

Lower Matrix Constraints:
      avec SBI SPI SII
lower  0.2   0   1   1

Upper Matrix Constraints:
      avec SBI SPI SII
upper  0.7   1   1   0

Cardinality Constraints:
      SBI SPI SII
Lower   0   0   0
Upper   1   1   1

```

Notice that the above conditions are reflected in the output of the function `portfolioConstraints()`. For example, the "eqsumW" constraint is shown in the "eqsumW" row of the Equal Matrix Constraints table. The value is give in the `ceq` column, and which assets this group constraint applies to is denoted by either 0 or 1. In this case, we can see that the constraint applies to the group comprising the SBI and the SII.

13.6 COVARIANCE RISK BUDGET CONSTRAINT STRINGS

By default, risk budgets are not included in the portfolio optimization. Covariance risk budgets have to be added explicitly, and have the following form:

LISTING 13.6: COVARIANCE RISK BUDGET CONSTRAINTS

```

Constraints Settings:
"minB"           lower bounds of the covariance risk budgets
"maxB"           upper bounds of the covariance risk budgets

```

"minB" and "maxB" have to be assigned to the lower and upper bounds of the covariance risk budgets. The assignments must be given to all assets. The following shows an example of covariance risk budget constraints for a portfolio with 3 assets:

```
> budget.1 <- "minB[1:nAssets]=-Inf"
```



```

> budget.2 <- "maxB[c(1, 2:nAssets)]=c(0.5, rep(0.6, times=2))"
> budgetConstraints <- c(budget.1, budget.2)
> budgetConstraints

[1] "minB[1:nAssets]=-Inf"
[2] "maxB[c(1, 2:nAssets)]=c(0.5, rep(0.6, times=2))"

```

and the portfolio constraints become

```

> portfolioConstraints(Data, Spec, budgetConstraints)

Title:
Portfolio Constraints

Lower/Upper Bounds:
      SBI SPI SII
Lower  0   0   0
Upper  1   1   1

Equal Matrix Constraints:
      ceq      SBI      SPI      SII
Return 0.036037 4.0663e-05 0.084175 0.023894
Budget -1.000000 -1.0000e+00 -1.000000 -1.000000

Lower/Upper Cov Risk Budget Bounds:
      SBI SPI SII
Lower -Inf -Inf -Inf
Upper  0.5  0.6  0.6

Cardinality Constraints:
      SBI SPI SII
Lower  0   0   0
Upper  1   1   1

```

Again, the variable *nAssets* is automatically recognized and the value of the total number of assets, here 3, will be assigned.

Risk budget constraints will enforce diversification at the expense of return generation. The resulting portfolios will thus lie below the unconstrained efficient frontier.

Note that adding risk budget constraints will modify the optimization problem since these constraints are quadratic, unlike the other constraints considered so far. This requires optimizers which can handle non-linear constraints⁴.

13.7 NON-LINEAR WEIGHT CONSTRAINT STRINGS

We can also make use of non-linear functional constraints. This can be achieved by using the following strings:

⁴Portfolio Solvers for quadratic constraints will be presented in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

LISTING 13.7: NON-LINEAR WEIGHT CONSTRAINTS

Constraints Settings:

```
"listF"          list of non-linear functions
"minF"           lower bounds of non-linear functions
"maxF"           upper bounds of non-linear functions
```

If, for example, we want to constrain our portfolio by a maximum draw-down, we first have to write a function to compute the maximum draw-down. We call the function `maxdd()`, and define it such that the data `x` will be passed to the function `min(drawdowns())`.

```
> maxdd <- function(x, ...) min(drawdowns(x, ...))
```

Next, we have to decide on the lower and upper bounds.

```
> nonlin.1 <- "listF=list(maxdd=maxdd)"
> nonlin.2 <- "minF=-0.04"
> nonlin.3 <- "maxF=0"
> nonlinConstraints <- c(nonlin.1, nonlin.2, nonlin.3)
> nonlinConstraints

[1] "listF=list(maxdd=maxdd)" "minF=-0.04"
[3] "maxF=0"
```

Then portfolio constraints become

```
> portfolioConstraints(Data, Spec, nonlinConstraints)

Title:
Portfolio Constraints

Lower/Upper Bounds:
      SBI SPI SII
Lower  0   0   0
Upper  1   1   1

Equal Matrix Constraints:
      ceq      SBI      SPI      SII
Return 0.036037 4.0663e-05 0.084175 0.023894
Budget -1.000000 -1.0000e+00 -1.000000 -1.000000

Non-Linear Function Constraints:
      maxdd
Lower -0.04
Upper  0.00

Cardinality Constraints:
      SBI SPI SII
Lower  0   0   0
Upper  1   1   1
```

These settings can be interpreted in the following way: "listF" lists all the function names of the non-linear constraints function, and "minF" and "maxF" hold the values of their lower and upper bounds. Notice that

if there is more than one non-linear constraints function, then "listF" holds the names of all functions, and "minF" and "maxF" are composed as vectors of the same length as the number of non-linear functions⁵.

13.8 CASE STUDY: HOW TO CONSTRUCT COMPLEX PORTFOLIO CONSTRAINTS

Box, group, risk budget and non-linear constraints can now be combined. First, let us create a larger data set than in the previous examples, using the "SBI", "SPI", "SII", "LMI", "MPI" and "ALT" columns from the LPP2005 returns data set.

```
> # create data set
> Data <- 100 * LPP2005.RET[,1:6]
> # create default portfolio spec
> Spec <- portfolioSpec()
> # set target return to grand mean of the data
> setTargetReturn(Spec) <- mean(Data)
```

Now that we have created our data set and portfolio specification, we can turn our attention to the individual constraint strings. These can all be created in the same way as before, but bear in mind that our data set now consists of six assets instead of three.

```
> box.1 <- "minW[1:6] = 0.1"
> box.2 <- "maxW[c(1:3, 5)] = c(rep(0.5, 3), 0.6)"
> box.3 <- "maxW[4] = 0.4"
> # combine individual strings
> boxConstraints <- c(box.1, box.2, box.3)
> boxConstraints

[1] "minW[1:6] = 0.1"
[2] "maxW[c(1:3, 5)] = c(rep(0.5, 3), 0.6)"
[3] "maxW[4] = 0.4"

> group.1 <- "eqsumW[c(2, 3, 5)]=0.2"
> group.2 <- "minsumW[c(2, 4)]=0.2"
> group.3 <- "maxsumW[c(4:6, 2)]=0.6"
> # combine individual strings
> groupConstraints <- c(group.1, group.2, group.3)
> groupConstraints

[1] "eqsumW[c(2, 3, 5)]=0.2" "minsumW[c(2, 4)]=0.2"
[3] "maxsumW[c(4:6, 2)]=0.6"

> budget.1 <- "minB[1:nAssets]= -Inf"
> budget.2 <- "maxB[c(1:3, 4:nAssets)]=rep(c(0.5, 0.6), each=3)"
> # combine individual strings
```

⁵Portfolio Solvers for non-linear constraints will be presented in the ebook *Advance Portfolio Optimization with R/Rmetrics*.

```

> budgetConstraints <- c(budget.1, budget.2)
> budgetConstraints
[1] "minB[1:nAssets]=-Inf"
[2] "maxB[c(1:3, 4:nAssets)]=rep(c(0.5, 0.6), each=3)"

```

The portfolio should be constrained by a maximum drawdown; therefore, we redefine the same function as used above to compute this:

```

> # create function to compute maximum drawdown
> maxdd <- function(x,...) min(drawdown(x,...))

> nonlin.1 <- "listF=list(maxdd=maxdd)"
> nonlin.2 <- "minF=-0.04"
> nonlin.3 <- "maxF=0"
> # combine individual strings
> nonlinConstraints <- c(nonlin.1, nonlin.2, nonlin.3)
> nonlinConstraints
[1] "listF=list(maxdd=maxdd)" "minF=-0.04"
[3] "maxF=0"

```

The constraint string can now be constructed from all the individual constraints in the following manner:

```

> Constraints <- c(boxConstraints, groupConstraints,
  budgetConstraints, nonlinConstraints)
> Constraints
[1] "minW[1:6] = 0.1"
[2] "maxW[c(1:3, 5)] = c(rep(0.5, 3), 0.6)"
[3] "maxW[4] = 0.4"
[4] "eqsumW[c(2, 3, 5)]=0.2"
[5] "minsumW[c(2, 4)]=0.2"
[6] "maxsumW[c(4:6, 2)]=0.6"
[7] "minB[1:nAssets]=-Inf"
[8] "maxB[c(1:3, 4:nAssets)]=rep(c(0.5, 0.6), each=3)"
[9] "listF=list(maxdd=maxdd)"
[10] "minF=-0.04"
[11] "maxF=0"

```

Finally, we can now create the portfolio constraints by passing the complex constraints as an argument of the `portfolioConstraints()` function:

```

> portfolioConstraints(Data, Spec, Constraints)
Title:
Portfolio Constraints

Lower/Upper Bounds:
      SBI SPI SII LMI MPI ALT
Lower 0.1 0.1 0.1 0.1 0.1 0.1
Upper 0.5 0.5 0.5 0.4 0.6 1.0

Equal Matrix Constraints:
      ceq      SBI      SPI      SII      LMI      MPI
Return 0.043077 4.0663e-05 0.084175 0.023894 0.0055315 0.059052

```

```

Budget -1.000000 -1.0000e+00 -1.000000 -1.000000 -1.0000000 -1.000000
eqsumW  0.200000  0.0000e+00  1.000000  1.000000  0.0000000  1.000000

```

```

    ALT

```

```

Return  0.085768

```

```

Budget -1.000000

```

```

eqsumW  0.000000

```

```

Lower Matrix Constraints:

```

```

    avec SBI SPI SII LMI MPI ALT

```

```

lower  0.2  0  1  0  1  0  0

```

```

Upper Matrix Constraints:

```

```

    avec SBI SPI SII LMI MPI ALT

```

```

upper  0.6  0  1  0  1  1  1

```

```

Lower/Upper Cov Risk Budget Bounds:

```

```

    SBI SPI SII LMI MPI ALT

```

```

Lower -Inf -Inf -Inf -Inf -Inf -Inf

```

```

Upper  0.5  0.5  0.5  0.6  0.6  0.6

```

```

Non-Linear Function Constraints:

```

```

    maxdd

```

```

Lower -0.04

```

```

Upper  0.00

```

```

Cardinality Constraints:

```

```

    SBI SPI SII LMI MPI ALT

```

```

Lower  0  0  0  0  0  0

```

```

Upper  1  1  1  1  1  1

```


CHAPTER 14

PORTFOLIO FUNCTIONS

```
> library(fPortfolio)
```

After we have loaded the data, specified the portfolio settings, and defined the constraints, we are ready to optimize the portfolio. The portfolio optimization functions take the data, the specifications and the constraints as inputs. The returned value is an S4 object of class `fPORTFOLIO` which can be used to print reports and/or to display graphs.

The usage of the functions is described in detail in [Part IV](#) for mean-variance portfolios and in [Part V](#) for mean-CVaR portfolios.

14.1 S4 CLASS REPRESENTATION

In Rmetrics, portfolios are represented by S4 `fPORTFOLIO` objects. An S4 object of class `fPORTFOLIO` has the following representation

```
> showClass("fPORTFOLIO")

Class "fPORTFOLIO" [package "fPortfolio"]

Slots:

Name:      call      data      spec constraints  portfolio
Class:     call fPFOLIOWORK fPFOLIOSPEC fPFOLIOCON  fPFOLIOVAL

Name:      title description
Class:     character  character
```

These objects are returned by the computation and optimization of any portfolio.

Functions to compute and optimize portfolios

The functions to compute or optimize a single portfolio are:

LISTING 14.1: FUNCTIONS FOR COMPUTING AND OPTIMIZING PORTFOLIOS

Portfolio Functions:

<code>feasiblePortfolio</code>	returns a feasible portfolio given the vector of portfolio weights
<code>efficientPortfolio</code>	returns the portfolio with the lowest risk for a given target return
<code>maxratioPortfolio</code>	returns the portfolio with the highest return/risk ratio
<code>tangencyPortfolio</code>	synonym for <code>maxratioPortfolio</code>
<code>minriskPortfolio</code>	returns a portfolio with the lowest risk at all
<code>minvariancePortfolio</code>	synonym for <code>minriskPortfolio</code>
<code>maxreturnPortfolio</code>	returns the portfolio with the highest return for a given target risk
<code>portfolioFrontier</code>	computes portfolios on the efficient frontier and/or on the minimum covariance locus.

If the weights for the function `feasiblePortfolio()` are not specified in the portfolio specification, then the function assumes that the portfolio should be an equal-weights portfolio.

The portfolio functions are called with the following arguments

```
> args(feasiblePortfolio)
function (data, spec = portfolioSpec(), constraints = "LongOnly")
NULL
```

The portfolio specification and the constraints specification have defaults, which are those for a mean-variance portfolio with long-only constraints. The returned values are those from an S4 object of class `fPORTFOLIO` with the following slots:

LISTING 14.2: FPORTFOLIO SLOTS

Returned Portfolio Values:

<code>call</code>	function call
<code>data</code>	data, an object of class <code>fPFOLIODATA</code>
<code>spec</code>	specification, an object of class <code>fPFOLIOSPEC</code>
<code>constraints</code>	constraints, an object of class <code>fPFOLIOCON</code>
<code>portfolio</code>	the portfolio result as returned by the implied solver
<code>title</code>	an optional title, by default the portfolio function name
<code>description</code>	an optional description, by default time and name of the user

How to display the structure of a portfolio object

The structure of an S4 object of class `fPORTFOLIO` is quite comprehensive, because it contains the previously presented data, spec, and constraints objects. To explore the whole structure of an S4 `fPFOLIOCON` object, type:

```
> tgPortfolio <- tangencyPortfolio(100 * LPP2005.RET[, 1:6])
> str(tgPortfolio, width = 65, strict.width = "cut")

Formal class 'fPORTFOLIO' [package "fPortfolio"] with 7 slots
..@ call      : language maxratioPortfolio(data = data, spec..
..@ data      :Formal class 'fPFOLIODATA' [package "fPortfo"..
.. .. ..@ data      :List of 3
.. .. .. ..$ series :Time Series:
      Name:      object
Data Matrix:
Dimension:      377 6
Column Names:   SBI SPI SII LMI MPI ALT
Row Names:      2005-11-01 ... 2007-04-11
Positions:
Start:          2005-11-01
End:            2007-04-11
With:
Format:         %Y-%m-%d
FinCenter:      GMT
Units:          SBI SPI SII LMI MPI ALT
Title:          Time Series Object
Documentation:   Tue Jan 20 17:49:06 2009 by user:
.. .. .. ..$ nAssets: int 6
.. .. .. ..$ names   : chr [1:6] "SBI" "SPI" "SII" "LMI" ...
.. .. .. ..@ statistics:List of 5
.. .. .. .. ..$ mean      : Named num [1:6] 4.07e-05 8.42e-02 2.3..
.. .. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII"..
.. .. .. .. ..$ Cov       : num [1:6, 1:6] 0.0159 -0.0127 0.0018 ..
.. .. .. .. ..- attr(*, "dimnames")=List of 2
.. .. .. .. .. ..$ : chr [1:6] "SBI" "SPI" "SII" "LMI" ...
.. .. .. .. .. ..$ : chr [1:6] "SBI" "SPI" "SII" "LMI" ...
.. .. .. .. ..$ estimator: chr "covEstimator"
.. .. .. .. ..$ mu        : Named num [1:6] 4.07e-05 8.42e-02 2.3..
.. .. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII"..
.. .. .. .. ..$ Sigma     : num [1:6, 1:6] 0.0159 -0.0127 0.0018 ..
.. .. .. .. ..- attr(*, "dimnames")=List of 2
.. .. .. .. .. ..$ : chr [1:6] "SBI" "SPI" "SII" "LMI" ...
.. .. .. .. .. ..$ : chr [1:6] "SBI" "SPI" "SII" "LMI" ...
.. .. .. ..@ tailRisk   : list()
..@ spec             :Formal class 'fPFOLIOSPEC' [package "fPortfo"..
.. .. .. ..@ model      :List of 5
.. .. .. .. ..$ type     : chr "MV"
.. .. .. .. ..$ optimize : chr "minRisk"
.. .. .. .. ..$ estimator: chr "covEstimator"
.. .. .. .. ..$ tailRisk : list()
.. .. .. .. ..$ params   :List of 2
.. .. .. .. .. ..$ alpha : num 0.05
.. .. .. .. .. ..$ a     : num 1
.. .. .. ..@ portfolio:List of 6
.. .. .. .. ..$ weights   : atomic [1:6] 0 0.000476 0.18239..
```

```

.. .. .. attr(*, "invest")= num 1
.. .. .. ..$ targetReturn : logi NA
.. .. .. ..$ targetRisk : logi NA
.. .. .. ..$ riskFreeRate : num 0
.. .. .. ..$ nFrontierPoints: num 50
.. .. .. ..$ status : num 0
.. .. ..@ optim :List of 5
.. .. .. ..$ solver : chr "solveRquadprog"
.. .. .. ..$ objective: chr [1:3] "portfolioObjective" "port"..
.. .. .. ..$ options :List of 1
.. .. .. .. ..$ meq: num 2
.. .. .. ..$ control : list()
.. .. .. ..$ trace : logi FALSE
.. .. ..@ messages :List of 2
.. .. .. ..$ messages: logi FALSE
.. .. .. ..$ note : chr ""
.. .. ..@ ampl :List of 5
.. .. .. ..$ ampl : logi FALSE
.. .. .. ..$ project : chr "ampl"
.. .. .. ..$ solver : chr "ipopt"
.. .. .. ..$ protocol: logi FALSE
.. .. .. ..$ trace : logi FALSE
..@ constraints:Formal class 'fPFOLIOCON' [package "fPortfol"..
.. .. ..@ stringConstraints : chr "LongOnly"
.. .. ..@ minWConstraints : Named num [1:6] 0 0 0 0 0 0
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
.. .. ..@ maxWConstraints : Named num [1:6] 1 1 1 1 1 1
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
.. .. ..@ eqsumWConstraints : num [1, 1:7] -1 -1 -1 -1 -1 -1..
.. .. .. ..- attr(*, "dimnames")=List of 2
.. .. .. .. ..$ : chr "Budget"
.. .. .. .. ..$ : chr [1:7] "ceq" "SBI" "SPI" "SII" ...
.. .. .. ..- attr(*, "na.action")=Class 'omit' Named num 1
.. .. .. .. ..- attr(*, "names")= chr "Return"
.. .. ..@ minsumWConstraints : logi [1, 1] NA
.. .. ..@ maxsumWConstraints : logi [1, 1] NA
.. .. ..@ minBConstraints : Named num [1:6] -Inf -Inf -Inf..
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
.. .. ..@ maxBConstraints : Named num [1:6] 1 1 1 1 1 1
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
.. .. ..@ listFConstraints : list()
.. .. ..@ minFConstraints : num(0)
.. .. ..@ maxFConstraints : num(0)
.. .. ..@ minBuyinConstraints: Named num [1:6] 0 0 0 0 0 0
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
.. .. ..@ maxBuyinConstraints: Named num [1:6] 1 1 1 1 1 1
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
.. .. ..@ nCardConstraints : int 6
.. .. ..@ minCardConstraints : Named num [1:6] 0 0 0 0 0 0
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
.. .. ..@ maxCardConstraints : Named num [1:6] 1 1 1 1 1 1
.. .. .. ..- attr(*, "names")= chr [1:6] "SBI" "SPI" "SII" ""..
..@ portfolio :Formal class 'fPFOLIOVAL' [package "fPortfol"..
.. .. ..@ portfolio:List of 6
.. .. .. ..$ weights : Named num [1:6] 0 0.000476 0.182..

```

```

.. .. .. attr(*, "names")= chr [1:6] "SBI" "SPI" "SII"..
.. .. .. ..$ covRiskBudgets: Named num [1:6] 0 0.00141 0.1538..
.. .. .. .. attr(*, "names")= chr [1:6] "SBI" "SPI" "SII"..
.. .. .. ..$ targetReturn : Named num [1:2] 0.0283 0.0283
.. .. .. .. attr(*, "names")= chr [1:2] "mean" "mu"
.. .. .. ..$ targetRisk : Named num [1:4] 0.153 0.153 0.31..
.. .. .. .. attr(*, "names")= chr [1:4] "Cov" "Sigma" "C"..
.. .. .. ..$ targetAlpha : num 0.05
.. .. .. ..$ status : num 0
.. .. ..@ messages : list()
..@ title : chr "Tangency Portfolio"
..@ description: chr "Tue Jan 27 13:38:47 2015 by user: Rmet"..

```

The above shows the structure of a mean-variance tangency portfolio with long-only constraints.

How to print a portfolio object

A compact and nicely formatted printout of the most important information can be obtained by using the generic `print()` function. Let us do this for the mean-variance tangency portfolio that we optimized above.

```

> print(tgPortfolio)
Title:
MV Tangency Portfolio
Estimator:      covEstimator
Solver:         solveRquadprog
Optimize:       minRisk
Constraints:     LongOnly

Portfolio Weights:
  SBI   SPI   SII   LMI   MPI   ALT
0.0000 0.0005 0.1824 0.5753 0.0000 0.2418

Covariance Risk Budgets:
  SBI   SPI   SII   LMI   MPI   ALT
0.0000 0.0014 0.1539 0.1124 0.0000 0.7324

Target Returns and Risks:
  mean   Cov   CVaR   VaR
0.0283 0.1533 0.3096 0.2142

Description:
Tue Jan 27 13:38:47 2015 by user: Rmetrics

```


PART IV

MEAN-VARIANCE PORTFOLIOS

INTRODUCTION

Modern portfolio theory proposes how rational investors use diversification to optimize their portfolio(s) of risky assets. The basic concepts of the theory go back to [Markowitz \(1952\)](#)'s idea of diversification and the efficient portfolio frontier. His model considers asset returns as a random variable, and models a portfolio as a weighted combination of assets. Being a random variable, a portfolio's returns have an expected mean and variance. In this model, return and risk are estimated by the sample mean and the sample standard deviation of the asset returns.

In [chapter 15](#) we briefly describe the mean-variance portfolio theory and present its solution when no restrictions are set on the weights. This is the unlimited short selling case where an analytically closed form solution is possible. We derive the feasible set and the efficient frontier. Two special points on the frontier are discussed in detail: the minimum variance portfolio and the tangency portfolio. In the case of constraints we discuss the solutions for box and group constraints defining linear constraints. The case of the maximum return mean-variance portfolio and in the case of additional covariance risk budget constraints we have a new situation where quadratic forms of the constraints are becoming active.

In [chapter 16](#) we show how to specify a mean-variance portfolio which describes an S4 class in Rmetrics. We discuss the slots representing the portfolio data, the portfolio specification, and the portfolio constraints.

In [chapter 17](#) we show in several examples how to compute an optimize several type of mean-variance portfolios. These include feasible portfolios, portfolios with the lowest risk for a given return, the global minimum variance portfolio, the tangency portfolio, and portfolios with the highest return for a given risk. We also show how to handle non-linear constraints, such as the maximum drawdown constrained portfolio, or the 130/30 extension strategy constrained portfolio.

In [chapter 18](#) we present how to compute and graphically display the whole efficient frontier of a portfolio. As special examples we consider the portfolio with long-only constraints, the unlimited short selling portfolio, box and/or group constrained portfolios, and covariance risk budget constrained portfolios. In addition we show how to create different re-

ward/risk views on the efficient frontier.

In [chapter 19](#) we present a case study.

In [chapter 20](#) we discuss how to robustify portfolios using alternative covariance estimators. The standard estimator used in the mean-variance portfolio optimization is the sample covariance estimator. Here we show the influence on the portfolio if we use robust and related statistical estimators. These include the minimum covariance determinant estimator, the minimum volume ellipsoid estimator, the orthogonalized Gnanadesikan-Kettenring estimator for large covariance matrixes, and the shrinkage covariance estimator.

CHAPTER 15

MARKOWITZ PORTFOLIO THEORY

```
> library(fPortfolio)
```

In this chapter we define the original mean-variance portfolio optimization problem and several related problems. The problem of minimizing the covariance risk for a given target return with optional box and group constraints is a quadratic programming problem with linear constraints. We call it QP1 or *minimum risk mean-variance portfolio*. The opposite case, fixing the risk and maximizing the return, has a linear objective function with quadratic constraints. We call this programming problem QP2 or *maximum return mean-variance portfolio* problem. The QP2 problem is much more complex than QP1. If we have even more complex constraints, i.e. nonlinear constraints, we need a new class of solvers, which we call NL1, or *non-linear constrained portfolio* problem. This allows us to handle the case of linear and quadratic objective functions with non-linear constraints. In all three cases we speak of Markowitz' portfolio optimization problem, although they require different classes of solvers with increasing complexity.

15.1 THE MINIMUM RISK MEAN-VARIANCE PORTFOLIO

Following [Markowitz \(1952\)](#) we define the problem of portfolio selection as follows:

$$\begin{aligned}
& \min_w w^T \hat{\Sigma} w \\
& s.t. \\
& w^T \hat{\mu} = \bar{r} \\
& w^T \mathbf{1} = 1
\end{aligned}$$

The formula expresses that we minimize the variance-covariance risk $\bar{\sigma}^2 = w^T \hat{\Sigma} w$, where the matrix $\hat{\Sigma}$ is an estimate of the covariance of the assets. The vector w denotes the individual investments subject to the condition $w^T \mathbf{1} = 1$ that the available capital is fully invested. The expected or target return \bar{r} is expressed by the condition $w^T \hat{\mu} = \bar{r}$, where the p -dimensional vector $\hat{\mu}$ estimates the expected mean of the assets. Markowitz' portfolio model¹ has a unique solution:

$$w^* = \hat{\mu} w_0^* + w_1^*$$

where

$$\begin{aligned}
w_0^* &= \frac{1}{\Delta} (B \hat{\Sigma}^{-1} \hat{\mu} - C \hat{\Sigma}^{-1} \mathbf{1}) \\
w_1^* &= \frac{1}{\Delta} (C \hat{\Sigma}^{-1} \hat{\mu} - A \hat{\Sigma}^{-1} \mathbf{1}) \\
\Delta &= AB - C^2
\end{aligned}$$

with

$$\begin{aligned}
A &= \hat{\mu}^T \hat{\Sigma}^{-1} \hat{\mu} \\
B &= \mathbf{1}^T \hat{\Sigma}^{-1} \mathbf{1} \\
C &= \mathbf{1}^T \hat{\Sigma}^{-1} \hat{\mu} .
\end{aligned}$$

15.2 THE FEASIBLE SET AND THE EFFICIENT FRONTIER

The corresponding standard deviation $\bar{\sigma}$ for the optimal portfolio with weights w^* is

$$\begin{aligned}
\bar{\sigma} &= \sqrt{\frac{1}{\Delta} (\hat{\mu} B - 2\hat{\mu} C + A)} \\
\bar{r} &= w^T \hat{\mu} .
\end{aligned}$$

¹For a detailed listing of Markowitz' assumptions and technical conditions underlying his approach we refer to [Vanini & Vignola \(2001\)](#).




FIGURE 15.1: Risk versus return view of the mean-variance portfolio: Mean-variance portfolio illustrations in the $(\sigma(r), r)$ space. The minimum variance point r_* separates the efficient frontier $\partial_+ \mathcal{A}$ from the minimum variance locus $\partial_- \mathcal{A}$.

The locus of this set in the $\{\bar{\sigma}, \bar{r}\}$ -space are hyperbolas. The set inside the hyperbola is the *feasible set* of mean/standard deviation portfolios, and the borders are the *efficient frontier* (upper border), and the minimum variance locus (lower border). Here, r_* is the return of the minimum variance portfolio.

15.3 THE MINIMUM VARIANCE PORTFOLIO

The point with the smallest risk on the efficient frontier is called the *global minimum variance portfolio*, MVP. The MVP represents just the minimum risk point on the efficient frontier. The set of weight is:

$$w_* = \frac{\Sigma^{-1} \mathbf{1}}{\mathbf{1}^T \Sigma^{-1} \mathbf{1}}$$

15.4 THE CAPITAL MARKET LINE AND TANGENCY PORTFOLIO

Reward/risk profiles of different combinations of a risky portfolio with a riskless asset, with expected return r_f , can be represented as a straight line, the so called capital market line, CML. The point where the CML touches the efficient frontier corresponds to the optimal risky portfolio. Mathematically, this can be expressed as the portfolio that maximizes the quantity

$$\begin{aligned} \max_w \quad & h(w) = \frac{\hat{\mu}^T w - r_f}{w^T \hat{\Sigma} w} \\ \text{s.t.} \quad & \\ & w^T \hat{\mu} = \bar{r} \\ & w^T \mathbf{1} = 1 \end{aligned}$$

among all w . This quantity is precisely the *Sharpe ratio* introduced by [Sharpe \(1994\)](#).

15.5 BOX AND GROUP CONSTRAINED MEAN-VARIANCE PORTFOLIOS

As shown above, the unlimited short selling portfolio can be solved analytically. However, if the weights are bounded by zero, which forbids short selling, then the optimization has to be done numerically. The structure of the portfolio problems is quadratic and thus we can use a quadratic solver to compute the weights of the portfolio. In the following we consider as the standard Markowitz portfolio problem a portfolio which sets box and group constraints on the weights:

$$\begin{aligned} \min_w \quad & w^T \Sigma w \\ \text{s.t.} \quad & \\ & Aw \leq b \end{aligned}$$

It can be shown that, if Σ is a positive definite matrix, the Markowitz portfolio problem is a convex optimization problem. As such, its local optimal solutions are also global optimal solutions.

The contributed R package `quadprog` ([Weingessel, 2004](#)) provides the function `solve.QP()`, which interfaces a FORTRAN subroutine. This subroutine implements the dual method of [Goldfarb & Idnani \(1982, 1983\)](#) for solving quadratic programming problems of the form $\min(-c^T x + 1/2 x^T C x)$ with the constraints $A^T x \geq b$. The Rmetrics solver function `solveRQuadprog(data, spec, constraints)` provides a direct interface to the FORTRAN subroutine provided in the `quadprog` package. The desired solver is selected through the specification structure, which means

that the user need not interact with the underlying FORTRAN subroutine. If necessary, the setting can be modified by calling the function `setSolver()`. This also allows you to supply an alternative solver providing access to another algorithm or to write your own code.

15.6 MAXIMUM RETURN MEAN-VARIANCE PORTFOLIOS

In contrast to minimum risk portfolios, where we minimize the risk for a given target return, maximum return portfolios work the opposite way: Maximize the return for a given target risk.

$$\begin{aligned} \max_w \quad & w^T \hat{\mu} \\ \text{s.t.} \quad & Aw \leq b \\ & w^T \hat{\Sigma} w \leq \sigma \end{aligned}$$

Note that now we are concerned with a linear programming problem and quadratic constraints. This can be solved in Rmetrics using either the second order cone programming solver from the R package `Rsocp` or the less efficient non-linear programming solver from the R package `Rdonlp2`.

15.7 COVARIANCE RISK BUDGETS CONSTRAINTS

Risk budgeting is a way of taking a finite risk resource, and deciding how best to allocate it. In a mean-variance world, this defaults to Markowitz' portfolio optimization, where results are not only in terms of weights and monetary allocations but also in terms of risk contributions ([Scherer & Martin, 2005](#)). In order to quantify risk contributions, we address the questions of how the portfolio risk changes if we increase or decrease holdings in a set of assets. This change for a given asset i can be computed from the derivative

$$\sigma = \sqrt{w^T \hat{\Sigma} w} = \sum_i w_i \frac{d\sigma}{dw_i} .$$

To make the interpretation easier, we divide through σ and arrive at normalized risk budgets that sum up to 100%, i.e. to 1.

$$1 = \sum_i \mathcal{B}_i = \sum_i w_i \frac{w_i}{\sigma} \frac{d\sigma}{dw_i} .$$

Now, adding risk budgeting constraints in portfolio optimization

$$\begin{aligned}
& \min w^T \hat{\Sigma} w \\
& s.t. \\
& \quad w^T \hat{\mu} = \bar{r} \\
& \quad w^T \mathbf{1} = 1 \\
& \quad \mathcal{B}_i^{lower} \leq \frac{w_i}{\sigma} \frac{d\sigma}{dw_i} \leq \mathcal{B}_i^{upper} \\
& \quad \dots
\end{aligned}$$

allows us to limit the maximum and minimum risk contributions arising from individual positions.

Covariance risk budgeting can be formulated for a risk minimizing portfolio as a portfolio with a quadratic objective function and quadratic constraints, in addition to the common linear constraints. The problem is discussed in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

CHAPTER 16

MEAN-VARIANCE PORTFOLIO SETTINGS

```
> library(fPortfolio)
```

Like all portfolios in Rmetrics, mean-variance portfolios are defined by the time series data set, the portfolio specification object, and the constraint strings. Specifying a portfolio thus requires three steps.

16.1 STEP 1: PORTFOLIO DATA

The portfolio functions expect S4 `timeSeries` objects. You can create them from scratch using one of the functions from the Rmetrics `timeSeries` package for time series generation. Alternatively, you can load a data set from the demo examples provided in the `fPortfolio` package. Note that the portfolio functions expect time-ordered data records. To sort S4 `timeSeries` objects, use the generic function `sort()`. To align time series objects and to manage missing values use the function `align()`. If you want to bind and merge several `timeSeries` to a data set of assets, you can use the functions `cbind()`, `rbind()` and `merge()`. These functions are explained in detail in [chapter 1](#)¹.

16.2 STEP 2: PORTFOLIO SPECIFICATION

The representation of the portfolio specification and how to manage the slots is discussed in detail in [chapter 11](#). For Markowitz' mean-variance portfolio we can just use the default settings

¹For further details please see the ebook *Chronological Objects with R/Rmetrics*

```

> mvSpec <- portfolioSpec()
> print(mvSpec)
Model List:
Type:                MV
Optimize:             minRisk
Estimator:           covEstimator
Params:              alpha = 0.05 a = 1

Portfolio List:
Target Weights:      NULL
Target Return:       NULL
Target Risk:         NULL
Risk-Free Rate:      0
Number of Frontier Points: 50
Status:              NA

Optim List:
Solver:              solveRquadprog
Objective:           portfolioObjective portfolioReturn portfolioRisk
Options:             meq = 2
Trace:               FALSE

```

The printout tells us that the portfolio type is concerned with the mean-variance portfolio "MV", that we want to optimize (minimize) the risk "minrisk" using the quadprog solver "solveRquadprog", and that the sample covariance estimator "covEstimator" will be applied. The other two parameters shown are the risk-free rate and the number of frontier points. The first will only be used when we calculate the tangency portfolio and the Sharpe ratio, and the second when we calculate the whole efficient frontier.

16.3 STEP 3: PORTFOLIO CONSTRAINTS

In many cases we will work with long-only portfolios. Specifying

```
> constraints <- "LongOnly"
```

will force the lower and upper bounds for the weights to zero and one, respectively.

Many alternative constraints have already been implemented in `fPortfolio`. These include unlimited short selling, lower and upper bounds, linear equality and inequality constraints, covariance risk budget constraints, and non-linear function constraints. For a full list, see [chapter 11](#). The solver for dealing with these constraints has to be selected by the user and assigned by the function `setSolver()`.

CHAPTER 17

MINIMUM RISK MEAN-VARIANCE PORTFOLIOS

```
> library(fPortfolio)
```

The following examples show how to compute the properties of a minimum risk mean-variance portfolio. These portfolios have a quadratic objective function defined by the covariance matrix of the financial assets and a fixed target return. Included are feasible, efficient, tangency and global minimum risk portfolios. We consider the case of linear constraints, as well as long-only, short selling, box and group constraints¹.

17.1 HOW TO COMPUTE A FEASIBLE PORTFOLIO

A *feasible portfolio* is an ‘existing’ portfolio described by the settings of the portfolio specification. ‘Existing’ means that the portfolio was specified by its parameters in such a way that in a risk versus return plot the portfolio has a solution and is part of the feasible set (including the efficient frontier and the minimum variance locus).

The generic way to define a feasible portfolio is to define the portfolio weights. For example, the *equal weights portfolio* is such a portfolio. To specify the equal weights portfolio for the LPP2005² data set, we first list the names of the instruments which are part of this data set, and then we subset those which we want to include in the portfolio.

LISTING 17.1: THE TABLE LISTS FUNCTIONS TO OPTIMIZE LINEARLY CONSTRAINED MEAN-VARIANCE PORTFOLIOS AND TO PLOT THE RESULTS

¹The case of non-linear constraints and the use of alternative solvers is described in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

²The LPP2005 data set has nine columns, the first six columns are domestic and foreign assets, the last three columns are the benchmark indices with increasing risk profiles.

Functions:

<code>feasiblePortfolio</code>	feasible portfolio given the weights
<code>efficientPortfolio</code>	minimum risk portfolio for given return
<code>tangencyPortfolio</code>	portfolio with highest Sharpe ratio
<code>minvariancePortfolio</code>	global minimum risk portfolio
<code>weightsPie</code>	weights pie plot
<code>weightedReturnsPie</code>	weighted returns pie plot
<code>covRiskBudgetsPie</code>	covariance risk budget pie plot

```
> colnames(LPP2005REC)
[1] "SBI" "SPI" "SII" "LMI" "MPI" "ALT" "LPP25" "LPP40" "LPP60"
> lppData <- 100 * LPP2005REC[, 1:6]
```

Next we add the vector of weights³ to the specification `spec`, using the function `setWeights()`. In this case, are using equal weights.

```
> ewSpec <- portfolioSpec()
> nAssets <- ncol(lppData)
> setWeights(ewSpec) <- rep(1/nAssets, times = nAssets)
```

Now we are ready to calculate the properties of this portfolio. To do so, we call the function `feasiblePortfolio()`

```
> ewPortfolio <- feasiblePortfolio(
  data = lppData,
  spec = ewSpec,
  constraints = "LongOnly")
> print(ewPortfolio)
```

Title:
MV Feasible Portfolio

Estimator: covEstimator
Solver: solveRquadprog
Optimize: minRisk
Constraints: LongOnly

Portfolio Weights:

SBI	SPI	SII	LMI	MPI	ALT
0.1667	0.1667	0.1667	0.1667	0.1667	0.1667

Covariance Risk Budgets:

SBI	SPI	SII	LMI	MPI	ALT
-0.0039	0.3526	0.0431	-0.0079	0.3523	0.2638

Target Returns and Risks:

mean	Cov	CVaR	VaR
0.0431	0.3198	0.7771	0.4472

Description:
Tue Jan 27 13:40:53 2015 by user: Rmetrics

³The default settings do not specify a weights vector; by default it is set to NULL. We therefore have to supply the portfolio weights explicitly

The output first reports the settings, then the portfolio weights, then the covariance risk budgets, and finally the target returns and risks. This includes the portfolio mean, and several portfolio risk measures, including the variance computed from the covariance matrix, the conditional value-at-risk, and the value-at-risk. Note that since we have specified no alternative covariance estimator, μ and Σ are the same as mean and Cov.

Now let us display the results from the equal weights portfolio, the assignment of weights, and the attribution of returns and risk.

```
> col <- divPalette(ncol(lppData), "RdBu")
> weightsPie(ewPortfolio, radius = 0.7, col = col)
> mtext(text = "Equally Weighted MV Portfolio", side = 3, line = 1.5,
      font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(ewPortfolio, radius = 0.7, col = col)
> mtext(text = "Equally Weighted MV Portfolio", side = 3, line = 1.5,
      font = 2, cex = 0.7, adj = 0)
> covRiskBudgetsPie(ewPortfolio, radius = 0.7, col = col)
> mtext(text = "Equally Weighted MV Portfolio", side = 3, line = 1.5,
      font = 2, cex = 0.7, adj = 0)
```

The pie plots are shown in the left-hand column of [Figure 17.1](#). We have created a view of the pies with a legend listing the asset names and the percentual part of the pies. All pie plots are surrounded by a box, which is the default. The colours are taken from a red to blue diverging palette.

17.2 HOW TO COMPUTE A MINIMUM RISK EFFICIENT PORTFOLIO

A *minimum risk efficient portfolio* is a portfolio with the lowest risk for a given target return. As a first example for an efficient portfolio, we calculate the efficient mean-variance portfolio with the same target return as the equal weights portfolio, but with the lowest possible risk. Since the default settings of the `portfolioSpec()` function does not define a target return we should not forget to explicitly add the target return to the portfolio specification.

```
> minriskSpec <- portfolioSpec()
> targetReturn <- getTargetReturn(ewPortfolio@portfolio)["mean"]
> setTargetReturn(minriskSpec) <- targetReturn
```

The next step is then to optimize the portfolio for the specified target return.

```
> minriskPortfolio <- efficientPortfolio(
  data = lppData,
  spec = minriskSpec,
  constraints = "LongOnly")
> print(minriskPortfolio)
```

Title:

MV Efficient Portfolio

Estimator: covEstimator

```
Solver:          solveRquadprog
Optimize:        minRisk
Constraints:      LongOnly
```

Portfolio Weights:

```
  SBI   SPI   SII   LMI   MPI   ALT
0.0000 0.0086 0.2543 0.3358 0.0000 0.4013
```

Covariance Risk Budgets:

```
  SBI   SPI   SII   LMI   MPI   ALT
0.0000 0.0184 0.1205 -0.0100 0.0000 0.8711
```

Target Returns and Risks:

```
  mean   Cov   CVaR   VaR
0.0431 0.2451 0.5303 0.3412
```

Description:

```
Tue Jan 27 13:40:53 2015 by user: Rmetrics
```

The weights and related pie plots are generated in the same way as for the equal weights portfolio shown in the previous section.

```
> col <- qualiPalette(ncol(lppData), "Dark2")
> weightsPie(minriskPortfolio, radius = 0.7, col = col)
> mtext(text = "Minimal Risk MV Portfolio", side = 3, line = 1.5,
        font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(minriskPortfolio, radius = 0.7, col = col)
> mtext(text = "Minimal Risk MV Portfolio", side = 3, line = 1.5,
        font = 2, cex = 0.7, adj = 0)
> covRiskBudgetsPie(minriskPortfolio, radius = 0.7, col = col)
> mtext(text = "Minimal Risk MV Portfolio", side = 3, line = 1.5,
        font = 2, cex = 0.7, adj = 0)
```

The pie plots are shown in the right-hand column of [Figure 17.1](#). We have created a view of the pies with the same settings as in the previous example, except for the colours; these are taken from a dark qualitative palette.

17.3 HOW TO COMPUTE THE GLOBAL MINIMUM VARIANCE PORTFOLIO

The *global minimum variance portfolio* is the efficient portfolio with the lowest possible risk. The global minimum variance point is thus the point which separates the efficient frontier from the minimum variance locus.

```
> globminSpec <- portfolioSpec()
> globminPortfolio <- minvariancePortfolio(
  data = lppData,
  spec = globminSpec,
  constraints = "LongOnly")
> print(globminPortfolio)
```

Title:

```
MV Minimum Variance Portfolio
Estimator:      covEstimator
Solver:         solveRquadprog
```

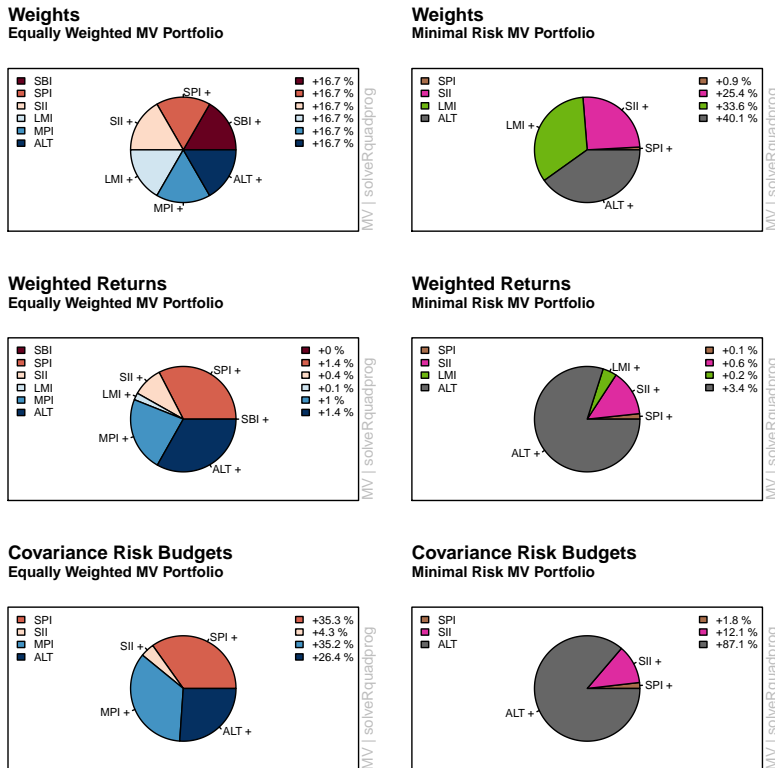


FIGURE 17.1: Weights, weighted returns, and covariance risk budgets plots for an equal weighted and a minimum variance portfolio: The equally weighted portfolio is shown to the left and the efficient portfolio with the same target return to the right. We have reduced the radius of the pies to 70% since we have a legend to the right and left. The legend to the left lists the assets and the legend to the right the percentual parts of the pie. The text to the right margin denotes the portfolio type, MV, and the solver, solveRquadprog, used for optimizing the portfolio.

```
Optimize:      minRisk
Constraints:    LongOnly
```

```
Portfolio Weights:
  SBI  SPI  SII  LMI  MPI  ALT
0.3555 0.0000 0.0890 0.4893 0.0026 0.0636
```

```
Covariance Risk Budgets:
  SBI  SPI  SII  LMI  MPI  ALT
0.3555 0.0000 0.0890 0.4893 0.0026 0.0636
```

```
Target Returns and Risks:
  mean  Cov  CVaR  VaR
0.0105 0.0986 0.2020 0.1558
```

Description:

Tue Jan 27 13:40:54 2015 by user: Rmetrics

Internally, the global minimum mean-variance portfolio is calculated by minimizing the efficient portfolio with respect to the target risk. This is a quadratic optimization problem with linear constraints.

```
> col <- seqPalette(ncol(lppData), "YlGn")
> weightsPie(globminPortfolio, box = FALSE, col = col)
> mtext(text = "Global Minimum Variance MV Portfolio", side = 3,
  line = 1.5, font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(globminPortfolio, box = FALSE, col = col)
> mtext(text = "Global Minimum Variance MV Portfolio", side = 3,
  line = 1.5, font = 2, cex = 0.7, adj = 0)
> covRiskBudgetsPie(globminPortfolio, box = FALSE, col = col)
> mtext(text = "Global Minimum Variance MV Portfolio", side = 3,
  line = 1.5, font = 2, cex = 0.7, adj = 0)
```

The pie plots for the global minimum mean-variance portfolio are shown in the left-hand column of [Figure 17.2](#). Compared to the previous pie plots in [Figure 17.1](#), we have chosen a different layout. The colours are taken from yellow to green sequential palettes, and the boxes around the pies have been suppressed.

17.4 HOW TO COMPUTE THE TANGENCY PORTFOLIO

The *tangency portfolio* is calculated by minimizing the Sharpe Ratio for a given risk-free rate. The Sharpe ratio is the ratio of the target return lowered by the risk-free rate and the covariance risk. The default risk-free rate is zero and can be reset to another value by modifying the portfolio's specification.

```
> tgSpec <- portfolioSpec()
> setRiskFreeRate(tgSpec) <- 0
```

The tangency portfolio is then obtained by calling the function `tangencyPortfolio()`

```
> tgPortfolio <- tangencyPortfolio(
  data = lppData,
  spec = tgSpec,
  constraints = "LongOnly")
> print(tgPortfolio)
```

Title:

MV Tangency Portfolio

Estimator: covEstimator

Solver: solveRquadprog

Optimize: minRisk

Constraints: LongOnly

Portfolio Weights:

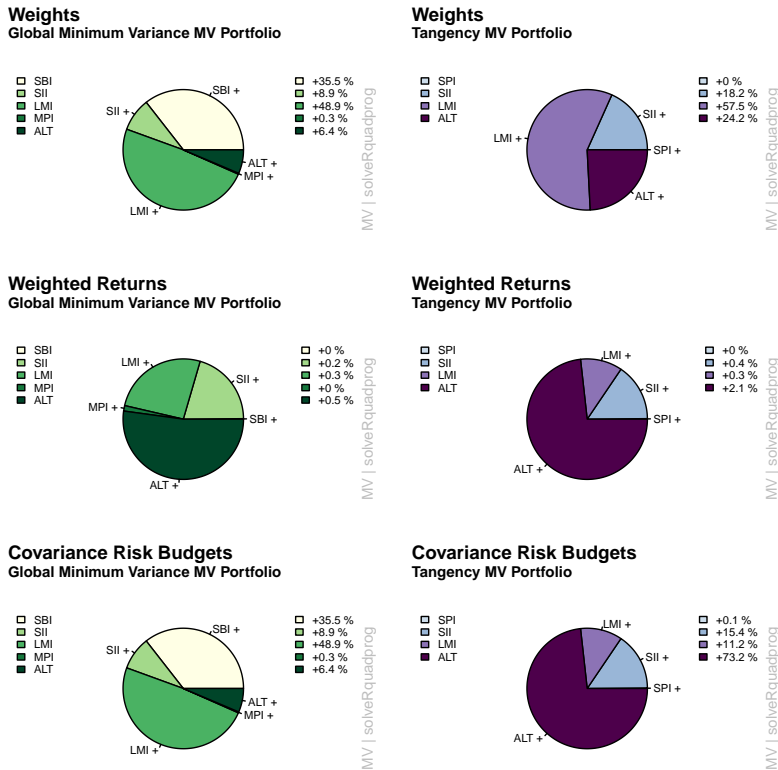


FIGURE 17.2: Weights, weighted returns, and covariance risk budget plots for the global minimum risk and the tangency portfolios: The global minimum risk portfolio is shown to the left and the tangency portfolio to the right. In this graph we have chosen an orange-to-red colour palette and removed the boxes from the pie charts.

```
SBI  SPI  SII  LMI  MPI  ALT
0.0000 0.0005 0.1824 0.5753 0.0000 0.2418
```

```
Covariance Risk Budgets:
SBI  SPI  SII  LMI  MPI  ALT
0.0000 0.0014 0.1539 0.1124 0.0000 0.7324
```

```
Target Returns and Risks:
mean  Cov  CVaR  VaR
0.0283 0.1533 0.3096 0.2142
```

```
Description:
Tue Jan 27 13:40:54 2015 by user: Rmetrics
```

and the pie plots are generated as in the previous examples, but this time using a blue to purple sequential palette.

```
> col <- seqPalette(ncol(lppData), "BuPu")
```

```

> weightsPie(tgPortfolio, box = FALSE, col = col)
> mtext(text = "Tangency MV Portfolio", side = 3, line = 1.5,
      font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(tgPortfolio, box = FALSE, col = col)
> mtext(text = "Tangency MV Portfolio", side = 3, line = 1.5,
      font = 2, cex = 0.7, adj = 0)
> covRiskBudgetsPie(tgPortfolio, box = FALSE, col = col)
> mtext(text = "Tangency MV Portfolio", side = 3, line = 1.5,
      font = 2, cex = 0.7, adj = 0)

```

The pie plots are shown in the right-hand column of [Figure 17.2](#).

17.5 HOW TO CUSTOMIZE A PIE PLOT

The functions `weightsPie()`, `weightedReturnsPie()`, and `covRiskBudgetsPie()` have with several arguments that allow you to customize the plots.

LISTING 17.2: FUNCTIONS TO PLOT PIE CHARTS OF PORTFOLIO WEIGHTS

Functions:	
<code>weightsPie</code>	displays the weights composition
<code>weightedReturnsPie</code>	displays weighted returns, the investment
<code>covRiskBudgetsPie</code>	displays the covariance risk budgets
Arguments:	
<code>object</code>	an S4 object of class <code>fPORTFOLIO</code>
<code>pos</code>	the point position on a whole frontier
<code>labels</code>	should the graph be labelled?
<code>col</code>	selects colour from a colour palette
<code>box</code>	should a box drawn around the pies?
<code>legend</code>	should a legend added to the pies?
<code>radius</code>	the radius of the pie
<code>...</code>	arguments to be passed

If you prefer a bar plot instead of a pie chart you can easily create it with R's base function `barplot()`. Here is an example of how to display the weights of the tangency portfolio optimized above as a horizontal bar chart.

```

> par(mfrow = c(2, 2))
> col <- rampPalette(ncol(lppData), "purple2green")
> weights <- 100 * as.vector(getWeights(tgPortfolio))
> weightedReturns <- weights * getMean(tgPortfolio)
> covRiskBudgets <- getCovRiskBudgets(tgPortfolio)
> names <- colnames(lppData)
> barplot(height = weights, names.arg = names, horiz = TRUE, las = 1, col = col)
> title(main = "Portfolio Weights", xlab = "Weights %")
> barplot(height = weightedReturns, names.arg = names, horiz = TRUE, las = 1, col = col)
> title(main = "Weighted Portfolio Returns", xlab = "Weighted Returns %")
> barplot(height = weights, names.arg = names, las = 1, col = col)

```





FIGURE 17.3: Weights, weighted returns, and covariance risk budget bar plots for the long-only tangency portfolio with zero risk-free rate. In this graph we have chosen a purple to green colour ramp palette and a horizontal display for the first two and a vertical display (the default) for the remaining two graphs.

```
> title(main = "Portfolio Weights", xlab = "Weights %")
> barplot(height = covRiskBudgets, names.arg = names, las = 1, col = col)
> title(main = "Covariance Risk Budgets", xlab = "Risk Budgets")
```

The bar plots of weights, weighted returns, and covariance risk budgets are shown in Figure 17.3. It is left to the reader to write his or her own functions for customized bar plots with the same argument list as for the pie plots.

CHAPTER 18

MEAN-VARIANCE PORTFOLIO FRONTIERS

```
> library(fPortfolio)
```

The efficient frontier together with the minimum variance locus form the ‘upper border’ and ‘lower border’ lines of the feasible set. To the right the feasible set is determined by the envelope of all pairwise asset frontiers. The region outside of the feasible set is unachievable by holding risky assets alone. No portfolios can be constructed corresponding to the points in this region. Points below the frontier are suboptimal. Thus, a rational investor will hold a portfolio only on the frontier. In this chapter we show how to compute the whole efficient frontier and minimum variance locus of a mean-variance portfolio with linear constraints¹ and show which functions can be used to display the results.

18.1 FRONTIER COMPUTATION AND GRAPHICAL DISPLAYS

The Rmetrics function `portfolioFrontier()` allows you to calculate optimized portfolios along the efficient frontier and the minimum variance locus. For the default settings with long-only constraints, the range spans all values equidistantly ranging from the asset with the lowest return up to the asset with the highest return. Allowing for box, group and other more complex constraints the range of the frontier will be downsized, i.e. the length of the frontier becomes shorter and shorter. Bear in mind that it is possible for the constraints to be too strong, and that a frontier might not even exist at all.

¹The case of non-linear constraints and the use of alternative solvers is described in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.

Many additional parameters can be set by the portfolio specification function, such as the number of frontier points.

LISTING 18.1: THIS TABLE LISTS FUNCTIONS TO COMPUTE THE EFFICIENT FRONTIER OF LINEARLY CONSTRAINED MEAN-VARIANCE PORTFOLIOS AND TO PLOT THE RESULTS.

Functions:	
portfolioFrontier	efficient portfolios on the frontier
frontierPoints	extracts risk/return frontier points
frontierPlot	creates an efficient frontier plot
cmlPoints	adds market portfolio
cmlLines	adds capital market line
tangencyPoints	adds tangency portfolio point
tangencyLines	adds tangency line
equalWeightsPoints	adds point of equal weights portfolio
singleAssetPoints	adds points of single asset portfolios
twoAssetsLines	adds frontiers of two assets portfolios
sharpeRatioLines	adds Sharpe ratio line
monteCarloPoints	adds randomly feasible portfolios
weightsPlot	weights bar plot along the frontier
weightedReturnsPlot	weighted returns bar plot
covRiskBudgetsPlot	covariance risk budget bar plot

How to compute the efficient frontier

The computation of the efficient frontier for the default MV portfolio just requires a few function calls. As a first example, we compute the efficient frontier for the six assets included in Pictet's pension fund benchmark portfolio LPP2005.RET. We multiply the series by 100 to convert them to returns in percentages.

```
> lppData <- 100 * LPP2005.RET[, 1:6]
> colnames(lppData)
[1] "SBI" "SPI" "SII" "LMI" "MPI" "ALT"
```

Let us compute the efficient frontier for example on 5 points.

```
> lppSpec <- portfolioSpec()
> setNFrontierPoints(lppSpec) <- 5
> longFrontier <- portfolioFrontier(lppData, lppSpec)
```

How to print an efficient frontier report

The printout of the S4 frontier object returned by the portfolioFrontier() function lists the major parameter settings, the portfolio weights, the covariance risk budgets, and the target return and risk values for each point along the frontier.

```
> print(longFrontier)
Title:
MV Portfolio Frontier
Estimator:      covEstimator
Solver:         solveRquadprog
Optimize:       minRisk
Constraints:     LongOnly
Portfolio Points: 5 of 5

Portfolio Weights:
      SBI   SPI   SII   LMI   MPI   ALT
1 1.0000 0.0000 0.0000 0.0000 0.0000 0.0000
2 0.0193 0.0000 0.1481 0.6665 0.0000 0.1661
3 0.0000 0.0085 0.2535 0.3386 0.0000 0.3994
4 0.0000 0.0210 0.3458 0.0000 0.0000 0.6332
5 0.0000 0.0000 0.0000 0.0000 0.0000 1.0000

Covariance Risk Budgets:
      SBI   SPI   SII   LMI   MPI   ALT
1 1.0000 0.0000 0.0000 0.0000 0.0000 0.0000
2 0.0064 0.0000 0.1593 0.3359 0.0000 0.4984
3 0.0000 0.0183 0.1208 -0.0097 0.0000 0.8707
4 0.0000 0.0286 0.0890 0.0000 0.0000 0.8824
5 0.0000 0.0000 0.0000 0.0000 0.0000 1.0000

Target Returns and Risks:
      mean   Cov   CVaR   VaR
1 0.0000 0.1261 0.2758 0.2177
2 0.0215 0.1214 0.2362 0.1760
3 0.0429 0.2439 0.5275 0.3392
4 0.0643 0.3939 0.8822 0.5886
5 0.0858 0.5684 1.3343 0.8978

Description:
Tue Jan 27 13:40:38 2015 by user: Rmetrics
```

How to interactively plot the efficient frontier

Several plotting facilities are available to display the efficient frontier. For a quick interactive overview you can use the generic `plot()` function offering the following selections:

```
> longFrontier <- portfolioFrontier(lppData)
> plot(longFrontier)
```

Make a plot selection (or 0 to exit):

```
1: Plot Efficient Frontier
2: Add Minimum Risk Portfolio
3: Add Tangency Portfolio
4: Add Risk/Return of Single Assets
5: Add Equal Weights Portfolio
```

```

6:  Add Two Asset Frontiers [0-1 PF Only]
7:  Add Wheel Pie of Weights
8:  Add Monte Carlo Portfolios
9:  Add Sharpe Ratio [MV PF Only]

```

Selection:

As an example, recalculate the frontier with the default number of frontier points (50), then call the interactive `plot()` function and add some optional graphs.

How to create a customized efficient frontier plot

For customized plots the function `frontierPlot()` with several add-on plot functions can be used.

LISTING 18.2: FUNCTIONS TO DISPLAY THE EFFICIENT FRONTIER.

Functions:	
<code>frontierPlot</code>	efficient frontier plot
<code>cmlPoints</code>	adds market portfolio
<code>cmlLines</code>	adds capital market line
<code>tangencyPoints</code>	adds tangency portfolio point
<code>tangencyLines</code>	adds tangency line
<code>equalWeightsPoints</code>	adds point of equal weights portfolio
<code>singleAssetPoints</code>	adds points of single asset portfolios
<code>twoAssetsLines</code>	adds frontiers of two assets portfolios
<code>sharpeRatioLines</code>	adds Sharpe ratio line
<code>monteCarloPoints</code>	adds randomly feasible portfolios
Arguments:	
<code>object</code>	an S4 object of class <code>fPORTFOLIO</code>
<code>frontier</code>	which frontier part should be plotted?
<code>col</code>	colours for the EF and the MV locus
<code>add</code>	should another frontier added?
<code>return</code>	select from 'mean' or 'mu'
<code>risk</code>	select from 'cov', 'sigma', 'VaR', 'CVaR'
<code>auto</code>	automatic risk/return selection
<code>labels</code>	should the plot be labelled?
<code>title</code>	should a default title be added?
<code>mcSteps</code>	number of Monte Carlo steps
<code>xlim, ylim</code>	set the plot range
<code>mText</code>	marginal text to be added
<code>...</code>	optional arguments to be passed

These functions allow you to create your own customized plotting function, allowing you to create your own presentation of the frontier. The Rmetrics `fPortfolio` package already has such a function, named `taioredFrontierPlot()`, which can be used as a starting point for your customized frontier plot function.

```
> args(tailoredFrontierPlot)
function (object, return = c("mean", "mu"), risk = c("Cov", "Sigma",
  "CVaR", "VaR"), mText = NULL, col = NULL, xlim = NULL, ylim = NULL,
  twoAssets = FALSE, sharpeRatio = TRUE, title = TRUE, ...)
NULL
```

If you want to write your own display function, it is a good idea to also inspect the code of the function.

```
> tailoredFrontierPlot
```

The result of calling the `tailoredFrontierPlot()` is shown in [Figure 18.1](#), now re-calculated on 25 frontier points.

```
> setNFrontierPoints(lppSpec) <- 25
> longFrontier <- portfolioFrontier(lppData, lppSpec)
> tailoredFrontierPlot(object = longFrontier, mText = "MV Portfolio - LongOnly Constraints",
  risk = "Cov")
```

How to create weights and related plots

Furthermore, bar and line plots for the weights, `weightsPlot()`, the performance attribution, `weightedReturnsPlot()`, and the covariance risk budgets, `covRiskBudgetsPlot()`, are available. Note that these plots can also be displayed as line plots.

```
> weightsPlot(longFrontier)
> text <- "Mean-Variance Portfolio - Long Only Constraints"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(longFrontier)
> covRiskBudgetsPlot(longFrontier)
```

18.2 THE 'LONG-ONLY' PORTFOLIO FRONTIER

The long-only constraints are the default constraints for the mean-variance portfolios. Remember that in this case all the weights are bounded between zero and one. In the previous example we optimized the default portfolio, and the results are shown in [Figure 18.1](#) for the frontier and in [Figure 18.2](#) for the weights.

Now we want to explore in more detail the feasible set of the long-only constrained mean-variance portfolio. For this, we plot the frontier, add randomly generated portfolios from a Monte Carlo simulation, and add the frontier lines of all two-asset portfolios.

```
> par(mfrow = c(1, 1))
> set.seed(1953)
> frontierPlot(object = longFrontier, pch = 19, xlim = c(0.05,
```




FIGURE 18.1: Efficient frontier of a long-only constrained mean-variance portfolio: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

```
0.85), cex = 0.5)
> monteCarloPoints(object = longFrontier, mcSteps = 1000, pch = 19,
  cex = 0.5)
> twoAssetsLines(object = longFrontier, col = "orange", lwd = 2)
> frontier <- frontierPoints(object = longFrontier)
> lines(frontier, col = "red", lwd = 2)
```

Note that the Monte Carlo simulation provided by the function `monteCarloPoints()` is only meaningful for long-only constraints.

18.3 THE UNLIMITED 'SHORT' PORTFOLIO FRONTIER

If all the weights are not restricted, we have the case of unlimited short selling. Since unlimited short selling portfolios can be solved analytically, we

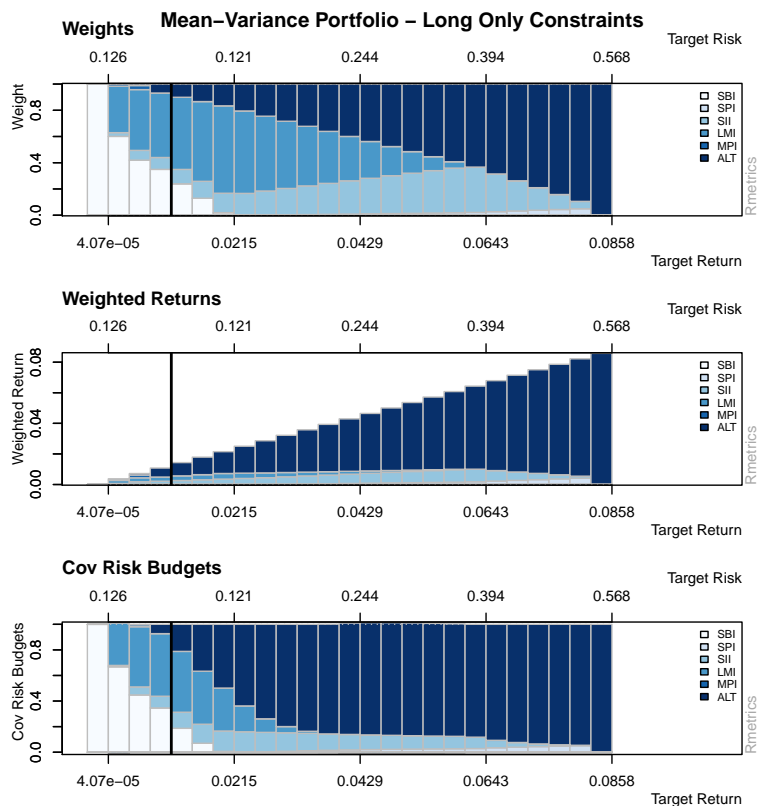


FIGURE 18.2: Weights along the efficient frontier of a long-only constrained mean-variance portfolio: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars.

can replace the solver "solveRquadprog" with the solver "solveRshort-Exact"

```
> shortSpec <- portfolioSpec()
> setNFrontierPoints(shortSpec) <- 5
> setSolver(shortSpec) <- "solveRshortExact"
> shortFrontier <- portfolioFrontier(
  data = lppData,
  spec = shortSpec,
  constraints = "Short")
> print(shortFrontier)
```

```
Title:
MV Portfolio Frontier
Estimator:      covEstimator
Solver:         solveRshortExact
```




FIGURE 18.3: The feasible set for a long-only constrained mean-variance portfolio: The graph shows the risk/return plot for 1000 randomly generated mean-variance portfolios with long-only constraints. The plot is overlayed by the efficient frontier, the minimum variance locus, and the pairwise frontier lines of all combinations of two-asset portfolios. The corners of the lines coincide with the risk/return values for the six assets.

Optimize: minRisk
Constraints: Short
Portfolio Points: 5 of 5

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.5348	-0.0310	0.0493	0.4245	0.1054	-0.0830
2	0.1560	0.0210	0.1337	0.5648	-0.1013	0.2258
3	-0.2227	0.0730	0.2181	0.7051	-0.3081	0.5346
4	-0.6014	0.1250	0.3025	0.8453	-0.5149	0.8435
5	-0.9801	0.1769	0.3869	0.9856	-0.7216	1.1523

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.5348	0.0267	0.0233	0.3730	-0.0322	0.0744
2	0.0788	0.0513	0.1412	0.3569	-0.1893	0.5610
3	-0.0038	0.1420	0.1230	0.1007	-0.4222	1.0602

```

4 0.0360 0.1658 0.1008 0.0260 -0.4700 1.1414
5 0.0674 0.1734 0.0885 -0.0003 -0.4813 1.1523

```

Target Returns and Risks:

```

      mean    Cov   CVaR   VaR
1 0.0000 0.1121 0.2374 0.1921
2 0.0215 0.1144 0.2187 0.1710
3 0.0429 0.1962 0.3790 0.2712
4 0.0643 0.2979 0.5912 0.4078
5 0.0858 0.4048 0.8123 0.5374

```

Description:

Tue Jan 27 13:40:40 2015 by user: Rmetrics

For the plot of the frontier we reset the number of frontier points to 20 and recalculate the frontier.

```

> setNFrontierPoints(shortSpec) <- 20
> shortFrontier <- portfolioFrontier(data = lppData, spec = shortSpec,
  constraints = "Short")
> tailoredFrontierPlot(object = shortFrontier, mText = "MV Portfolio - Short Constraints",
  risk = "Cov")

> weightsPlot(shortFrontier)
> text <- "MV Portfolio - Short Constrained Portfolio"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(shortFrontier)
> covRiskBudgetsPlot(shortFrontier)

```

The results are shown in [Figure 18.4](#) and [Figure 18.5](#).

18.4 THE BOX-CONSTRAINED PORTFOLIO FRONTIER

A box-constrained portfolio is a portfolio where the weights are constrained by lower and upper bounds, e.g. we want to invest at least in each asset 1% and no more than 50%.

```

> boxSpec <- portfolioSpec()
> setNFrontierPoints(boxSpec) <- 15
> boxConstraints <- c(
  "minW[1:6]=0.01",
  "maxW[1:6]=0.5")
> boxFrontier <- portfolioFrontier(
  data = lppData,
  spec = boxSpec,
  constraints = boxConstraints)
> print(boxFrontier)

```

Title:

MV Portfolio Frontier

Estimator: covEstimator

Solver: solveRquadprog

Optimize: minRisk




FIGURE 18.4: Efficient frontier of an unlimited short selling constrained mean-variance portfolio: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

Constraints: minW maxW
Portfolio Points: 5 of 13

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.4986	0.0100	0.0538	0.4128	0.0148	0.0100
4	0.1114	0.0100	0.1824	0.5000	0.0100	0.1862
7	0.0100	0.0100	0.2540	0.3242	0.0100	0.3919
10	0.0100	0.1081	0.3619	0.0100	0.0100	0.5000
13	0.0100	0.4128	0.0572	0.0100	0.0100	0.5000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.5513	0.0042	0.0344	0.4120	-0.0007	-0.0012
4	0.0249	0.0329	0.1882	0.1353	0.0360	0.5828
7	-0.0003	0.0213	0.1191	-0.0104	0.0248	0.8455

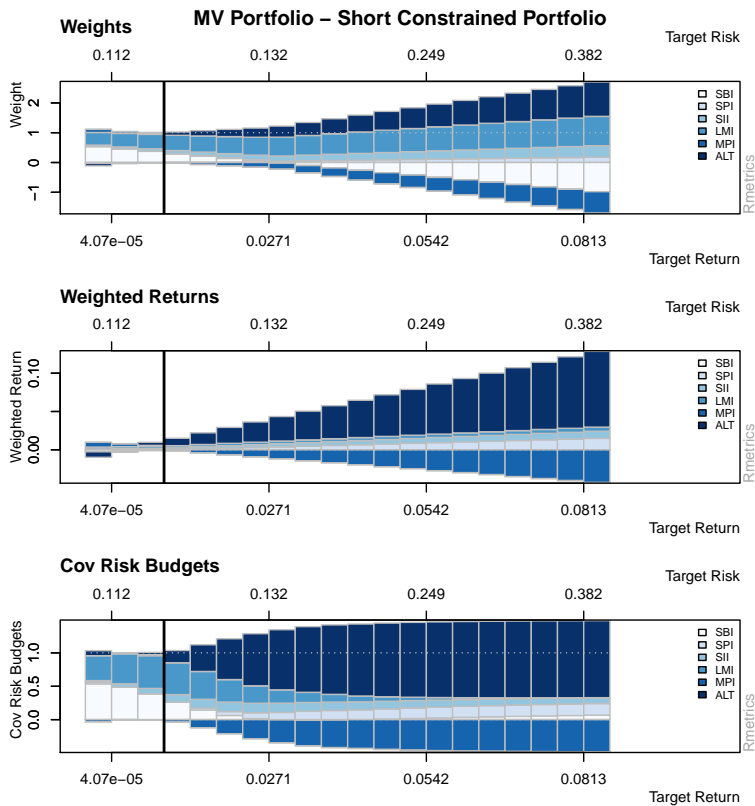


FIGURE 18.5: Weights along the efficient frontier of an unlimited short selling constrained mean-variance portfolio: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars.

```
10 -0.0005 0.1712 0.1063 -0.0007 0.0168 0.7069
13 -0.0004 0.5228 0.0047 -0.0005 0.0115 0.4619
```

Target Returns and Risks:

	mean	Cov	CVaR	VaR
1	0.0062	0.1023	0.2141	0.1653
4	0.0245	0.1380	0.2780	0.1928
7	0.0429	0.2469	0.5373	0.3395
10	0.0613	0.3787	0.8674	0.5227
13	0.0796	0.5593	1.3953	0.7938

Description:

Tue Jan 27 13:40:40 2015 by user: Rmetrics




FIGURE 18.6: Efficient frontier of a box-constrained mean-variance portfolio: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP; all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

For the plot of the frontier we reset the number of frontier points to 25 and recalculate the frontier.

```
> setNFrontierPoints(boxSpec) <- 25
> boxFrontier <- portfolioFrontier(data = lppData, spec = boxSpec,
  constraints = boxConstraints)
> tailoredFrontierPlot(object = boxFrontier, mText = "MV Portfolio - Box Constraints",
  risk = "Cov")
```

The efficient frontier of the box-constrained MV portfolio is shown in [Figure 18.6](#). The weights, weighted returns and covariance risk budgets are shown in the left-hand column of [Figure 18.8](#).

```
> weightsPlot(boxFrontier)
> weightedReturnsPlot(boxFrontier)
> covRiskBudgetsPlot(boxFrontier)
```

18.5 THE GROUP-CONSTRAINED PORTFOLIO FRONTIER

A group-constrained portfolio is a portfolio where the weights of groups of selected assets are constrained by lower and upper bounds for the total weights of the groups, e.g. we want to invest at least in the group of bonds 30% and no more than 50% in the groups of assets.

```
> groupSpec <- portfolioSpec()
> setNFrontierPoints(groupSpec) <- 7
> groupConstraints <- c("minsumW[c(1,4)]=0.3",
                        "maxsumW[c(2,5)]=0.5")
> groupFrontier <- portfolioFrontier(
  data = lppData,
  spec = groupSpec,
  constraints = groupConstraints)
> print(groupFrontier)
```

Title:

MV Portfolio Frontier

Estimator: covEstimator

Solver: solveRquadprog

Optimize: minRisk

Constraints: minsumW maxsumW

Portfolio Points: 5 of 5

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.2394	0.0000	0.1097	0.5500	0.0000	0.1009
3	0.0000	0.0006	0.1838	0.5705	0.0000	0.2450
4	0.0000	0.0085	0.2535	0.3386	0.0000	0.3994
5	0.0000	0.0284	0.0721	0.3000	0.0000	0.5995

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.1869	0.0000	0.1258	0.4757	0.0000	0.2117
3	0.0000	0.0019	0.1532	0.1061	0.0000	0.7388
4	0.0000	0.0183	0.1208	-0.0097	0.0000	0.8707
5	0.0000	0.0445	0.0097	-0.0150	0.0000	0.9609

Target Returns and Risks:

	mean	Cov	CVaR	VaR
1	0.0000	0.1261	0.2758	0.2177
2	0.0143	0.1016	0.2002	0.1534
3	0.0286	0.1549	0.3136	0.2170
4	0.0429	0.2439	0.5275	0.3392
5	0.0572	0.3516	0.8182	0.5552

Description:

Tue Jan 27 13:40:40 2015 by user: Rmetrics

For the plot of the frontier we reset the number of frontier points to 25 and recalculate the frontier.

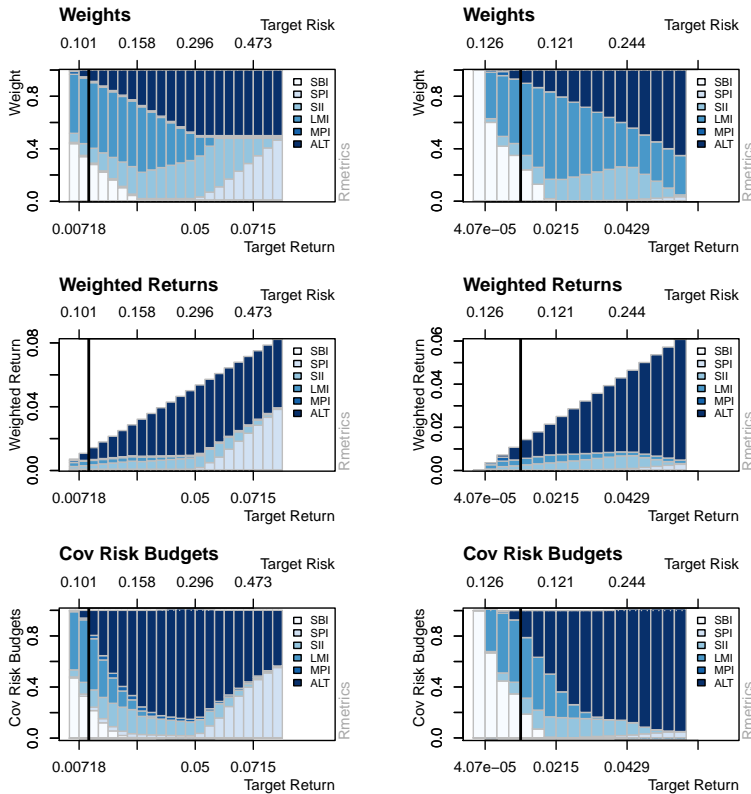


FIGURE 18.8: MV box (left) and group (right) constrained weights, weighted returns, and covariance risk budgets plots along the frontier.

18.6 THE BOX/GROUP-CONSTRAINED PORTFOLIO FRONTIER

Box and group constraints can be combined

```
> boxgroupSpec <- portfolioSpec()
> setNFrontierPoints(boxgroupSpec) <- 15
> boxgroupConstraints <- c(boxConstraints,
                           groupConstraints)
> boxgroupFrontier <- portfolioFrontier(
  data = lppData,
  spec = boxgroupSpec,
  constraints = boxgroupConstraints)
> print(boxgroupFrontier)
```

Title:

MV Portfolio Frontier

Estimator: covEstimator

Solver: solveRquadprog

Optimize: minRisk

Constraints: minW maxW minsumW maxsumW

Portfolio Points: 5 of 9

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.4986	0.0100	0.0538	0.4128	0.0148	0.0100
3	0.2126	0.0100	0.1411	0.5000	0.0100	0.1263
5	0.0102	0.0100	0.2236	0.5000	0.0100	0.2462
7	0.0100	0.0100	0.2540	0.3242	0.0100	0.3919
9	0.0100	0.0918	0.0982	0.2900	0.0100	0.5000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.5513	0.0042	0.0344	0.4120	-0.0007	-0.0012
3	0.1090	0.0329	0.1692	0.2783	0.0343	0.3762
5	0.0007	0.0290	0.1867	0.0552	0.0325	0.6959
7	-0.0003	0.0213	0.1191	-0.0104	0.0248	0.8455
9	-0.0004	0.1630	0.0165	-0.0141	0.0191	0.8159

Target Returns and Risks:

	mean	Cov	CVaR	VaR
1	0.0062	0.1023	0.2141	0.1653
3	0.0184	0.1132	0.2262	0.1632
5	0.0307	0.1692	0.3474	0.2245
7	0.0429	0.2469	0.5373	0.3395
9	0.0552	0.3403	0.8076	0.5211

Description:

Tue Jan 27 13:40:41 2015 by user: Rmetrics

For the plot of the frontier we reset the number of frontier points to 25 and recalculate the frontier.

```
> boxgroupSpec <- portfolioSpec()
> setNFrontierPoints(boxgroupSpec) <- 25
> boxgroupFrontier <- portfolioFrontier(
  data = lppData,
  spec = boxgroupSpec,
  constraints = boxgroupConstraints)
> tailoredFrontierPlot(
  object = boxgroupFrontier,
  mText = "MV Portfolio - Box/Group Constraints",
  risk = "Cov")
```

The results of combining box and group constraints are shown in [Figure 18.9](#) and [Figure 18.10](#).

```
> weightsPlot(boxgroupFrontier)
> text <- "MV Portfolio - Box/Group Constrained Portfolio"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(boxgroupFrontier)
> covRiskBudgetsPlot(boxgroupFrontier)
```




FIGURE 18.9: Efficient frontier of a box/group constrained mean-variance portfolio: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

18.7 CREATING DIFFERENT 'REWARD/RISK VIEWS' ON THE EFFICIENT FRONTIER

In the efficient frontier plots we have plotted the target return as a function of the covariance risk, expressed as the standard deviation. We can now ask what the efficient frontier looks like when we plot the sample mean versus the conditional Value-at-Risk. The `frontierPlot()` and `add-on` functions allow you to change the view specifying the arguments for the return and the risk in the `frontierPlot()` function.

As an example let us plot the efficient frontier for the sample mean return versus the covariance, the CVaR and VaR risk measures. Note that if we specify a risk/reward measure, we have to set the argument `auto=FALSE`, explicitly.

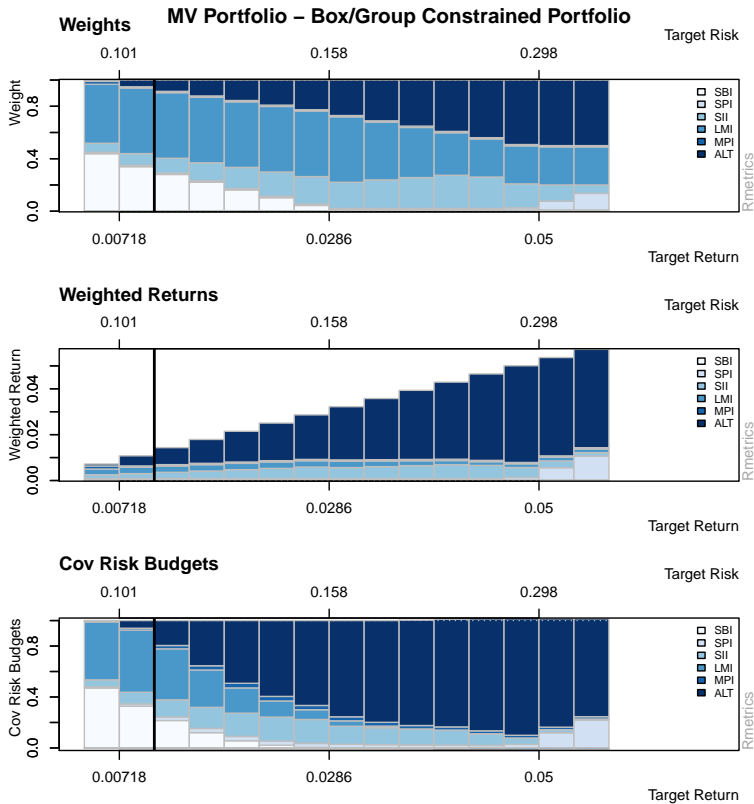


FIGURE 18.10: Weights along the efficient frontier of a mixed box/group constrained mean-variance portfolio: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars.

```
> frontierPlot(longFrontier, auto = TRUE)
> frontierPlot(longFrontier, return = "mean", risk = "Cov",
  auto = FALSE)
> frontierPlot(longFrontier, return = "mean", risk = "CVaR",
  auto = FALSE)
> frontierPlot(longFrontier, return = "mean", risk = "VaR",
  auto = FALSE)
```

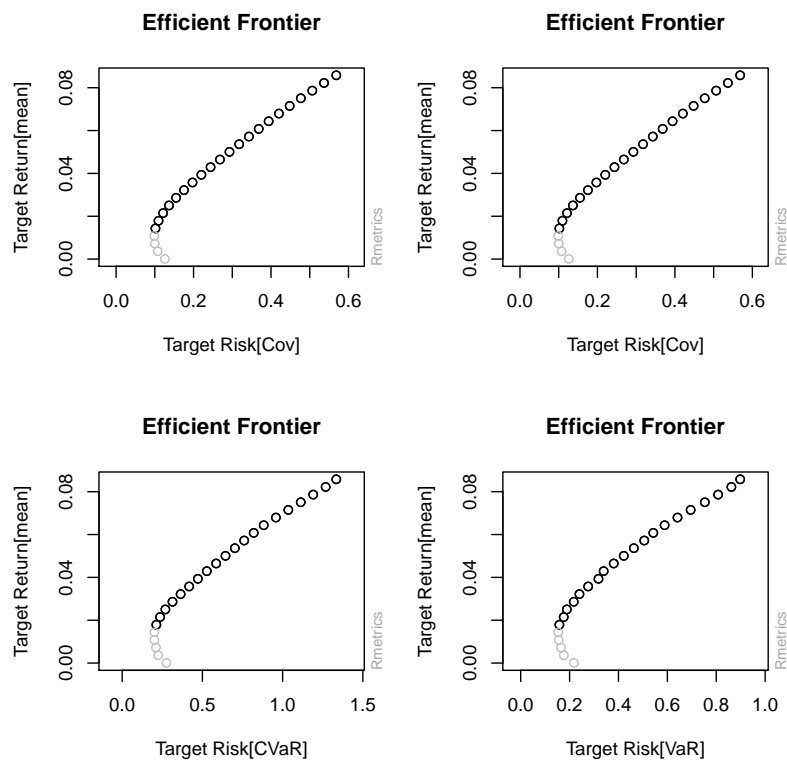


FIGURE 18.11: MV long-only constrained frontier plots for different risk measures: *Upper left*, the default, risk and reward type are selected automatically from the portfolio specifications, *upper right*, the same graph, the mean plotted versus the covariance risk, *lower left*, now the mean return plotted versus the conditional value at risk, and *lower right*, now the mean return plotted versus the value-at-risk.

CHAPTER 19

CASE STUDY: DOW JONES INDEX

```
> library(fPortfolio)
```

In this chapter we have prepared a real-world case study for optimizing a portfolio with the 30 shares given in the Dow Jones Index. We use the DowJones30 data set provided in the fBasics Package.

1. Load the data set

```
> djiData <- as.timeSeries(DowJones30)
> djiData.ret <- 100 * returns(djiData)
> colnames(djiData)
 [1] "AA"   "AXP"  "T"    "BA"   "CAT"  "C"    "KO"   "DD"   "EK"   "XOM"
[11] "GE"   "GM"   "HWP"  "HD"   "HON"  "INTC" "IBM"  "IP"   "JPM"  "JNJ"
[21] "MCD"  "MRK"  "MSFT" "MMM"  "MO"   "PG"   "SBC"  "UTX"  "WMT"  "DIS"
> c(start(djiData), end(djiData))
GMT
[1] [1990-12-31] [2001-01-02]
```

The data cover 10 years of daily data. If you would like to use more recent data, please feel free to update the data from [Yahoo Finance](http://finance.yahoo.com/)¹.

2. Perform an exploratory data analysis Explore the returns series, and the series of share prices. Then investigate pairwise dependencies between the asset returns, including correlations and distributional properties from star plots. Which of the shares are similar or dissimilar? Use hierarchical clustering and a PCA analysis of the equities.

¹<http://finance.yahoo.com/>

```

> for (i in 1:3) plot(djiData.ret[, (10 * i - 9):(10 * i)])
> for (i in 1:3) plot(djiData[, (10 * i - 9):(10 * i)])
> assetsCorImagePlot(djiData.ret)
> plot(assetsSelect(djiData.ret))
> assetsCorEigenPlot(djiData.ret)

```

3. Find the optimal weights for a long-only MV portfolio Apply the mean-variance portfolio approach to explore the efficient frontier and to display the weights along the frontier.

```

> frontier <- portfolioFrontier(djiData.ret)
> tailoredFrontierPlot(frontier)
> weightsPlot(frontier)

```

4. Find the optimal weights for a group-constrained MV portfolio Perform a clustering of the equities, grouping the data into 5 clusters. Limit the investment for each cluster to a maximum of 50%.

```

> selection <- assetsSelect(djiData.ret, method = "kmeans")
> cluster <- selection$cluster
> cluster[cluster == 1]

BA  EK  HON  MMM  UTX
1   1   1   1   1

> cluster[cluster == 2]

T  KO  XOM  GE  JNJ  MCD  MRK  MO  PG  SBC  DIS
2   2   2   2   2   2   2   2   2   2   2

> cluster[cluster == 3]

HWP  INTC  IBM  MSFT
3   3   3   3

> cluster[cluster == 4]

AXP  C  HD  JPM  WMT
4   4   4   4   4

> cluster[cluster == 5]

AA  CAT  DD  GM  IP
5   5   5   5   5

> constraints <- c(
  'maxsumW[c("BA", "DD", "EK", "XOM", "GM", "HON", "MMM", "UTX")] = 0.30',
  'maxsumW[c("T", "KO", "GE", "HD", "JNJ", "MCD", "MRK", "MO", "PG", "SBC", "WMT", "DIS")] = 0.30',
  'maxsumW[c("AXP", "C", "JPM")] = 0.30',
  'maxsumW[c("AA", "CAT", "IP")] = 0.30',
  'maxsumW[c("HWP", "INTC", "IBM", "MSFT")] = 0.30')

```

Estimate the covariance matrix using the shrinkage estimator and compute the weights along the frontier. The weights are shown in [Figure 19.1](#).




FIGURE 19.1: The graph shows the weights along the DJI mean-variance frontier. The legend to the right links the equity names to colour of the bars.

```
> djiSpec <- portfolioSpec()
> setNFrontierPoints(djiSpec) <- 25
> setEstimator(djiSpec) <- "shrinkEstimator"
> djiFrontier <- portfolioFrontier(djiData.ret, djiSpec)
> col = seqPalette(30, "YlOrRd")
> weightsPlot(djiFrontier, col = col)
```


CHAPTER 20

ROBUST PORTFOLIOS AND COVARIANCE ESTIMATION

```
> library(fPortfolio)
```

Mean-variance portfolios constructed using the sample mean and covariance matrix of asset returns often perform poorly out-of-sample due to estimation errors in the mean vector and covariance matrix. As a consequence, minimum-variance portfolios may yield unstable weights that fluctuate substantially over time. This loss of stability may also lead to extreme portfolio weights and dramatic swings in weights with only minor changes in expected returns or the covariance matrix. Consequentially, we observe frequent re-balancing and excessive transaction costs.

To achieve better stability properties compared to traditional minimum-variance portfolios, we try to reduce the estimation error using robust methods to compute the mean and/or covariance matrix of the set of financial assets. Two different approaches are implemented: robust mean and covariance estimators, and the shrinkage estimator¹.

If the number of time series records is small and the number of considered assets increases, then the sample estimator of covariance becomes more and more unstable. Specifically, it is possible to provide estimators that improve considerably upon the maximum likelihood estimate in terms of mean-squared error. Moreover, when the number of records is smaller than the number of assets, the empirical estimate of the covariance matrix becomes singular.

¹For further information, we recommend the text book by [Marazzi \(1993\)](#)

20.1 ROBUST MEAN AND COVARIANCE ESTIMATORS

In the mean-variance portfolio approach, the sample mean and sample covariance estimators are used by default to estimate the mean vector and covariance matrix.

This information, i.e. the name of the covariance estimator function, is kept in the specification structure and can be shown by calling the function `getEstimator()`. The default setting is

```
> getEstimator(portfolioSpec())
[1] "covEstimator"
```

There are many different implementations of robust and related estimators for the mean and covariance in R's base packages and in contributed packages. The estimators listed below can be accessed by the portfolio optimization program.

LISTING 20.1: RMETRICS FUNCTIONS TO ESTIMATE ROBUST COVARIANCES FOR PORTFOLIO OPTIMIZATION

Functions:	
<code>covEstimator</code>	Covariance sample estimator
<code>kendallEstimator</code>	Kendall's rank estimator
<code>spearmanEstimator</code>	Spearman's rank estimator
<code>mcdEstimator</code>	MCD, minimum covariance determinant estimator
<code>mveEstimator</code>	MVE, minimum volume ellipsoid estimator
<code>covMcdEstimator</code>	Minimum covariance determinant estimator
<code>covOGKEstimator</code>	Orthogonalized Gnanadesikan-Kettenring estimator
<code>shrinkEstimator</code>	Shrinkage covariance estimator
<code>baggedEstimator</code>	Bagged covariance estimator

20.2 THE MCD ROBUSTIFIED MEAN-VARIANCE PORTFOLIO

The *minimum covariance determinant*, MCD, estimator of location and scatter looks for the $h > n/2$ observations out of n data records whose classical covariance matrix has the lowest possible determinant. The raw MCD estimate of location is then the average of these h points, whereas the raw MCD estimate of scatter is their covariance matrix, multiplied by a consistency factor and a finite sample correction factor (to make it consistent with the normal model and unbiased for small sample sizes). The algorithm from the MASS library is quite slow, whereas the one from contributed package `robustbase` (Rousseeuw et al., 2008) is much more time-efficient. The implementation in `robustbase` uses the fast MCD algorithm of Rousseeuw & Van Driessen (1999). To optimize a Markowitz mean-variance portfolio, we just have to specify the name of the mean/covariance estimator function. Unfortunately, this can take some time since

we have to apply the MCD estimator in every instance when we call the function `covMcdEstimator()`. To circumvent this, we perform the covariance estimation only once at the very beginning, store the value globally, and use its estimate in the new function `fastCovMcdEstimator()`.

```
> lppData <- 100 * LPP2005.RET[, 1:6]
> covMcdEstimate <- covMcdEstimator(lppData)
> fastCovMcdEstimator <-
  function(x, spec = NULL, ...)
    covMcdEstimate
```

Next we define the portfolio specification

```
> covMcdSpec <- portfolioSpec()
> setEstimator(covMcdSpec) <- "fastCovMcdEstimator"
> setNFrontierPoints(covMcdSpec) <- 5
```

and optimize the MCD robustified portfolio (with long-only default constraints).

```
> covMcdFrontier <- portfolioFrontier(
  data = lppData, spec = covMcdSpec)
> print(covMcdFrontier)
```

```
Title:
MV Portfolio Frontier
Estimator:      fastCovMcdEstimator
Solver:         solveRquadprog
Optimize:       minRisk
Constraints:     LongOnly
Portfolio Points: 5 of 5
```

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.1379	0.0377	0.1258	0.5562	0.0000	0.1424
3	0.0000	0.0998	0.2088	0.3712	0.0000	0.3202
4	0.0000	0.1661	0.2864	0.0430	0.0000	0.5046
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0492	0.1434	0.1209	0.2452	0.0000	0.4413
3	0.0000	0.2489	0.0878	-0.0071	0.0000	0.6704
4	0.0000	0.2624	0.0660	-0.0027	0.0000	0.6743
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Target Returns and Risks:

	mean	mu	Cov	Sigma	CVaR	VaR
1	0.0000	0.0000	0.1261	0.1304	0.2758	0.2177
2	0.0215	0.0215	0.1242	0.1153	0.2552	0.1733
3	0.0429	0.0429	0.2493	0.2117	0.5698	0.3561
4	0.0643	0.0643	0.4023	0.3363	0.9504	0.5574
5	0.0858	0.0858	0.5684	0.5016	1.3343	0.8978

Description:

Tue Jan 27 13:40:58 2015 by user: Rmetrics

Note that for the Swiss Pension Fund benchmark data set the "covMcdEstimator" is about 20 time slower than the sample covariance estimator, and the "mcdEstimator" is even slower by a factor of about 300.

For the plot we recalculate the frontier on 20 frontier points.

```
> setNFrontierPoints(covMcdSpec) <- 20
> covMcdFrontier <- portfolioFrontier(
  data = lppData, spec = covMcdSpec)
> tailoredFrontierPlot(
  covMcdFrontier,
  mText = "MCD Robustified MV Portfolio",
  risk = "Sigma")
```

The frontier plot is shown in [Figure 20.1](#).

To display the weights, risk attributions and covariance risk budgets for the MCD robustified portfolio in the left-hand column and the same plots for the sample covariance MV portfolio in the right-hand column of a figure:

```
> ## MCD robustified portfolio
> par(mfcol = c(3, 2), mar = c(3.5, 4, 4, 3) + 0.1)
> col = qualiPalette(30, "Dark2")
> weightsPlot(covMcdFrontier, col = col)
> text <- "MCD"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(covMcdFrontier, col = col)
> covRiskBudgetsPlot(covMcdFrontier, col = col)
> ## Sample covariance MV portfolio
> longSpec <- portfolioSpec()
> setNFrontierPoints(longSpec) <- 20
> longFrontier <- portfolioFrontier(data = lppData, spec = longSpec)
> col = qualiPalette(30, "Set1")
> weightsPlot(longFrontier, col = col)
> text <- "COV"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(longFrontier, col = col)
> covRiskBudgetsPlot(longFrontier, col = col)
```

The weights, risk attributions and covariance risk budgets are shown in [Figure 20.2](#).

20.3 THE MVE ROBUSTIFIED MEAN-VARIANCE PORTFOLIO

[Rousseeuw & Leroy \(1987\)](#) proposed a very robust alternative to classical estimates of mean vectors and covariance matrices, the Minimum Volume Ellipsoid, MVE. Samples from a multivariate normal distribution form ellipsoid-shaped 'clouds' of data points. The MVE corresponds to the smallest point cloud containing at least half of the observations, the




FIGURE 20.1: Efficient frontier of a long-only constrained mean-variance portfolio with robust MCD covariance estimates: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

uncontaminated portion of the data. These ‘clean’ observations are used for preliminary estimates of the mean vector and the covariance matrix. Using these estimates, the program computes a robust Mahalanobis distance for every observation vector in the sample. Observations for which the robust Mahalanobis distances exceed the 97.5% significance level for the chi-square distribution are flagged as probable outliers.

Rmetrics provides a function, `mveEstimator()`, to compute the MVE estimator; it is based on the `cov.rob()` estimator from the MASS package. We define a function called `fastMveEstimator()`

```
> mveEstimate <- mveEstimator(lppData)
> fastMveEstimator <- function(x, spec = NULL, ...) mveEstimate
```

and set the portfolio specifications

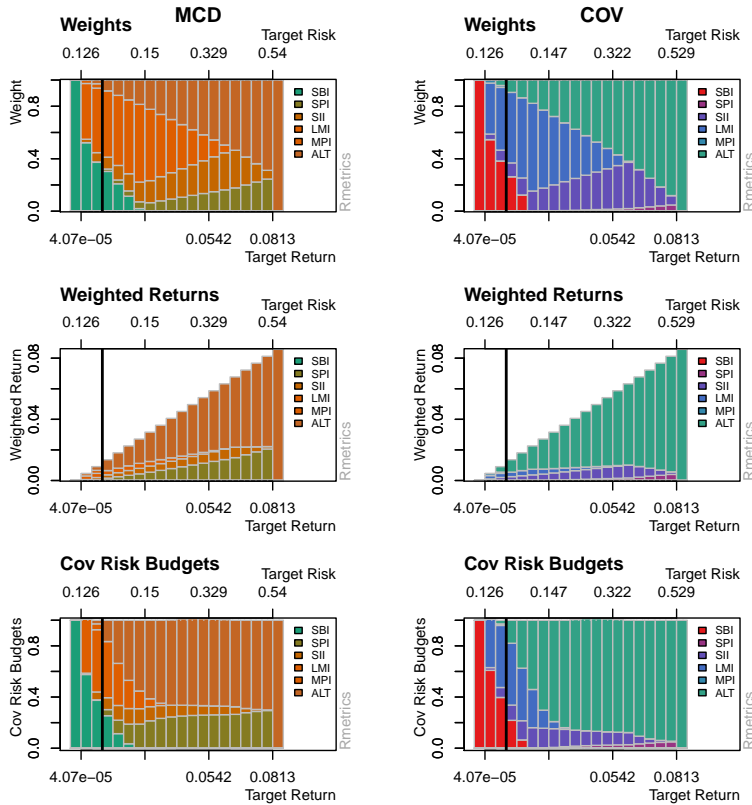


FIGURE 20.2: Weights plot for MCD robustified and COV MV portfolios. Weights along the efficient frontier of a long-only constrained mean-variance portfolio with robust MCD (left) and sample (right) covariance estimates: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets, which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars.


```

> mveSpec <- portfolioSpec()
> setEstimator(mveSpec) <- "fastMveEstimator"
> setNFrontierPoints(mveSpec) <- 5

```

Then we compute the MVE robustified efficient frontier

```

> mveFrontier <- portfolioFrontier(
  data = lppData,
  spec = mveSpec,
  constraints = "LongOnly")
> print(mveFrontier)

```

Title:

```

MV Portfolio Frontier
Estimator:      fastMveEstimator
Solver:         solveRquadprog
Optimize:       minRisk
Constraints:     LongOnly
Portfolio Points: 5 of 5

```

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.1244	0.0298	0.1414	0.5587	0.0000	0.1456
3	0.0000	0.0766	0.2496	0.3402	0.0000	0.3336
4	0.0000	0.1270	0.3431	0.0000	0.0000	0.5299
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0442	0.1094	0.1485	0.2490	0.0000	0.4489
3	0.0000	0.1850	0.1191	-0.0085	0.0000	0.7044
4	0.0000	0.1943	0.0893	0.0000	0.0000	0.7163
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Target Returns and Risks:

	mean	mu	Cov	Sigma	CVaR	VaR
1	0.0000	0.0000	0.1261	0.1230	0.2758	0.2177
2	0.0215	0.0215	0.1233	0.1086	0.2493	0.1715
3	0.0429	0.0429	0.2468	0.2003	0.5526	0.3428
4	0.0643	0.0643	0.3983	0.3190	0.9219	0.5517
5	0.0858	0.0858	0.5684	0.4743	1.3343	0.8978

Description:

```

Tue Jan 27 13:40:58 2015 by user: Rmetrics

```

For the frontier plot, we recompute the robustified frontier on 20 points.

```

> setNFrontierPoints(mveSpec) <- 20
> mveFrontier <- portfolioFrontier(
  data = lppData, spec = mveSpec)
> tailoredFrontierPlot(
  mveFrontier,
  mText = "MVE Robustified MV Portfolio",
  risk = "Sigma")

```




FIGURE 20.3: Efficient frontier of a long-only constrained mean-variance portfolio with robust MVE covariance estimates: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

The frontier plot is shown in [Figure 20.3](#).

To complete this section, we will show the weights and the performance and risk attribution plots (left-hand column of [Figure 20.4](#)).

```
> col = divPalette(6, "RdBu")
> weightsPlot(mveFrontier, col = col)
> boxL()
> text <- "MVE Robustified MV Portfolio"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(mveFrontier, col = col)
> boxL()
> covRiskBudgetsPlot(mveFrontier, col = col)
> boxL()
```

For the colours we have chosen a diverging red to blue palette. The `boxL()`

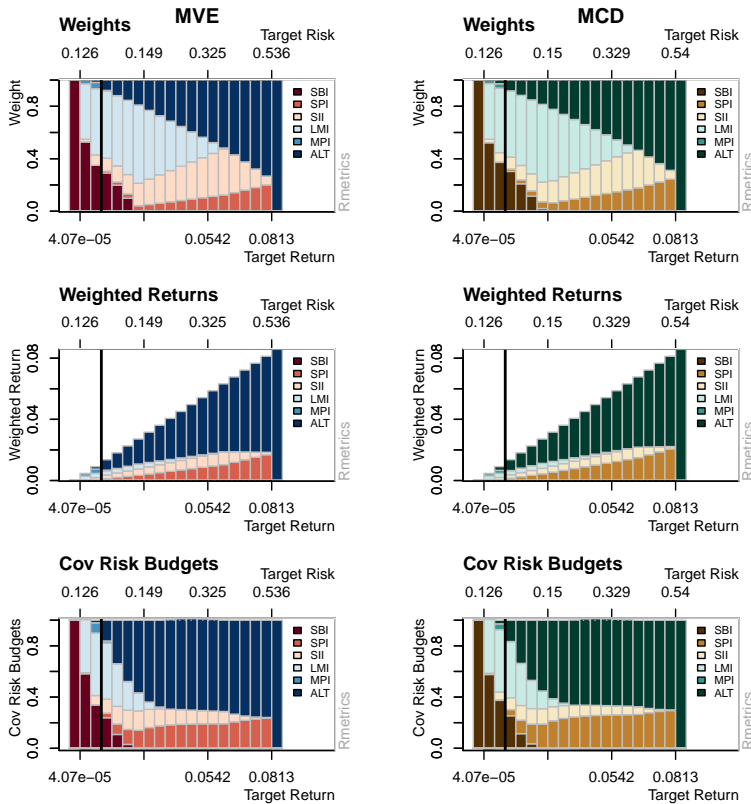


FIGURE 20.4: Weights along the efficient frontier of a long-only constrained mean-variance portfolio with robust MVE (left) and MCD (right) covariance estimates: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars. Note that the comparison of weights between the MVE and MCD with sample covariance estimates shows a much better diversification of the portfolio weights and also leads to a better diversification of the covariance risk budgets.

function draws an alternative frame around the graph with axes to the left and bottom.

20.4 THE OGK ROBUSTIFIED MEAN-VARIANCE PORTFOLIO

The Orthogonalized Gnanadesikan-Kettenring (OGK) estimator computes the orthogonalized pairwise covariance matrix estimate described in [Maronna & Zamar \(2002\)](#). The pairwise proposal goes back to [Gnanadesikan & Kettenring \(1972\)](#).

We first write a fast estimator function, `fastCovOGKEstimator()`

```
> covOGKEstimate <- covOGKEstimator(lppData)
> fastCovOGKEstimator <- function(x, spec = NULL, ...) covOGKEstimate
```

then we set the portfolio specification

```
> covOGKSpec <- portfolioSpec()
> setEstimator(covOGKSpec) <- "fastCovOGKEstimator"
> setNFrontierPoints(covOGKSpec) <- 5
```

and finally we compute the OGK robustified frontier

```
> covOGKFrontier <- portfolioFrontier(
  data = lppData, spec = covOGKSpec)
> print(covOGKFrontier)
```

Title:

```
MV Portfolio Frontier
Estimator:      fastCovOGKEstimator
Solver:         solveRquadprog
Optimize:       minRisk
Constraints:     LongOnly
Portfolio Points: 5 of 5
```

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0990	0.0171	0.1593	0.5723	0.0000	0.1522
3	0.0000	0.0650	0.2661	0.3277	0.0000	0.3411
4	0.0000	0.1179	0.3433	0.0000	0.0000	0.5388
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0347	0.0583	0.1827	0.2605	0.0000	0.4639
3	0.0000	0.1540	0.1329	-0.0089	0.0000	0.7221
4	0.0000	0.1790	0.0895	0.0000	0.0000	0.7315
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Target Returns and Risks:

	mean	mu	Cov	Sigma	CVaR	VaR
1	0.0000	0.0000	0.1261	0.1270	0.2758	0.2177
2	0.0215	0.0215	0.1223	0.1197	0.2419	0.1741
3	0.0429	0.0429	0.2460	0.2222	0.5450	0.3418
4	0.0643	0.0643	0.3976	0.3532	0.9175	0.5523
5	0.0858	0.0858	0.5684	0.5236	1.3343	0.8978

Description:

```
Tue Jan 27 13:40:59 2015 by user: Rmetrics
```

```
> setNFrontierPoints(covOGKSpec) <- 20
> covOGKFrontier <- portfolioFrontier(
  data = lppData, spec = covOGKSpec)
> tailoredFrontierPlot(
  covOGKFrontier,
  mText = "OGK Robustified MV Portfolio",
```




FIGURE 20.5: Efficient frontier of a long-only constrained mean-variance portfolio with robust OGK covariance estimates: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

```
risk = "Sigma")
```

The frontier plot is shown in [Figure 20.5](#).

The weights, and the performance and risk attributions are shown in the left-hand column of [Figure 20.6](#).

```
> col = divPalette(6, "RdYlGn")
> weightsPlot(covOGKFrontier, col = col)
> text <- "OGK Robustified MV Portfolio"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(covOGKFrontier, col = col)
> covRiskBudgetsPlot(covOGKFrontier, col = col)
```

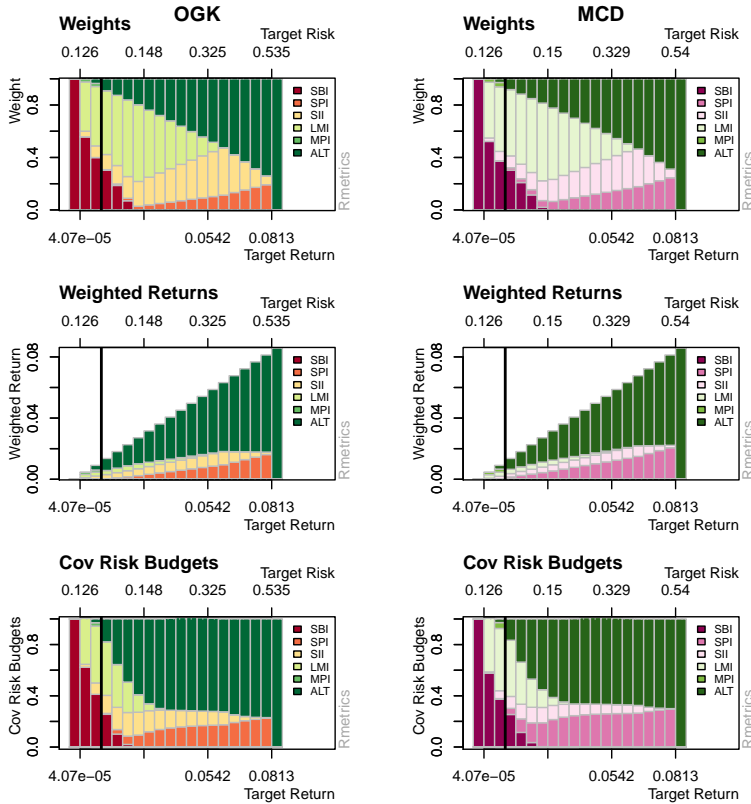


FIGURE 20.6: Weights along the efficient frontier of a long-only constrained mean-variance portfolio with robust OGK (left) and MCD (right) covariance estimates: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars. Note that both estimators result in a similar behaviour concerning the diversification of the weights. A remark, for larger data sets of assets the OGK estimator becomes favourable since it is more computation efficient.

20.5 THE SHRUNKED MEAN-VARIANCE PORTFOLIO

A simple version of a shrinkage estimator of the covariance matrix is constructed as follows. We consider a convex combination of the empirical estimator with some suitable chosen target, e.g., the diagonal matrix. Subsequently, the mixing parameter is selected to maximize the expected accuracy of the shrunked estimator. This can be done by cross-validation, or by using an analytic estimate of the shrinkage intensity. The resulting regularized estimator can be shown to outperform the maximum likelihood estimator for small samples. For large samples, the shrinkage intensity will reduce to zero, therefore in this case the shrinkage estimator will be identical to the empirical estimator. Apart from increased efficiency, the shrinkage estimate has the additional advantage that it is always positive definite and well conditioned, (Schäfer & Strimmer, 2005)².

```
> shrinkSpec <- portfolioSpec()
> setEstimator(shrinkSpec) <- "shrinkEstimator"
> setNFrontierPoints(shrinkSpec) <- 5
> shrinkFrontier <- portfolioFrontier(
  data = lppData, spec = shrinkSpec)
> print(shrinkFrontier)
```

Title:

MV Portfolio Frontier

Estimator: shrinkEstimator

Solver: solveRquadprog

Optimize: minRisk

Constraints: LongOnly

Portfolio Points: 5 of 5

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0328	0.0020	0.1498	0.6507	0.0000	0.1647
3	0.0000	0.0193	0.2550	0.3372	0.0000	0.3885
4	0.0000	0.0378	0.3454	0.0000	0.0000	0.6168
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0111	0.0062	0.1629	0.3231	0.0000	0.4968
3	0.0000	0.0422	0.1227	-0.0096	0.0000	0.8447
4	0.0000	0.0527	0.0893	0.0000	0.0000	0.8580
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Target Returns and Risks:

	mean	mu	Cov	Sigma	CVaR	VaR
1	0.0000	0.0000	0.1261	0.1261	0.2758	0.2177
2	0.0215	0.0215	0.1214	0.1217	0.2368	0.1810

² The covariance shrinkage estimator we use here is implemented in the R package `corpcor` (Schaefer et al., 2008).




FIGURE 20.7: Efficient frontier of a long-only constrained mean-variance portfolio with shrunk covariance estimates: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

```
3 0.0429 0.0429 0.2440 0.2438 0.5305 0.3382
4 0.0643 0.0643 0.3940 0.3932 0.8881 0.5834
5 0.0858 0.0858 0.5684 0.5684 1.3343 0.8978
```

Description:

Tue Jan 27 13:40:59 2015 by user: Rmetrics

The results are shown in [Figure 20.7](#) and [Figure 20.8](#).

20.6 HOW TO WRITE YOUR OWN COVARIANCE ESTIMATOR

Since we have just to set the name of the mean/covariance estimator function calling the function `setEstimator()` it becomes straightforward to add user-defined covariance estimators.

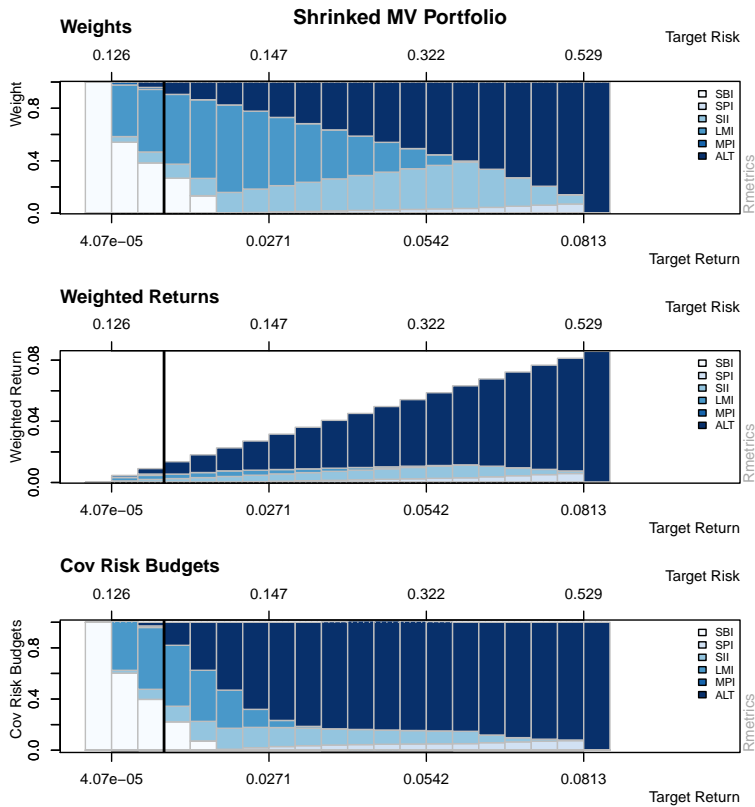


FIGURE 20.8: Weights along the efficient frontier of a long-only constrained mean-variance portfolio with shrunked covariance estimates: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars.

Let us show an example. In R's recommended package MASS there is a function (`cov.trob()`) which estimates a covariance matrix assuming the data come from a multivariate Student's *t* distribution. This approach provides some degree of robustness to outliers without giving a high breakdown point³.

```
> covtEstimator <- function (x, spec = NULL, ...) {
  x.mat = as.matrix(x)
  list(mu = colMeans(x.mat), Sigma = MASS::cov.trob(x.mat)$cov) }
> covtSpec <- portfolioSpec()
> setEstimator(covtSpec) <- "covtEstimator"
> setNFrontierPoints(covtSpec) <- 5
> covtFrontier <- portfolioFrontier(
  data = lppData, spec = covtSpec)
> print(covtFrontier)
```

Title:

```
MV Portfolio Frontier
Estimator:      covtEstimator
Solver:         solveRquadprog
Optimize:       minRisk
Constraints:     LongOnly
Portfolio Points: 5 of 5
```

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0749	0.0156	0.1490	0.6061	0.0000	0.1544
3	0.0000	0.0517	0.2479	0.3420	0.0000	0.3583
4	0.0000	0.0896	0.3441	0.0000	0.0000	0.5663
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0260	0.0527	0.1627	0.2873	0.0000	0.4714
3	0.0000	0.1205	0.1179	-0.0089	0.0000	0.7706
4	0.0000	0.1326	0.0897	0.0000	0.0000	0.7777
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Target Returns and Risks:

	mean	mu	Cov	Sigma	CVaR	VaR
1	0.0000	0.0000	0.1261	0.1109	0.2758	0.2177
2	0.0215	0.0215	0.1220	0.1043	0.2420	0.1741
3	0.0429	0.0429	0.2451	0.2006	0.5424	0.3432
4	0.0643	0.0643	0.3958	0.3217	0.9066	0.5645
5	0.0858	0.0858	0.5684	0.4697	1.3343	0.8978

Description:

```
Tue Jan 27 13:41:00 2015 by user: Rmetrics
```

³Intuitively, the breakdown point of an estimator is the proportion of incorrect observations an estimator can handle before giving an arbitrarily unreasonable result




FIGURE 20.9: Efficient frontier of a long-only constrained mean-variance portfolio with Student's t estimated covariance estimates: The plot includes the efficient frontier, the tangency line and tangency point for a zero risk-free rate, the equal weights portfolio, EWP, all single assets risk vs. return points. The line of Sharpe ratios is also shown, with its maximum coinciding with the tangency portfolio point. The range of the Sharpe ratio is printed on the right hand side axis of the plot.

```
> setNFrontierPoints(covtSpec) <- 20
> covtFrontier <- portfolioFrontier(
  data = lppData, spec = covtSpec)
> tailoredFrontierPlot(
  shrinkFrontier,
  mText = "Student's t MV Portfolio",
  risk = "Sigma")
```

The frontier plot is shown in Figure 20.9. The weights and related plots are computed in the usual way, and presented in Figure 20.10.

```
> par(mfrow = c(3, 1), mar = c(3.5, 4, 4, 3) + 0.1)
> weightsPlot(covtFrontier)
> text <- "Student's t"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(covtFrontier)
```

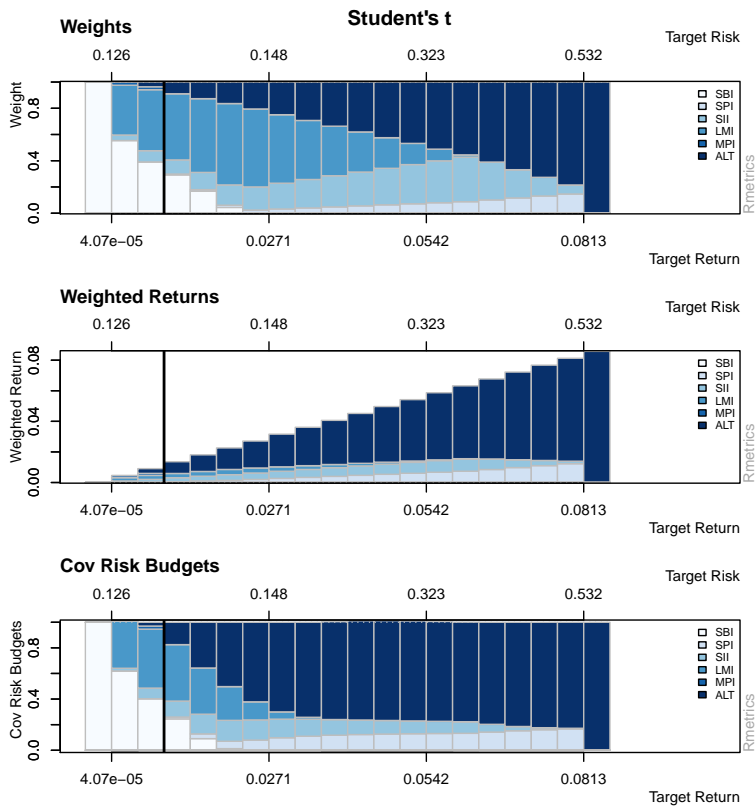


FIGURE 20.10: Weights along the efficient frontier of a long-only constrained mean-variance portfolio with robust Student's t covariance estimates: The graphs from top to bottom show the weights, the weighted returns or in other words the performance attribution, and the covariance risk budgets which are a measure for the risk attribution. The upper axis labels the target risk, and the lower labels the target return. The thick vertical line separates the efficient frontier from the minimum variance locus. The risk axis thus increases in value to both sides of the separator line. The legend to the right links the assets names to colour of the bars.

```
> covRiskBudgetsPlot(covtFrontier)
```

PART V

MEAN-CVAR PORTFOLIOS

INTRODUCTION

An alternative risk measure to the covariance is the Conditional Value at Risk, CVaR, which is also known as mean excess loss, mean shortfall or tail Value at Risk, VaR. For a given time horizon and confidence level, CVaR is the conditional expectation of the loss above VaR for the time horizon and the confidence level under consideration.

Pflug (2000) was the first to show that CVaR is a coherent risk measure and Rockafellar & Uryasev (2000) has shown that CVaR has other attractive properties including convexity.

We briefly describe the mathematical formulation of mean-CVaR optimization problems, which can be formulated as an equivalent linear programming problem and can be solved using standard linear programming solvers.

In chapter 21 we briefly describe mean-CVaR portfolio theory and present its solution. We derive the feasible set and the efficient frontier. Two special points on the frontier are discussed in detail.

In chapter 23 we present examples of how to compute feasible mean-CVaR portfolios and efficient mean-CVaR portfolios. These include not only the general cases, i.e. computing the portfolio with the lowest risk for a given return, or the portfolio with the highest return for a given risk, but also the special cases of the global minimum risk portfolio and the portfolio with the highest return/risk ratio.

In chapter 24 we explore the efficient frontier of mean-CVaR portfolios. We proceed in the same way as for the mean-variance portfolios. We consider the case of long-only, short, box, and group constrained efficient frontiers of mean-CVaR portfolios.

CHAPTER 21

MEAN-CVaR PORTFOLIO THEORY

In this chapter we formulate and solve the mean-CVaR portfolio model, where covariance risk is now replaced by the conditional Value at Risk as the risk measure. In contrast to the mean-variance portfolio optimization problem, we are no longer restrict the set of assets to have a multivariate elliptically contoured distribution.

We consider a portfolio of assets with random returns. We denote the portfolio vector of weights with w and the random events by the vector r . Let $f(w, r)$ denote the loss function when we choose the portfolio W from a set X of feasible portfolios and let r be the realization of the random events. We assume that the random vector r has a probability density function denoted by $p(r)$. For a fixed decision vector w , we compute the cumulative distribution function of the loss associated with that vector w .

$$\Psi(w, \gamma) = \int_{f(w, r) \leq \gamma} p(r) dr$$

Then, for a given confidence level α , the VaR_α associated with portfolio W is given as

$$\text{VaR}_\alpha(w) = \min\{\gamma \in \Re : \Psi(w, \gamma) \geq \alpha\}$$

Similarly, we define the CVaR_α associated with portfolio W

$$\text{CVaR}_\alpha(w) = \frac{1}{1-\alpha} \int_{f(w, r) \leq \text{VaR}_\alpha(w)} f(w, r) p(r) dr$$

We then define the problem of mean-CVaR portfolio selection as follows:

$$\begin{aligned}
& \min_w \text{CVaR}_\alpha(w) \\
& \text{s.t.} \\
& \quad w^T \hat{\mu} = \bar{r} \\
& \quad w^T \mathbf{1} = 1
\end{aligned}$$

21.1 SOLUTION OF THE MEAN-CVaR PORTFOLIO

In general, minimizing CVaR_α and VaR_α are not equivalent. Since the definition of CVaR_α involves the VaR_α function explicitly, it is difficult to work with and optimize this function. Instead, we consider the following simpler auxiliary function:

$$F_\alpha(w, \gamma) = \gamma + \frac{1}{1-\alpha} \int_{f(w, r) \geq \gamma} (f(w, r) - \gamma) p(r) dr$$

Alternatively, we can write $F_\alpha(w, \gamma)$ as follows:

$$F_\alpha(w, \gamma) = \gamma + \frac{1}{1-\alpha} \int (f(w, r) - \gamma)^+ p(r) dr$$

where $z^+ = \max(z, 0)$. This final function of γ has the following important properties that make it useful for the computation of VaR_α and CVaR_α :

- $F_\alpha(w, \gamma)$ is a convex function of γ ,
- $\text{VaR}_\alpha(w)$ is a minimizer of $F_\alpha(w, \gamma)$,
- the minimum value of the function $F_\alpha(w, \gamma)$ is $\text{CVaR}_\alpha(w)$.

As a consequence, we deduce that CVaR_α can be optimized via optimization of the function $F_\alpha(w, \gamma)$ with respect to the weights w and $\text{VaR } \gamma$. If the loss function $f(w, r)$ is a convex function of the portfolio variables w , then $F_\alpha(w, \gamma)$ is also a convex function of w . In this case, provided the feasible portfolio set W is also convex, the optimization problems are smooth convex optimization problems that can be solved using well-known optimization techniques for such problems.

21.2 DISCRETIZATION

Often it is not possible or desirable to compute/determine the joint density function $p(r)$ of the random events in our formulation. Instead, we may have a number of scenarios, say r_s for $s = 1, \dots, S$, which may represent some historical values of the returns. In this case, we obtain the following approximation to the function $F_\alpha(w, \gamma)$ by using the empirical distribution of the random returns based on the available scenarios:




FIGURE 21.1: Efficient mean-variance frontier with the global minimum variance portfolio, the global minimum Value at Risk (5%) portfolio and the global minimum Conditional Value at Risk (5%) portfolio. The efficient frontiers under the various measures, are the subset of boundaries above the corresponding minimum global risk portfolios. We see that under 5% VaR and 5% CVaR the set of efficient portfolios is reduced with respect to the variance. *Source De Giorgi (2002).*




FIGURE 21.2: Mean-VaR(5%)-boundary with the global minimum variance portfolio. Portfolios on the mean-VaR(5%)-boundary between the global minimum VaR(5%) portfolio and the global minimum variance portfolio, are mean-variance efficient. The VaR constraint (vertical line) could force mean-variance investors with high variance to reduce the variance, and mean-variance investors with low variance to increase the variance, in order to be on the left side of the VaR constraint. *Source De Giorgi (2002).*

$$\hat{F}_\alpha(w, \gamma) = \gamma + \frac{1}{(1-\alpha)S} \sum_{s=1}^S (f(w, r_s) - \gamma)^+.$$

Now, the problem $\min_{w \in W} CVaR_\alpha(w)$ can be approximated by

$$\min_{w, \gamma} = \gamma + \frac{1}{(1-\alpha)S} \sum_{s=1}^S (f(w, r_s) - \gamma)^+.$$

To solve this optimization problem, we introduce artificial variables z_s to replace $(f(w, r_s) - \gamma)^+$. This is achieved by imposing the constraints $z_s \geq f(w, r_s) - \gamma$ and $z_s \geq 0$

$$\begin{aligned} \min \quad & \gamma + \frac{1}{(1-\alpha)S} \sum_{s=1}^S z_s \\ \text{s.t.} \quad & \\ & z_s \geq f(w, r_s) - \gamma \\ & z_s \geq 0. \end{aligned}$$

Note that the constraints $z_s \geq f(w, r_s) - \gamma$ and $z_s \geq 0$ alone cannot ensure that $z_s = (f(w, r_s) - \gamma)^+$, since z_s can be larger than both right-hand term and still be feasible. However, since we are minimizing the objective function, which involves a positive multiple of z_s , it will never be optimal to assign z_s a value larger than the maximum of the two quantities $f(w, r_s) - \gamma$ and 0, and therefore, in an optimal solution z_s will be precisely $(f(w, r_s) - \gamma)^+$, thus justifying our substitution (Tütüncü, Toh & Todd, 2003).

21.3 LINEARIZATION

In the case that $f(w, r_s)$ is linear in w , all the expressions $z_s \geq f(w, r_s) - \gamma$ represent linear constraints and therefore the optimization problem becomes a linear programming problem that can be solved using the simplex method or alternative linear programming algorithms.

CHAPTER 22

MEAN-CVaR PORTFOLIO SETTINGS

```
> library(fPortfolio)
```

Like all portfolios in Rmetrics, mean-CVaR portfolios are defined by the time series data set, the portfolio specification object, and the constraint strings. Specifying a mean-CVaR portfolio thus requires the three steps already familiar from the mean-variance portfolio approach.

22.1 STEP 1: PORTFOLIO DATA

The input data for the portfolio is an S4 "timeSeries" object.

22.2 STEP 2: PORTFOLIO SPECIFICATION

As in the case of the mean-variance portfolio, the portfolio specification manages all the settings which characterize the mean-CVaR portfolio.

It is important to note that in contrast to the mean-variance portfolio specification, the type of the portfolio always has to be specified in the case of CVaR portfolios. The significance level of α is 0.05 by default, but can be modified by the user. The default solver is the LP solver from the GLPK, `Rglpk()`. Alternative solvers are the solvers from the contributed R packages `lpSolveAPI` and `Rsymphony`. The following is an example of how to modify the default specifications to use them together with the mean-CVaR portfolios:

```
> cvarSpec <- portfolioSpec()
> setType(cvarSpec) = "CVaR"
> setAlpha(cvarSpec) = 0.05
> setSolver(cvarSpec) = "solveRglpk"
```

```

> print(cvarSpec)
Model List:
  Type:          CVaR
  Optimize:       minRisk
  Estimator:      covEstimator
  Params:         alpha = 0.05 a = 1

Portfolio List:
  Target Weights: NULL
  Target Return:  NULL
  Target Risk:    NULL
  Risk-Free Rate: 0
  Number of Frontier Points: 50
  Status:         NA

Optim List:
  Solver:          solveRglpk
  Objective:       portfolioObjective portfolioReturn portfolioRisk
  Trace:          FALSE

```

22.3 STEP 3: PORTFOLIO CONSTRAINTS

In many cases we will work with long-only mean-CVaR portfolios. Specifying `constraints="LongOnly"` will force the lower and upper bounds for the weights to zero and one, respectively.

However, `fPortfolio` provides many alternative constraints. These include unlimited short-selling, lower and upper bounds, as well as linear equality and inequality constraints. The solver for dealing with these constraints has to be selected by the user and assigned by the function `setSolver()`.

CHAPTER 23

MEAN-CVAR PORTFOLIOS

```
> library(fPortfolio)
```

The following examples show how to compute feasible mean-CVaR portfolios and efficient CVaR portfolios. These include not only the general cases, i.e. computing the portfolio with the lowest risk for a given return, or the portfolio with the highest return for a given risk, but also the special cases of the global minimum-risk portfolio and the portfolio with the highest return/risk ratio.

23.1 HOW TO COMPUTE A FEASIBLE MEAN-CVaR PORTFOLIO

As a first example we consider the equal weights *feasible portfolio* with "LongOnly" constraints, which is the default case.

```
> lppData <- 100 * LPP2005.RET[, 1:6]
> cvarSpec <- portfolioSpec()
> setType(cvarSpec) <- "CVAR"
> nAssets <- ncol(lppData)
> setWeights(cvarSpec) <- rep(1/nAssets, times = nAssets)
> setSolver(cvarSpec) <- "solveRglpk.CVAR"
> ewPortfolio <- feasiblePortfolio(
  data = lppData,
  spec = cvarSpec,
  constraints = "LongOnly")
> print(ewPortfolio)
```

```
Title:
CVAR Feasible Portfolio
Estimator:      covEstimator
Solver:         solveRglpk.CVAR
Optimize:       minRisk
```

```

Constraints:      LongOnly

Portfolio Weights:
  SBI   SPI   SII   LMI   MPI   ALT
0.1667 0.1667 0.1667 0.1667 0.1667 0.1667

Covariance Risk Budgets:
  SBI   SPI   SII   LMI   MPI   ALT
-0.0039 0.3526 0.0431 -0.0079 0.3523 0.2638

Target Returns and Risks:
  mean   Cov   CVaR   VaR
0.0431 0.3198 0.7771 0.4472

Description:
Tue Jan 27 13:40:07 2015 by user: Rmetrics

```

To display the results let us write a customized function to plot the weights, the performance attribution, and the risk attribution expressed by the covariance risk budgets.

```

> weightsPie(ewPortfolio, radius = 0.7)
> text <- "Equal Weights Man-CVaR Portfolio"
> mtext(text, side = 3, line = 1.5, font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(ewPortfolio, radius = 0.8, legend = FALSE)
> covRiskBudgetsPie(ewPortfolio, radius = 0.9, legend = FALSE)

```

The result is shown in [Figure 23.1](#).

Now let us observe how the results change if we change the CVaR confidence level from $\alpha = 0.05$ to $\alpha = 0.10$

```

> setAlpha(cvarSpec) = 0.10
> ew10Portfolio <- feasiblePortfolio(
  data = lppData,
  spec = cvarSpec,
  constraints = "LongOnly")
> print(ew10Portfolio)

Title:
CVAR Feasible Portfolio
Estimator:      covEstimator
Solver:         solveRglpk.CVAR
Optimize:       minRisk
Constraints:     LongOnly

Portfolio Weights:
  SBI   SPI   SII   LMI   MPI   ALT
0.1667 0.1667 0.1667 0.1667 0.1667 0.1667

Covariance Risk Budgets:
  SBI   SPI   SII   LMI   MPI   ALT
-0.0039 0.3526 0.0431 -0.0079 0.3523 0.2638

Target Returns and Risks:
  mean   Cov   CVaR   VaR

```





FIGURE 23.1: Weights plots for an equal-weights CVaR portfolio: Although we invest the same amount in each asset, the major contribution comes from the Swiss and foreign equities and alternative instruments. The same holds for the covariance risk budgets and the weighted returns.

0.0431 0.3198 0.5858 0.3215

Description:
Tue Jan 27 13:40:08 2015 by user: Rmetrics

23.2 HOW TO COMPUTE THE MEAN-CVaR PORTFOLIO WITH THE LOWEST RISK FOR A GIVEN RETURN

Specifying the target return, we can compute an optimized efficient portfolio which has the lowest risk for a given return. In this example, we start from the equal weights portfolio, and search for a portfolio with the same returns, but a lower covariance risk.

```
> minriskSpec <- portfolioSpec()  
> setType(minriskSpec) <- "CVaR"
```

```

> setAlpha(minriskSpec) <- 0.05
> setSolver(minriskSpec) <- "solveRglpk.CVAR"
> setTargetReturn(minriskSpec) <- getTargetReturn(ewPortfolio@portfolio)["mean"]
> minriskPortfolio <- efficientPortfolio(data = lppData, spec = minriskSpec,
  constraints = "LongOnly")
> print(minriskPortfolio)

Title:
CVaR Efficient Portfolio
Estimator:      covEstimator
Solver:         solveRglpk.CVAR
Optimize:       minRisk
Constraints:     LongOnly
VaR Alpha:      0.05

Portfolio Weights:
   SBI   SPI   SII   LMI   MPI   ALT
0.0000 0.0000 0.3848 0.2354 0.0000 0.3799

Covariance Risk Budgets:
   SBI   SPI   SII   LMI   MPI   ALT
0.0000 0.0000 0.2425 -0.0102 0.0000 0.7677

Target Returns and Risks:
   mean   Cov   CVaR   VaR
0.0431 0.2484 0.5101 0.3353

Description:
Tue Jan 27 13:40:08 2015 by user: Rmetrics

```

The covariance risk of the optimized portfolio has been lowered from 0.32 to 0.25 for the same target return.

```

> weightsPie(minriskPortfolio, radius = 0.7)
> text <- "Minimum Risk CVaR Portfolio"
> mtext(text, side = 3, line = 1.5, font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(minriskPortfolio, radius = 0.8, legend = FALSE)
> covRiskBudgetsPie(minriskPortfolio, radius = 0.9, legend = FALSE)

```

The plots are shown in [Figure 23.1](#).

23.3 HOW TO COMPUTE THE GLOBAL MINIMUM MEAN-CVaR PORTFOLIO

The global *minimum risk portfolio* is the efficient portfolio with the lowest possible risk.

```

> globminSpec <- portfolioSpec()
> setType(globminSpec) <- "CVaR"
> setAlpha(globminSpec) <- 0.05
> setSolver(globminSpec) <- "solveRglpk.CVAR"
> setTargetReturn(globminSpec) <- getTargetReturn(ewPortfolio@portfolio)["mean"]
> globminPortfolio <- minriskPortfolio(data = lppData, spec = globminSpec,
  constraints = "LongOnly")
> print(globminPortfolio)

```




FIGURE 23.2: Weights plots for a minimum risk CVaR portfolio: Optimizing the risk for the target return of the equal weights portfolio leads to badly diversified portfolio, dominated by the risky alternative instruments.

```
Title:
CVaR Minimum Risk Portfolio
Estimator:      covEstimator
Solver:         solveRglpk.CVaR
Optimize:       minRisk
Constraints:     LongOnly
VaR Alpha:      0.05

Portfolio Weights:
SBI  SPI  SII  LMI  MPI  ALT
0.1846 0.0000 0.1432 0.5952 0.0000 0.0770

Covariance Risk Budgets:
SBI  SPI  SII  LMI  MPI  ALT
0.1436 0.0000 0.1986 0.5500 0.0000 0.1078

Target Returns and Risks:
mean  Cov  CVaR  VaR
```




Weights**Global Minimum Risk Portfolio****Covariance Risk Budgets****Weighted Returns**

FIGURE 23.3: Weights plots for a global minimum risk CVaR portfolio: As expected, the global minimum risk portfolio is dominated by the low-risk Swiss and foreign bond assets.

```
0.0133 0.1015 0.1964 0.1523
```

Description:

Tue Jan 27 13:40:08 2015 by user: Rmetrics

As expected, the portfolio is now dominated by the Swiss and foreign equities, which contribute 78% to the weights of the optimized portfolio. Internally, the global minimum mean-CVaR portfolio is calculated by minimizing the efficient portfolio with respect to the target risk.

```
> weightsPie(globminPortfolio, radius = 0.7)
> text <- "Global Minimum Risk Portfolio"
> mtext(text, side = 3, line = 1.5, font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(globminPortfolio, radius = 0.8, legend = FALSE)
> covRiskBudgetsPie(globminPortfolio, radius = 0.9, legend = FALSE)
```

The plots are shown in [Figure 23.3](#).

How to Compute the Max Return/Risk Ratio Mean-CVaR Portfolio

The *Max Return/Risk portfolio* is calculated by minimization of the ‘Sortino Ratio’ for a given risk-free rate. The Sortino ratio is the ratio of the target return lowered by the risk-free rate and the CvaR risk. The risk-free rate in the default specification is zero and can be set to another value by using the function `setRiskFreeRate<-`.

```
> scData <- SMALLCAP.RET[, c("BKE", "GG", "GYMB", "KRON")]
> ratioSpec <- portfolioSpec()
> setType(ratioSpec) <- "CVaR"
> setAlpha(ratioSpec) <- 0.05
> setSolver(ratioSpec) <- "solveRglpk.CVAR"
> setRiskFreeRate(ratioSpec) <- 0
> ratioPortfolio <- maxratioPortfolio(data = scData, spec = ratioSpec,
  constraints = "LongOnly")
> print(ratioPortfolio)
```

Title:

CVaR Max Return/Risk Ratio Portfolio	
Estimator:	covEstimator
Solver:	solveRglpk.CVAR
Optimize:	minRisk
Constraints:	LongOnly
VaR Alpha:	0.05

Portfolio Weights:

BKE	GG	GYMB	KRON
0.3468	0.2551	0.0000	0.3981

Covariance Risk Budgets:

BKE	GG	GYMB	KRON
0.3613	0.1385	0.0000	0.5002

Target Returns and Risks:

mean	Cov	CVaR	VaR
0.0294	0.0984	0.1269	0.1139

Description:


Tue Jan 27 13:40:08 2015 by user: Rmetrics

```
> weightsPie(ratioPortfolio, radius = 0.7)
> text <- "Maximum Return/Risk Portfolio"
> mtext(text, side = 3, line = 1.5, font = 2, cex = 0.7, adj = 0)
> weightedReturnsPie(ratioPortfolio, radius = 0.8, legend = FALSE)
> covRiskBudgetsPie(ratioPortfolio, radius = 0.9, legend = FALSE)
```


The plots are shown in [Figure 23.4](#).

Weights

Maximum Return/Risk Portfolio



Covariance Risk Budgets



Weighted Returns




FIGURE 23.4: Weights, covariance risk budgets and weighthd returns plots for a max/risk ratio mean-CVaR portfolio.

CHAPTER 24

MEAN-CVaR PORTFOLIO FRONTIERS

```
> library(fPortfolio)
```

In this section we explore the efficient frontier, EF, and the minimum variance locus, MVL, of mean-CVaR portfolios. We proceed in the same way as for mean-variance portfolios: We select the two assets which lead to the smallest and largest returns and divide their range into equidistant parts which determine the target returns for which we try to find the efficient portfolios. We compute the global minimum risk portfolio and start from the closest returns to this point in both directions of the EF and the MVL. Note that only in the case of the long-only portfolio constraints do we reach both ends of the EF and the MVL. Usually, constraints will shorten the EF and MVL, and may even happen, that the constraints were so strong that do not find any solution at all.

In the following we compute and compare long-only, unlimited short, box, and group constrained efficient frontiers of mean-CVaR portfolios¹.

24.1 THE LONG-ONLY PORTFOLIO FRONTIER

The long-only mean-variance portfolios. In this case all the weights are bounded between zero and one.

LISTING 24.1: THE TABLE LISTS FUNCTIONS TO COMPUTE THE EFFICIENT FRONTIER OF LINEARLY CONSTRAINED MEAN-CVaR PORTFOLIOS AND TO PLOT THE RESULTS.

Functions:

<code>portfolioFrontier</code>	efficient portfolios on the frontier
--------------------------------	--------------------------------------

¹Note that throughout this section we set the portfolio type to CVaR and the solver function to `solveRglpk.CVAR()`.

frontierPoints	extracts risk/return frontier points
frontierPlot	creates an efficient frontier plot
cmlPoints	adds market portfolio
cmlLines	adds capital market line
tangencyPoints	adds tangency portfolio point
tangencyLines	adds tangency line
equalWeightsPoints	adds point of equal weights portfolio
singleAssetPoints	adds points of single asset portfolios
twoAssetsLines	adds frontiers of two assets portfolios
sharpeRatioLines	adds Sharpe ratio line
monteCarloPoints	adds randomly feasible portfolios
weightsPlot	weights bar plot along the frontier
weightedReturnsPlot	weighted returns bar plot
covRiskBudgetsPlot	covariance risk budget bar plot

```
> lppData <- 100 * LPP2005.RET[, 1:6]
> longSpec <- portfolioSpec()
> setType(longSpec) <- "CVaR"
> setAlpha(longSpec) <- 0.05
> setNFrontierPoints(longSpec) <- 5
> setSolver(longSpec) <- "solveRglpk.CVAR"
> longFrontier <- portfolioFrontier(data = lppData, spec = longSpec,
  constraints = "LongOnly")
> print(longFrontier)
```

Title:

CVaR Portfolio Frontier

```
Estimator:      covEstimator
Solver:         solveRglpk.CVAR
Optimize:       minRisk
Constraints:     LongOnly
Portfolio Points: 5 of 5
VaR Alpha:      0.05
```

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0000	0.0000	0.1988	0.6480	0.0000	0.1532
3	0.0000	0.0000	0.3835	0.2385	0.0000	0.3780
4	0.0000	0.0000	0.3464	0.0000	0.0000	0.6536
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.0000	0.0000	0.2641	0.3126	0.0000	0.4233
3	0.0000	0.0000	0.2432	-0.0101	0.0000	0.7670
4	0.0000	0.0000	0.0884	0.0000	0.0000	0.9116
5	0.0000	0.0000	0.0000	0.0000	0.0000	1.0000

Target Returns and Risks:

	mean	Cov	CVaR	VaR
1	0.0000	0.1261	0.2758	0.2177


```

2 0.0215 0.1224 0.2313 0.1747
3 0.0429 0.2472 0.5076 0.3337
4 0.0643 0.3941 0.8780 0.5830
5 0.0858 0.5684 1.3343 0.8978

```

Description:

Tue Jan 27 13:38:57 2015 by user: Rmetrics

To shorten the output in the example above, we have lowered the number of frontier points to 5 points. The printout lists the weights, the covariance risk budgets and the target return and risk values along the minimum variance locus and the efficient frontier starting with the portfolio with the lowest return and ending with the portfolio with the highest achievable return at the end of the efficient frontier.

To plot the efficient frontier we repeat the optimization with 25 points at the frontier and plot the result using the function `tailoredFrontierPlot()`

```

> setNFrontierPoints(longSpec) <- 25
> longFrontier <- portfolioFrontier(data = lppData, spec = longSpec,
  constraints = "LongOnly")
> tailoredFrontierPlot(object = longFrontier, mText = "Mean-CVaR Portfolio - Long Only Constraints",
  risk = "CVaR")

```

The function `tailoredFrontierPlot()` displays, as the name says, a customized plot with fixed colour, font and symbol settings and several selected add-ons including, single assets points, tangency line, and Sharpe ratio line.

Figure 24.1 and Figure 24.2 show the results for the weights, the weighted returns and the covariance risk budgets along the minimum variance locus and the efficient frontier.

```

> par(mfrow = c(3, 1), mar = c(3.5, 4, 4, 3) + 0.1)
> weightsPlot(longFrontier)
> text <- "Mean-CVaR Portfolio - Long Only Constraints"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(longFrontier)
> covRiskBudgetsPlot(longFrontier)

```

24.2 THE UNLIMITED 'SHORT' PORTFOLIO FRONTIER

When all weights are not restricted we have the case of unlimited short selling. Unlike in the mean-variance portfolio, we cannot optimize the portfolio analytically. To circumvent this we define box constraints with large lower and upper bounds.

```

> shortSpec <- portfolioSpec()
> setType(shortSpec) <- "CVaR"
> setAlpha(shortSpec) <- 0.05

```




FIGURE 24.1: The graph shows the minimum variance locus and the efficient frontier for 25 equidistant return points. Added are the risk-return points for the individual assets and the equal weights portfolio. The line through the origin is the tangency line for a zero risk-free rate. The curved line with the maximum at the tangency point is the Sharpe ratio along the frontier.

```
> setNFrontierPoints(shortSpec) <- 5
> setSolver(shortSpec) <- "solveRglpk.CVAR"
> shortConstraints <- c("minW[1:6]=-999", "maxW[1:6]=+999")
> shortFrontier <- portfolioFrontier(data = lppData, spec = shortSpec,
  constraints = shortConstraints)
> print(shortFrontier)
```

Title:

CVaR Portfolio Frontier

Estimator: covEstimator

Solver: solveRglpk.CVAR

Optimize: minRisk

Constraints: minW maxW

Portfolio Points: 5 of 5

VaR Alpha: 0.05

Portfolio Weights:

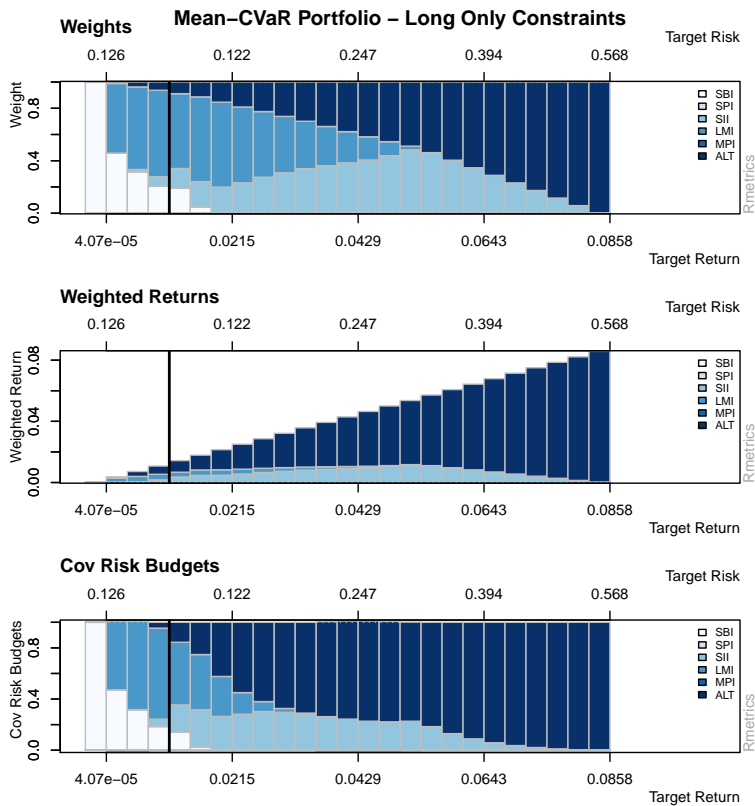


FIGURE 24.2: The graph shows for the weights, weighted returns, and covariance risk budgets 25 equidistant return points, along the minimum variance locus and the efficient frontier. Note that the strong separation line marks the position between the minimum variance locus and the efficient frontier. Target returns are increasing from left to right, whereas target risks, are increasing to the left and right with respect to the separation line.

	SBI	SPI	SII	LMI	MPI	ALT
1	0.4257	-0.0242	0.0228	0.5661	0.0913	-0.0816
2	-0.0201	-0.0101	0.1746	0.7134	-0.0752	0.2174
3	-0.3275	-0.0196	0.4318	0.6437	-0.2771	0.5486
4	-0.8113	0.0492	0.5704	0.8687	-0.5273	0.8503
5	-1.6975	0.0753	0.6305	1.5485	-0.6683	1.1115

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.4056	0.0256	0.0062	0.5384	-0.0730	0.0972
2	-0.0080	-0.0173	0.2124	0.4559	-0.1256	0.4825
3	0.0054	-0.0204	0.3674	0.0592	-0.2787	0.8671
4	0.0572	0.0409	0.2901	0.0223	-0.2984	0.8880
5	0.1513	0.0510	0.1966	0.0215	-0.3235	0.9031

Target Returns and Risks:

mean Cov CVaR VaR




FIGURE 24.3: The graph shows for 25 equidistant return points the minimum variance locus and the efficient frontier when short selling is allowed. The major difference to the previous long-only is the fact, that MVL and EF do not end at the assets with the lowest and highest risks: For the same return the risk has lowered through short selling.

```
1 0.0000 0.1136 0.2329 0.1859
2 0.0215 0.1172 0.2118 0.1733
3 0.0429 0.2109 0.3610 0.2923
4 0.0643 0.3121 0.5570 0.4175
5 0.0858 0.4201 0.7620 0.5745
```

Description:

Tue Jan 27 13:39:10 2015 by user: Rmetrics

```
> setNFrontierPoints(shortSpec) <- 25
> shortFrontier <- portfolioFrontier(data = lppData, spec = shortSpec,
  constraints = shortConstraints)
> tailoredFrontierPlot(object = shortFrontier, mText = "Mean-CVaR Portfolio - Short Constraints",
  risk = "CVaR")

> par(mfrow = c(3, 1), mar = c(3.5, 4, 4, 3) + 0.1)
> weightsPlot(shortFrontier)
```

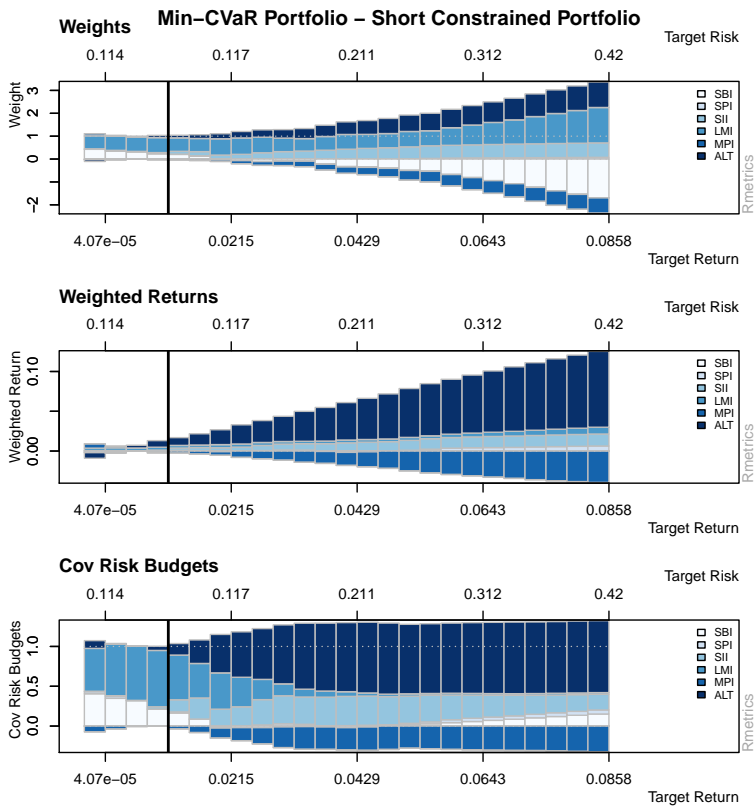


FIGURE 24.4: The graph shows for equidistant return points the weights, weighted returns, and covariance risk budgets, along the minimum variance locus and the efficient frontier.

```
> text <- "Min-CVaR Portfolio - Short Constrained Portfolio"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(shortFrontier)
> covRiskBudgetsPlot(shortFrontier)
```

24.3 THE BOX-CONSTRAINED PORTFOLIO FRONTIER

A box-constrained portfolio is a portfolio where the weights are constrained by lower and upper bounds, e.g. we want to invest at least 10% and no more than 50% in each asset.

```
> boxSpec <- portfolioSpec()
> setType(boxSpec) <- "CVaR"
> setAlpha(boxSpec) <- 0.05
> setNFrontierPoints(boxSpec) <- 15
> setSolver(boxSpec) <- "solveRglpk.CVAR"
> boxConstraints <- c("minW[1:6]=0.05", "maxW[1:6]=0.66")
> boxFrontier <- portfolioFrontier(data = lppData, spec = boxSpec,
```

```
constraints = boxConstraints)
> print(boxFrontier)
```

Title:

```
CVaR Portfolio Frontier
Estimator:      covEstimator
Solver:         solveRglpk.CVaR
Optimize:       minRisk
Constraints:     minW maxW
Portfolio Points: 5 of 9
VaR Alpha:      0.05
```

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.0526	0.0500	0.1370	0.6600	0.0500	0.0504
3	0.0500	0.0500	0.2633	0.4127	0.0500	0.1739
5	0.0500	0.0500	0.4109	0.1463	0.0500	0.2928
7	0.0500	0.0500	0.3378	0.0500	0.0500	0.4622
9	0.0500	0.0501	0.1399	0.0500	0.0500	0.6600

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.0189	0.1945	0.1435	0.3378	0.1742	0.1312
3	0.0023	0.1528	0.2209	0.0248	0.1582	0.4410
5	-0.0017	0.1042	0.2478	-0.0080	0.1136	0.5440
7	-0.0024	0.0823	0.1099	-0.0035	0.0931	0.7206
9	-0.0022	0.0650	0.0182	-0.0031	0.0752	0.8469

Target Returns and Risks:

	mean	Cov	CVaR	VaR
1	0.0184	0.1232	0.2604	0.1913
3	0.0307	0.1838	0.3999	0.2651
5	0.0429	0.2655	0.5787	0.3654
7	0.0552	0.3456	0.7832	0.4818
9	0.0674	0.4388	1.0382	0.6675

Description:

```
Tue Jan 27 13:39:23 2015 by user: Rmetrics
```

```
> setNFrontierPoints(boxSpec) <- 25
> boxFrontier <- portfolioFrontier(data = lppData, spec = boxSpec,
  constraints = boxConstraints)
> tailoredFrontierPlot(object = boxFrontier, mText = "Mean-CVaR Portfolio - Box Constraints",
  risk = "CVaR")

> weightsPlot(boxFrontier)
> text <- "Min-CVaR Portfolio - Box Constrained Portfolio"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(boxFrontier)
> covRiskBudgetsPlot(boxFrontier)
```




FIGURE 24.5: Box constrained Min-CVaR portfolio frontier plot.

24.4 THE GROUP-CONSTRAINED PORTFOLIO FRONTIER

A group-constrained portfolio is a portfolio where the weights of groups of selected assets are constrained by lower and upper bounds for the total weights of the groups, e.g. we want to invest at least 30% in the group of bounds and not more than 50% in the groups of assets.

```
> groupSpec <- portfolioSpec()
> setType(groupSpec) <- "CVaR"
> setAlpha(groupSpec) <- 0.05
> setNFrontierPoints(groupSpec) <- 10
> setSolver(groupSpec) <- "solveRglpk.CVAR"
> groupConstraints <- c("minsumW[c(1,4)]=0.3", "maxsumW[c(2:3,5:6)]=0.66")
> groupFrontier <- portfolioFrontier(data = lppData, spec = groupSpec,
  constraints = groupConstraints)
> print(groupFrontier)
```

Title:
CVaR Portfolio Frontier
Estimator: covEstimator
Solver: solveRglpk.CVAR

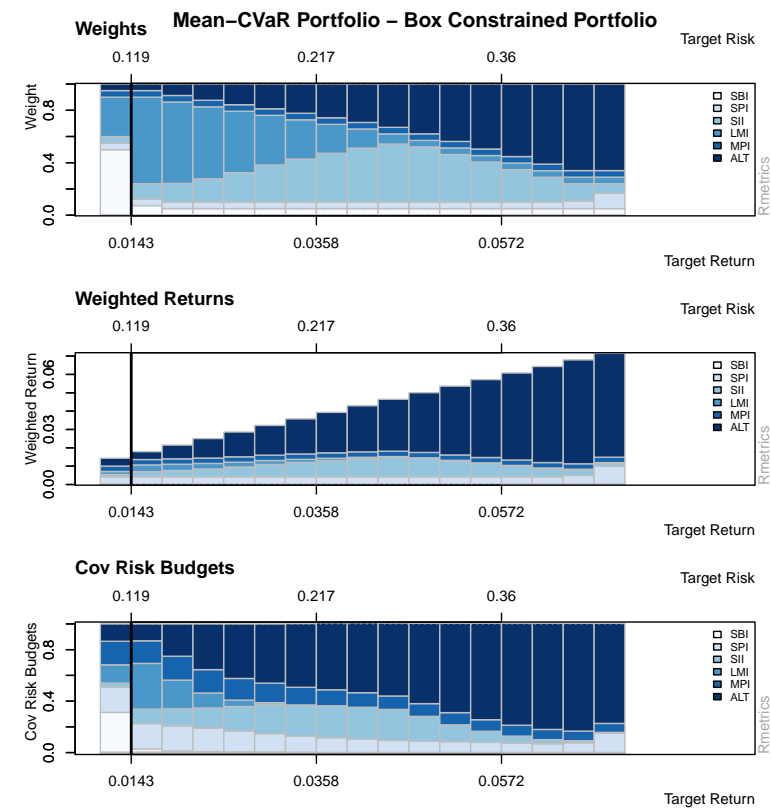


FIGURE 24.6: Mean-CVaR Portfolio - Box Constrained Weights Plot

Optimize: minRisk
Constraints: minsumW maxsumW
Portfolio Points: 5 of 7
VaR Alpha: 0.05

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.2440	0.0000	0.0410	0.6574	0.0000	0.0576
4	0.0000	0.0000	0.2744	0.5007	0.0000	0.2249
5	0.0000	0.0000	0.3288	0.3400	0.0000	0.3312
7	0.0000	0.0000	0.0209	0.3400	0.0000	0.6391

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	1.0000	0.0000	0.0000	0.0000	0.0000	0.0000
2	0.2294	0.0000	0.0218	0.7226	0.0000	0.0263
4	0.0000	0.0000	0.3024	0.0780	0.0000	0.6196
5	0.0000	0.0000	0.2388	-0.0024	0.0000	0.7636
7	0.0000	0.0000	0.0020	-0.0159	0.0000	1.0139




FIGURE 24.7: Group constrained Min-CVaR portfolio frontier plot.

Target Returns and Risks:

	mean	Cov	CVaR	VaR
1	0.0000	0.1261	0.2758	0.2177
2	0.0096	0.1012	0.2003	0.1622
4	0.0286	0.1575	0.3076	0.2178
5	0.0381	0.2147	0.4392	0.2783
7	0.0572	0.3559	0.8228	0.5517

Description:

Tue Jan 27 13:39:36 2015 by user: Rmetrics

```
> setNFrontierPoints(groupSpec) <- 25
> groupFrontier <- portfolioFrontier(data = lppData, spec = groupSpec,
  constraints = groupConstraints)
> tailoredFrontierPlot(object = groupFrontier, mText = "Mean-CVaR Portfolio - Group Constraints",
  risk = "CVaR")

> weightsPlot(groupFrontier)
> text <- "Min-CVaR Portfolio - Group Constrained Portfolio"
```

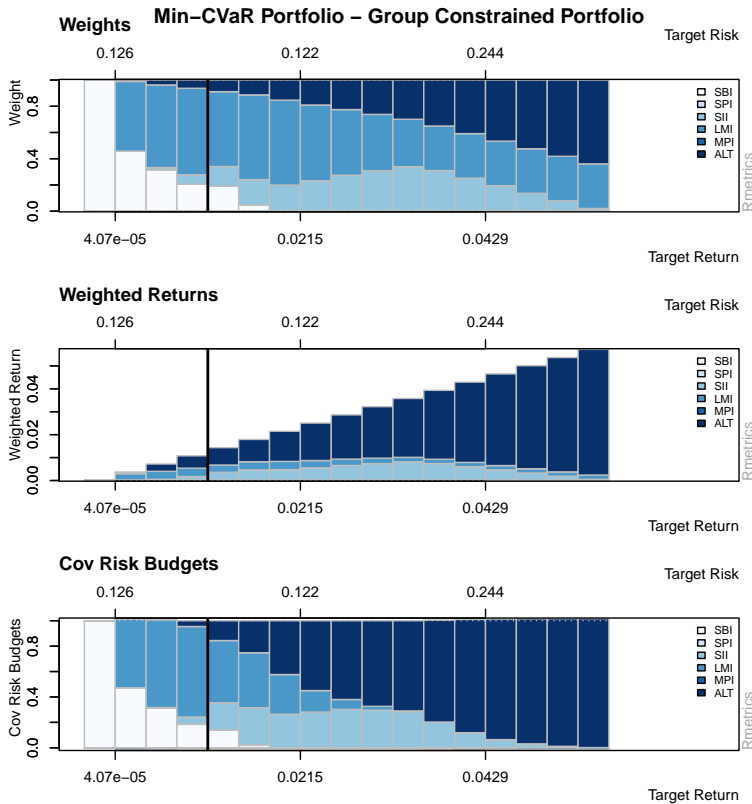


FIGURE 24.8: Mean-CVaR Portfolio - Group Constrained Weights Plot

```
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(groupFrontier)
> covRiskBudgetsPlot(groupFrontier)
```

24.5 THE BOX/GROUP-CONSTRAINED PORTFOLIO FRONTIER

Box and group constraints can be combined

```
> boxgroupSpec <- portfolioSpec()
> setType(boxgroupSpec) <- "CVaR"
> setAlpha(boxgroupSpec) <- 0.05
> setNFrontierPoints(boxgroupSpec) <- 5
> setSolver(boxgroupSpec) <- "solveRglpk.CVAR"
> boxgroupConstraints <- c(boxConstraints, groupConstraints)
> boxgroupFrontier <- portfolioFrontier(data = lppData, spec = boxgroupSpec,
  constraints = boxgroupConstraints)
> print(boxgroupFrontier)
```

Title:

CVaR Portfolio Frontier

```

Estimator:      covEstimator
Solver:         solveRglpk.CVAR
Optimize:       minRisk
Constraints:     minW maxW minsumW maxsumW
Portfolio Points: 2 of 2
VaR Alpha:      0.05

```

Portfolio Weights:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.0500	0.0500	0.1421	0.6207	0.0500	0.0872
2	0.0500	0.0500	0.2245	0.2900	0.0500	0.3355

Covariance Risk Budgets:

	SBI	SPI	SII	LMI	MPI	ALT
1	0.0125	0.1951	0.1321	0.2234	0.1858	0.2510
2	-0.0013	0.1117	0.0906	-0.0112	0.1232	0.6871

Target Returns and Risks:

	mean	Cov	CVaR	VaR
1	0.0215	0.1347	0.2852	0.2047
2	0.0429	0.2609	0.5953	0.3814

Description:

Tue Jan 27 13:39:50 2015 by user: Rmetrics

```

> setNFrontierPoints(boxgroupSpec) <- 25
> boxgroupFrontier <- portfolioFrontier(
  data = lppData,
  spec = boxgroupSpec,
  constraints = boxgroupConstraints)
> tailoredFrontierPlot(
  object = boxgroupFrontier,
  mText = "Mean-CVaR Portfolio - Box/Group Constraints",
  risk = "CVaR")

> weightsPlot(boxgroupFrontier)
> text <- "Mean-CVaR Portfolio - Box/Group Constrained Portfolio"
> mtext(text, side = 3, line = 3, font = 2, cex = 0.9)
> weightedReturnsPlot(boxgroupFrontier)
> covRiskBudgetsPlot(boxgroupFrontier)

```

24.6 OTHER CONSTRAINTS

Like in the case of the mean-variance portfolios quadratic and/or non-linear constraints complicate portfolio optimization. Those constraints will include for example quadratic covariance risk budget constraints and tail risk budget constraints, as well as non-linear function constraints, such as maximum drawdowns limit constraints or extension strategy constraints².

²These more complex constraints require quadratic and non-linear portfolio solvers which are considered in the ebook *Advanced Portfolio Optimization with R/Rmetrics*.




FIGURE 24.9: Box/Group constrained CVaR portfolio frontier plot.

24.7 MORE ABOUT THE FRONTIER PLOT TOOLS

Note that the default axis type of the frontier plot is automatically taken from the portfolio specification, here the "CVaR" axis was selected and thus displayed. The reason for this is that the function `frontierPlot()` inspects the type of the portfolio and then decides what type of axis to display.

The function `frontierPlot()` returns a two column matrix with the target risk and target return to be plotted. For the target return we can extract either the mean or the mu values, for the target risk we can select from four choices, "Cov", "Sigma", "CVaR", and "VaR". Furthermore, we can overwrite the risk choice, and allow for an automated selection, `auto=TRUE`, which is the default. The auto selection does the following:

```
Type - Risk Relationships:
Type <- getType(object)
if (Type == "MV") risk = "Cov"
```

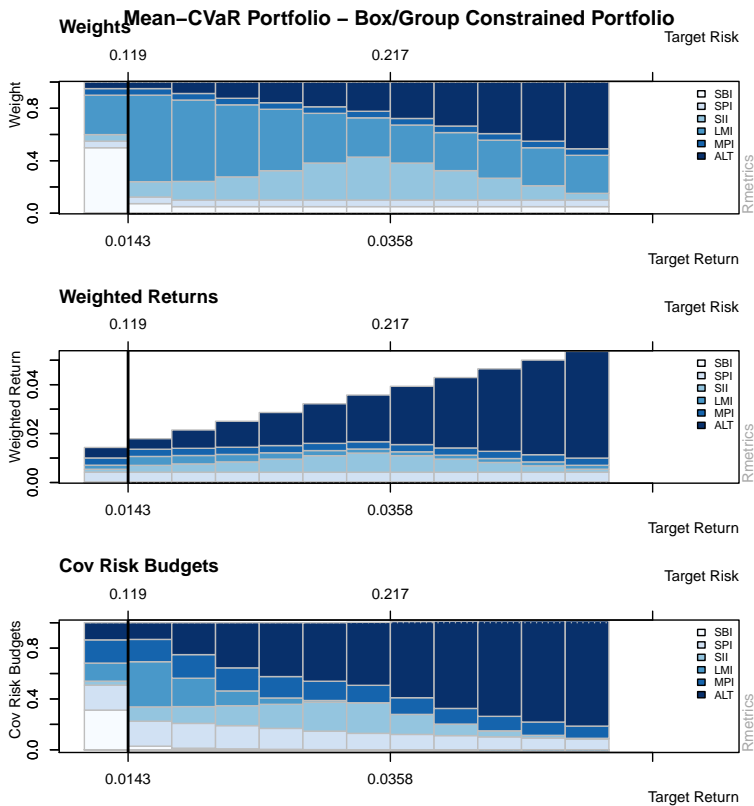


FIGURE 24.10: Mean-CVaR Portfolio - Box/Group Constrained Weights Plot

```
if (Type == "MV" & Estimator != "covEstimator") risk = "Sigma"
if (Type == "CVaR") risk = "CVaR"
```

Explicitly specifying the risk type in the function argument the function `frontierPlot()` allows us to display several views from the efficient frontier. Now let us plot the "covariance" frontier together with the "CVaR" frontier in the covariance risk view:

```
> longSpec <- portfolioSpec()
> setType(longSpec) <- "CVaR"
> setAlpha(longSpec) <- 0.1
> setNFrontierPoints(longSpec) <- 20
> setSolver(longSpec) <- "solveRglpk.CVAR"
> longFrontier <- portfolioFrontier(data = lppData, spec = longSpec,
  constraints = "LongOnly")
> par(mfrow = c(2, 2))
> frontierPlot(longFrontier, pch = 16, type = "b", cex = 0.7)
> frontierPlot(longFrontier, risk = "Cov", auto = FALSE, pch = 16,
  type = "b", cex = 0.7)
> frontierPlot(longFrontier, risk = "VaR", auto = FALSE, pch = 16,
```

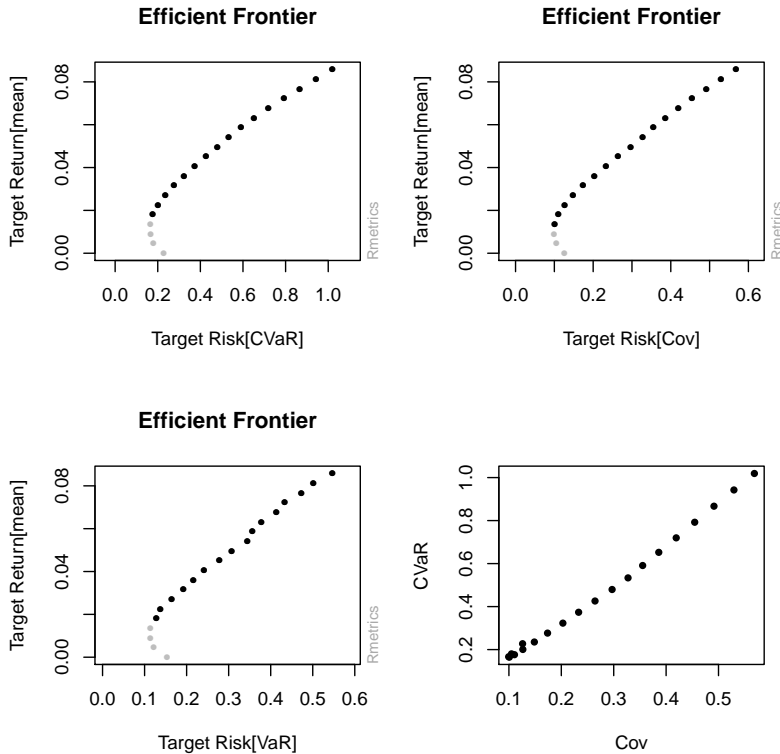


FIGURE 24.11: The first three charts show different views of return vs. risk plots. Here the risk is measured as covariance risk, as conditional value-at-risk, and as value-at-risk. Note the kinks in the VaR measure plot arise from the fact that VaR is not a coherent risk measure. The last graph plots the relationship between covariance risk and conditional value-at-risk.

```

type = "b", cex = 0.7)
> Cov <- frontierPoints(longFrontier, risk = "Cov", auto = FALSE)[,
  "targetRisk"]
> CVaR <- frontierPoints(longFrontier, risk = "CVaR", auto = FALSE)[,
  "targetRisk"]
> plot(Cov, CVaR, pch = 19, cex = 0.7)

```

The result is shown in Figure 24.11.

Interactively plotting the efficient frontier

The generic `plot()` function allows you to interactively display the efficient frontier with several add-on plots

```

> plot(frontier)
Make a plot selection (or 0 to exit):

```

- 1: Plot Efficient Frontier
- 2: Add Minimum Risk Portfolio
- 3: Add Tangency Portfolio
- 4: Add Risk/Return of Single Assets
- 5: Add Equal Weights Portfolio
- 6: Add Two Asset Frontiers [0-1 PF Only]
- 7: Add Wheel Pie of Weights
- 8: Add Monte Carlo Portfolios
- 9: Add Sharpe Ratio [MV PF Only]

Selection:

PART VI

PORTFOLIO BACKTESTING

INTRODUCTION

Backtesting is a key component of portfolio management and it is often used to assess and compare different statistical models. Portfolio backtesting is accomplished by reconstructing, with historical data, trades that would have occurred in the past using rules defined by a given strategy. The underlying theory is that any strategy that worked well in the past is likely to work well in the future.

The backtesting results offer statistics that can be used to gauge the effectiveness of the strategy. However, the results achieved from the backtest are highly dependent on the movements of the tested period. Therefore, it is often a good idea to backtest over a long time frame that encompasses several different types of market conditions.

In [chapter 25](#) we introduce a new S4 object class `fPFOLIOBACKTEST`, which represents a portfolio backtest. The four slots keep all information on the rolling windows, on the investment strategy portfolio, on the smoother computing the weights for their re-balancing, and on optional messages. We describe in detail the rolling analysis technique used to run a portfolio backtest.

In [chapter 26](#) we illustrate the usage of the `fPortfolio` package with a realistic portfolio. The first example shows how to re-balance the Swiss Performance Sector Indexes over time to outperform the SPI market index.

In [chapter 27](#) we present a second example which shows us the backtest results for the MSCI GCC Index, a market index traded in the gulf region which covers the economies Bahrain, Kuwait, Oman, Qatar and the United Arab Emirates.

CHAPTER 25

S4 PORTFOLIO BACKTEST CLASS

```
> library(fPortfolio)
```

In this chapter we introduce the S4 class `fPFOLIOWBACKTEST` for performing a rolling optimization and performance analysis on portfolios over time. We present the components which allow such an analysis. This includes a discussion of how to create rolling windows, how to formulate portfolio strategies, and how to smooth and post-process the optimal portfolio weights.

25.1 CLASS REPRESENTATION

All settings specifying the backtesting procedure are represented by an S4 class named `fPFOLIOWBACKTEST`.

```
> showClass("fPFOLIOWBACKTEST")
Class "fPFOLIOWBACKTEST" [package "fPortfolio"]

Slots:

Name:   windows strategy smoother messages
Class:  list      list      list      list
```

An object of class `fPFOLIOWBACKTEST` has four slots, named `@windows`, `@strategy`, `@smoother`, and `@messages`. The first slot, `@windows`, holds the rolling windows information, the second slot, `@strategy`, contains the portfolio strategy for backtesting, the `@smoother` slot contains information regarding weights smoothing and the last slot, named `@messages`, holds a list of optional messages.

How to create a portfolio backtest object

The function `portfolioBacktest()` allows the user to set backtest settings from scratch.

```
> formals(portfolioBacktest)

$windows
list(windows = "equidistWindows", params = list(horizon = "12m"))

$strategy
list(strategy = "tangencyStrategy", params = list())

$smoother
list(smoother = "emaSmoother", params = list(doubleSmoothing = TRUE,
  lambda = "3m", skip = 0, initialWeights = NULL))

$messages
list()
```

The default settings for rolling portfolios are composed of equidistant windows with a horizon of 12 months, a tangency portfolio strategy, and a double exponential moving average (EMA) smoother. The decay length, `lambda="3m"`, is three months and the computation starts with the first data point, i.e. `skip=0`. The messages list is empty.

To look inside the portfolio backtest structure you can call the function `str()`. This function compactly displays the internal structure of the portfolio backtest object. It can be considered as a diagnostic function and as a simple way to summarize the internal structure of the object. Let us create a new backtest object, and then print the structure for the default settings.

```
> backtest <- portfolioBacktest()
> str(backtest, width = 65, strict.width = "cut")

Formal class 'fPFOLIOBACKTEST' [package "fPortfolio"] with 4 slots:
 ..@ windows :List of 2
 .. ..$ windows: chr "equidistWindows"
 .. ..$ params :List of 1
 .. .. ..$ horizon: chr "12m"
 ..@ strategy:List of 2
 .. ..$ strategy: chr "tangencyStrategy"
 .. ..$ params   : list()
 ..@ smoother:List of 2
 .. ..$ smoother: chr "emaSmoother"
 .. ..$ params   :List of 4
 .. .. ..$ doubleSmoothing: logi TRUE
 .. .. ..$ lambda         : chr "3m"
 .. .. ..$ skip           : num 0
 .. .. ..$ initialWeights : NULL
 ..@ messages: list()
```

25.2 THE WINDOWS SLOT

The `@windows` slot is a list with two named entries. The first entry named `windows` holds the name of the rolling windows function that defines the ‘backtest’ windows, and the second slot, entitled `params`, holds the parameters, such as the horizon of the windows. The slot and its entries can be extracted and modified by the user through extractor and constructor functions.

LISTING 25.1: EXTRACTOR FUNCTIONS FOR THE `@windows` SLOT OF AN `fPFOLIOBACKTEST` OBJECT

```
Extractor Functions:
getWindows          gets windows slot
getWindowsFun       gets windows function
getWindowsParams    gets windows specific parameters
getWindowsHorizon   gets windows horizon
```

To modify the settings from a portfolio backtest specifications we use the constructor functions.

LISTING 25.2: CONSTRUCTOR FUNCTIONS FOR THE `@windows` SLOT OF AN `fPFOLIOBACKTEST` OBJECT

```
Constructor Functions:
setWindowsFun       sets the name of the windows function
setWindowsParams    sets parameters for the windows function
setWindowsHorizon   sets the windows horizon
```

Note that you can write your own windows function, and if required, you can add additional parameters to the parameter list `params`. This is explained in [section 25.2](#). The extractor and constructor functions can also be used to set the new parameters.

How to inspect the default rolling windows

The default rolling window function in the `fPortfolio` package is called `equidistWindows()` and as the name of the function implies, this function generates windows with fixed equidistant horizons. Their value is set in the `params` list under the name `horizon`. The following code shows how to inspect this value for the default settings.

```
> defaultBacktest <- portfolioBacktest()
> getWindowsFun(defaultBacktest)
[1] "equidistWindows"
> getWindowsParams(defaultBacktest)
```

```
$horizon
[1] "12m"

> getWindowsHorizon(defaultBacktest)

[1] "12m"
```

Bear in mind that the horizon is specified as a span string with integer length, here 12, and the unit, here "m" indicating that we measure spans in months. A windows functions has the two arguments:

```
> args(equidistWindows)
function (data, backtest = portfolioBacktest())
NULL
```

the data and the backtest object. Let us inspect the code of the function.

```
> equidistWindows
function (data, backtest = portfolioBacktest())
{
  horizon = getWindowsHorizon(backtest)
  ans = rollingWindows(x = data, period = horizon, by = "1m")
  ans
}
<environment: namespace:fPortfolio>
```

First, we use the function `getWindowsHorizon()` to extract the horizon, which is the length of the windows returned as a span value. Then we call the function `rollingWindows()` to create the windows. The result returned by the function `rollingWindows()` is a list with two entries named `from` and `to`. These two list entries give the start and the end date for each window. In addition, the `control` attribute holds information about the start and end dates of the whole series, the period length of the windows (also called horizon), and its regular time shift.

```
> swxData <- 100 * SWX.RET
> swxBacktest <- portfolioBacktest()
> setWindowsHorizon(swxBacktest) <- "24m"
> equidistWindows(data = swxData, backtest = swxBacktest)

$from
GMT
[1] [2000-01-01] [2000-02-01] [2000-03-01] [2000-04-01] [2000-05-01]
[6] [2000-06-01] [2000-07-01] [2000-08-01] [2000-09-01] [2000-10-01]
[11] [2000-11-01] [2000-12-01] [2001-01-01] [2001-02-01] [2001-03-01]
[16] [2001-04-01] [2001-05-01] [2001-06-01] [2001-07-01] [2001-08-01]
[21] [2001-09-01] [2001-10-01] [2001-11-01] [2001-12-01] [2002-01-01]
[26] [2002-02-01] [2002-03-01] [2002-04-01] [2002-05-01] [2002-06-01]
[31] [2002-07-01] [2002-08-01] [2002-09-01] [2002-10-01] [2002-11-01]
[36] [2002-12-01] [2003-01-01] [2003-02-01] [2003-03-01] [2003-04-01]
[41] [2003-05-01] [2003-06-01] [2003-07-01] [2003-08-01] [2003-09-01]
[46] [2003-10-01] [2003-11-01] [2003-12-01] [2004-01-01] [2004-02-01]
[51] [2004-03-01] [2004-04-01] [2004-05-01] [2004-06-01] [2004-07-01]
[56] [2004-08-01] [2004-09-01] [2004-10-01] [2004-11-01] [2004-12-01]
```



```

[61] [2005-01-01] [2005-02-01] [2005-03-01] [2005-04-01] [2005-05-01]
[66] [2005-06-01]

$to
GMT
[1] [2001-12-31] [2002-01-31] [2002-02-28] [2002-03-31] [2002-04-30]
[6] [2002-05-31] [2002-06-30] [2002-07-31] [2002-08-31] [2002-09-30]
[11] [2002-10-31] [2002-11-30] [2002-12-31] [2003-01-31] [2003-02-28]
[16] [2003-03-31] [2003-04-30] [2003-05-31] [2003-06-30] [2003-07-31]
[21] [2003-08-31] [2003-09-30] [2003-10-31] [2003-11-30] [2003-12-31]
[26] [2004-01-31] [2004-02-29] [2004-03-31] [2004-04-30] [2004-05-31]
[31] [2004-06-30] [2004-07-31] [2004-08-31] [2004-09-30] [2004-10-31]
[36] [2004-11-30] [2004-12-31] [2005-01-31] [2005-02-28] [2005-03-31]
[41] [2005-04-30] [2005-05-31] [2005-06-30] [2005-07-31] [2005-08-31]
[46] [2005-09-30] [2005-10-31] [2005-11-30] [2005-12-31] [2006-01-31]
[51] [2006-02-28] [2006-03-31] [2006-04-30] [2006-05-31] [2006-06-30]
[56] [2006-07-31] [2006-08-31] [2006-09-30] [2006-10-31] [2006-11-30]
[61] [2006-12-31] [2007-01-31] [2007-02-28] [2007-03-31] [2007-04-30]
[66] [2007-05-31]

attr("control")
attr("control")$start
GMT
[1] [2000-01-04]

attr("control")$end
GMT
[1] [2007-05-08]

attr("control")$period
[1] "24m"

attr("control")$by
[1] "1m"

```

Currently, the `backtest` function is fully tested only for end-of-month windows with varying horizons, shifted monthly. We are working on implementing shifts of arbitrary lengths, so that one can create rolling windows which depend on market volatility or may even be triggered by trading decisions.

How to modify rolling window parameters

The list entry `params` from the `@windows` slot is a list with additional parameters used in different situations. If required, you can enhance this.

LISTING 25.3: LIST ENTRY IN THE `@windows` SLOT AN `fPFOLIOBACKTEST` OBJECT

```

horizon      fixed horizon length used for the default
              windows function
...          your parameter settings for your
              custom windows functions (if required)

```

By default the window size is fixed at 12 months (`horizon = "12m"`). This fixed horizon can be changed with the function `setWindowsHorizon()`.

```
> setWindowsHorizon(backtest) <- "24m"
```

The entire list of parameters can be extracted with the function `getWindowsParams()` and you can add and modify parameter settings with the function `setWindowsParams()`. Note that you must take care not to omit the horizon parameter when setting windows parameters with the `setWindowsParams()` function.

```
> getWindowsParams(backtest)
$horizon
[1] "24m"
```

How to write your own windows function

If you want to add your own rolling windows function you should proceed in the following way

LISTING 25.4: EXAMPLE OF A ROLLING WINDOWS FUNCTION

```
# Set ANY additional windows parameters if required:
setWindowsParams(backtest) <- list(...)

# Template to create your own rolling windows function:
myRollingWindows <- function(data, backtest = portfolioBacktest())
{
  # Code:
  Params <- getWindowsParams(backtest)
  ...

  # Return:
  list(from = <...> , to = <...>)
}
```

In this function template, `data` is a multivariate `timeSeries` object, and `backtest` the portfolio backtest object with class `fPFOLIOBACKTEST`. Additional parameters required by the function `myRollingWindows()` can be passed in through the `list@windows$params` of the backtest object. With the function `getWindowsParams` we can extract its parameter list. Note that `myRollingWindows` must at least return a named list, with two named entries `$from` and `$to`, which give the start and end dates for each backtest window.

As an example we want to create rolling windows which depend on the volatility of the underlying returns. In this case, if the volatility is low, we want to use longer window horizons, and if the volatility increases, then we shorten the window horizons. The result are windows of varying length, dependent on the volatility.

25.3 THE STRATEGY SLOT

The `@strategy` slot is a list with two named entries. The first entry, `strategy`, holds the name of the strategy function that defines the backtest portfolio strategy, and the second, `params`, holds their parameters if required. The slot and its entries can be extracted and modified by through extractor and constructor functions.

LISTING 25.5: EXTRACTOR FUNCTIONS FOR THE `@strategy` SLOT OF AN `fPFOLIOBACKTEST` OBJECT

```
Extractor Functions:
getStrategy      gets strategy slot
getStrategyFun   gets the name of the strategy function
getStrategyParams gets strategy specific parameters
```

To modify the settings for a portfolio backtest strategy we use the constructor functions.

LISTING 25.6: CONSTRUCTOR FUNCTIONS FOR THE `@strategy` SLOT OF AN `fPFOLIOBACKTEST` OBJECT

```
Constructor Functions:
setStrategyFun   sets the name of the strategy function
setStrategyParams sets strategy specific parameters
```

How to inspect the default portfolio strategy

The default rolling portfolio strategy provided by the `fPortfolio` package is the `tangencyStrategy`.

```
> args(tangencyStrategy)
function (data, spec = portfolioSpec(), constraints = "LongOnly",
         backtest = portfolioBacktest())
NULL
```

The first argument expects the data as a `timeSeries` object, the second the portfolio specification as an object of class `fPFOLIOSPEC`, constraints as a string vector, and backtest information from an object of class `fPFOLIOBACKTEST`.

Let us take a look at the code of the very simple portfolio investment strategy defined by the tangency strategy.

```
> tangencyStrategy
function (data, spec = portfolioSpec(), constraints = "LongOnly",
         backtest = portfolioBacktest())
{
  strategyPortfolio <- try(tangencyPortfolio(data, spec, constraints))
}
```

```

    if (class(strategyPortfolio) == "try-error") {
      strategyPortfolio <- minvariancePortfolio(data, spec,
        constraints)
    }
    strategyPortfolio
  }
}
<environment: namespace:fPortfolio>

```

The `tangencyStrategy` invests in a portfolio that has the highest Sharpe ratio, and if such a portfolio does not exist, the minimum variance portfolio is taken instead. Strategy functions always have to return an object of class `fPORTFOLIO`.

The function name of the portfolio strategy can be extracted with the function `getStrategyFun()`.

```

> getStrategyFun(backtest)
[1] "tangencyStrategy"

```

and changed with the function `setStrategyFun()`.

```

> setStrategyFun(backtest) <- "tangencyStrategy"

```

where "tangencyStrategy" can be replaced with the name of the function you wish to use.

How to write your own strategy function

If you want to test your own portfolio strategies, you will need to write your own function. This is shown in the following example; here, we define a function called `myPortfolioStrategy()`.

```

> ## Add parameters needed in the function 'myPortfolioStrategy':
> setStrategyParams(backtest) <- list()
> ## Creating a new portfolio strategy function:
> myPortfolioStrategy <-
  function(data, spec, constraints, backtest)
  {
    ## Extract Parameters:
    Parameters <- getStrategyParams(backtest)

    ## Strategy Portfolio:
    strategyPortfolio <- tangencyPortfolio(data, spec, constraints)

    ## Return :
    strategyPortfolio
  }

```

Here, `data` is a multivariate time series object, `spec` the portfolio specification, `constraints` the string of portfolio constraints and `backtest` the portfolio backtest object. Additional parameters can be passed through

the backtest object with the function `setStrategyParams()` and can be extracted within the portfolio strategy function with `getStrategyParams()`. Note that `myPortfolioStrategy()` function must return an S4 object of class `fPORTFOLIO`.

25.4 THE SMOOTHER SLOT

The `@smoother` slot is a list with two named entries. The first entry named `smoother` holds the name of the smoother function that defines the backtest function to smooth the weights over time. The second, named `params`, holds the required parameters for the smoother function. The slot and its entries can be extracted and modified through extractor and constructor functions.

LISTING 25.7: EXTRACTOR FUNCTIONS FOR THE `@smoother` SLOT OF AN `fPFOLIOBACKTEST` OBJECT

Extractor Functions:	
<code>getSmoother</code>	gets smoother slot
<code>getSmootherFun</code>	gets the name of the smoother function
<code>getSmootherParams</code>	gets parameters for strategy function
<code>getSmootherLambda</code>	gets smoothing parameter lambda
<code>getSmootherDoubleSmoothing</code>	gets setting for double smoothing
<code>getSmootherInitialWeights</code>	gets initial weights used in smoothing
<code>getSmootherSkip</code>	gets number of skipped months

To modify the settings for a portfolio backtest strategy you can use the constructor functions.

LISTING 25.8: CONSTRUCTOR FUNCTIONS FOR THE `@smoother` SLOT OF AN `fPFOLIOBACKTEST` OBJECT

Constructor Functions:	
<code>setSmootherFun</code>	sets the name of the smoother function
<code>setSmootherParams</code>	sets parameters for strategy function
<code>setSmootherLambda</code>	sets smoothing parameter lambda
<code>setSmootherDoubleSmoothing</code>	sets setting for double smoothing
<code>setSmootherInitialWeights</code>	sets initial weights used in smoothing
<code>setSmootherSkip</code>	sets number of skipped months

How to inspect the default smoother function

The default smoother function provided by the `fPortfolio` package is the `emaSmoother`.

```
> args(emaSmoother)
function (weights, spec, backtest)
```

NULL

The first argument expects the weights, the second the portfolio specification spec as an object of class `fPFOLIOSPEC`, and backtest information from an object of class `fPFOLIOBACKTEST`.

Let us examine the code of the exponential moving average (EMA) smoothing approach implemented in the `emaSmoother()` function.

```
> emaSmoother
function (weights, spec, backtest)
{
  ema <- function(x, lambda) {
    x = as.vector(x)
    lambda = 2/(lambda + 1)
    xlam = x * lambda
    xlam[1] = x[1]
    ema = filter(xlam, filter = (1 - lambda), method = "rec")
    ema[is.na(ema)] <- 0
    as.numeric(ema)
  }
  lambda <- getSmootherLambda(backtest)
  lambdaLength <- as.numeric(substr(lambda, 1, nchar(lambda) -
    1))
  lambdaUnit <- substr(lambda, nchar(lambda), nchar(lambda))
  stopifnot(lambdaUnit == "m")
  lambda <- lambdaLength
  nAssets <- ncol(weights)
  initialWeights = getSmootherInitialWeights(backtest)
  if (!is.null(initialWeights))
    weights[1, ] = initialWeights
  smoothWeights1 = NULL
  for (i in 1:nAssets) {
    EMA <- ema(weights[, i], lambda = lambda)
    smoothWeights1 <- cbind(smoothWeights1, EMA)
  }
  doubleSmooth <- getSmootherDoubleSmoothing(backtest)
  if (doubleSmooth) {
    smoothWeights = NULL
    for (i in 1:nAssets) {
      EMA <- ema(smoothWeights1[, i], lambda = lambda)
      smoothWeights = cbind(smoothWeights, EMA)
    }
  }
  else {
    smoothWeights <- smoothWeights1
  }
  rownames(smoothWeights) <- rownames(weights)
  colnames(smoothWeights) <- colnames(weights)
  smoothWeights
}
<environment: namespace:fPortfolio>
```

The `emaSmoother()` applies a single or double EMA filter to a vector of weights. First we define the internal smoother function `ema` and retrieve

the decay parameter `lambda`. Then we apply single or double EMA smoothing to each series of asset weights. Finally, the smoothed weights are returned. Note that in each step you have to make sure that the weights add to one if you are fully invested.

How to modify rolling window parameters

The `emaSmoother()` function is controlled by four parameters, which are defined in the `params` entry of the smoother slot:

LISTING 25.9: PARAMETERS OF THE EMASMOOTHER FUNCTION

Smoother Control Parameters:

<code>doubleSmoothing</code>	logical, TRUE means the EMA filter is applied twice
<code>lambda</code>	character, the amount of smoothing - e.g. "3m", "6m", ...
<code>skip</code>	numeric value, number of months to skip - e.g. 12, 18, ...
<code>initialWeights</code>	numeric vector containing the initial weights

The default settings are:

```
> getSmootherParams(backtest)
$doubleSmoothing
[1] TRUE

$lambda
[1] "3m"

$skip
[1] 0

$initialWeights
NULL
```

To modify these control parameters individually, we use the `setSmoother*` functions.

To change to single smoothing, type

```
> setSmootherDoubleSmoothing(backtest) <- FALSE
```

To modify the smoother's decay length, type

```
> setSmootherLambda(backtest) <- "12m"
```

To start rebalancing 12 months after the original start date, type

```
> setSmootherSkip(backtest) <- "12m"
```

To use equal weights as starting points, type

```
> nAssets <- 5
> setSmootherInitialWeights(backtest) <- rep(1/nAssets, nAssets)
```

After you have made changes, check if your settings are active.

```
> getSmootherParams(backtest)

$doubleSmoothing
[1] FALSE

$lambda
[1] "12m"

$skip
[1] "12m"

$initialWeights
[1] 0.2 0.2 0.2 0.2 0.2
```

How to write your own smoother function

If you want to apply your own weight-smoothing style, you will need to write a custom function. Note that additional smoothing parameters can be passed through the backtest object with the `setSmootherParams()` function, and they can be extracted within the function with a call to `getSmootherParams()`.

The following is an example of how to implement your own smoother function:

LISTING 25.10: CUSTOM SMOOTHER FUNCTION

```
# Add additional parameters used by 'mySmoother':
setSmootherParams(backtest) <- list(...)

# Creating a new smoother function:
mySmoother <- function(weights, spec, backtest)
{
  # Code:
  Params <- getSmootherParams(backtest)
  ... Code to return smoothed Weights ...

  # Return:
  smoothedWeights
}
```

Here, the `weights` are a numeric vector of weights, `spec` is the portfolio specification and `backtest` is the portfolio backtest object.

25.5 ROLLING ANALYSIS

Rolling portfolio analysis is commonly used to backtest the outcome of portfolio optimization over time. A rolling backtest on historical data can give us insights into the stability and performance of a given strategy. The

strategy can then be optimized and improved before applying it to real markets. The function available in the `fPortfolio` package for portfolio backtesting is `portfolioBacktesting()`.

How to backtest a portfolio

The `portfolioBacktesting()` function has the following arguments:

```
> args(portfolioBacktesting)
function (formula, data, spec = portfolioSpec(),
  constraints = "LongOnly", backtest = portfolioBacktest(),
  trace = TRUE)
NULL
```

`portfolioBacktesting()` requires a formula expression that tells us which assets from the data set have to be analyzed against a given benchmark. The second argument, `data`, requests a multivariate timeSeries object, which contains at least the columns referenced in the formula expression. Portfolio specifications are given by the argument `spec` and portfolio constraints are given by `constraints`. The last argument `backtest` takes a portfolio backtest object, as described above.

Suppose we want to backtest the simple tangency portfolio strategy with the following three assets: the SBI (Swiss Bond Index), the SPI (Swiss Performance Index) and the SII (Swiss Immofunds Index), and we want to backtest this strategy against the Pictet Benchmark Index LP40.

```
> swxData <- 100 * SWX.RET
> swxSpec <- portfolioSpec()
> swxConstraints <- "LongOnly"
> swxBacktest <- portfolioBacktest()
> setWindowsHorizon(swxBacktest) <- "18m"
> setSmootherLambda(swxBacktest) <- "6m"
> swxFormula <- LP40 ~ SBI + SPI + SII

> swxPortfolios <- portfolioBacktesting(formula = swxFormula,
  data = swxData, spec = swxSpec, constraints = swxConstraints,
  backtest = swxBacktest, trace = FALSE)
```

The output on the screen is too lengthy to show, so here is a limited trace:

```
Portfolio Backtesting:
Assets:          SBI SPI SII
Benchmark:      LP40
Start Series:   2000-01-04
End Series:    2007-05-08
Type:          MV
Cov Estimator:  covEstimator
Solver:         solveRquadprog
Portfolio Windows:
  Horizon:      18m
Portfolio Strategy: tangencyStrategy
```

```
Portfolio Smoother:  emaSmoother
doubleSmoothing:    TRUE
Lambda:             6m
```

Portfolio Optimization:

Optimization Period	Target	Benchmark	Weights				
2000-01-01 2001-06-30	0	0.002	1	0	0	*	1
2000-02-01 2001-07-31	0.003	0.001	0.975	0.025	0	*	1
2000-03-01 2001-08-31	0.006	-0.006	0.893	0	0.107	*	1
2000-04-01 2001-09-30	0.01	-0.023	0.983	0	0.017	*	1
2000-05-01 2001-10-31	0.015	-0.017	0.956	0	0.044	*	1
2000-06-01 2001-11-30	0.013	-0.012	0.905	0	0.095	*	1
2000-07-01 2001-12-31	0.009	-0.011	0.84	0	0.16	*	1
2000-08-01 2002-01-31	0.011	-0.013	0.727	0	0.273	*	1
2000-09-01 2002-02-28	0.007	-0.02	0.81	0	0.19	*	1
2000-10-01 2002-03-31	0.01	-0.011	0.625	0	0.375	*	1
...							

The `portfolioBacktesting()` function returns a list with the following named elements:

LISTING 25.11: NAMED ELEMENTS OF THE LIST RETURNED BY `portfolioBacktesting()`

<code>portfolioBacktesting</code> Values:	
<code>formula</code>	backtest formula expression
<code>data</code>	multivariate returns
<code>spec</code>	portfolio specifications
<code>constraints</code>	portfolio constraints
<code>backtest</code>	portfolio backtest specifications
<code>benchmarkName</code>	name of the benchmark
<code>assetsNames</code>	names of the assets
<code>weights</code>	a matrix of portfolio weights
<code>strategyList</code>	a list of the invested portfolios

How to smooth the weights from a backtest

As you can see from the above R output, the portfolio weights may vary strongly as market conditions change from one optimization period to the next. It is often a reasonable idea to smooth the weights first before applying them to the portfolio performance analysis. The function provided in the `fPortfolio` package for weights smoothing is `portfolioSmoothing()`.

```
> args(portfolioSmoothing)
function (object, backtest = NULL, trace = TRUE)
NULL
```

The first argument `object` is simply the list returned by the `portfolioBacktesting()` function.

An advantage of separating the process of backtesting from that of smoothing is that once we have backtested a portfolio, we can investigate the effect of alternative weights smoothers on the performance of the portfolio, without backtesting again and again for different types of smoothers. In `fPortfolio` the default smoother function is the `emaSmoother()`, which applies an exponential moving average, EMA, filter to the portfolio weights. There are several control parameters for the EMA smoother function, but for the following example we will use the default settings, except that we start with an equally weighted portfolio and set the decay parameter for the EMA to `lambda = "12m"`.

```
> setSmootherInitialWeights(swxPortfolios$backtest) <- rep(1/3,
3)
> setSmootherLambda(swxPortfolios$backtest) <- "12m"

> swxSmooth <- portfolioSmoothing(swxPortfolios)
```

The `portfolioSmoothing()` function returns a list with the following named elements:

LISTING 25.12: NAMED ELEMENTS OF THE LIST RETURNED BY `portfolioSmoothing()`

```
portfolioSmoothing - Values:
smoothWeights      matrix of smoothed portfolio weights
monthlyAssets      timeSeries of monthly asset returns
monthlyBenchmark    timeSeries of monthly benchmark returns
portfolioReturns    timeSeries of cumulated monthly portfolio returns
benchmarkReturns    timeSeries of cumulated monthly benchmark returns
P                   timeSeries of monthly portfolio returns
B                   timeSeries of monthly benchmark returns
stats               matrix of basic backtest statistics
```

The list returned by `portfolioSmoothing()` also inherits some of the named elements from the list returned by `portfolioBacktesting()`.

How to plot backtesting results

The function `backtestPlot()` provides graphs and statistics to summarize the backtesting results:

```
> backtestPlot
function (object, which = "all", labels = TRUE, legend = TRUE,
  at = NULL, format = NULL, cex = 0.6, font = 1, family = "mono")
{
  if (any(which == "all"))
    par(mfrow = c(3, 2), mar = c(1.5, 4, 5, 2), oma = c(5,
      1, 0, 1))
  if (any(which == "1") || which == "all")
    backtestAssetsPlot(object, labels, legend, at, format)
```

```

if (any(which == "2") || which == "all")
  backtestWeightsPlot(object, labels, legend, at, format)
if (any(which == "3") || which == "all")
  backtestRebalancePlot(object, labels, legend, at, format)
if (any(which == "4") || which == "all")
  backtestPortfolioPlot(object, labels, legend, at, format)
if (any(which == "5") || which == "all")
  backtestDrawdownPlot(object, labels, legend, at, format)
if (any(which == "6") || which == "all")
  backtestReportPlot(object, cex = cex, font = font, family = family)
invisible()
}
<environment: namespace:fPortfolio>

```

The first argument object is the list returned by the `portfolioSmoothing()` function followed by which, either a numeric vector from 1 to 6 or "all" which plots all 6 graphs. For a plot of the current example, see Figure Figure 25.1.

```
> backtestPlot(swxSmooth, cex = 0.6, font = 1, family = "mono")
```

How to print backtest statistics

`backtestStats()` is a wrapper function that gathers rolling statistics over each of the backtest windows.

```

> args(backtestStats)
function (object, FUN = "rollingSigma", ...)
NULL

```

The function requires an object returned by the function `portfolioSmoothing()` and the name of the `stats()` function given as a character string that computes the statistics. By default, the portfolio risk, σ , is calculated with the function `rollingSigma()`.

```

> defaultStats <- backtestStats(swxSmooth)
> head(defaultStats)

GMT
      Sigma
2001-06-29 0.099987
2001-07-31 0.093903
2001-08-31 0.094450
2001-09-28 0.093175
2001-10-31 0.085676
2001-11-30 0.092309

```

Bear in mind that there are other predefined rolling statistics functions available in `fPortfolio`:

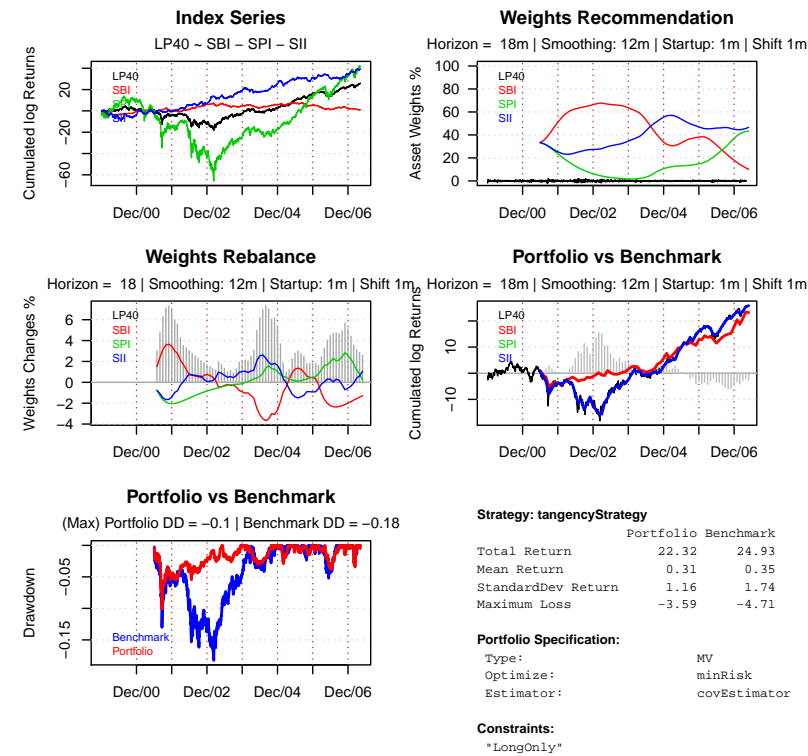


FIGURE 25.1: Portfolio backtesting for major Swiss indices: The backtesting is performed for a portfolio composed of three Swiss market indexes, equities, bonds and reits. The *Series* graph shows the time series indexes, the *Weights Recommendation* graph shows the smoothed recommended weights every month for the investment of the next month, the *Weights Rebalance* graph shows to which amount the weights were rebalanced, and the *Drawdowns* graph shows the drawdowns of the optimized portfolio at the end of each month in comparison to the benchmark. The table gives some characterizing numbers for the optimized portfolio and benchmark.

LISTING 25.13: PREDEFINED ROLLING STATISTICS FUNCTION IN FPORTFOLIOBACKTEST

rollingSigma	portfolio risk Sigma over time
rollingVaR	rolling Value at Risk
rollingCVaR	rolling Conditional Value at Risk
rollingDaR	rolling Drawdowns at Risk
rollingCDaR	rolling Conditional Drawdowns at Risk
rollingRiskBudgets	rolling Risk Budget

How to write your own statistics functions

You can write your own stats functions to compute other statistics such as the Drawdown at Risk (DaR), Conditional Drawdown at Risk (CDaR), Shapiro-Wilk's test statistic, etc. Note that the function must take a list of `fPORTFOLIO` objects¹ as an argument, and additional arguments are passed in through the `...` argument in `backtestStats(object, stats, ...)`.

As an example, the following function, `rollingCDaR()`, returns a rolling estimate of the Conditional Drawdown at Risk.

```
> rollingCDaR
function (object)
{
  .cdar <- function(x) {
    alpha <- getAlpha(x)
    R <- as.numeric(getSeries(x) %**% getWeights(x))
    dd <- 100 * drawdowns(as.timeSeries(R)/100)
    z <- quantile(as.numeric(dd), probs = alpha)
    mean(dd[dd <= z])
  }
  portfolios <- object$strategyList
  ans <- sapply(portfolios, FUN = .cdar)
  dates <- sapply(portfolios, function(x) rev(rownames(getSeries(x)))[1])
  alpha <- getAlpha(portfolios[[1]])
  timeSeries(ans, charvec = dates, units = paste("CDaR", alpha,
    sep = "."))
}
<environment: namespace:fPortfolio>
```

The CDaR for the current example is:

```
> CDaRstats <- backtestStats(swxSmooth, FUN = "rollingCDaR")
> head(CDaRstats)

GMT                CDaR.0.05
2001-06-29    -3.7913
2001-07-31    -2.7909
2001-08-31    -2.3463
```

¹For example like the list returned by the `portfolioBacktesting()` function under the name `$strategyList`




FIGURE 25.2: Rolling Analysis of the Swiss Performance Index: This plot shows the covariance, Value-At-Risk, Conditional Value-At-Risk and Conditional Drawdowns-At-Risk over time. The latter three all have a confidence level (α) of 0.05.

2001-09-28 -1.7138
2001-10-31 -1.0517
2001-11-30 -1.2134

It can be very helpful to plot various risk measures; the covariance, Value-At-Risk, Conditional Value-At-Risk and Conditional Drawdowns-At-Risk are displayed in [Figure 25.2](#).

CHAPTER 26

CASE STUDY: SPI SECTOR ROTATION

```
> library(fPortfolio)
```

In this chapter we will demonstrate how to backtest portfolio strategies on realistic portfolios. The first case study concerns the sector rotation of Swiss equities.

26.1 SPI PORTFOLIO BACKTESTING

The data we will be using for this case study are the returns from the Swiss Performance Index, SPI. This data set, `SPISECTOR.RET`, is comprised of 9 sectors; basic materials, industrials, consumer goods, health care, consumer service, telecommunications, utilities, finance and technology. For further details, see [section B.3](#).

To view the names of the individual sectors, type:

```
> colnames(SPISECTOR.RET)
[1] "SPI" "BASI" "INDU" "CONG" "HLTH" "CONS" "TELE" "UTIL" "FINA"
[10] "TECH"
```

The data set contains historical asset returns from January 2000 to June 2008, and the strategy we are going to backtest is to invest in the tangency portfolio

strategy with a fixed rolling window of 12 months shifted monthly. We call this investment type in the following *tangency strategy*. The portfolios will be optimized with the mean-variance Markowitz approach and we will use the sample covariance as the risk measure. First we specify the settings for the portfolio data, for the specification, for the constraints and finally for the portfolio backtest.

```
> spiData <- SPISECTOR.RET
> spiSpec <- portfolioSpec()
> spiConstraints <- "LongOnly"
> spiBacktest <- portfolioBacktest()
```

Then we specify the assets which should be used for backtesting.

```
> spiFormula <- SPI ~ BASI + INDU + CONG + HLTH + CONS + TELE +
  UTIL + FINA + TECH
```

Next, we optimize the rolling portfolios and perform the backtests.

```
> spiPortfolios <- portfolioBacktesting(formula = spiFormula,
  data = spiData, spec = spiSpec, constraints = spiConstraints,
  backtest = spiBacktest, trace = FALSE)
```

The weights of the first 12 months for the portfolios which are rebalanced every month are given by

```
> Weights <- round(100 * spiPortfolios$weights, 2)[1:12, ]
> Weights
```

	BASI	INDU	CONG	HLTH	CONS	TELE	UTIL	FINA	TECH
2000-12-31	0	0	48.08	0	0	0	27.54	16.06	8.33
2001-01-31	0	0	22.25	0	0	0	28.62	49.13	0.00
2001-02-28	0	0	31.15	0	0	0	32.80	36.05	0.00
2001-03-31	0	0	51.92	0	0	0	48.08	0.00	0.00
2001-04-30	0	0	39.77	0	0	0	46.70	13.53	0.00
2001-05-31	0	0	35.16	0	0	0	64.68	0.16	0.00
2001-06-30	0	0	47.84	0	0	0	52.16	0.00	0.00
2001-07-31	0	0	27.19	0	0	0	72.81	0.00	0.00
2001-08-31	0	0	0.00	0	0	0	100.00	0.00	0.00
2001-09-30	0	0	0.00	0	0	100	0.00	0.00	0.00
2001-10-31	0	0	0.00	0	0	100	0.00	0.00	0.00
2001-11-30	0	0	0.00	0	0	100	0.00	0.00	0.00

We see that the weights are fluctuating significantly and thus we smooth them to prevent to reduce a cost effective monthly rebalancing.

26.2 SPI PORTFOLIO WEIGHTS SMOOTHING

We set the smoothing parameter lambda to 12 months to increase the smoothing effect and we have taken the recommended initial weights from the portfolio backtesting function as opposed to setting our own.

```
> setSmootherLambda(spiPortfolios$backtest) <- "12m"
> spiSmoothPortfolios <- portfolioSmoothing(object = spiPortfolios,
  trace = FALSE)
```

The weights during the first 12 months are now

```
> smoothWeights <- round(100 * spiSmoothPortfolios$smoothWeights,
  2)[1:12, ]
> smoothWeights
```

	BASI	INDU	CONG	HLTH	CONS	TELE	UTIL	FINA	TECH
2000-12-31	0	0	48.08	0	0	0.00	27.54	16.06	8.33
2001-01-31	0	0	47.47	0	0	0.00	27.56	16.84	8.13
2001-02-28	0	0	46.64	0	0	0.00	27.70	17.86	7.80
2001-03-31	0	0	46.18	0	0	0.00	28.29	18.16	7.37
2001-04-30	0	0	45.70	0	0	0.00	29.14	18.27	6.89
2001-05-31	0	0	45.10	0	0	0.00	30.59	17.92	6.39
2001-06-30	0	0	44.74	0	0	0.00	32.15	17.24	5.88
2001-07-31	0	0	44.06	0	0	0.00	34.22	16.35	5.37
2001-08-31	0	0	42.54	0	0	0.00	37.26	15.32	4.88
2001-09-30	0	0	40.44	0	0	2.37	38.55	14.23	4.41
2001-10-31	0	0	37.98	0	0	6.37	38.57	13.11	3.98
2001-11-30	0	0	35.32	0	0	11.46	37.67	11.99	3.57

Note that the process of the re-balancing now appears in a much smoother way.

26.3 SPI PORTFOLIO BACKTEST PLOTS

Figure 26.1 summarizes the backtest results.

```
> backtestPlot(spiSmoothPortfolios, cex = 0.6, font = 1, family = "mono")
```

The function `backtestPlot()` shows five different views including the series, the recommended weights, the rebalancing of weights, the performance graphs and a drawdown plot.

26.4 SPI PERFORMANCE REVIEW

As we can see from Figure 26.1, if we had started the tangency strategy in January 2001, the total return for the portfolio would have been around 79.60% compared to 24.71% if we had invested in the SPI. Furthermore, during 2003, when the market was on a downward trend, the portfolio was able to absorb the losses better than the SPI, with the portfolio's maximum drawdown during that period being only -0.28 compared to -0.54 for the SPI. The amount of re-balancing seems reasonable as total weight changes per month were at most around 9%.

```
> netPerformance(spiSmoothPortfolios)

Net Performance % to 2008-10-31:
      1 mth 3 mths 6 mths  1 yr 3 yrs 5 yrs 3 yrs p.a. 5 yrs p.a.
Portfolio -0.21 -0.25 -0.31 -0.44 -0.26  0.49      -0.09   0.10
Benchmark -0.13 -0.17 -0.22 -0.38 -0.29  0.29      -0.10   0.06

Net Performance % Calendar Year:
      2001 2002 2003 2004 2005 2006 2007  YTD Total
Portfolio -0.10 -0.09 0.25 0.25 0.22 0.29 0.05 -0.39  0.48
Benchmark -0.25 -0.30 0.20 0.07 0.30 0.19 0.00 -0.32 -0.11
```

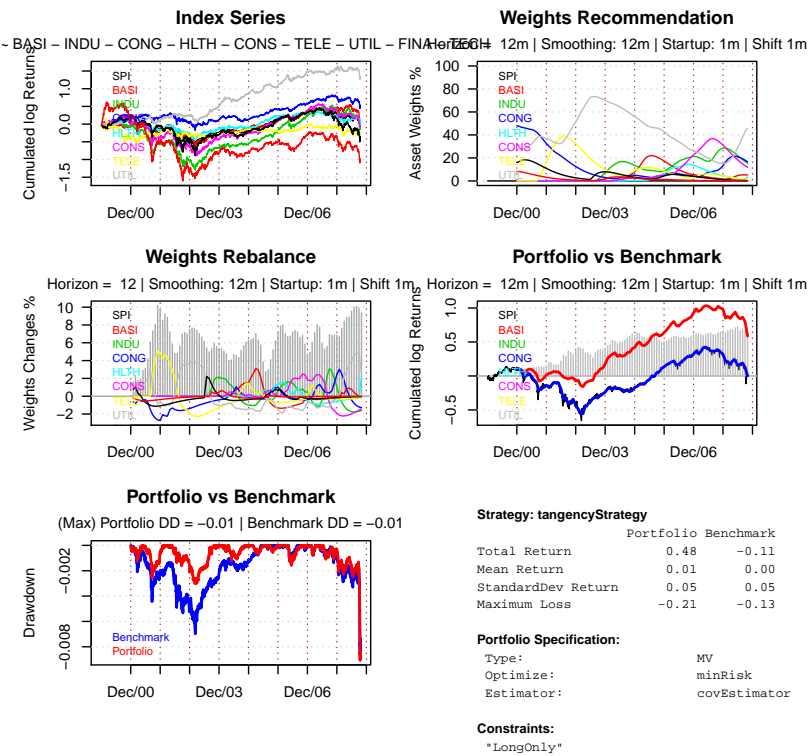


FIGURE 26.1: The five graphs show the results from portfolio backtesting with instruments from the SPI Subsectors. The graph to the upper left shows the subsector time series. To the upper right we see the weights recommendation over time. We have rolled a window with a 12 month horizon every month. The weights are smoothed with a double exponential EMA smoother with a time decay of 6 months. The end-of-month rebalancing is derived from the weights recommendation shown in the left graph in the middle of the graph sheets. The next graph to the right shows the development of the monthly returns over time. The last plot shows the rolling drawdowns for the optimized portfolio and the benchmark index.

Over the last year the portfolio lost 13.04% in returns, 7% less than the SPI. For the past seven calendar years the portfolio strategy seemed to perform much better than the SPI, except for 2005. The portfolio strategy gave us relatively small losses in 2001 and 2002, and yielded significant gains in years 2003 to 2007.

CHAPTER 27

CASE STUDY: GCC INDEX ROTATION

```
> library(fPortfolio)
```

The MSCI GCC Countries Indices was launched in January 2006 to reflect growing investor interest in the Gulf region. The GCC index is a comprehensive family of equity markets traded in the GCC region, which covers Bahrain, Kuwait, Oman, Qatar and the United Arab Emirates. The index excludes Saudi Arabia because it is not open to foreign investment. In this case study, we will backtest the five indices open to foreign investment against the MSCI GCC index.

All of the indices have daily history going back to May 31, 2005 and the most recent data we have is to July 28, 2008. Again, the tangency strategy will be applied with a fixed rolling window of 12 months. However, this time the portfolios will be optimized with the Conditional Value at Risk (CVaR) approach with $\alpha = 0.05$.

First, let us look at the column names of the GCC Index data set

```
> colnames(GCCINDEX.RET)
[1] "BAHDSC" "BAHSC" "KUWDSC" "OMADSC" "OMASC"
[6] "KSADSC" "UAEDSC" "UAESC" "QATSC" "GCCXSASC"
[11] "GCCSC"
```

For further information on this data set, see [section B.5](#). Next, we define the specification

```
> gccData <- GCCINDEX.RET
> gccSpec <- portfolioSpec()
> setType(gccSpec) <- "CVaR"
> setSolver(gccSpec) <- "solveRglpk.CVAR"
```

and then set the constraints.

```
> gccConstraints <- "LongOnly"
```

Let us use the default settings for the backtest, and select the instruments from which we build the portfolio, using the formula notation.

```
> gccBacktest <- portfolioBacktest()
> gccFormula <- GCCSC ~ BAHDSC + KUWDSC + OMADSC + UAEDSC +
  QATSC
```

Now, we are ready to combine all of the above:

```
> gccPortfolios <- portfolioBacktesting(formula = gccFormula,
  data = gccData, spec = gccSpec, constraints = gccConstraints,
  backtest = gccBacktest, trace = FALSE)
```

27.1 GCC PORTFOLIO WEIGHTS SMOOTHING

Since we only have three years of historical data, we will change the smoothing parameter to `lambda="6m"`. Shortening the smoothing parameter (`lambda`) means that the portfolio is more responsive to changes in the market, which sounds like a good portfolio strategy. However, it does come with higher transaction costs, because the portfolio is likely to be rebalanced whenever the market situation changes¹.

```
> setSmootherLambda(gccPortfolios$backtest) <- "6m"
> gccSmooth <- portfolioSmoothing(gccPortfolios)
```

27.2 GCC PERFORMANCE REVIEW

The backtest plots suggest that the tangency strategy could also be an effective portfolio strategy for investing in the Gulf region. The results show that the total return for the portfolio over the two years was 43.07%, which is considerably greater than the benchmark at 5.39%. Maximum drawdown for the portfolio was about half that of the benchmark over the two years.

```
> backtestPlot(gccSmooth, cex = 0.6, font = 1, family = "mono")
```

However, the backtest identified a few caveats with the strategy, the main one being diversification or the lack thereof, see [Figure 27.1](#). The result of this poorly diversified portfolio is clustering of risk. In this case, the majority of the portfolio risk came from the Kuwait index. Furthermore, during the months where more than 15% of the portfolio was rebalanced, the transaction costs would reduce the actual return from the portfolio.

¹Note that the rolling windows are generated by the `equidistWindows()` function and that portfolios are re-balanced monthly

Note that a risk-seeking investor may well justify these concerns with having higher returns, but a strategy that operates with a more diversified portfolio may be easier to market.

```
> netPerformance(gccSmooth)
Net Performance % to 2008-07-31:
      1 mth 3 mths 6 mths 1 yrs 2 yrs 2 yrs p.a.
Portfolio -0.03 -0.02  0.03  0.14  0.47      0.23
Benchmark -0.05 -0.09 -0.05  0.11  0.10      0.05

Net Performance % Calendar Year:
      2006 2007   YTD Total
Portfolio  0.06 0.28  0.09  0.43
Benchmark -0.19 0.38 -0.14  0.05
```

27.3 ALTERNATIVE STRATEGY

An alternative strategy would be to start with an equally weighted portfolio. In order to reduce overall weight changes we increase `lambda` from "6m" to "12m".

```
> setSmootherLambda(gccPortfolios$backtest) <- "12m"
> setSmootherInitialWeights(gccPortfolios$backtest) <- rep(1/5,
5)

> gccSmoothAlt <- portfolioSmoothing(gccPortfolios)
```

Notice how overall weight changes have dropped to below 8% per month and the portfolio is much more diversified, see [Figure 27.2](#).

```
> backtestPlot(gccSmoothAlt, cex = 0.6, font = 1, family = "mono")
```

Total return over the two years is lower than the previous strategy (43.07% to 31.63%) but, surprisingly, it is still considerably greater than the benchmark.

```
> netPerformance(gccSmoothAlt)
Net Performance % to 2008-07-31:
      1 mth 3 mths 6 mths 1 yrs 2 yrs 2 yrs p.a.
Portfolio -0.03 -0.02  0.01  0.14  0.34      0.17
Benchmark -0.05 -0.09 -0.05  0.11  0.10      0.05

Net Performance % Calendar Year:
      2006 2007   YTD Total
Portfolio  0.01 0.24  0.07  0.32
Benchmark -0.19 0.38 -0.14  0.05
```

From the above output we can see that the defining point for this strategy was in 2006, where it avoided losses even when the majority of the market

was in decline. The ability to minimize losses when the market conditions are poor seems to be a recurring feature of this tangency strategy. However, in saying that, we have only backtested this strategy with the two examples, further investigations are required to support this statement.

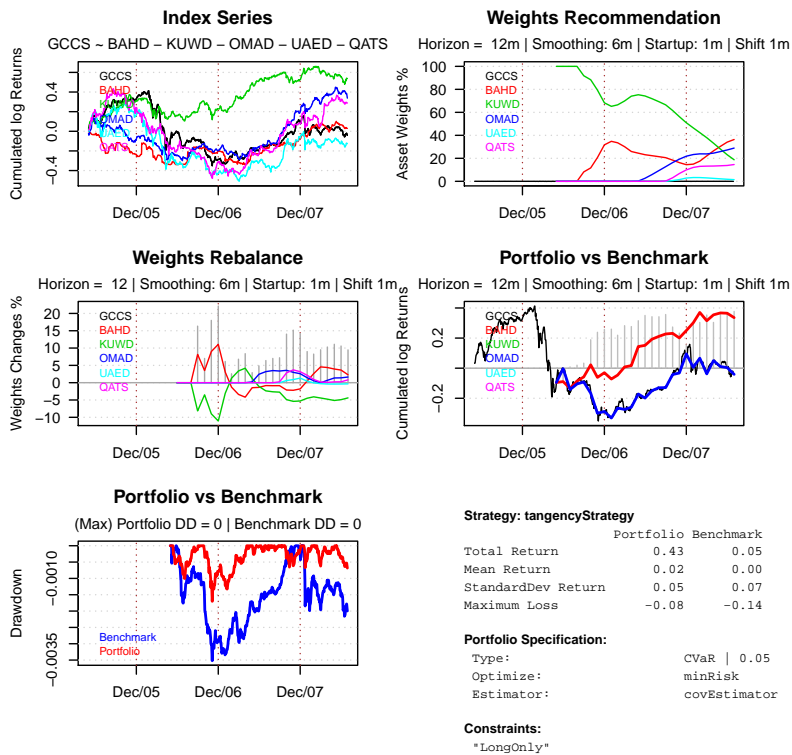


FIGURE 27.1: The five graphs show the results from portfolio backtesting with instruments from the GCC index. The graph to the upper left shows the five series which we have selected from the GCC indexes. To the upper right we see the weights recommendation over time. We have rolled a window with a 12 month horizon every month. The weights are smoothed with a double exponential EMA smoother with a time decay of 6 months. The end-of-month re-balancing is derived from the weights recommendation shown in the left graph in the middle of the graph sheets. The next graph to the right shows the development of the monthly returns over time. The last plot shows the rolling drawdowns for the optimized portfolio and the benchmark index, the GCC Market Index.

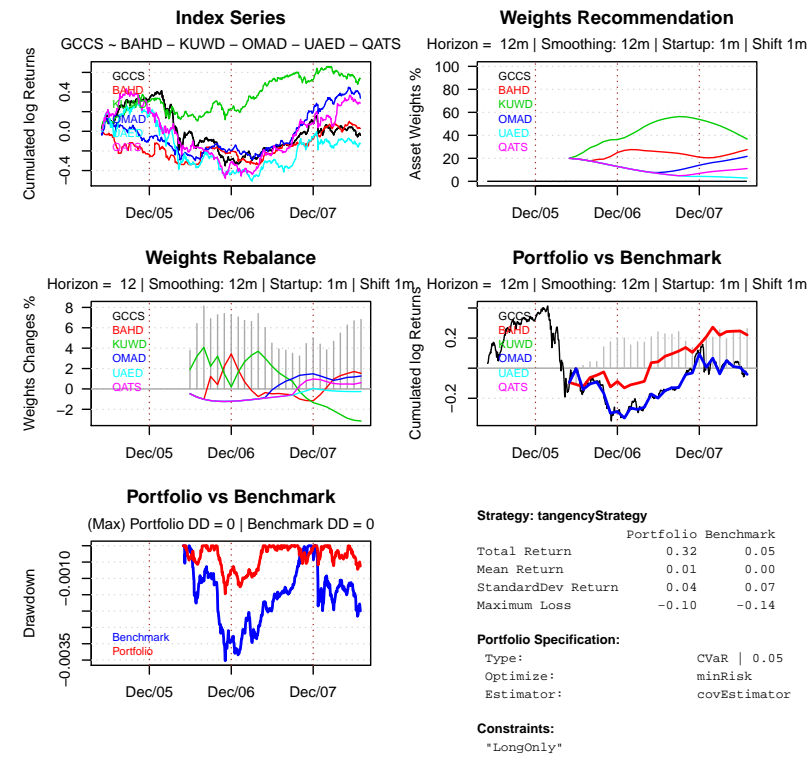


FIGURE 27.2: Backtesting for the GCC Index with alternative strategy: The graph to the upper left shows the five series which we have selected from the GCC indexes. To the upper right we see the weights recommendation over time. We have rolled a window with a 12 month horizon every month. The weights are smoothed with a double exponential EMA smoother with a time decay of 12 months, as opposed to 6 months in the previous figure. The end-of-month re-balancing is derived from the weights shown in the left graph in the middle row. The next graph to the right shows the development of the monthly returns over time. The last plot shows the rolling drawdowns for the optimized portfolio and the benchmark index, the GCC Market Index.

PART VII

APPENDIX

APPENDIX A

PACKAGES REQUIRED FOR THIS EBOOK

```
> library(fPortfolio)
```

In the following we briefly describe the packages required for this ebook. There is one major package named `fPortfolio`. It allows us to model mean-variance and mean-CVaR portfolios with linear constraints and to analyze the data sets of assets used in the portfolios. It also adds additional functionality, including backtesting functions over rolling windows.

A.1 RMETRICS PACKAGE: FPORTFOLIO

`fPortfolio` (Würtz & Chalabi, 2009a) contains the R functions for solving mean-variance and mean-CVaR portfolio problems with linear constraints. The package depends on the contributed R packages `quadprog` (Weingesel, 2004) for quadratic programming problems and `Rglpk` (Theussl & Hornik, 2009) with the appropriate solvers for quadratic and linear programming problems.

```
> listDescription(fPortfolio)
Package: fPortfolio
Title: Rmetrics - Portfolio Selection and Optimization
Date: 2014-10-30
Version: 3011.81
Author: Rmetrics Core Team, Diethelm Wuertz [aut], Tobias Setz
       [cre], Yohan Chalabi [ctb]
Maintainer: Tobias Setz <tobias.setz@rmetrics.org>
Description: Environment for teaching "Financial Engineering and
            Computational Finance".
Depends: R (>= 2.15.1), methods, timeDate, timeSeries, fBasics,
        fAssets
```

```
Imports: fCopulae, robustbase, MASS, Rglpk, slam, Rsymphony,
        Rsolnp, kernlab, quadprog, rneos
Suggests: Rsocp, Rnlnmb2, Rdonlp2, dplR, bcp, fGarch, mvoutlier
Additional_repositories: http://r-forge.r-project.org/
Note: SEVERAL PARTS ARE STILL PRELIMINARY AND MAY BE CHANGED IN
      THE FUTURE. THIS TYPICALLY INCLUDES FUNCTION AND ARGUMENT
      NAMES, AS WELL AS DEFAULTS FOR ARGUMENTS AND RETURN VALUES.
LazyData: yes
License: GPL (>= 2)
URL: https://www.rmetrics.org
Packaged: 2014-10-30 14:13:59 UTC; Tobi
NeedsCompilation: no
Repository: CRAN
Date/Publication: 2014-10-30 15:54:08
Built: R 3.1.2; ; 2015-01-26 00:50:01 UTC; windows
```

A.2 RMETRICS PACKAGE: `timeDate`

`timeDate` (Würtz & Chalabi, 2009b) contains R functions to handle time, date and calendar aspects. The S4 `timeDate` class is used in Rmetrics for financial data and time management together with the management of public and ecclesiastical holidays. The class fulfils the conventions of the ISO 8601 standard as well as of the ANSI C and POSIX standards. Beyond these standards, Rmetrics has added the ‘Financial Center’ concept, which allows you to handle data records collected in different time zones and combine them with the proper time stamps of your personal financial center, or, alternatively, to the GMT reference time. The S4 class can also handle time stamps from historical data records from the same time zone, even if the financial centers changed daylight saving times at different calendar dates. Moreover, `timeDate` is almost compatible with Insightful’s SPlus `timeDate` class. If you move between the two worlds of R and SPlus, you will not have to rewrite your code. This is important for many business applications. The class offers not only date and time functionality, but also sophisticated calendar manipulations for business days, weekends, public and ecclesiastical holidays. `timeSeries` can be downloaded from the CRAN server. Development versions are also available from the R-forge repository.

```
> listDescription(timeDate)

Package: timeDate
Title: Rmetrics - Chronological and Calendar Objects
Date: 2015-01-22
Version: 3012.100
Author: Rmetrics Core Team, Diethelm Wuertz [aut], Tobias Setz
       [cre], Yohan Chalabi [ctb], Martin Maechler [ctb], Joe W.
       Byers [ctb]
Maintainer: Tobias Setz <tobias.setz@rmetrics.org>
Description: Environment for teaching "Financial Engineering and
```



```

Computational Finance". Managing chronological and
calendar objects.
Depends: R (>= 2.15.1), graphics, utils, stats, methods
Suggests: date, RUnit
Note: SEVERAL PARTS ARE STILL PRELIMINARY AND MAY BE CHANGED IN
      THE FUTURE. THIS TYPICALLY INCLUDES FUNCTION AND ARGUMENT
      NAMES, AS WELL AS DEFAULTS FOR ARGUMENTS AND RETURN VALUES.
LazyData: yes
License: GPL (>= 2)
URL: https://www.rmetrics.org
Packaged: 2015-01-23 00:36:51 UTC; Tobi
NeedsCompilation: no
Repository: CRAN
Date/Publication: 2015-01-23 09:30:30
Built: R 3.1.2; ; 2015-01-23 23:20:50 UTC; windows

```

A.3 RMETRICS PACKAGE: timeSeries

timeSeries (Würtz & Chalabi, 2009c) is the Rmetrics package that allows us to work very efficiently with S4 timeSeries objects. Let us briefly summarize the major functions available in this package. You can create timeSeries objects in several different ways, i.e. you can create them from scratch or you can read them from a file. You can print and plot these objects, and modify them in many different ways. Rmetrics provides functions that compute financial returns from price/index series or the cumulated series from returns. Further modifications deal with drawdowns, durations, spreads, midquotes and may other special series. timeSeries objects can be subset in several ways. You can extract time windows, or you can extract start and end data records, and you can aggregate the records on different time scale resolutions. Time series can be ordered and resampled, and can be grouped according to statistical approaches. You can apply dozens of math operations on time series. timeSeries can also handle missing values.

```

> listDescription(timeSeries)

Package: timeSeries
Title: Rmetrics - Financial Time Series Objects
Date: 2015-01-22
Version: 3012.99
Author: Rmetrics Core Team, Diethelm Wuertz [aut], Tobias Setz
       [cre], Yohan Chalabi [ctb]
Maintainer: Tobias Setz <tobias.setz@rmetrics.org>
Description: Environment for teaching "Financial Engineering and
            Computational Finance". Managing financial time series
            objects.
Depends: R (>= 2.10), graphics, grDevices, stats, methods, utils,
         timeDate (>= 2150.95)
Suggests: RUnit, robustbase, xts, PerformanceAnalytics, fTrading
Note: SEVERAL PARTS ARE STILL PRELIMINARY AND MAY BE CHANGED IN

```

```

    THE FUTURE. THIS TYPICALLY INCLUDES FUNCTION AND ARGUMENT
    NAMES, AS WELL AS DEFAULTS FOR ARGUMENTS AND RETURN VALUES.
LazyData: yes
License: GPL (>= 2)
URL: http://www.rmetrics.org
Packaged: 2015-01-23 00:37:25 UTC; Tobi
NeedsCompilation: no
Repository: CRAN
Date/Publication: 2015-01-23 09:55:08
Built: R 3.1.2; ; 2015-01-25 23:38:38 UTC; windows

```

A.4 RMETRICS PACKAGE: fBasics

fBasics (Würtz, 2009a) provides basic functions to analyze and to model data sets of financial time series. The topics from this package include distribution functions for the generalized hyperbolic distribution, the stable distribution, and the generalized lambda distribution. Beside the functions to compute density, probabilities, and quantiles, you can find there also random number generators, functions to compute moments and to fit the distributional parameters. Matrix functions, functions for hypothesis testing, general utility functions and plotting functions are further important topics of the package.

```

> listDescription(fBasics)

Package: fBasics
Title: Rmetrics - Markets and Basic Statistics
Date: 2014-10-29
Version: 3011.87
Author: Rmetrics Core Team, Diethelm Wuertz [aut], Tobias Setz
       [cre], Yohan Chalabi [ctb]
Maintainer: Tobias Setz <tobias.setz@rmetrics.org>
Description: Environment for teaching "Financial Engineering and
             Computational Finance".
Depends: R (>= 2.15.1), timeDate, timeSeries
Imports: gss, stabledist, MASS
Suggests: methods, spatial, RUnit, tcltk, akima
Note: SEVERAL PARTS ARE STILL PRELIMINARY AND MAY BE CHANGED IN
      THE FUTURE. THIS TYPICALLY INCLUDES FUNCTION AND ARGUMENT
      NAMES, AS WELL AS DEFAULTS FOR ARGUMENTS AND RETURN VALUES.
LazyData: yes
License: GPL (>= 2)
URL: https://www.rmetrics.org
Packaged: 2014-10-29 17:34:48 UTC; Tobi
NeedsCompilation: yes
Repository: CRAN
Date/Publication: 2014-10-29 20:07:26
Built: R 3.1.2; x86_64-w64-mingw32; 2015-01-26 00:10:28 UTC;
       windows

```

A.5 RMETRICS PACKAGE: FASSETS

fAssets (Würtz, 2009a) provides functions to analyze and to model multivariate data sets of financial asset returns. The package depends on R's recommended packages methods and MASS (Venables & Ripley, 2008). It also depends on the contributed R packages sn (which depends on mnormt), and robustbase.

```
> listDescription(fAssets)
Package: fAssets
Title: Rmetrics - Analysing and Modelling Financial Assets
Date: 2014-10-30
Version: 3011.83
Author: Rmetrics Core Team, Diethelm Wuertz [aut], Tobias Setz
       [cre], Yohan Chalabi [ctb]
Maintainer: Tobias Setz <tobias.setz@rmetrics.org>
Description: Environment for teaching "Financial Engineering and
             Computational Finance".
Depends: R (>= 2.15.1), timeDate, timeSeries, fBasics
Imports: fMultivar, robustbase, MASS, sn, ecodist, mvnrmtest,
         energy
Suggests: methods, mnormt, RUnit
Note: SEVERAL PARTS ARE STILL PRELIMINARY AND MAY BE CHANGED IN
      THE FUTURE. THIS TYPICALLY INCLUDES FUNCTION AND ARGUMENT
      NAMES, AS WELL AS DEFAULTS FOR ARGUMENTS AND RETURN VALUES.
LazyData: yes
License: GPL (>= 2)
URL: https://www.rmetrics.org
Packaged: 2014-10-30 11:52:35 UTC; Tobi
NeedsCompilation: no
Repository: CRAN
Date/Publication: 2014-10-30 13:38:28
Built: R 3.1.2; ; 2015-01-26 00:33:51 UTC; windows
```

A.6 CONTRIBUTED R PACKAGE: QUADPROG

quadprog implements the dual method of Goldfarb & Idnani (1982, 1983) for solving quadratic programming problems with linear constraints. The original S package was written by Turlach, the R port was done by Weingesel (2004), who also maintains the package. The contributed R package quadprog is the default solver in Rmetrics for quadratic programming problems.

```
> listDescription(quadprog)
Package: quadprog
Type: Package
Title: Functions to solve Quadratic Programming Problems.
Version: 1.5-5
Date: 2013-04-17
Author: S original by Berwin A. Turlach <Berwin.Turlach@gmail.com>
```

```

R port by Andreas Weingessel
<Andreas.Weingessel@ci.tuwien.ac.at>
Maintainer: Berwin A. Turlach <Berwin.Turlach@gmail.com>
Description: This package contains routines and documentation for
              solving quadratic programming problems.
Depends: R (>= 2.15.0)
License: GPL (>= 2)
Packaged: 2013-04-17 08:35:43 UTC; berwin
NeedsCompilation: yes
Repository: CRAN
Date/Publication: 2013-04-17 13:42:49
Built: R 3.1.2; x86_64-w64-mingw32; 2014-11-01 02:06:38 UTC;
       windows

```

A.7 CONTRIBUTED PACKAGE: RGLPK

Rglpk is the R interface to the GNU Linear Programming Kit, GLPK version 4.33, written and maintained by [Makhorin \(2008\)](#). GLPK is open source software for solving large-scale linear programming, mixed integer linear programming, and other related problems. The R port provides a high level interface to the low level C interface of the C solver. The interface was written by [Theussl & Hornik \(2009\)](#), the former author is also the maintainer of the package. The contributed R package Rglpk is Rmetrics' default solver for linear programming problems.

```

> listDescription(Rglpk)
Package: Rglpk
Version: 0.6-0
Title: R/GNU Linear Programming Kit Interface
Description: R interface to the GNU Linear Programming Kit. GLPK
             is open source software for solving large-scale linear
             programming (LP), mixed integer linear programming (MILP)
             and other related problems.
Authors@R: c(person("Stefan", "Theussl", role = c("aut", "cre"),
                   email = "Stefan.Theussl@R-project.org"), person("Kurt",
                   "Hornik", role = "aut"), person("Christian", "Buchta", role
                   = "ctb"), person("Andrew", "Makhorin", role = "cph"),
                   person("Timothy A.", "Davis", role = "cph"),
                   person("Niklas", "Sorensson", role = "cph"), person("Mark",
                   "Adler", role = "cph"), person("Jean-loup", "Gailly", role
                   = "cph"))
Depends: slam (>= 0.1-9)
SystemRequirements: GLPK library package (e.g., libglpk-dev on
                   Debian/Ubuntu)
License: GPL-2 | GPL-3
URL: http://R-Forge.R-project.org/projects/rglp/,
     http://www.gnu.org/software/glpk/
Packaged: 2014-06-15 20:23:05 UTC; theussl
Author: Stefan Theussl [aut, cre], Kurt Hornik [aut], Christian
        Buchta [ctb], Andrew Makhorin [cph], Timothy A. Davis
        [cph], Niklas Sorensson [cph], Mark Adler [cph], Jean-loup

```

```
Gailly [cph]
Maintainer: Stefan Theussl <Stefan.Theussl@R-project.org>
NeedsCompilation: yes
Repository: CRAN
Date/Publication: 2014-06-16 07:41:04
Built: R 3.1.2; x86_64-w64-mingw32; 2014-11-01 02:48:39 UTC;
      windows
```

A.8 RECOMMENDED PACKAGES FROM BASE R

methods (R Development Core Team, 2009a), and MASS (Venables & Ripley, 2008) are used by Rmetrics. The two packages are recommended R packages, which means that they are installed with the base R environment.

A.9 CONTRIBUTED RPACKAGES

sn (Azzalini, Azzalini) and robustbase (Rousseeuw et al., 2008) are two contributed packages used by Rmetrics. The package sn comes with functions for manipulating skew-normal and skew-t probability distributions, and for fitting them to data, in the scalar and in the multivariate case. sn itself depends on the package mnormt (Azzalini, 2009), which provides functions for computing the density and the distribution function of, and for generating random vectors from, the multivariate normal and multivariate t distributions. The package robustbase provides ‘essential’ robust statistics. The goal of the package is to provide tools allowing to analyze data with robust methods. This includes regression methodology including model selections and multivariate statistics where the authors strive to cover the book ‘Robust Statistics, Theory and Methods’ by Maronna, Martin & Yohai (2006).

A.10 RMETRICS PACKAGE: FPORTFOLIOBACKTEST

fPortfolioBacktest is used to perform portfolio backtests together with a performance analysis for rolling portfolios. The functionality of this package is now integrated into package fPortfolio and is no longer required to be loaded.

APPENDIX B

DESCRIPTION OF DATA SETS

```
> library(fPortfolio)
```

In the following be briefly describe the data sets used in this ebook.

B.1 DATA SET: SWX

SWX stands for the Swiss Exchange in Zurich. The SWX provides downloads for historical time series including equities, bonds, reits, their indices, and many other financial instruments. The data set SWX, which we provide here, contains three daily SWX market indices, the Swiss Performance Index SPI, the Swiss Bond Index SBI, and the Swiss ImmoFunds Index SII. In addition, the data set contains Pictet's Pension Fund Indices LP25, LP40, LP60 from the LPP2000 index family.

```
> colnames(SWX)
[1] "SBI" "SPI" "SII" "LP25" "LP40" "LP60"
> range(time(SWX))
GMT
[1] [2000-01-03] [2007-05-08]
> nrow(SWX)
[1] 1917
```

B.2 DATA SET: LPP2005

Pictet is one of Switzerland's largest private banks. The bank is well known for its Swiss Pension Fund Benchmarks LPP2000 and LPP2005. The family

of Pictet LPP indices was created in 1985 with the introduction of new regulations in Switzerland governing the investment of pension fund assets. Since then it has established itself as the authoritative pension fund index for Switzerland. In 2000, a family of three indices, called LP25, LP40, LP60, where the number denotes increasing risk profiles, was introduced to provide a suitable benchmark for Swiss pension funds. During the last years, new investment instruments have become available for alternative asset classes. With Pictet's LPP2005 indices the bank took this new situation into consideration. The LPP2005 keeps the family of the three indices now named LPP25, LPP40, and LPP60 by adding 'plus' to the name represented by the second P in the index names.

```
> colnames(LPP2005)
[1] "SBI" "SPI" "SII" "LMI" "MPI" "ALT" "LPP25" "LPP40" "LPP60"

> range(time(LPP2005))

GMT
[1] [2005-11-01] [2007-04-11]

> nrow(LPP2005)

[1] 377
```

B.3 DATA SET: SPISECTOR

The SPISECTOR data set also provides data from the Swiss exchange. It covers the *Swiss Performance Index*, SPI, and nine of its sectors. The sectors include basic materials, industrials, consumer goods, health care, consumer services, telecommunications, utilities, financial, and technology. The oil and gas sector index is not included since it was introduced later than the others.

```
> colnames(SPISECTOR)
[1] "SPI" "BASI" "INDU" "CONG" "HLTH" "CONS" "TELE" "UTIL" "FINA" "TECH"

> range(time(SPISECTOR))

GMT
[1] [1999-12-30] [2008-10-17]

> nrow(SPISECTOR)

[1] 2216
```


B.4 DATA SET: SMALLCAP

Scherer & Martin (2005) comes with several CRSP (Center for Research in Security Prices) data sets used in the book in examples. These data sets contain monthly data records of market-cap-weighted equities recorded between 1997 and 2001. One of the data sets, with 20 small cap equities, the MARKET index and the T90 rates are made available in the Rmetrics data file SMALLCAP¹.

```
> colnames(SMALLCAP)
[1] "MODI" "MGF" "MEE" "FCEL" "OII" "SEB" "RML"
[8] "AEOS" "BRC" "CTC" "TNL" "IBC" "KWD" "TOPP"
[15] "RARE" "HAR" "BKE" "GG" "GYMB" "KRON" "MARKET"
[22] "T90"

> range(time(SMALLCAP))
GMT
[1] [1997-01-31] [2001-12-31]

> nrow(SMALLCAP)
[1] 60
```

B.5 DATA SET: GCCINDEX

The Gulf Cooperation Council, GCC, is an organization of six Arab states which share many social and economic objectives. These states are Saudi Arabia, Bahrain, Oman, Qatar, United Arab Emirate, and Kuwait. The index was launched by MSCI Barra in January 2006 to reflect growing investor interest in this region. The GCC Countries Indices offer broad coverage (up to 99 MSCI Barra maintains two series indices for the GCC and Arabian Markets. One is applicable to international investors, while the domestic series is aimed at investors not constrained by foreign ownership limits). The indices have daily history back to May 31, 2002.

```
> colnames(GCCINDEX)
[1] "BAHDSC" "BAHSC" "KUWDSC" "OMADSC" "OMASC"
[6] "KSADSC" "UAEDSC" "UAESC" "QATSC" "GCCXSASC"
[11] "GCCSC"

> range(time(GCCINDEX))
GMT
[1] [2005-05-31] [2008-07-28]

> nrow(GCCINDEX)
[1] 825
```

¹Please note that the data provided in the Rmetrics data file are not those from the CRSP data base. The data records were obtained from free sources, such as Yahoo, amongst others. Therefore, the SMALLCAP data records are exactly the same as those from the CRSP database.

APPENDIX C

R MANUALS ON CRAN

The R core team creates several manuals for working with R¹. The platform dependent versions of these manuals are part of the respective R installations. They can be downloaded as PDF files from the URL given below or directly browsed as HTML.

<http://cran.r-project.org/manuals.html>

The following manuals are available:

- An Introduction to R is based on the former "Notes on R", gives an introduction to the language and how to use R for doing statistical analysis and graphics.
- R Data Import/Export describes the import and export facilities available either in R itself or via packages which are available from CRAN.
- R Installation and Administration.
- Writing R Extensions covers how to create your own packages, write R help files, and the foreign language (C, C++, Fortran, ...) interfaces.
- A draft of The R language definition documents the language per se. That is, the objects that it works on, and the details of the expression evaluation process, which are useful to know when programming R functions.
- R Internals: a guide to the internal structures of R and coding standards for the core team working on R itself.
- The R Reference Index: contains all help files of the R standard and recommended packages in printable form.

¹The manuals are created on Debian Linux and may differ from the manuals for Mac or Windows on platform-specific pages, but most parts will be identical for all platforms.

The \LaTeX or texinfo sources of the latest version of these documents are contained in every R source distribution. Have a look in the subdirectory `doc/manual` of the extracted archive.

The HTML versions of the manuals are also part of most R installations. They are accessible using function `help.start()`.

APPENDIX D

RMETRICS ASSOCIATION

Rmetrics is a non-profit taking association founded in 2007 in Zurich working in the public interest. Regional bodies include the Rmetrics Asia Chapter. Rmetrics provides support for innovations in statistical computing. Starting with the Rmetrics Open Source code libraries which have become a valuable tool for education and business Rmetrics has developed a wide variety of activities.

- **Rmetrics Research:** supporting research activities done by the Econophysics group at the Institute for Theoretical Physics at ETH Zurich.
- **Rmetrics Software:** maintaining high quality open source code libraries.
- **Rmetrics Publishing:** publication of various Rmetrics books as well as from contributed authors.
- **Rmetrics Events:** organizing lectures, trainings and workshops on various topics.
- **Rmetrics Juniors:** helping companies to find students for interim jobs such as reviewing and checking code for higher quality or statistical analyses of various problems.
- **Rmetrics Stability:** licensing of stability signals and indicators to describe changing environments.

RMETRICS RESEARCH

The Rmetrics Association is mainly run by the researchers working at the Econophysics group at the Institute for Theoretical Physics at ETH Zurich. Research activities include:

- PhD, Master, Bachelor and Semester Theses

- Papers and Articles
- Presentations on international conferences
- Sponsored and tailored theses for companies
- Paid student internships at ETH Zurich

Rmetrics SOFTWARE

Without the Rmetrics Open Source Software Project it wouldn't be possible to realize all the research projects done in the Econophysics Group at ETH in such a short time. But it is not only the Econophysics group who has profited from the Open Source Rmetrics Software, there are worldwide many other research institutes and companies that are using Rmetrics Software.

The Rmetrics Software environment provides currently more than 40 R packages authored and maintained by 22 developers from all over the world. Amongst others it includes topics about basic statistics, portfolio optimization, option pricing as well as ARMA and GARCH processes.

An "ohloh" evaluation in 2012 of the Rmetrics Software concluded with the following results:

- Mature, well-established codebase
- Large, active development team
- Extremely well-documented source
- Cocomo project cost estimation
 - Codebase: 367'477 Lines
 - Effort: 97 Person Years
 - Estimated Cost: USD 5'354'262

This powerful software environment is freely available for scientific research and even for commercial applications.

Rmetrics PUBLISHING

Rmetrics Publishing is an electronic publishing project with a [bookstore](http://finance.e-bookshelf.ch/)¹ offering textbooks, handbooks, conference proceedings, software user guides and manuals related to R in finance and insurance. Most of the

¹<http://finance.e-bookshelf.ch/>

books can be ordered and downloaded for free. The bookstore is sponsored by the Rmetrics Association and the ETH spin-off company Finance Online. For contributed authors our bookstore offers a peer-reviewing process and a free platform to publish and to distribute books without transferring their copyright to the publisher. You can find a list of our books on [page ii](#).

Rmetrics EVENTS

Trainings and Seminars are offered by Rmetrics for the most recent developments in R. Topics include all levels of knowledge:

- Basic R programming
- Advanced R project management
- Efficiently debugging code
- Speeding up code by e.g. byte compiling or using foreign language interfaces
- Managing big data
- Professional reporting by e.g. using R Sweave, knitr, Markdown or interactive R Shiny web applications and presentations

There also exists an Rmetrics Asia Chapter for teaching and training R with its home in Mumbai, India.

Besides that Rmetrics organizes a yearly international summer school together with a workshop for users and developers.

Rmetrics JUNIORS

The Rmetrics Juniors initiative helps companies to find students for interim jobs. This ranges from reviewing and checking code for higher quality, to building R projects from scratch, to statistical analyses of inhouse problems and questions. The work is done by experienced Rmetrics Juniors members, usually Master or PhD thesis students. This is an advisory concept quite similar to that offered by ETH Juniors.

Rmetrics STABILITY

Analyzing and enhancing the research results from the Econophysics Group at ETH Zurich and other research institutions worldwide, the Rmetrics Association implements stability and thresholding indicators. These indicators can then be licensed by industry.

In this context it is important to keep in mind that Rmetrics is an independent non-profit taking association. With the money we earn from the stability project, we support open source software projects, student internships, summer schools, and PhD student projects.

APPENDIX E

RMETRICS TERMS OF LEGAL USE

Grant of License

Rmetrics Association (Zurich) and *Finance Online* (Zurich) have authorized you to download one copy of this electronic book (eBook). The service includes free updates for the period of one year. *Rmetrics Association* (Zurich) and *Finance Online* (Zurich) grant you a nonexclusive, nontransferable license to use this eBook according to the terms and conditions specified in the following. This License Agreement permits you to install and use the eBook for your personal use only.

Restrictions

You shall not resell, rent, assign, timeshare, distribute, or transfer all or any part of this eBook (including code snippets and functions) or any rights granted hereunder to any other person.

You shall not duplicate this eBook, except for a single backup or archival copy. You shall not remove any proprietary notices, labels, or marks from this eBook and transfer to any other party.

The code snippets and functions provided in this book are for teaching and educational research, i.e. for non commercial use. It is not allowed to use the provided code snippets and functions for any commercial use. This includes workshops, seminars, courses, lectures, or any other events. The unauthorized use or distribution of copyrighted or other proprietary content from this eBook is illegal.

Intellectual Property Protection

This eBook is owned by the *Rmetrics Association* (Zurich) and *Finance Online* (Zurich) and is protected by international copyright and other intellectual property laws.

Rmetrics Association (Zurich) and *Finance Online* (Zurich) reserve all rights in this eBook not expressly granted herein. This license and your right to use this eBook terminates automatically if you violate any part of this agreement. In the event of termination, you must destroy the original and all copies of this eBook.

General

This agreement constitutes the entire agreement between you and *Rmetrics Association* (Zurich) and *Finance Online* (Zurich) and supersedes any prior agreement concerning this eBook. This Agreement is governed by the laws of Switzerland.

(C) 2007-2015 Rmetrics Association Zurich. All rights reserved.

BIBLIOGRAPHY

- Aragon, T. (2008). *EpiTools: R Package for Epidemiologic Data and Graphics*. cran.r-project.org.
- Azzalini, A. *The sn package: The skew-normal and skew-t distributions*. cran.r-project.org.
- Azzalini, A. (2009). *The mnormt package: The multivariate normal and t distributions*. cran.r-project.org.
- Azzalini, A. & Capitanio, A. (1999). Statistical applications of the multivariate skew-normal distribution. *J.Roy.Statist.Soc*, 61, 579–602.
- Azzalini, A. & Capitanio, A. (2003). Distributions generated by perturbation of symmetry with emphasis on a multivariate skew-t distribution. *J.Roy.Statist.Soc*, 65, 367–389.
- Azzalini, A. & Dalla Valle, A. (1996). The multivariate skew-normal distribution. *Biometrika*, 83(4), 715–726.
- Bacon, C. R. (2008). *Practical Portfolio Performance Measurement and Attribution* (2 ed.). John Wiley & Sons.
- Carl, P. & Peterson, B. G. (2008). *PerformanceAnalytics*. cran.r-project.org.
- Chambers, J. M., Cleveland, W. S., Kleiner, B., & Tukey, P. A. (1983). *Graphical Methods for Data Analysis*. Wadsworth & Brooks/Cole.
- Cleveland, W. S. (1979). Robust locally weighted regression and smoothing scatterplots. *J. Amer. Statist. Assoc.*, 74, 829–836.
- Cleveland, W. S. (1981). Lowess: A program for smoothing scatterplots by robust locally weighted regression. *The American Statistician*, 35, 54.
- De Giorgi, E. G. (2002). A note on portfolio selection under various risk measures. Working Paper 122.
- Filzmoser, P., Garrett, R. G., & Reimann, C. (2005). Multivariate outlier detection in exploration geochemistry. *Computers & Geosciences*, 31(5), 579–587.

- Forgy, E. W. (1965). Cluster analysis of multivariate data: Efficiency vs interpretability of classifications. *Biometrics*, 21, 768–769.
- Friendly, M. (2002). Corrrgrams: Exploratory displays for correlation matrices. *The American Statistician*, 56, 316–324. working paper.
- Gnanadesikan, R. & Kettenring, J. R. (1972). Robust estimates, residuals and outlier detection with multiresponse data. *Biometrics*, 28, 81–124.
- Goldfarb, D. & Idnani, A. (1982). Dual and primal-dual methods for solving strictly convex quadratic programs. In J. Hennart (Ed.), *Numerical Analysis* (pp. 226–239). Springer.
- Goldfarb, D. & Idnani, A. (1983). A numerically stable dual method for solving strictly convex quadratic programs. *Mathematical Programming*, (27), 1–33.
- Gschwandtner, M. & Filzmoser, P. (2009). *mvoutlier: Multivariate Outlier Detection Based on Robust Methods*. cran.r-project.org.
- Hartigan, J. A. & Wong, M. A. (1979). A k-means clustering algorithm. *Applied Statistics*, 28, 100–108.
- Hyndman, R. J. & Fan, Y. (1996). Sample quantiles in statistical packages. *The American Statistician*, 50, 361–365.
- Ibanez, F., Grosjean, P., & Etienne, M. (2009). *pastecs: Package for Analysis of Space-Time Ecological Series*. cran.r-project.org.
- Jarek, S. (2009). *mvnormtest: Normality test for multivariate variables*. cran.r-project.org.
- Kaufman, L. & Rousseeuw, P. J. (1990). *Finding Groups in Data: An Introduction to Cluster Analysis*. New York: Wiley.
- Keitt, T. (2007). *colorRamps*. cran.r-project.org.
- Kotsiantis, S. & Pintelas, P. (2004). Combining bagging and boosting. *International Journal of Computational Intelligence*, 1(4), 324–333.
- Lloyd, S. P. (1982). Least squares quantization in pcm. *IEEE Transactions on Information Theory*, 28, 128–137.
- MacQueen, J. B. (1967). Some methods for classification and analysis of multivariate observations. In *Proceedings of 5th Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, (pp. 281–297)., Berkeley. University of California Press.
- Maechler, M., Rousseeuw, P., Struyf, A., & Hubert, M. (2009). *The cluster Package*. cran.r-project.org.

- Makhorin, A. (2008). *GNU Linear Programming Kit Reference Manual*. www.gnu.org/software/glpk.
- Marazzi, A. (1993). *Algorithms, Routines and S Functions for Robust Statistics* (2 ed.). Chapman & Hall.
- Markowitz, H. (1952). Portfolio selection. *Journal of Finance*, 7, 77–91.
- Maronna, R. & Zamar, R. (2002). Robust estimates of location and dispersion of high-dimensional datasets. *Technometrics*, 44(4), 307–317.
- Maronna, R. A., Martin, D. R., & Yohai, V. J. (2006). *Robust Statistics: Theory and Method* (1 ed.). Wiley.
- Neuwirth, E. (2007). *RColorBrewer*. cran.r-project.org.
- Opgen-Rhein, R. & Strimmer, K. (2007). Accurate ranking of differentially expressed genes by a distribution-free shrinkage approach. *Statistical Applications in Genetics and Molecular Biology*, 6(1).
- Pflug, G. (2000). Some remarks on the value-at-risk and the conditional value-at-risk. In S. Uryasev (Ed.), *Probabilistic Constrained Optimization*, volume 38 (pp. 272–281). Dordrecht: Kluwer.
- R Development Core Team (2009a). *methods*. cran.r-project.org.
- R Development Core Team (2009b). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing.
- Rizzo, M. L. (2002). *A New Rotation Invariant Goodness-of-Fit Test*. PhD thesis, Bowling Green State University.
- Rizzo, M. L. & Szekely, G. J. (2008). *energy: E-statistics*. cran.r-project.org.
- Rockafellar, R. & Uryasev, S. (2000). Optimization of conditional value-at-risk. *The Journal of Risk*, 2(3), 21–41.
- Rousseeuw, P., Croux, C., Todorov, V., Ruckstuhl, A., Salibian-Barrera, M., Verbeke, T., & Maechler, M. (2008). *The robustbase package: Basic Robust Statistics*. cran.r-project.org.
- Rousseeuw, P. & Van Driessen, K. (1999). A fast algorithm for the minimum covariance determinant estimator. *Technometrics*, 41, 212–223.
- Rousseeuw, P. J. (1985). Multivariate estimation with high breakdown point. In W. e. a. Grossmann (Ed.), *Mathematical Statistics and Applications*, volume B (pp. 283–297). Budapest: Akademiai Kiado.

- Rousseeuw, P. J. & Leroy, A. M. (1987). *Robust Regression and Outlier Detection*. Wiley.
- Royston, P. (1982). An extension of shapiro and wilk's w test for normality to large samples. *Applied Statistics*, 31, 115–124.
- Ryan, J. A. (2008). *IBrokers: R API to Interactive Brokers Trader Workstation*. cran.r-project.org.
- Sams, R. (2009). *RBloomberg*. cran.r-project.org.
- Schaefer, J., Opgen-Rhein, R., & Strimmer, K. (2008). *The corpcor package: Efficient Estimation of Covariance and (Partial) Correlation*. cran.r-project.org.
- Schäfer, J. & Strimmer, K. (2005). A shrinkage approach to large-scale covariance matrix estimation and implications for functional genomics. *Statistical Applications in Genetics and Molecular Biology*, 4(1).
- Scherer, B. & Martin, R. D. (2005). *Introduction to Modern Portfolio optimization with NUOPT and S-PLUS* (1 ed.). New York: Springer New York.
- Sharpe, W. F. (1994). The sharpe ratio. *Journal of Portfolio Management*, 21(1), 49–58.
- Szekely, G. & Rizzo, M. (2005). A new test for multivariate normality. *Journal of Multivariate Analysis*, 93(1), 58–80.
- Szekely, G. J. (1989). Potential and kinetic energy in statistics. Lecture notes, Budapest Institute of Technology (Technical University).
- Szekely, G. J., Rizzo, M. L., & Bakirov, N. K. (2007). Measuring and testing independence by correlation of distances. *Annals of Statistics*, 35(6), 2769–2794.
- Theussl, S. & Hornik, K. (2009). *The Rglpk package: R/GNU Linear Programming Kit Interface*. cran.r-project.org.
- Tukey, J. W. (1977). *Exploratory Data Analysis*. Addison-Wesley.
- Tütüncü, R. H., Toh, K. C., & Todd, M. J. (2003). Solving semidefinite-quadratic-linear programs using sdpt3. *Mathematical Programming*, 95(2), 189–217.
- Vanini, P. & Vignola, L. (2001). Markowitz model.
- Venables, W. N. & Ripley, B. D. (2008). *The MASS Package*. cran.r-project.org.

- Wang, N. & Raftery, A. (2002). Nearest neighbor variance estimation (nnve): Robust covariance estimation via nearest neighbor cleaning. *Journal of the American Statistical Association*, 97, 994–1019.
- Wang, N., Raftery, A., & Fraley, C. (2008). *The covRobust package: Robust Covariance Estimation via Nearest Neighbor Cleaning*. cran.r-project.org.
- Weingessel, A. (2004). *quadprog - Functions to Solve Quadratic Programming Problems*. cran.r-project.org.
- Wright, K. (2009). *corrgram: Plot a correlogram*. cran.r-project.org.
- Würtz, D. (2009a). *The fBasics Package*. cran.r-project.org.
- Würtz, D. (2009b). *fCopulae: Rmetrics - Dependence Structures with Copulae*. cran.r-project.org.
- Würtz, D. (2009c). *The fImport Package*. cran.r-project.org.
- Würtz, D. & Chalabi, Y. (2009a). *The fPortfolio Package*. cran.r-project.org.
- Würtz, D. & Chalabi, Y. (2009b). *The timeDate Package*. cran.r-project.org.
- Würtz, D. & Chalabi, Y. (2009c). *The timeSeries Package*. cran.r-project.org.

INDEX

- assets
 - arrange, 108
 - correlation tests, 122
 - correlations, 115
 - correlogram plots, 118
 - fitting parameters, 95
 - image plots, 122
 - modelling, 93
 - normality tests, 93
 - scatterplots, 115
 - simulating assets sets, 96
- backtesting, 305
- bagged estimator, 139
- bar chart, 194
- bar plot, 194
- bootstrap aggregation, 49
- breakdown point, 236
- chi-square distribution, 225
- classes
 - S4, 133
- conditional Value at Risk, 241
- constraints
 - box constraints, 160
 - class representation, 155
 - constraint strings, 155
 - group constraints, 161
 - long-only, 158
 - non-linear constraints, 163
 - risk budget constraints, 162
 - short selling, 159
- copula, 96
- covariance, 41
- CSV, 9
- cumulated column statistics, 40
- data
 - class representation, 149
 - Internet downloads, 10
- data sets, 9
- daylight saving time, 7
- drawdown, 27
- efficient, 197
- efficient frontier, 190
- efficient frontier, 187, 257, 259
- estimation errors, 221
- extrema, 30
- extreme value theory, 39
- fast Fourier transform, 72
- feasible portfolio, 187
- generalized hyperbolic Student's
 - t, 74
- generalized hyperbolic Student's
 - t, 75
- global minimum variance portfolio, 190
- GNU linear programming kit, 145
- hard-rejection weights, 50
- intervals, 29
- keywords
 - abbreviate, 123
 - abline, 78

- add, 67, 80, 200
- aggregate, 20
- all, 20, 39, 159, 170
- any, 90
- as, 8, 14, 51, 67, 122, 158, 159, 170
- axis, 79, 80, 82
- box, 63, 67, 69, 78, 79, 110, 156, 158–160, 194
- by, 13, 14, 17, 20, 22, 51, 59, 63, 69, 141, 145, 170
- c, 14
- call, 102, 170
- case, 63
- cbind, 14
- character, 13, 14, 146
- class, 11, 13, 14, 18, 26, 170, 194, 200
- cm, 86
- coefficients, 140
- col, 63, 67, 69, 78, 194, 200
- colors, 63
- colours, 63, 69, 200
- control, 100, 103, 135, 147
- cor, 51, 123
- cov, 33, 36, 51
- data, 10, 20, 34, 63, 94, 100, 106, 110, 116, 124, 170, 292
- date, 17, 21
- default, 63, 78, 81–83, 91, 146, 170, 200, 283
- density, 71
- description, 170
- deviance, 138
- distance, 100, 102
- do, 90
- double, 287
- end, 18
- equal, 198, 200, 258
- example, 63
- expression, 292
- file, 8
- For, 14
- for, 14, 26, 29, 34, 38, 63, 78, 80, 82, 94, 100, 102, 103, 105, 106, 138, 146, 147, 159, 160, 170, 188, 200, 281, 283, 287
- formula, 292
- function, 8, 20, 67, 78–83, 124, 135, 138, 140, 145, 146, 153, 157, 170, 281, 283–285, 287, 290
- grid, 63, 67, 78, 79
- heat, 63, 86
- hist, 71
- if, 283
- image, 116, 123
- index, 26, 60, 63
- integer, 13, 14, 82, 105, 135, 144, 145
- invisible, 67
- is, 14, 20, 21, 63, 105, 289
- labels, 63, 67, 79, 80, 82, 102, 123, 194, 200
- legend, 194
- length, 13, 14, 22, 123, 283
- lines, 59, 78, 79, 81
- list, 100, 103, 135, 137, 140, 147, 157, 164, 284, 290, 292
- log, 71, 80
- logical, 8, 17, 63, 135, 146
- lower, 157–162, 164
- lowess, 26, 29, 67, 122
- margin, 79
- matrix, 51, 105, 140, 292, 293
- max, 121, 137
- mean, 33, 36, 51, 137, 153
- merge, 14, 102
- min, 121, 137
- missing, 14
- model, 134, 137
- mtext, 79
- names, 80, 121, 135, 137, 152, 292
- ncol, 63

nrow, 67
 numeric, 135, 140–142, 144, 289
 objects, 8, 13, 14, 18, 102
 offset, 14
 on, 8, 170, 194, 198, 257
 optim, 135
 optimize, 134, 137, 138
 options, 135, 147
 or, 11, 14, 17, 20, 26, 60, 78, 80, 123, 137, 146, 170, 200
 order, 11, 102
 pairs, 116
 palette, 69, 78, 86, 194
 panel, 121, 122
 plot, 8, 41, 59, 63, 69, 71, 74, 78–83, 110, 116, 121, 124, 188, 198, 200, 258
 points, 59, 78, 81, 105, 116, 122, 124, 141, 142, 144, 198, 200, 258
 pos, 194
 quantile, 34, 51, 74
 R, 86, 90
 rainbow, 86
 range, 80, 200
 rank, 140, 222
 rbind, 14
 return, 51, 110, 137, 138, 141, 142, 170, 188, 198, 200, 258, 290
 rev, 11
 rug, 63, 67, 78
 sample, 11, 33, 34, 36, 42, 139, 153, 222
 segments, 78
 seq, 8
 sequence, 20
 set, 34, 94, 100, 200
 show, 123
 single, 198, 200, 258
 sort, 11
 stars, 110
 start, 18

status, 135, 141, 142, 145
 sub, 82
 subset, 17
 sum, 105
 summary, 20, 33
 symbol, 80, 91
 symbols, 78, 82
 system, 90
 table, 91
 terrain, 86
 text, 79, 80, 82, 91, 121, 200
 time, 17, 59, 170, 296
 title, 63, 67, 78–80, 170, 200
 topo, 86
 trace, 135, 146
 upper, 156–162, 164
 var, 36
 vector, 38, 51, 67, 105, 135, 142, 146, 170, 289
 weighted, 29, 188, 194, 198, 258
 weights, 38, 135, 141, 142, 158–160, 170, 188, 194, 198, 200, 258, 287, 289, 292, 293
 which, 63, 80, 82, 102, 105, 200
 window, 17
 kurtosis, 37
 linear programming solver named, 145
 long-only, 201
 Mahalanobis distance, 50, 225
 maximum likelihood estimator, 233
 mean, 41
 mean excess loss, 241
 mean shortfall, 241
 mean-variance portfolio, 201
 minimum variance locus, 197
 minimum variance locus, 187, 190, 257, 259

- minimum volume ellipsoid, 224, 227
- Monte Carlo simulation, 201
- multivariate normal distribution, 224
- multivariate Student's t distribution, 236
- NAs, 16
- nearest neighbour variance estimation, 47, 139
- normal inverse Gaussian, 74, 75
- out-of-sample, 221
- outlier, 225
- outliers, 50
- peaks, 30
- pits, 30
- plot
 - covariance ellipse plot, 41
- portfolio
 - analysis, 1
 - specification, 135
- principal component analysis, 108
- quadratic programming solver, 145
- quantile-quantile plot, 73, 74
- quantiles, 38
- R classes
 - data.frame, 30
 - fHTEST, 94
 - fPFOLIOBACKTEST, 277, 279, 281, 283–285, 287, 288
 - fPFOLIOCON, 155–157, 171
 - fPFOLIODATA, 149
 - fPFOLIOSPEC, 133, 134, 285, 288
 - fPORTFOLIO, 169–171, 286, 287
 - hclust, 101
 - hexBinning, 124
 - kmeans, 104
 - matrix, 36
 - squareBinning, 124
 - timeDate, 7, 8, 314
 - timeSeries, 5, 7–13, 18, 20, 28–30, 33–36, 39, 40, 50, 51, 59–61, 68, 69, 71, 73, 94, 95, 97, 105, 115, 125, 131, 149, 150, 156, 284, 285, 315
 - ts, 30, 59
- R data
 - data, 284, 286
 - DowJones30, 217
 - LPP2005, 108
 - LPP2005.RET, 44, 150
 - SPISECTOR.RET, 299
 - SWX, 11, 26, 30, 63
- R functions
 - .cov.bagged, 49
 - .cov.nnve, 47, 48
 - .help, 71
 - .hist, 71
 - aggregate, 20
 - align, 13, 185
 - applySeries, 22
 - as.timeSeries, 8
 - as.timeSeries.foo, 8
 - assetsArrange, 108
 - assetsBasicStatsPlot, 110
 - assetsBoxStatsPlot, 110
 - assetsCorEigenPlot, 105
 - assetsCorgramPlot, 118, 119
 - assetsCorImagePlot, 122, 124
 - assetsCorTestPlot, 122
 - assetSelect, 107
 - assetsFit, 95, 97
 - assetsMeanCov, 41–43, 50
 - assetsMomentsPlot, 110
 - assetsOutliers, 51
 - assetsPairsPlot, 115, 116
 - assetsSelect, 99, 100, 102, 106
 - assetsSim, 96
 - assetsStarsPlot, 110
 - assetsTest, 93, 94

backtestPlot, 293, 301
backtestStats, 294
baggedEstimator, 139
barplot, 194
basicStats, 33–35, 110
box_, 66
boxL, 66, 228
boxPercentilePlot, 69
boxPlot, 69, 70, 113
boxplot, 69, 70
cbind, 15, 185
ceiling, 19
characterTable, 90
cm.colors, 86
cmPalette, 86
colCumstats, 40
colMeans, 36
colorLocator, 83
colorMatrix, 84
colorTable, 83
colQuantiles, 38
colStats, 40
copyright, 66
cor, 122
cor.test, 122
cov, 36, 37
cov.nnve, 47
cov.rob, 44, 225
cov.trob, 236
covEllipsesPlot, 43, 44
covEstimator, 139
covMcd, 44
covMcdEstimator, 139, 223
covOGKEstimator, 139
covRiskBudgetsPie, 194
covRiskBudgetsPlot, 201
cumulatedPlot, 63, 64
data, 8
density, 71–73
densityPlot, 71, 72
divPalette, 87
drawdowns, 27
drawdownsPlot, 67
drawdownsStats, 33, 35
drawdownStats, 27
durations, 28
ellipsePanel, 120, 121
emaSmoother, 288, 289, 293
end, 17–19
equidistWindows, 281, 306
fastCovMcdEstimator, 223
fastCovOGKEstimator, 229
fastMveEstimator, 225
feasiblePortfolio, 170, 188
floor, 19
focusPalette, 88
fooPalette, 86
fooPlot, 71, 74
formals, 134
frontierPlot, 200, 213, 270, 271
getCenterRob, 43
getControl, 147
getCovRob, 43
getData, 152
getEstimator, 140, 222
getModelParams, 141
getOptimize, 138, 143, 144
getRiskFreeRate, 144
getSmootherParams, 290
getSolver, 144
getStatistics, 153
getStrategyFun, 286
getStrategyParams, 287
getTailRisk, 140
getType, 138
getWindowsHorizon, 282
getWindowsParams, 284
grayPalette, 86
hclust, 100, 101
heat.colors, 86
heatPalette, 86
help, 71
help.start, 326
hexBinning, 124
hgrid, 66
hist, 71
histPlot, 71
isUnivariate, 30

- kurtosis, 37
- lines, 59
- logDensityPlot, 71, 73
- lowess, 29, 31, 67, 119–123
- lowessPanel, 120, 121
- maxdd, 164
- mcdEstimator, 139
- mean, 36–38
- merge, 16, 185
- monoPalette, 88
- monteCarloPoints, 202
- mtext, 78
- mveEstimator, 139, 225
- myEstimator, 140
- myPortfolioStrategy, 286, 287
- myRollingWindows, 284
- nnveEstimator, 139
- pairs, 115
- par, 78, 81
- periods, 23
- piePanel, 120, 121
- plot, 59, 68, 78, 92, 199, 200, 272
- plot.default, 81
- plot.ts, 59
- points, 59
- pointsPanel, 120, 121
- portfolioBacktest, 280
- portfolioBacktesting, 291–293, 296
- portfolioConstraints, 155, 156, 159, 162, 166
- portfolioData, 149, 150
- portfolioFrontier, 197, 198
- portfolioSmoothing, 292–294
- portfolioSpec, 134, 135, 139, 189
- print, 136, 151, 158, 173
- qqline, 73, 74
- qqnigPlot, 75
- qqnorm, 73
- qqplot, 73, 74
- qualiPalette, 87
- quantile, 38
- rainbow, 86
- rainbowPalette, 86
- rampPalette, 87
- rbind, 15, 16, 185
- readSeries, 10
- returnPlot, 63, 64
- returns, 26, 149
- rev, 11, 12
- Rglpk, 247
- rollapply, 22
- rollingCDaR, 296
- rollingSigma, 294
- rollingWindows, 282
- round, 19
- sample, 11
- seqPalette, 87
- series, 97, 115
- seriesPlot, 63, 64, 66, 67, 77
- setControl, 147
- setEstimator, 140, 152, 234
- setMethod, 29
- setModelParams, 141
- setNFrontierPoints, 144
- setObjective, 146
- setOptimize, 138
- setOptions, 147
- setRiskFreeRate, 144
- setSmootherParams, 290
- setSolver, 146, 183, 186, 248
- setStatus, 145
- setStrategyFun, 286
- setStrategyParams, 287
- setTailRisk, 140
- setTargetReturn, 142
- setTargetRisk, 142
- setTrace, 146
- setType, 138
- setWeights, 142, 188
- setWindowsHorizon, 284
- setWindowsParams, 284
- shrinkage, 48
- shrinkEstimator, 139
- signif, 19
- skewness, 37

- solve.QP, 182
- solveRglpk, 145
- solveRglpk.CVAR, 257
- solveRquadprog, 145
- sort, 11, 12, 185
- squareBinning, 124
- stars, 110, 114
- start, 17, 18
- stats, 294
- str, 135, 150, 280
- summary, 33, 34
- symbolTable, 92
- symnum, 122
- tailoredFrontierPlot, 200, 201, 259
- terrain.colors, 86
- terrainPalette, 86
- text, 78
- timPalette, 87
- title, 66, 78
- topo.colors, 86
- topoPalette, 86
- truncate, 19
- turnpoints, 30, 31
- var, 36
- vgrid, 66
- weightedReturnsPie, 194
- weightedReturnsPlot, 201
- weightsPie, 194
- weightsPlot, 201
- window, 17
- write.csv, 9
- R packages
 - cluster, vii, 105, 106
 - corpcor, 48, 49, 139, 233
 - corrgram, 118
 - covRobust, 47, 139
 - energy, 93, 95
 - epitools, 83
 - fAssets, 317
 - fBasics, 70, 217, 316
 - fCopulae, 96
 - fExtremes, 39
 - fImport, 10
 - fPortfolio, vii, 1, 8, 148, 185, 186, 200, 248, 277, 281, 285, 287, 291–294, 313, 319
 - fPortfolioBacktest, 319
 - graphics, 69, 70
 - grDevices, 86
 - IBrokers, 25
 - lpSolveAPI, 247
 - MASS, 43, 44, 50, 139, 222, 225, 236, 317, 319
 - methods, 29, 317, 319
 - mnormt, 319
 - mnormt), 317
 - mvnormtest, 93, 94
 - mvoutlier, vii, 50
 - pastecs, vii, 30, 31
 - PerformanceAnalytics, 27
 - periods, 22
 - quadprog, 182, 313, 317
 - RBloomberg, 25
 - Rdonlp2, 183
 - Rglpk, 313, 318
 - robustbase, 44, 47, 139, 222, 317, 319
 - Rsocp, 183
 - Rsymphony, 247
 - sn, 317, 319
 - stats, 20, 29, 59, 100
 - timeDate, 22, 314
 - timeSeries, 10, 22, 185, 314, 315
 - zoo, 8
- R slots
 - smoother, 289
- R variable
 - nAssets, 161, 163
- rate, 192
- re-weighted estimator, 50
- Real Estate Investment Trust, 321
- returns
 - compound, 26
 - simple, 26
- risk-free rate, 192

robust estimators, 41, 43, 221
robust locally weighted regression,
29

Robust statistics, 41

robustness to outliers, 236

scatter plot, 115, 116

Sharpe Ratio, 192

Sharpe ratio, 192

short selling, 259

shrinkage, 48

shrinkage estimator, 221, 233

skewness, 37

smoothing, 29, 30

specification

consistency checks, 147

model slot, 137

modifying, 137, 141, 145

optim slot, 145

portfolio slot, 141

stability, 221

Student's t distribution, 50

tangency portfolio, 192

transaction costs, 221

turnpoints, 29–31

Value at Risk, 241

ABOUT THE AUTHORS

Diethelm Würtz is professor at the Institute for Theoretical Physics, ITP, and for the Curriculum Computational Science and Engineering, CSE, at the Swiss Federal Institute of Technology in Zurich. He teaches Econophysics at ITP and supervises seminars in Financial Engineering at CSE. Diethelm is senior partner of Finance Online, an ETH spin-off company in Zurich, and co-founder of the Rmetrics Association.

Tobias Setz has a Bachelor and Master degree in Computational Science and Engineering from ETH in Zurich and has contributed with his thesis projects on wavelet and Bayesian change point analytics to this ebook. He is now a PhD student in the Econophysics group at ETH Zurich at the Institute for Theoretical Physics and maintainer of the Rmetrics packages.

Yohan Chalabi has a PhD in Physics from the Swiss Federal Institute of Technology in Zurich. After his PhD he left the Econophysics group at ETH Zurich at the Institute for Theoretical Physics. Up to that time Yohan was a co-maintainer of the Rmetrics packages.

William Chen has a master in statistics from University of Auckland in New Zealand. In the summer of 2008, he did a Student Internship in the Econophysics group at ETH Zurich at the Institute for Theoretical Physics. During his three months internship, William contributed to the portfolio backtest package.

Andrew Ellis read neuroscience and mathematics at the University in Zurich. He worked for Finance Online and did an internship in the Econophysics group at ETH Zurich at the Institute for Theoretical Physics. Andrew co-authored this ebook about portfolio optimization.

